RESKILL: Explicit Failure Attribution and Structured Repair for Interactive Language Agents

Mengyi Deng[1], **Xin Li**[1], **Duyi Pan**[1], **Zilin Wang**[1], **Zhiwei Li**[1], **Zhijiang Guo**[1,2,†], **Wei Wang**[1,2,†]
[1]Information Hub, The Hong Kong University of Science and Technology (Guangzhou), China
[2]The Hong Kong University of Science and Technology, Hong Kong SAR
{mdeng974, xli420, dpan457, zwang374, zli404}@connect.hkust-gz.edu.cn
zhijiangguo@hkust-gz.edu.cn, weiwcs@ust.hk

arXiv:2609.15684v1 [cs.CL] 14 Sep 2026

Abstract

Language agents increasingly rely on reusable skills, but post-failure repair is often handled by opaque one-shot reflection: a model generates a skill patch without explicitly maintaining how failure explanations relate to candidate repairs or how unsuccessful retests should influence later edits. We introduce RESKILL, a structured repair framework that maintains an explicit repair state across repair rounds. Given a failed rollout, the framework links failure hypotheses to candidate skill patches, selects local repairs through coverage-based attribution, retests the edited skill set in the environment, and uses retest outcomes to guide subsequent repair updates. The language model supplies structured repair factors, while the repair procedure records them, compares local skill patches by how well they address active failure explanations, and carries unsuccessful retest outcomes into later repair rounds. We evaluate RESKILL on ALFWorld and TextCraft across three model sizes under fixed repair budgets. RESKILL obtains the strongest final success in all six benchmark–model settings, improving average final success by 3.7 percentage points over direct repair and 3.3 points over hypothesis-conditioned repair. These results suggest that explicit attribution alone is insufficient; durable improvement emerges when attribution is integrated with repair selection and persistent retest-conditioned update.

1 Introduction

Interactive language agents increasingly rely on reusable skills and behavioral rules for long-horizon tool-use and embodied interaction (Wang et al., 2023a; Yang et al., 2026b). Prior work has studied reusable skill acquisition (Wang et al., 2023a; Chen et al., 2026), iterative self-improvement, and continual skill evolution (Liu et al., 2026a) from reflection and interaction experience (Zheng et al.; Yang et al., 2026a). Despite this progress, reusable skills remain brittle, and repairing them after failure is difficult because a single failed rollout may point to multiple plausible, potentially co-occurring causes, such as a missing precondition, incorrect object handling, or an invalid action. Because candidate edits may address different subsets of these causes, repair requires deciding which edit to apply and, after an unsuccessful retest, which explanations should guide the next repair round.

However, existing repair approaches often handle post-failure correction by asking a language model to rewrite the skill set directly from a failed trajectory (Shinn et al., 2023; Madaan et al., 2023a; Liu et al., 2026a; Patel et al., 2024). While such approaches can produce useful local patches and progressively improve reusable skills (Shen et al., 2026; Zhang et al., 2026; Yang et al., 2026a), they typically conflate diagnosis, repair generation, and repair selection into a single generation step. As a result, post-failure repair becomes a repair credit-assignment problem: a failed trajectory may support several plausible explanations, but an effective edit may address only a subset of them. If the system does not record which explanations each patch is meant to address, or how retest feedback changes support for those explanations, iterative repair can repeatedly apply ineffective edits or overfit to a spurious diagnosis. Effective repair therefore requires reducing failure ambiguity, attributing observed outcomes to the explanations under repair, and carrying credit across repair rounds.

We formulate post-failure skill repair as hypothesis-conditioned repair selection under environment feedback. Given a failed trajectory, a language model generates a support set of failure hypotheses, candidate skill patches, and a coverage matrix describing which patches address which hypotheses. The repair algorithm tracks prioritized failure hypotheses, compares local patches by how well they cover the active hypotheses, applies the selected patch to the skill set, and retests the task

in the environment. If a retest fails, the resulting post-repair trace is converted into updated hypotheses for the next repair round. Figure 1(b) illustrates these components in a concrete ALFWorld example. The agent is asked to cool a mug and place it on the coffee machine, but instead cools an egg and repeatedly examines the mug without picking it up. From this evidence, RESKILL forms two active hypotheses: incorrect object selection and a missing pickup action. It selects a patch that requires picking up the mug before cooling it, addressing both hypotheses, and validates the repair through a successful retest.

We evaluate RESKILL in two interactive settings that stress different forms of skill use: grounded household interaction in ALFWorld (Shridhar et al., 2020) and symbolic crafting in TextCraft from AgentGym (Xi et al., 2024). Across the resulting six benchmark–model settings, RESKILL improves final repair success under fixed repair budgets, with average gains of 3.7 and 3.3 percentage points over two competitive repair baselines. The results suggest that grounding repair decisions in hypothesized failure causes, comparing candidate edits against competing explanations, and revising repair credit through environment retesting improve post-failure skill repair across model scales. In summary, our contributions are:

- We formulate post-failure skill repair as a repair credit-assignment problem over ambiguous failure explanations, candidate edits, and retest outcomes.
- We introduce RESKILL, a structured framework that links failure hypotheses, candidate repairs, coverage-based selection, environment retesting, and outcome-conditioned repair updates.
- We instantiate RESKILL in ALFWorld and TextCraft without changing model weights, showing improvements across three model sizes and providing implementation-grounded analyses of repair traces.

2 Related Work

Reusable skill for agents. Recent language-agent systems increasingly externalize reusable behavior into skills, tool calls, or programmatic action interfaces. In robotics and tool use, language models can be grounded through affordances, APIs, or code-like policies (Ahn et al., 2022b; Huang et al., 2022; Liang et al., 2023; Huang et al., 2024; Xu et al., 2026; Deng et al., 2026; Zhang et al., 2025). In open-ended embodied environments, agents accumulate competence through executable skill libraries, retrieved memories, or structured action knowledge (Wang et al., 2023a; Zhu et al., 2023; Wang et al., 2024; Liu et al., 2024b; Fan et al., 2022; Park et al., 2023; Liu et al., 2026b).

Recent skill-centric work makes this substrate more explicit: SkillAct studies prompting agents with reusable skill abstractions (Liu et al., 2024a); SKILL0 studies internalizing skill context through an in-context reinforcement-learning curriculum (Lu et al., 2026); SkillX constructs transferable skill knowledge bases from trajectories (Wang et al., 2026); and SkillGen and SkillOS study skill synthesis and curation from execution experience (Ma et al., 2026; Ouyang et al., 2026). These works establish reusable skills as a practical interface for long-horizon behavior. RESKILL builds on this view by studying the post-failure repair stage: how failed interactions can be attributed to competing skill-level explanations and converted into targeted skill updates.

Feedback-driven self-improvement. Reflection methods show that feedback can improve future behavior. ReAct interleaves reasoning and acting during interaction (Yao et al., 2022); DEPS uses language-model explanation and selection for open-world planning (Wang et al., 2023b); Reflexion stores verbal reflections after failed trials (Shinn et al., 2024); and Self-Refine iteratively improves outputs using self-feedback (Madaan et al., 2023b).

Related work also separates generation from selection. Tree of Thoughts expands multiple reasoning paths before choosing among them (Yao et al., 2023); SayCan combines language scores with affordance values for action selection (Ahn et al., 2022a); and tool-use methods select or invoke external capabilities as part of solving a task (Schick et al., 2023). These methods demonstrate the value of separating proposal generation from downstream selection. However, they do not extend this separation to post-failure skill repair, where candidate edits must be linked to competing failure hypotheses. RESKILL makes this link explicit through hypothesis–patch coverage and environment retesting.

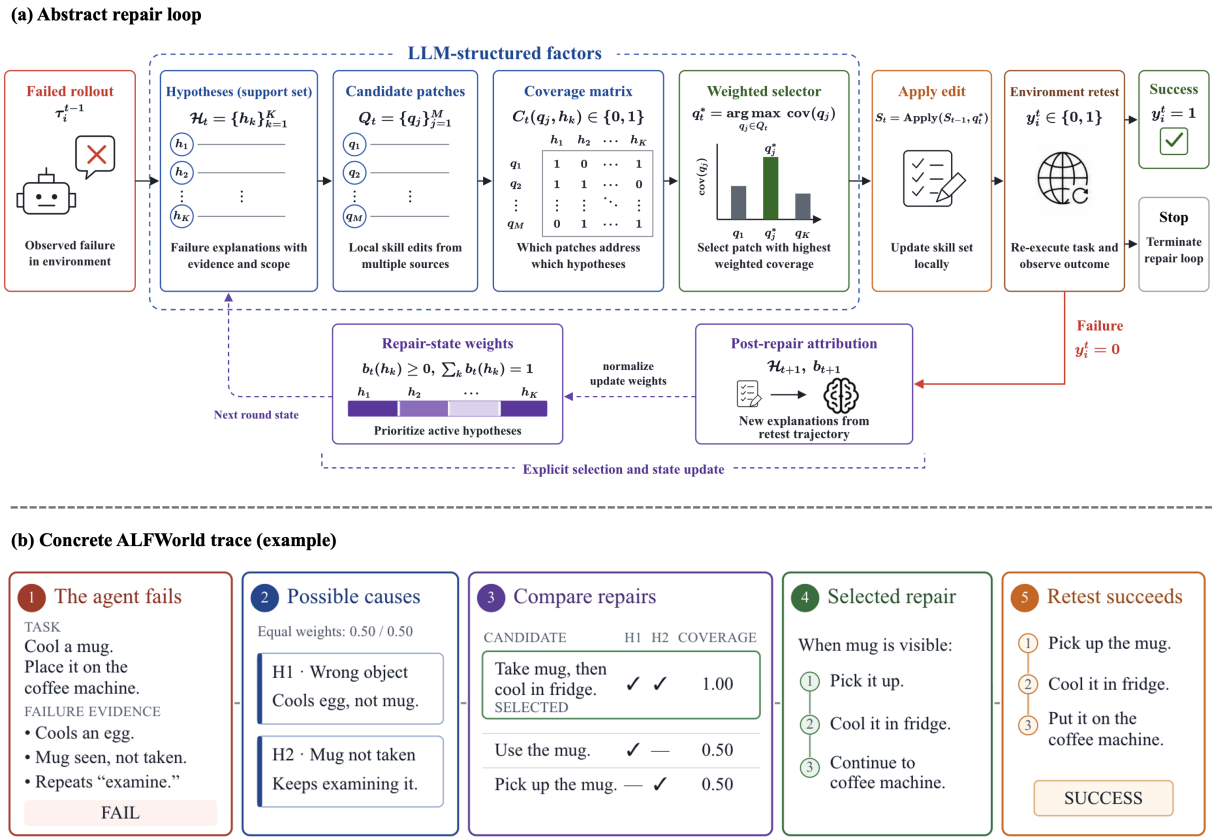


Figure 1: Overview of RESKILL for structured post-failure skill repair. (a) Given a failed rollout, the framework constructs a set of failure hypotheses, generates candidate skill patches, selects a patch by weighted hypothesis coverage, and updates the hypothesis state through environment retesting. (b) A concrete ALFWorld example showing how failure evidence gives rise to multiple hypotheses and candidate patches. The repair that best addresses the active hypotheses is then selected and validated through environment retesting.

3 Method

RESKILL is a post-failure repair framework for editable agent skills. As shown in Figure 1, given a failed trajectory, the language model proposes multiple failure hypotheses and local candidate patches, and identifies which hypotheses each patch addresses. RESKILL maintains weights over the active hypotheses, selects a patch based on its weighted coverage, and retests the edited skill set in the environment. If the retest succeeds, repair stops; otherwise, the failed retest trajectory is used to update the hypotheses and their weights for the next round.

3.1 Post-Failure Skill Repair

Let x_i denote an interactive task, S_0 an initial skill set, and $\pi_\theta(\cdot \mid S)$ the frozen language-agent policy induced by prompting the model with skill set S. Model parameters θ remain fixed throughout repair; repair modifies only the external skill set used to condition the policy. The same underlying frozen language model is used for task execution and structured repair generation through separate prompts and calls. The first pass executes a complete rollout under S_0, producing an interaction trajectory

$$\tau_i^0 = (o_0, a_0, o_1, \ldots, a_{K_i^0-1}, o_{K_i^0})$$

and a binary success indicator $y_i^0 \in \{0, 1\}$. Here o_k is the environment observation after k actions, a_k is the action emitted after observing o_k, and K_i^0 is the number of agent actions before the episode terminates or reaches the task budget. Repair is triggered only after a completed unsuccessful rollout, not during intermediate action steps.

At repair round $t \geq 1$, provided that $y_i^{t-1} = 0$, RESKILL uses the failed trajectory τ_i^{t-1} as evidence to select a patch q_t. Applying the patch changes the skill set locally from S_{t-1} to:

$$S_t = \mathrm{Apply}(S_{t-1}, q_t).$$

The operator Apply inserts the patch into the skill section identified by its target scope and leaves un-

related sections unchanged. The agent then retests the task under S_t, producing a new trajectory τ_i^t and success indicator y_i^t. The repair budget R limits the number of post-failure retests. For an evaluation set of N tasks and any budget $r \in \{0, \ldots, R\}$, we report cumulative task success, including both first-pass successes and tasks recovered within r repair rounds:

$$\mathrm{Succ@}r = \frac{1}{N} \sum_{i=1}^{N} \mathbb{I}\left[y_i^0 = 1 \vee \exists t \leq r : \; y_i^t = 1\right].$$

Tasks solved on the first pass count at every budget, while initially failed tasks can be recovered by any later repair retest. Each attempted patch is evaluated by re-executing the triggering task under the edited skill set. The selected patch, retest trajectory, and outcome are recorded for the next repair round.

3.2 Structured Repair State

At repair round t, RESKILL organizes the information needed for the next repair decision into a structured state rather than directly rewriting the skill set from the failed trajectory. This state contains the current skill set S_{t-1}, the latest failed trajectory τ_i^{t-1}, the active hypothesis set $\mathcal{H}_t$ with weights $b_t(h)$, and the history of previously attempted patches and their retest outcomes. Using this state, the language model generates candidate local patches and coverage relations indicating which active hypotheses each patch addresses. RESKILL selects the patch $q_t^\star$ with the highest weighted coverage, applies it to the skill set, and retests the task in the environment. If the retest fails, the new trajectory is used to update the hypothesis set and its weights, while the attempted patch and its outcome are recorded for the next repair round.

3.3 Hypothesis State

RESKILL represents a failed trajectory as a finite support of possible failure explanations. At the first repair round, the model produces an initial hypothesis set $\mathcal{H}_1$. After an unsuccessful retest, post-repair attribution produces the next set $\mathcal{H}_{t+1}$. At round t, the active support is

$$\mathcal{H}_t = \{h_1, \ldots, h_K\},$$

where each hypothesis contains a concise failure explanation, supporting evidence from the trajectory, the affected skill scope, and a repair hint. Across rounds, hypotheses are aligned when they represent the same underlying failure cause within the same skill scope; hypotheses that cannot be aligned with an earlier explanation are treated as newly introduced. RESKILL maintains a normalized weight $b_t(h)$ for each $h \in \mathcal{H}_t$. The first repair round starts from uniform weights,

$$b_1(h) = \frac{1}{|\mathcal{H}_1|}, \qquad h \in \mathcal{H}_1,$$

and later rounds update the weights from the retest outcome. These weights determine which active explanations receive priority when the selector compares candidate edits. A high-weight hypothesis contributes to every candidate that covers it, allowing the selector to compare narrower and broader local edits under the same coverage-based comparison.

3.4 Repair Candidates

Given the latest failed trajectory τ_i^{t-1}, the current skill set S_{t-1}, the active hypotheses $\mathcal{H}_t$, and the patches attempted in earlier rounds, RESKILL builds a finite candidate pool $\mathcal{Q}_t$:

$$\mathcal{Q}_t = \mathcal{Q}_t^{\mathrm{trace}} \cup \mathcal{Q}_t^{\mathrm{hyp}} \cup \mathcal{Q}_t^{\mathrm{prop}} \cup \mathcal{Q}_t^{\mathrm{guard}}.$$

The superscripts index candidate sources used within a single RESKILL round. $\mathcal{Q}_t^{\mathrm{trace}}$ contains a trajectory anchor patch written from the observed failure evidence. $\mathcal{Q}_t^{\mathrm{hyp}}$ contains a hypothesis-conditioned anchor patch written with access to the active support $\mathcal{H}_t$. $\mathcal{Q}_t^{\mathrm{prop}}$ contains additional targeted proposals that cover different parts of $\mathcal{H}_t$, including edits that intentionally address overlapping hypotheses. $\mathcal{Q}_t^{\mathrm{guard}}$ contains interface-preserving guard candidates used when a local grounding, admissibility, or executable-format correction can be derived from the observed trace.

Each candidate $q \in \mathcal{Q}_t$ is represented as a local skill patch with an intended scope, the evidence it claims to address, and the skill text to be inserted or revised. Candidates are screened before selection: the patch must be reusable beyond the current episode, grounded in the observed trajectory, compatible with the agent's action interface, and narrow enough to avoid rewriting unrelated skills. The pool therefore gives the selector alternatives to compare while keeping repair local and bounded.

3.5 Coverage-Based Selection

For each candidate patch $q \in \mathcal{Q}_t$ and each hypothesis $h \in \mathcal{H}_t$, the model assigns a binary coverage

indicator,

$$C_t(q, h) \in \{0, 1\}.$$

The model assigns $C_t(q, h) = 1$ only when the proposed edit directly addresses the failure cause represented by h and applies to the relevant skill scope. Before selection, candidate patches are required to specify a local, executable edit grounded in the failed trajectory. The selector scores each candidate by the active hypothesis weight it covers:

$$\mathrm{cov}_t(q) = \sum_{h \in \mathcal{H}_t} b_t(h) C_t(q, h).$$

A patch can therefore receive a high score by addressing either one high-weight hypothesis or several lower-weight hypotheses. The selected patch is:

$$q_t^\star = \arg\max_{q \in \mathcal{Q}_t} \mathrm{cov}_t(q),$$

with ties resolved in favor of more local and evidence-grounded edits. This design separates patch generation from the final repair decision: the model proposes candidate edits and their coverage relations, while RESKILL applies an explicit coverage-based selection rule.

3.6 Outcome-Conditioned Update

After applying $q_t^\star$, the agent retests the task in the environment. If $y_i^t = 1$, repair terminates. Otherwise, the model uses the failed retest trajectory τ_i^t, the previous hypotheses, and the attempted patch to construct $\mathcal{H}_{t+1}$ for the next repair round. If no post-repair hypothesis supported by the retest trajectory is obtained, RESKILL retains $\mathcal{H}_t$ and b_t while continuing from the latest failed trajectory. RESKILL aligns recurring explanations with their counterparts from the previous round and treats unmatched explanations as newly introduced. The total active hypothesis weight covered by the selected patch is

$$m_t = \sum_{h \in \mathcal{H}_t} b_t(h) C_t(q_t^\star, h).$$

This quantity determines the initial weight assigned to explanations newly exposed by an unsuccessful retest. Intuitively, when a patch addressing active hypotheses still fails, the retest increases the relative priority of explanations newly exposed after that intervention. Recurring hypotheses retain their previous weights, while m_t is divided evenly among newly introduced hypotheses. The resulting weights are then normalized over $\mathcal{H}_{t+1}$. If all assigned weights are zero, RESKILL uses uniform weights. The next candidate pool $\mathcal{Q}_{t+1}$ is generated from the failed retest trajectory, the updated hypotheses, and the patches attempted in earlier rounds.

Algorithm 1 RESKILL for Post-Failure Skill Repair

Input: task x_i, initial skill set S_0
repair budget R, frozen policy π_θ
Output: final skill set, success flag, and repair trace
Roll out $\pi_\theta(\cdot \mid S_0)$ to obtain τ_i^0 and y_i^0
if $y_i^0 = 1$ **then**
  return S_0, success, first-pass trace
end if
Generate initial hypotheses $\mathcal{H}_1$ from τ_i^0
Initialize $b_1(h) = 1/|\mathcal{H}_1|$ for $h \in \mathcal{H}_1$
for $t = 1, \ldots, R$ **do**
  Build candidate repairs $\mathcal{Q}_t$ from the current trace and repair state
  Screen candidates for locality, grounding, validity, and non-redundancy
  Construct repair–hypothesis coverage relations $C_t(q, h)$ for $q \in \mathcal{Q}_t$ and $h \in \mathcal{H}_t$
  Select $q_t^\star$ using covered weight and fixed locality/-grounding criteria
  Apply the patch: $S_t \leftarrow \mathrm{Apply}(S_{t-1}, q_t^\star)$
  Retest with S_t to obtain τ_i^t and y_i^t
  if $y_i^t = 1$ **then**
    return S_t, success, repair trace
  end if
  Generate post-repair hypotheses $\mathcal{H}_{t+1}$
  if $\mathcal{H}_{t+1}$ contains usable hypotheses **then**
    Update b_{t+1} over $\mathcal{H}_{t+1}$ using the outcome-conditioned update
  else
    Carry forward the previous repair state
  end if
end for
return S_R, failure, repair trace

Algorithm 1 summarizes the repair loop for one task. Each attempted repair records the active hypotheses, candidate patches, coverage relations, selected patch, retest outcome, and next-round state. When a retest fails, these records guide the next repair decision by preserving persistent explanations and incorporating newly exposed ones.

4 Experiments

4.1 Experimental Setting

Our experiments separate first-pass skill use, one-shot repair, and structured multi-round repair. All methods start from the same failure pool, edit the same skill-file format, and use frozen language models; improvements come only from external skill edits and retesting. We evaluate on two interactive benchmarks with complementary failure

Model	Method	ALFWorld						TextCraft		
		Overall	Pick	Look	Clean	Heat	Cool	Overall	D1	D2
Qwen3-1.7B	No skills	7.7	9.0	19.4	5.2	2.6	4.3	26.3	70.9	5.6
Qwen3-1.7B	With skills	16.8	9.0	35.5	27.6	12.8	10.9	21.1	74.6	5.2
Qwen3-1.7B	Direct	30.7 (+13.9)	22.0 (+13.0)	48.4 (+12.9)	46.6 (+19.0)	**30.8** (+18.0)	**17.4** (+6.5)	28.5 (+7.4)	85.1 (+10.5)	13.9 (+8.7)
Qwen3-1.7B	Hypothesis	27.7 (+10.9)	18.0 (+9.0)	48.4 (+12.9)	48.3 (+20.7)	23.1 (+10.3)	13.0 (+2.1)	28.9 (+7.8)	86.6 (+12.0)	14.2 (+9.0)
Qwen3-1.7B	RESKILL	**33.2** (+16.4)	**25.0** (+16.0)	**58.1** (+22.6)	**53.4** (+25.8)	23.1 (+10.3)	**17.4** (+6.5)	**30.7** (+9.6)	**91.8** (+17.2)	**15.3** (+10.1)
Qwen2.5-3B	No skills	9.9	9.0	25.8	6.9	5.1	8.7	13.0	29.9	5.2
Qwen2.5-3B	With skills	22.6	22.0	45.2	17.2	17.9	19.6	9.9	29.9	4.9
Qwen2.5-3B	Direct	53.3 (+30.7)	**50.0** (+28.0)	74.2 (+29.0)	**67.2** (+50.0)	41.0 (+23.1)	39.1 (+19.5)	17.5 (+7.6)	50.7 (+20.8)	9.4 (+4.5)
Qwen2.5-3B	Hypothesis	52.2 (+29.6)	47.0 (+25.0)	**77.4** (+32.2)	51.7 (+34.5)	51.3 (+33.4)	**47.8** (+28.2)	22.8 (+12.9)	61.2 (+31.3)	14.2 (+9.3)
Qwen2.5-3B	RESKILL	**54.7** (+32.1)	47.0 (+25.0)	**77.4** (+32.2)	58.6 (+41.4)	**59.0** (+41.1)	**47.8** (+28.2)	**29.4** (+19.5)	**85.8** (+55.9)	**15.6** (+10.8)
Qwen3-4B	No skills	54.7	73.0	48.4	37.9	64.1	32.6	42.4	81.3	24.3
Qwen3-4B	With skills	80.7	82.0	93.5	89.7	61.5	73.9	47.4	87.3	28.8
Qwen3-4B	Direct	93.8 (+13.1)	96.0 (+14.0)	**100.0** (+6.5)	96.6 (+6.9)	**92.3** (+30.8)	82.6 (+8.7)	48.1 (+0.7)	89.6 (+2.3)	28.8 (+0.0)
Qwen3-4B	Hypothesis	93.4 (+12.7)	**98.0** (+16.0)	**100.0** (+6.5)	94.8 (+5.1)	87.2 (+25.7)	82.6 (+8.7)	49.3 (+1.9)	91.0 (+3.7)	29.9 (+1.1)
Qwen3-4B	RESKILL	**95.6** (+14.9)	97.0 (+15.0)	**100.0** (+6.5)	**98.3** (+8.6)	**92.3** (+30.8)	**89.1** (+15.2)	**50.7** (+3.3)	**92.5** (+5.2)	**31.2** (+2.4)

Table 1: First-pass and final repair success rates in percent. Overall reports aggregate success; ALFWorld columns report task families and TextCraft columns report visible recipe depth. No skills and With skills are unrepaired first-pass rows; for TextCraft, With skills is the repair-pool R0 used for repair deltas. Direct, Hypothesis, and RESKILL report final cumulative success after R7 for ALFWorld and R4 for TextCraft. Gray parentheses show change relative to the same model's With-skills/R0 row.

modes:

- ALFWorld contains 274 household tasks, split into 140 seen and 134 unseen tasks. For reporting, we group the two object-placement variants into a single Pick family and retain Look, Clean, Heat, and Cool as separate task families. The initial ALFWorld skill file is adapted from the ALFWorld skill substrate used by SKILL0 (Lu et al., 2026): each rollout receives general skills plus the mapped task-family section, and repair notes are appended to the corresponding local section.
- TextCraft uses the AgentGym TextCraft task set, where agents must follow visible recipe graphs and emit exact executable crafting actions. Because TextCraft has no external task-family skill file, we initialize it with a compact LLM-authored, recipe-free skill file containing only response-format, observation-grounding, exact-name, recipe-chain, and count-discipline constraints. We report TextCraft by visible recipe depth: D1 tasks require direct recipe use, while D2 tasks require one additional intermediate dependency.

Models and scale. Our primary evaluation uses Qwen3-1.7B, Qwen2.5-3B, and Qwen3-4B. Because repair is triggered only after an initial failure, informative evaluation requires a sufficiently large and diverse residual failure pool. As first-pass performance approaches ceiling, fewer tasks remain eligible for repair and the headroom for comparing repair methods decreases. We therefore focus on open-weight models in the 1.7B–4B range, which represent locally deployable agents while providing sufficient repair opportunities across capability levels.

We compare RESKILL with two single-patch repair procedures. Direct Repair generates a single local patch at each repair round from the current failed trajectory before retesting. Hypothesis Repair first generates failure hypotheses at each round, then writes one local patch conditioned on them. Neither baseline builds a multi-candidate repair pool, a coverage matrix, or an outcome-conditioned repair state. In contrast, RESKILL keeps multiple hypotheses and candidate edits, selects by weighted coverage, and carries failed retest evidence into later rounds. ALFWorld is tracked through seven repair rounds, while TextCraft is tracked through four. In Table 1, No skills (Li et al., 2026) and With skills are unrepaired first-pass results; the three repair rows report final cumulative success after the full repair budget, i.e., R7 for ALFWorld and R4 for TextCraft. Figure 2 and Appendix B show the intermediate budgets.

4.2 Skill Use and Repair Effects

Skills provide useful but incomplete guidance. Table 1 first separates skill use from post-failure repair. In ALFWorld, the With skills row improves most task families for all three model sizes, with especially large gains on Look and Clean for the smaller models. TextCraft is more uneven: skills can improve depth-1 behavior, but the repair-pool R0 remains low for smaller models, and depth-2 tasks remain difficult. This pattern motivates post-failure editing: skills provide useful structure, but a failed rollout can still expose missing conditions, overly broad rules, or brittle surface-form behavior.

One-shot repair works when the trace exposes a single correction. Direct Repair is strongest when the failure feedback already points to a concrete local edit. In ALFWorld, it matches or exceeds RESKILL on several family columns: it is strongest on Qwen3-1.7B Heat, Qwen2.5-3B Pick and Clean, and ties Qwen3-4B Heat. These cases often involve a specific procedural omission or action-interface mismatch, where a single patch can be sufficient. This pattern identifies a natural strength of one-shot repair: when the evidence localizes a single correction, direct editing can recover quickly. The complementary regime is where RESKILL is most useful: the failed trace supports several plausible edits, or an unsuccessful retest should redirect the next repair rather than trigger another unconstrained rewrite.

Structured attribution helps distinguish between plausible repairs. RESKILL obtains the best full-suite ALFWorld score for all three model sizes while also giving the best TextCraft D1 and D2 results for all three models. Its advantage is clearest when the repair choice is underdetermined. In ALFWorld, this appears in families where a failed trajectory can implicate object search, precondition handling, action ordering, or family scope. In TextCraft, surface form and planning interact: a wrong action may reveal an ingredient-string mismatch, but a useful skill edit may also need to preserve counts and follow the visible recipe chain. Direct Repair often captures one constraint, and Hypothesis Repair can name the failure more explicitly; RESKILL compares candidate edits by the hypotheses they cover before committing to a patch.

Model strength changes the repair regime. The smaller models leave larger recoverable failure pools, so the benefit of repair is more visible across both domains. The largest TextCraft margin appears for Qwen2.5-3B, where RESKILL improves both D1 and D2 substantially over the With skills row and over the two single-patch baselines. Qwen3-4B starts from stronger first-pass behavior, so the remaining repair pool is smaller and family-level differences are often close. Even in this higher-ceiling regime, RESKILL remains best on TextCraft D1/D2 and reaches the highest ALFWorld full-suite repair result, with gains concentrated in harder families such as Clean and Cool.

4.3 Repair-Budget Dynamics

Figure 2 plots cumulative success as the repair budget R increases. The curves reinforce the boundary suggested by Table 1. Direct Repair is often competitive in the first one or two rounds because many traces expose a concrete surface error that a single patch can fix. When that first patch is correct, the direct baseline recovers quickly. When it is incomplete, however, the direct and hypothesis-conditioned curves more often flatten, because the next prompt has little structured record of which explanation was tried and how the retest changed the evidence.

The trend is especially visible for TextCraft Qwen2.5-3B, where RESKILL jumps sharply after the first repair round and keeps a large margin through R4. TextCraft failures often combine recipe-chain selection, exact surface-form copying, and count discipline, so recording which explanation a repair targeted makes a failed retest useful: the next round can emphasize remaining explanations instead of rewriting the same general rule. For Qwen3-4B TextCraft, stronger first-pass performance leaves a smaller residual failure pool and less room for improvement, resulting in closer repair curves. Even when there is limited room for improvement, RESKILL achieves the highest final success, indicating that structured selection remains useful for difficult residual failures.

ALFWorld shows the complementary pattern. Smaller models benefit steadily across later rounds, while Qwen3-4B approaches saturation by the middle of the budget. The important point is not only the final percentage, but the shape of the curve: RESKILL is less dependent on the first patch being correct, because the post-repair trajectory can redirect later edits toward persistent or newly exposed failures. The budget curves also provide a practical reference for choosing skill-repair strategies under

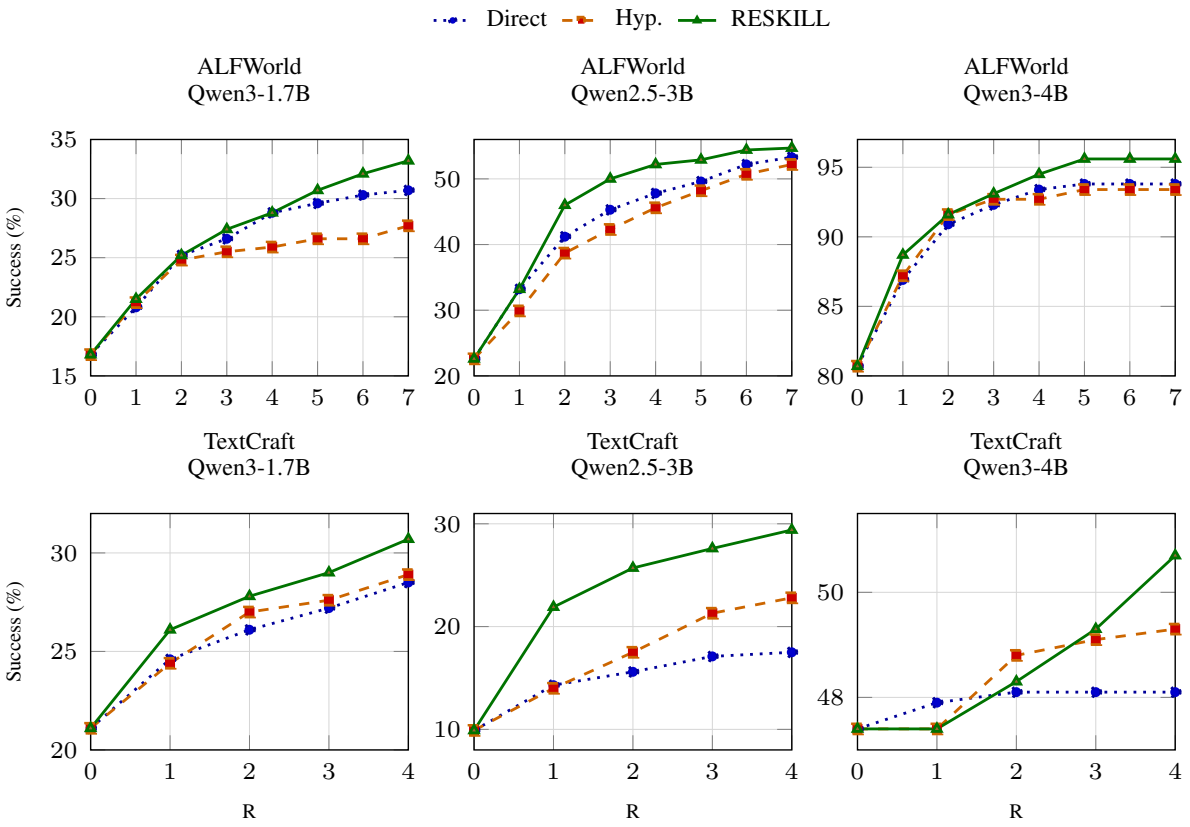


Figure 2: Cumulative repair-budget trends. Each panel fixes one benchmark–model setting and plots success rate as the repair budget R increases. ALFWorld panels run through R7; TextCraft panels run through R4.

different repair-round budgets. Appendix B reports the full budget tables, separated by benchmark because ALFWorld is tracked through R7 whereas TextCraft is tracked through R4.

4.4 Failure Modes and Repair Behavior

The repair traces help explain why the same framework behaves differently across domains. Direct Repair and Hypothesis Repair both commit to a model-selected edit at each round. RESKILL instead exposes the decision as a comparison among local repairs and carries retest feedback into the next state. This matters most when early diagnoses are noisy: for Qwen2.5-3B TextCraft, all methods begin from the same 9.9% first-pass success rate, but RESKILL reaches 21.9% after one repair round, while Direct and Hypothesis Repair reach 14.3% and 14.0%. The gap persists through R4, suggesting that structured selection and outcome-conditioned update help the agent use post-failure feedback more effectively.

TextCraft depth stratification also clarifies where the gains come from. Most recoveries occur in depth-1 and depth-2 tasks, where the visible observation exposes a short recipe chain and small surface mistakes can invalidate an otherwise plausible plan. These tasks reward repairs that preserve exact item names, numeric counts, and command form. Appendix E gives a representative TextCraft trace: the initial agent rewrites a listed recipe token, while RESKILL separates the surface-form hypothesis from broader planning hypotheses and selects a patch that directly targets the exposed mismatch.

For ALFWorld, trace inspection points to a different pattern. Direct patches can work well when the failure is a single procedural omission, such as not checking the right receptacle or not applying a family-specific action. RESKILL is most useful when a failed household trajectory leaves several explanations plausible: object search, precondition handling, action ordering, or family scope. In those cases, the coverage relation gives the selector a reason to prefer a patch that addresses more of the active explanation mass, and an unsuccessful retest changes the next-round state rather than simply asking for another unconstrained rewrite.

5 Conclusion

We presented RESKILL, a structured framework for post-failure skill repair. By tracking competing failure hypotheses, targeted patches, retest outcomes, and repair credit across rounds, RESKILL turns failed repairs into feedback for later edits. Experiments on ALFWorld and TextCraft show gains over direct and hypothesis-conditioned repair across three model sizes, with the benefits when failures admit multiple plausible explanations.

6 Acknowledgements

This work is supported by Advanced Materials-National Science and Technology Major Project (Grant No. 2025ZD0620100), National Key R&D Program of China (No. 2024YFA1012700), and Guangdong Provincial Key Lab of Integrated Communication, Sensing and Computation for Ubiquitous Internet of Things (No. 2023B1212010007).

7 Limitations

Our evaluation studies interactive environments where failures produce observable action feedback and repaired skills can be retested under matched budgets. This scope supports controlled comparison of repair decisions and makes it possible to inspect how hypotheses, selected patches, and retest outcomes affect later rounds. Broader tool-use settings may introduce delayed feedback, noisier failure signals, or edits whose effects appear across longer interaction histories; extending structured repair to those settings is an important direction for future work.

References

Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, Alex Herzog, Daniel Ho, Jasmine Hsu, Julian Ibarz, Brian Ichter, Alex Irpan, Eric Jang, and 1 others. 2022a. Do as i can, not as i say: Grounding language in robotic affordances. In *Conference on Robot Learning*.

Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, and 1 others. 2022b. Do as i can, not as i say: Grounding language in robotic affordances. *arXiv preprint arXiv:2204.01691*.

Shiqi Chen, Jingze Gai, Ruochen Zhou, Jinghan Zhang, Tongyao Zhu, Junlong Li, Kangrui Wang, Zihan Wang, Zhengyu Chen, Klara Kaleb, and 1 others. 2026. Skillcraft: Can llm agents learn to use tools skillfully? *arXiv preprint arXiv:2603.00718*.

Mengyi Deng, Zhiwei Li, Xin Li, Tingyu Zhu, Ying Zhao, Zhijiang Guo, and Wei Wang. 2026. Uncertainty-aware clarification in llm agents with information gain. *arXiv preprint arXiv:2606.03135*.

Linxi Fan, Guanzhi Wang, Yunfan Jiang, Ajay Mandlekar, Yuncong Yang, Haoyi Zhu, Andrew Tang, De-An Huang, Yuke Zhu, and Anima Anandkumar. 2022. Minedojo: Building open-ended embodied agents with internet-scale knowledge. *Advances in Neural Information Processing Systems*, 35:18343–18362.

Dong Huang, Jianbo Dai, Han Weng, Puzhen Wu, Yuhao Qing, Heming Cui, Zhijiang Guo, and Jie Zhang. 2024. Effilearner: Enhancing efficiency of generated code via self-optimization. In *Advances in Neural Information Processing Systems*, volume 37, pages 84482–84522. Curran Associates, Inc.

Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, Yevgen Chebotar, and 1 others. 2022. Inner monologue: Embodied reasoning through planning with language models. *arXiv preprint arXiv:2207.05608*.

Xiangyi Li, Wenbo Chen, Yimin Liu, Shenghan Zheng, Xiaokun Chen, Yifeng He, Yubo Li, Bingran You, Haotian Shen, Jiankai Sun, and 1 others. 2026. Skillsbench: Benchmarking how well agent skills work across diverse tasks. *arXiv preprint arXiv:2602.12670*.

Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. 2023. Code as policies: Language model programs for embodied control. In *2023 IEEE International conference on robotics and automation (ICRA)*, pages 9493–9500. IEEE.

Anthony Zhe Liu, Jongwook Choi, Sungryull Sohn, Yao Fu, Jaekyeom Kim, Dong-Ki Kim, Xinhe Wang, Jaewon Yoo, and Honglak Lee. 2024a. Skillact: Using skill abstractions improves llm agents. In *ICML 2024 Workshop on LLMs and Cognition*.

Shunyu Liu, Yaoru Li, Kongcheng Zhang, Zhenyu Cui, Wenkai Fang, Yuxuan Zheng, Tongya Zheng, and Mingli Song. 2024b. Odyssey: Empowering minecraft agents with open-world skills. *arXiv preprint arXiv:2407.15325*.

Xingyan Liu, Xiyue Luo, Linyu Li, Ganghong Huang, Jianfeng Liu, and Honglin Qiao. 2026a. Skillforge: Forging domain-specific, self-evolving agent skills in cloud technical support. *arXiv preprint arXiv:2604.08618*.

Yujian Liu, Jiabao Ji, Li An, Tommi Jaakkola, Yang Zhang, and Shiyu Chang. 2026b. How well do agentic skills work in the wild: Benchmarking llm

skill usage in realistic settings. *arXiv preprint arXiv:2604.04323*.

Zhengxi Lu, Zhiyuan Yao, Jinyang Wu, Chengcheng Han, Qi Gu, Xunliang Cai, Weiming Lu, Jun Xiao, Yueting Zhuang, and Yongliang Shen. 2026. Skill0: In-context agentic reinforcement learning for skill internalization. *arXiv preprint arXiv:2604.02268*.

Yuchen Ma, Yue Huang, Han Bao, Haomin Zhuang, Swadheen Shukla, Michel Galley, Xiangliang Zhang, and Stefan Feuerriegel. 2026. Skillgen: Verified inference-time agent skill synthesis. *arXiv preprint arXiv:2605.10999*.

Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. 2023a. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems*.

Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, and 1 others. 2023b. Self-refine: Iterative refinement with self-feedback. *Advances in neural information processing systems*, 36:46534–46594.

Siru Ouyang, Jun Yan, Yanfei Chen, Rujun Han, Zifeng Wang, Bhavana Dalvi Mishra, Rui Meng, Chun-Liang Li, Yizhu Jiao, Kaiwen Zha, and 1 others. 2026. Skillos: Learning skill curation for self-evolving agents. *arXiv preprint arXiv:2605.06614*.

Joon Sung Park, Joseph O'Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S Bernstein. 2023. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th annual acm symposium on user interface software and technology*, pages 1–22.

Ajay Patel, Markus Hofmarcher, Claudiu Leoveanu-Condrei, Marius-Constantin Dinu, Chris Callison-Burch, and Sepp Hochreiter. 2024. Large language models can self-improve at web agent tasks. *arXiv preprint arXiv:2405.20309*.

Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. 2023. Toolformer: Language models can teach themselves to use tools. *Advances in neural information processing systems*, 36:68539–68551.

Shuaike Shen, Wenduo Cheng, Mingqian Ma, Alistair Turcan, Martin Jinye Zhang, and Jian Ma. 2026. Skillfoundry: Building self-evolving agent skill libraries from heterogeneous scientific resources. *arXiv preprint arXiv:2604.03964*.

Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2024. Reflexion: Language agents with verbal reinforcement learning, 2023. *URL https://arxiv.org/abs/2303.11366*, 8.

Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language agents with verbal reinforcement learning. *Advances in neural information processing systems*, 36:8634–8652.

Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. 2020. Alfworld: Aligning text and embodied environments for interactive learning. *arXiv preprint arXiv:2010.03768*.

Chenxi Wang, Zhuoyun Yu, Xin Xie, Wuguannan Yao, Runnan Fang, Shuofei Qiao, Kexin Cao, Guozhou Zheng, Xiang Qi, Peng Zhang, and 1 others. 2026. Skillx: Automatically constructing skill knowledge bases for agents. *arXiv preprint arXiv:2604.04804*.

Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. 2023a. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*.

Zihao Wang, Shaofei Cai, Guanzhou Chen, Anji Liu, Xiaojian Ma, and Yitao Liang. 2023b. Describe, explain, plan and select: Interactive planning with large language models enables open-world multi-task agents. *arXiv preprint arXiv:2302.01560*.

Zihao Wang, Shaofei Cai, Anji Liu, Yonggang Jin, Jinbing Hou, Bowei Zhang, Haowei Lin, Zhaofeng He, Zilong Zheng, Yaodong Yang, and 1 others. 2024. Jarvis-1: Open-world multi-task agents with memory-augmented multimodal language models. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47(3):1894–1907.

Zhiheng Xi, Yiwen Ding, Wenxiang Chen, Boyang Hong, Honglin Guo, Junzhe Wang, Dingwen Yang, Chenyang Liao, Xin Guo, Wei He, Songyang Gao, Lu Chen, Rui Zheng, Yicheng Zou, Tao Gui, Qi Zhang, Xipeng Qiu, Xuanjing Huang, Zuxuan Wu, and Yu-Gang Jiang. 2024. AgentGym: Evolving large language model-based agents across diverse environments. *arXiv preprint arXiv:2406.04151*.

Minrui Xu, Zilin Wang, Mengyi Deng, Zhiwei Li, Zhicheng Yang, Xiao Zhu, Yinhong Liu, Boyu Zhu, Baiyu Huang, Chao Chen, and 1 others. 2026. Envfactory: Scaling tool-use agents via executable environments synthesis and robust rl. *arXiv preprint arXiv:2605.18703*.

Yutao Yang, Junsong Li, Qianjun Pan, Bihao Zhan, Yuxuan Cai, Lin Du, Jie Zhou, Kai Chen, Qin Chen, Xin Li, and 1 others. 2026a. Autoskill: Experience-driven lifelong learning via skill self-evolution. *arXiv preprint arXiv:2603.01145*.

Zhicheng Yang, Zhijiang Guo, Yifan Song, Minrui Xu, Yongxin Wang, Yiwei Wang, Xiaodan Liang, and

Jing Tang. 2026b. Prune-opd: Efficient and reliable on-policy distillation for long-horizon reasoning. *arXiv preprint arXiv:2605.07804*.

Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Tom Griffiths, Yuan Cao, and Karthik Narasimhan. 2023. Tree of thoughts: Deliberate problem solving with large language models. *Advances in neural information processing systems*, 36:11809–11822.

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2022. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*.

Duzhen Zhang, Zhong-Zhi Li, Ming-Liang Zhang, Jiaxin Zhang, Zengyan Liu, Yuxuan Yao, Haotian Xu, Junhao Zheng, Xiuyi Chen, Yingying Zhang, and 1 others. 2025. From system 1 to system 2: A survey of reasoning large language models. *IEEE Transactions on Pattern Analysis and Machine Intelligence*.

Hanrong Zhang, Shicheng Fan, Henry Peng Zou, Yankai Chen, Zhenting Wang, Jiayu Zhou, Chengze Li, Wei-Chieh Huang, Yifei Yao, Kening Zheng, and 1 others. 2026. Coevoskills: Self-evolving agent skills via co-evolutionary verification. *arXiv preprint arXiv:2604.01687*.

Boyuan Zheng, Michael Y Fatemi, Xiaolong Jin, Zora Zhiruo Wang, Apurva Gandhi, Yueqi Song, Yu Gu, Jayanth Srinivasa, Gaowen Liu, Graham Neubig, and 1 others. Skillweaver: Web agents can self-improve by discovering and honing skills, 2025. *URL https://arxiv. org/abs/2504.07079*.

Xizhou Zhu, Yuntao Chen, Hao Tian, Chenxin Tao, Weijie Su, Chenyu Yang, Gao Huang, Bin Li, Lewei Lu, Xiaogang Wang, and 1 others. 2023. Ghost in the minecraft: Generally capable agents for open-world environments via large language models with text-based knowledge and memory. *arXiv preprint arXiv:2305.17144*.

A Implementation Details

This appendix documents the configuration of every reported experimental setting. Table 4 maps each result to its task pool, model configuration, initial skill file, decoding parameters, context and action limits, transcript policy, repair budget, task-level logs, and reproduction command.

A.1 Generation Settings

All reported experiments use frozen model weights served through an OpenAI-compatible local inference endpoint. Environment rollouts are decoded greedily. For ALFWorld first-pass rollouts and repair retests, we set the LLM-agent decoding temperature to 0.0 and top-p to 0.7, with a 512-token action-generation cap, seed 0, and an 8192-token serving context. ALFWorld repair-stage generations use stage-specific settings: failure hypotheses use temperature 0.2, top-$p = 0.95$; repair proposals use temperature 0.6, top-$p = 0.95$; and coverage construction uses temperature 0.1, top-$p = 0.9$. Each ALFWorld repair-stage call uses a 2048-token completion cap in the reported runs.

For TextCraft, environment rollouts also use greedy decoding with top-$p = 1.0$ and at most ten interaction turns. The full TextCraft setting uses a 4096-token serving context and a 128-token action-generation cap, while the Qwen3-4B depth-1/depth-2 setting uses an 8192-token serving context and a 256-token action-generation cap. TextCraft structured repair calls use top-$p = 1.0$ and a JSON completion cap of 768 tokens. Hypothesis, coverage, and post-repair attribution calls use temperature 0.0 under the compact JSON profile. Single-patch proposal calls use temperature 0.2 for Qwen3 models and 0.3 for Qwen2.5-3B; multi-proposal RESKILL calls use temperature 0.2 for Qwen3 models and 0.5 for Qwen2.5-3B. Failed JSON-format attempts are retried with temperature 0.0.

A.2 ALFWorld Implementation

The ALFWorld evaluation contains 274 tasks, including 140 seen and 134 unseen tasks. Within each model setting, Direct Repair, Hypothesis Repair, and RESKILL share the same first-pass trajectories, initial skill file, evaluator, decoding configuration, and repair budget. The Qwen2.5-3B seen and unseen results are aggregated over the same task pools used by all three repair methods.

The shared repair loop builds failure packets, generates hypotheses and repairs, filters candidate edits, constructs coverage matrices, applies selected repair notes, retests failed tasks, runs post-repair attribution, and writes budget summaries.

Skill editing. ALFWorld skills are stored as task-family sections. Each rollout receives general guidance plus the section mapped to the current family. The two picking variants share one picking section, while inspection, cleaning, heating, and cooling use separate sections. Repairs are appended as local notes rather than rewriting the entire skill file.

A.3 TextCraft Implementation

The Qwen3-1.7B and Qwen2.5-3B TextCraft experiments use the full task pool, recipe-free skills, a 4096-token context, and a 128-token action limit. The Qwen3-4B experiment uses the depth-1 and

depth-2 task pool, an 8192-token context, and a 256-token action limit; its agent-visible trajectories contain only environment observations and actions. Within each model setting, Direct Repair, Hypothesis Repair, and RESKILL share the same task pool, initial skill file, evaluator, decoding configuration, transcript policy, and repair budget. The Qwen3-4B results therefore support controlled comparisons among repair methods within that setting, but are not used for strict cross-model comparisons with the 1.7B and 3B results. The corresponding configurations, task-level logs, and reproduction commands are listed in Table 4.

Recipe-free skills. The TextCraft skill file contains only general response constraints: follow environment instructions, output one action line, include numeric quantities, use visible item names and craft commands, avoid external Minecraft knowledge, and react to execution failures. It contains no task-specific recipes.

Visible transcript policy. In the final 4B setting, any model-generated analysis or rationale preceding the executable action is treated as internal inference and excluded from the trajectory passed to subsequent model calls or the environment. The visible transcript therefore contains only environment observations and cleaned action text. Hypothesis generation, proposal generation, and coverage scoring condition on this cleaned transcript. Repair retesting and post-repair attribution may still use internal reasoning, but only cleaned actions and structured repair outputs are retained as interaction evidence.

B Repair Budget Tables and Trends

R0 is the unrepaired with-skills first pass for each repair curve. For ALFWorld, this is the full-suite with-skills baseline in Table 1; for TextCraft, it is the recipe-free first pass in the corresponding full or depth-filtered repair pool. Table 2 reports ALFWorld through R7, while Table 3 reports TextCraft through R4.

C Evaluation Setting Summary

D Additional Notes on TextCraft Depth

The TextCraft depth diagnostic is computed from visible recipe lines in the initial observation. The full-setting analysis includes depths one through four for the 1.7B and 3B settings. The final 4B evaluation focuses on the depth-1 and depth-2 subset and uses the same visible-recipe rule.

E Case Study: Auditable TextCraft Repair

Table 5 shows a representative Qwen3-4B TextCraft example from the reported depth-1 and depth-2 setting. The goal is to craft spruce planks. Direct Repair and Hypothesis Repair remain unsuccessful after the repair budget on this item, while RESKILL succeeds after one repair round.

This case illustrates why the repair trace is useful beyond the final success flag. A direct patch can state a generic formatting rule without resolving the exact mismatch between "spruce log" and "spruce logs." RESKILL records the more specific surface-form hypothesis, compares candidate repairs through a binary coverage matrix, and preserves the retest outcome for the selected patch.

Model	Method	R0	R1	R2	R3	R4	R5	R6	R7
Qwen3-1.7B	Direct	16.8	20.8 (+4.0)	25.2 (+8.4)	26.6 (+9.8)	28.8 (+12.0)	29.6 (+12.8)	30.3 (+13.5)	30.7 (+13.9)
Qwen3-1.7B	Hypothesis	16.8	21.2 (+4.4)	24.8 (+8.0)	25.5 (+8.7)	25.9 (+9.1)	26.6 (+9.8)	26.6 (+9.8)	27.7 (+10.9)
Qwen3-1.7B	ReSkill	16.8	**21.5** (+4.7)	**25.2** (+8.4)	**27.4** (+10.6)	**28.8** (+12.0)	**30.7** (+13.9)	**32.1** (+15.3)	**33.2** (+16.4)
Qwen2.5-3B	Direct	22.6	33.2 (+10.6)	41.2 (+18.6)	45.3 (+22.7)	47.8 (+25.2)	49.6 (+27.0)	52.2 (+29.6)	53.3 (+30.7)
Qwen2.5-3B	Hypothesis	22.6	29.9 (+7.3)	38.7 (+16.1)	42.3 (+19.7)	45.6 (+23.0)	48.2 (+25.6)	50.7 (+28.1)	52.2 (+29.6)
Qwen2.5-3B	ReSkill	22.6	**33.2** (+10.6)	**46.0** (+23.4)	**50.0** (+27.4)	**52.2** (+29.6)	**52.9** (+30.3)	**54.4** (+31.8)	**54.7** (+32.1)
Qwen3-4B	Direct	80.7	86.9 (+6.2)	90.9 (+10.2)	92.3 (+11.6)	93.4 (+12.7)	93.8 (+13.1)	93.8 (+13.1)	93.8 (+13.1)
Qwen3-4B	Hypothesis	80.7	87.2 (+6.5)	**91.6** (+10.9)	92.7 (+12.0)	92.7 (+12.0)	93.4 (+12.7)	93.4 (+12.7)	93.4 (+12.7)
Qwen3-4B	ReSkill	80.7	**88.7** (+8.0)	**91.6** (+10.9)	**93.1** (+12.4)	**94.5** (+13.8)	**95.6** (+14.9)	**95.6** (+14.9)	**95.6** (+14.9)

Table 2: ALFWorld cumulative repair-budget success rates in percent. Parentheses show the gain relative to R0 within the same row; shaded rows mark ReSkill.

Model	Method	R0	R1	R2	R3	R4
Qwen3-1.7B	Direct	21.1	24.6 (+3.5)	26.1 (+5.0)	27.2 (+6.1)	28.5 (+7.4)
Qwen3-1.7B	Hypothesis	21.1	24.4 (+3.3)	27.0 (+5.9)	27.6 (+6.5)	28.9 (+7.8)
Qwen3-1.7B	ReSkill	21.1	**26.1** (+5.0)	**27.8** (+6.7)	**29.0** (+7.9)	**30.7** (+9.6)
Qwen2.5-3B	Direct	9.9	14.3 (+4.4)	15.6 (+5.7)	17.1 (+7.2)	17.5 (+7.6)
Qwen2.5-3B	Hypothesis	9.9	14.0 (+4.1)	17.5 (+7.6)	21.3 (+11.4)	22.8 (+12.9)
Qwen2.5-3B	ReSkill	9.9	**21.9** (+12.0)	**25.7** (+15.8)	**27.6** (+17.7)	**29.4** (+19.5)
Qwen3-4B	Direct	47.4	**47.9** (+0.5)	48.1 (+0.7)	48.1 (+0.7)	48.1 (+0.7)
Qwen3-4B	Hypothesis	47.4	47.4 (+0.0)	**48.8** (+1.4)	49.1 (+1.7)	49.3 (+1.9)
Qwen3-4B	ReSkill	47.4	47.4 (+0.0)	48.3 (+0.9)	**49.3** (+1.9)	**50.7** (+3.3)

Table 3: TextCraft cumulative repair-budget success rates in percent. Parentheses show the gain relative to R0 within the same row; shaded rows mark ReSkill. TextCraft is tracked through R4, so this table intentionally omits later repair budgets.

Reported setting	Evaluation configuration
ALFWorld Qwen3-1.7B and Qwen3-4B	274 tasks, including 140 seen and 134 unseen tasks; 8192-token context; 512-token action limit; repair budget $R = 7$. Repair methods share the same first-pass trajectories and initial skill file within each model setting.
ALFWorld Qwen2.5-3B	The same 140 seen and 134 unseen tasks; 8192-token context; 512-token action limit; repair budget $R = 7$. Repair methods share the same first-pass trajectories within each split.
TextCraft Qwen3-1.7B and Qwen2.5-3B	Full TextCraft task pool with recipe-free skills, which contain general rules for reading recipes and executing actions but no task-specific crafting recipes; 4096-token context; 128-token action limit; at most ten interaction turns; repair budget $R = 4$.
TextCraft Qwen3-4B	Depth-1 and depth-2 task pool with the same recipe-free skill design; 8192-token context; 256-token action limit; at most ten interaction turns; agent-visible trajectories containing only environment observations and actions; repair budget $R = 4$.
ALFWorld skills ablation	Matched first-pass evaluations with and without the initial skill file over the same 140 seen and 134 unseen tasks for each model.
TextCraft skills ablation	Matched first-pass evaluations with and without the recipe-free initial skill file over the same depth-1 and depth-2 task order for each model setting.

Table 4: Evaluation configurations for the reported settings. All repair methods are compared under matched conditions within each model and benchmark setting.

Field	Trace summary
Task	Goal: spruce planks; visible recipe: “craft 4 spruce planks using 1 spruce logs”; recipe depth: 1.
Trajectory excerpt	The observation lists the recipe as “craft 4 spruce planks using 1 spruce logs.” The failed action rewrites the ingredient as “craft 4 spruce planks using 1 spruce log.” The environment rejects the action as an invalid recipe. The important evidence is not that the agent chose the wrong goal, but that it changed the visible surface form of the ingredient.
Initial hypotheses	The first hypothesis states that the agent singularized a listed ingredient token instead of copying the recipe surface form. Other hypotheses cover generic recipe-format mismatch, direct-goal execution without checking prerequisites, and count discipline. The first-round weights are uniform because no repair has been tested yet.
Candidate repairs	One repair targets exact surface-form copying. Another targets visible recipe-chain execution. A third focuses on count preservation. These repairs overlap with different hypotheses, so the coverage matrix distinguishes a patch that fixes the exposed token mismatch from patches that only address broader recipe planning.
Selected repair	The selected patch tells the agent to treat listed TextCraft commands as exact surface forms: fetch missing ingredients using the exact listed item text and count, and craft by copying the visible recipe command rather than singularizing or pluralizing item names.
Retest outcome	The repaired skill snapshot succeeds on the retest after one repair round. The trace records both the chosen patch and the competing candidates, so the final success can be inspected as a consequence of the surface-form hypothesis and its selected repair.

Table 5: A TextCraft case study illustrating the audit trail produced by RESKILL. The example is drawn from the final Qwen3-4B depth-1 and depth-2 run. It exposes the surface-form hypothesis, competing coverage-based repair choices, and the selected successful patch.

F Prompt Templates

This appendix specifies the model inputs, instructions, and structured outputs used at each stage of the repair process.

F.1 Shared Evidence Bundle

All repair-stage calls are grounded in a common set of observable task evidence, supplemented with stage-specific inputs such as active hypotheses, candidate patches, or the selected repair. The shared evidence includes:

- the task goal and the initial environment observation visible to the agent;
- the relevant portion of the current skill set;
- a compact trajectory excerpt containing the agent's actions and the corresponding environment responses;
- benchmark-specific interface information available to the agent, such as visible recipes in TextCraft or valid action syntax in ALFWorld.

Only observable interaction content is treated as trajectory evidence. The prompts require each hypothesis and candidate patch to be grounded in this evidence, and require patches to express reusable local skill edits rather than episode-specific action sequences.

F.2 Repair Prompt Contracts

Failure Attribution Prompt Template

Role
You are an analyst identifying concrete reasons why an interactive rollout failed.
Input Evidence: {task goal}, {current skill excerpt}, {failed trajectory}, {environment feedback}.
Instructions: List a small set of distinct failure explanations. Each explanation must cite observable evidence from the trajectory and indicate the skill scope it may affect. Avoid vague summaries and avoid proposing a repair in this step.
Expected Output: Hypotheses with stable identifiers, evidence, affected scope, and repair direction.

Direct Repair Prompt Template

Role
You are editing a reusable skill file after a failed rollout.
Input Evidence: {task goal}, {current skill excerpt}, {failed trajectory}, {environment feedback}.
Instructions: Write one local skill patch that could prevent the same type of failure while remaining compatible with the environment interface. Prefer reusable rules over one-episode action plans.
Expected Output: One concise skill edit with target scope and appended rule.

Hypothesis Repair Prompt Template

Role
You are editing a reusable skill file using an explicit failure explanation.
Input Evidence: {shared evidence bundle}, {current hypotheses $\mathcal{H}_t$}.
Instructions: Choose one useful hypothesis from the list and write a local patch that targets it. Keep the edit small, grounded, and compatible with the observed interface.
Expected Output: Selected hypothesis identifier and one reusable skill edit.

Symbol	Object	Role in RESKILL
S_t	Current skill set	Provides the editable skill context; selected repairs append local skill notes.
τ_i^t	Failed or retest trajectory	Provides evidence for hypotheses, repair hints, and post-repair attribution.
$\mathcal{H}_t$	Hypothesis list	Finite support over candidate failure explanations.
b_t	State weights	Maintained by the algorithm over the current explanations.
$\mathcal{Q}_t$	Repair candidates	The finite set of reusable repairs considered at round t.
$C_t(q, h)$	Coverage relation	Binary relation indicating whether repair q addresses hypothesized failure explanation h.
$q_t^\star$	Selected repair	Chosen by weighted coverage plus fixed locality and grounding preferences.
b_{t+1}	Updated state	Computed from post-repair explanations and carried into the next repair round.

Table 6: Method variables used in the structured repair loop.

Post-Repair Attribution Prompt Template

Role
You are updating the repair state after a selected patch was retested and the task still failed.
Input Evidence: {previous hypotheses}, {selected repair $q_t^\star$}, {failed retest trajectory}.
Instructions: Explain whether the repair missed the issue, partially helped, or exposed a new failure. Reuse stable identifiers when possible and introduce a new hypothesis only when supported by the retest trace.
Expected Output: Next-round hypotheses $\mathcal{H}_{t+1}$ and a concise state-update rationale.

G Information About Use Of AI Assistants

This manuscript uses Ai Assistants strictly for the purpose of language editing and textual polishing to enhance presentation quality. We declare that the novel ideas, methodological framework, experimental execution, and data analysis are the original work of the authors. All content modified by AI tools has been carefully reviewed and validated by the authors to ensure accuracy.