How Many Thoughts Can a Vector Hold? The Capacity of Reasoning by Superposition

Hongyu Gu[1]
University of Science and Technology of China
Hefei, China
ustc_23ghy@mail.ustc.edu.cn

Chang Liu[2]
Zhongguancun Academy
Beijing, China
liuchang@bza.edu.cn

Jingwen Fu[2*]
Zhongguancun Academy
Beijing, China
jwfu99@gmail.com

arXiv:2609.13747v1 [cs.AI] 12 Sep 2026

Abstract

Large language models solve difficult problems by carrying intermediate computation across many reasoning steps. Conventional chain-of-thought writes this computation as tokens; recent continuous and recurrent approaches instead move part of it into fixed-dimensional latent states, where one thought can superpose several alternatives. This shift raises a basic design question: as reasoning proceeds, what should a continuous thought preserve? The natural strategy is to discard earlier computation and keep only the current frontier: storing more objects appears to dilute the state and waste its limited capacity. We show that this intuition can be exactly backwards. When the same downstream computation is applied to both states, a cumulative superposition that retains the reached history can require less representation dimension than a frontier superposition that stores only the current alternatives. At a fixed hidden width, this advantage lets a latent reasoner retain more usable evidence, distinguish more possible downstream outcomes, and delay the point at which its compressed state becomes unreliable. The reversal occurs because useful historical components reinforce one another coherently, while unrelated alternatives remain random interference. This perspective also resolves a practical design question: how should a model weight the memories accumulated inside a latent state when their future use is not yet known? Across the full family of reusable weighted superpositions, concentrating on a few recent or salient objects creates a weakly represented memory that bottlenecks later attention. Uniform cumulative weighting avoids this failure, and we prove that it is minimax-optimal for robust future reasoning. Our results turn superposition from an observed latent phenomenon into a design principle: balanced cumulative memory can make a fixed representational budget support more reliable and reusable computation.

1 Introduction

What should a continuous thought preserve as reasoning unfolds? This question is becoming increasingly relevant as language-model reasoning moves from explicit text into continuous computation. Continuous chain-of-thought feeds hidden states directly back into the model (Hao et al., 2024); looped and recurrent Transformers repeatedly update a latent workspace (Xu & Sato, 2024; Fu et al., 2026); and reasoning by superposition shows that one continuous thought can encode several search alternatives and expand them in parallel (Zhu et al., 2025b). These results establish that superposed latent reasoning is possible. They leave open a consequential question for model design: among the many superpositions that can represent a computation, which one uses a limited hidden dimension most effectively?

We study three structures that isolate this choice. A *frontier superposition* stores the alternatives active at the current step, a *cumulative superposition* stores the full reached history, and a general *weighted superposition* distributes representational energy arbitrarily across that history. All three

*Corresponding author.

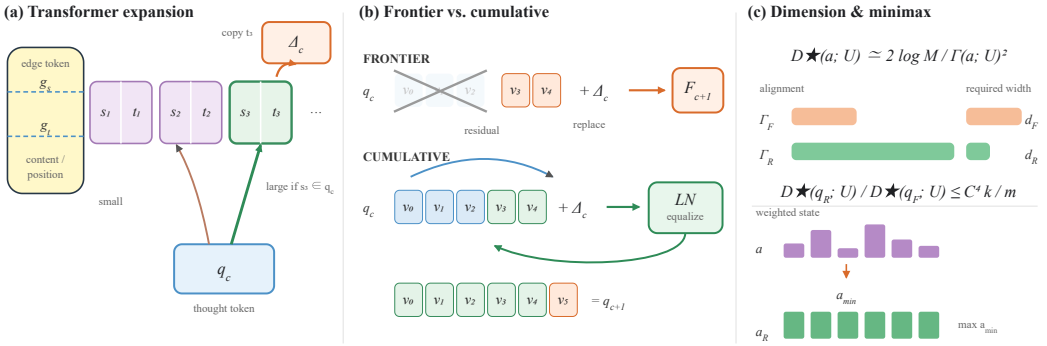

Figure 1: Our theory in one picture. Attention identifies active edge sources and copies their destinations into the next reasoning state. Frontier superposition overwrites the state and discards earlier reached objects, whereas cumulative superposition preserves them through the residual path. The retained history supplies coherent evidence to later computation, allowing cumulative superposition to remain usable at a smaller hidden dimension. For general weighted superpositions, uneven coefficients create a weakest-object routing bottleneck; uniform cumulative weighting removes this bottleneck.

use the same graph, downstream query, candidate bank, and decoder. The comparison therefore changes only what the continuous thought preserves, not the task it is asked to solve.

Ordinary capacity intuition favors frontier superposition: storing more objects in one normalized vector dilutes their individual coefficients and should make them harder to recover (Elhage et al., 2022). We show that this intuition can be exactly backwards for reasoning. Cumulative superposition can remain usable at a smaller hidden dimension despite storing strictly more objects, because useful historical components reinforce one another coherently while unrelated alternatives remain random interference. Superposition capacity is therefore governed not by support size alone, but by how the stored information supports future computation.

This reversal holds for a broad, state-independent class of downstream queries, which may be continuous, correlated, asymmetric, and identity-dependent. The condition captures bounded-influence computations over accumulated objects, including whole-set aggregation, additive route objectives, algorithmic execution traces, and intermediate-step evaluation (Zaheer et al., 2017; Lee et al., 2019; Kool et al., 2019; Veličković et al., 2022; Lightman et al., 2024).

The comparison also has a direct interpretation inside a Transformer: attention expands the current computation, while residual writing determines whether earlier objects are discarded or preserved. For general weighted superpositions, concentrating weight on a few objects leaves other memories too weak for later attention to reuse. When their future use is unknown, uniform cumulative weighting solves the resulting robust design problem:

$$\boldsymbol{a}_{\mathrm{cum}} \in \arg\min_{\boldsymbol{a}} \mathfrak{D}(\boldsymbol{a}; \mathcal{U}_C), \tag{1}$$

where $\mathfrak{D}$ is the worst-case end-to-end critical dimension. This suggests a simple principle for latent reasoning: expand the current computation while preserving its history with balanced representational weight.

Contributions. Our main contributions are:

1. **Critical-dimension theory.** We derive an exact finite-dimensional score law and show that cumulative superposition requires fewer dimensions than frontier superposition for a broad class of nonlocalized queries. This result shifts the relevant notion of capacity from how many objects are stored to how effectively they support future computation, explaining when retaining more history benefits a fixed-width reasoner.
2. **End-to-end Transformer mechanism.** We realize both superpositions through the same attention–residual construction and identify the weakest retained coefficient as the routing bot-

tleneck. This connects the abstract dimension advantage to multi-step Transformer computation and provides a concrete diagnostic for failures of reusable latent memory.

3. **A minimax design principle.** Across reusable weighted superpositions, uniform cumulative weighting is sharply near-optimal for fixed nonlocalized query classes and exactly minimax when future queries are unknown. This provides a robust rule for allocating latent memory, supported by width-scaling and token-level experiments in both lightweight and GPT-2 models.

2 Related Work

Explicit and continuous reasoning. Chain-of-thought externalizes intermediate computation as tokens (Wei et al., 2022; Merrill & Sabharwal, 2024), whereas COCONUT feeds continuous thoughts back into the model (Hao et al., 2024). Recent theory shows that continuous tokens can maintain and explore several reasoning traces in parallel (Zhu et al., 2025b; Gozeten et al., 2025); subsequent work studies how such superpositions emerge during training (Zhu et al., 2025a). Our question is different from existence or learnability: among frontier, cumulative, and weighted states, which representation requires the least embedding dimension for the same later computation?

Capacity of superposition. Feature-superposition analyses explain how non-orthogonal features share a low-dimensional space and how interference depends on feature sparsity, importance, and polysemanticity (Elhage et al., 2022; Foote, 2024; Gupta & Kumar, 2025; Bereska et al., 2025). These analyses primarily concern recovery or interpretation of individual features. Our result isolates a reasoning-specific reversal: when later computation aggregates several stored objects, additional components can contribute coherent signal, making query-aligned margin—rather than support size—the operative capacity.

Recurrent computation and latent memory. Universal and looped Transformers reuse computation depth while maintaining a latent workspace (Dehghani et al., 2019; Giannou et al., 2023; Xu & Sato, 2024; Fu et al., 2026), while recurrent-memory and buffer-based models carry information through dedicated latent tokens (Bulatov et al., 2022; Yang et al., 2024; Lin et al., 2026). Related work studies the difficulty of multi-hop reasoning in long contexts and without explicit intermediate steps (Li et al., 2024; Yao et al., 2025), as well as limits and representations of hidden reasoning (Brown-Cohen et al., 2026; Hu et al., 2024). Our dimension laws separate two roles of this latent budget: storing atomic states and keeping a compact state aligned with computations that must reuse them later.

3 Setup: Three Superposition Structures

Let $G = (V, E)$ be a directed graph with $n = |V|$ and root r. Coordinates of vectors in $\mathbb{R}^n$ are indexed by vertices $v \in V$ (also called objects or nodes). At a fixed depth c,

$$F := \{v : \mathrm{dist}(r, v) = c\}, \qquad R := \{v : \mathrm{dist}(r, v) \le c\}, \tag{2}$$

with $|F| = k$, $|R| = m$, $F \subseteq R$, and $1 \le k \le m$. Write D for the graph depth and, at layer c,

$$R_c := \{v : \mathrm{dist}(r, v) \le c\}, \qquad F_c := \{v : \mathrm{dist}(r, v) = c\}, \qquad N_c := R_{c+1} \setminus R_c, \tag{3}$$

with $r_c := |N_c|$. Write $\mathbf{1}_S$ for the 0–1 indicator of $S \subseteq V$ and $e_v := \mathbf{1}_{\{v\}}$.

Each vertex has an independent embedding $g_v \sim N(0, I_d/d)$. Equivalently, the columns of $G \in \mathbb{R}^{d\times n}$ are i.i.d. $N(0, 1/d)$. For $a \in \mathbb{R}^n$,

$$X_a := Ga = \sum_{v\in V} a_v g_v, \qquad q_a := \frac{X_a}{\|X_a\|}. \tag{4}$$

Since q_a is invariant to positive scaling we take $\|a\|_2 = 1$. Write χ_d (also C_d) for a chi random variable with d degrees of freedom.

The states under comparison are

$$a_F := \frac{\mathbf{1}_F}{\sqrt{k}}, \qquad a_R := \frac{\mathbf{1}_R}{\sqrt{m}}, \qquad a \in \mathcal{A}_S, \tag{5}$$

where $\mathcal{A}_S := \{a \in \mathbb{R}^n_+ : \mathrm{supp}(a) \subseteq S,\ \|a\|_2 = 1\}$ for $S \subseteq V$, and $q_F := q_{a_F}$, $q_R := q_{a_R}$. A weighted sequence is $\boldsymbol{a} = (a_c)_{c=0}^{D}$ with $a_c \in \mathcal{A}_{R_c}$, and $a_{\min,c} := \min_{v \in R_c} a_{c,v}$.

Fix, independently of G, unit directions $u_0, u_1, \ldots, u_M \in \mathbb{R}^n$, with u_0 useful and $u_1, \ldots, u_M$ competing; write $X_{u_j} := Gu_j$. Every state uses the same decoder

$$\widehat{j}(q_a) := \arg\max_{0 \le j \le M} \langle q_a, X_{u_j} \rangle, \tag{6}$$

and succeeds when $\widehat{j}(q_a) = 0$. When the bank is understood we write u for u_0. The coordinate u_v is query weight on vertex v, not a competing candidate; a node readout is $u = e_v$. The main text assumes isotropic competition:

$$u_1, \ldots, u_M \text{ are orthonormal and orthogonal to } \mathrm{span}\{a, u_0\}. \tag{7}$$

When several states supported on R are compared, one common bank is used: $u_1, \ldots, u_M$ are orthonormal in the orthogonal complement of the R-coordinates, which requires $M \le n - m$.

Alignment and the two uniform specializations are

$$\mu(a,u) := \langle a, u \rangle \in [0,1], \quad \mu_F(u) := \langle a_F, u \rangle = \frac{\sum_{v \in F} u_v}{\sqrt{k}}, \quad \mu_R(u) := \langle a_R, u \rangle = \frac{\sum_{v \in R} u_v}{\sqrt{m}}, \tag{8}$$

for nonnegative unit u supported on R. For $\delta \in (0,1)$,

$$d_\star(a, u; M, \delta) := \min\{d : \mathbb{P}(\widehat{j}(q_a) = 0) \ge 1 - \delta\}, \tag{9}$$

$$d_{\star,F}(u) := d_\star(a_F, u; M, \delta), \qquad d_{\star,R}(u) := d_\star(a_R, u; M, \delta), \tag{10}$$

under (7), with the value ∞ if $\mu(a,u) = 0$ and $1 - \delta > 1/(M+1)$.

A query class is a nonempty compact $\mathcal{U} \subseteq \mathcal{A}_S$; at a fixed depth, $S = R$. Set

$$\Gamma(a; \mathcal{U}) := \inf_{u \in \mathcal{U}} \langle a, u \rangle, \tag{11}$$

$$D_\star(a; \mathcal{U}, M, \delta) := \min\left\{ d : \inf_{u \in \mathcal{U}} \mathbb{P}(\widehat{j}(q_a) = 0) \ge 1 - \delta \right\}. \tag{12}$$

$$d_{1-\delta} := \min\{d \in \mathcal{D}_{\mathrm{test}} : \widehat{\mathrm{Acc}}(d) \ge 1 - \delta\}. \tag{13}$$

,which is the discrete version of the theoretical critical dimension $(d_\star(\delta))$ in experiments Call $\mathcal{U} \subseteq \mathcal{A}_R$ C-nonlocalized, for a constant $C \ge 1$, if

$$\|u\|_\infty \le \frac{C}{\sqrt{m}}, \qquad u \in \mathcal{U}. \tag{14}$$

Statements containing $o(1)$ are as $M \to \infty$ at fixed δ, with n large enough that $M \le n - m$; remainders in $d_\star$ are uniform over $\mu \in (0,1]$.

4 The Dimension Limit

We first derive the dimension law for an arbitrary weighted state. Frontier and cumulative comparisons will then follow by substitution, not by changing the readout.

Lemma 4.1 (Exact score competition)**.** *Fix unit a, u_0 with $\mu := \langle a, u_0 \rangle \in [0,1]$, and assume* (7)*. There exist mutually independent $Z_0, Z_1, \ldots, Z_M \sim N(0,1)$ and $C_d \sim \chi_d$ such that*

$$\sqrt{d}\,\langle q_a, X_{u_0} \rangle \overset{d}{=} \mu C_d + \sqrt{1 - \mu^2}\, Z_0, \tag{15}$$

$$\sqrt{d}\,\langle q_a, X_{u_j} \rangle \overset{d}{=} Z_j, \qquad 1 \le j \le M. \tag{16}$$

Consequently,

$$\mathbb{P}(\widehat{j}(q_a) = 0) = \mathbb{P}\left(\mu C_d + \sqrt{1 - \mu^2}\, Z_0 > \max_{1 \le j \le M} Z_j \right). \tag{17}$$

Lemma 4.1 already exposes the mechanism. Dimension increases only the aligned signal $\mu C_d \approx \mu\sqrt{d}$. The M unrelated candidates contribute a maximum of order $\sqrt{2\log M}$.

Theorem 4.2 (Dimension phase transition)**.** *Consider any sequence with $M \to \infty$, unit states $a = a_M$, useful directions $u_0 = u_{0,M}$, alignments $\mu_M \in (0, 1]$, and dimensions d_M. Under the fixed-bank condition* (7)*, if*

$$\frac{d_M \mu_M^2}{2\log M} \longrightarrow \lambda, \tag{18}$$

for $\lambda \in [0, \infty]$, then

$$\mathbb{P}(\widehat{j}(q_a) = 0) \longrightarrow \begin{cases} 0, & \lambda < 1, \\ 1, & \lambda > 1. \end{cases} \tag{19}$$

At the boundary $\lambda = 1$, the theorem makes no zero–one claim.

Remark 4.3. *The theorem turns superposition from a representational possibility into an operational criterion: a latent state is useful only when its aligned signal clears the competition created by the candidate set. Because every state is evaluated on the same task with the same decoder, the boundary isolates the effect of representation geometry and predicts a sharp failure when width is insufficient. It therefore provides the basis for asking whether retaining more reasoning history can make a fixed-width state more useful.*

Corollary 4.4 (Matching critical dimension)**.** *Fix $\delta \in (0, 1)$. For any sequence $M \to \infty$ and $\mu_M \in (0, 1]$, the critical dimension at which superposition reasoning succeeds with high probability is*

$$d_\star(a, u_0; M, \delta) = (2 + o(1)) \frac{\log M}{\mu_M^2}. \tag{20}$$

Moreover, let $N := n - m$ and define

$$M_{\max}(d, \mu, \delta, N) := \max\{1 \le M \le N : p_{d,M}(\mu) \ge 1 - \delta\}, \tag{21}$$

where $p_{d,M}(\mu)$ denotes the probability in (17)*. Along any sequence for which $\min\{d\mu^2, \log N\} \to \infty$,*

$$\log M_{\max} = (1 + o(1)) \min\left\{ \frac{d\mu^2}{2}, \log N \right\}. \tag{22}$$

Remark 4.5. *The matching law makes the phase transition actionable as a model-sizing rule: it specifies the hidden width required for a target reliability and the number of competing continuations that a fixed-width state can distinguish. For a continuous-reasoning model, it separates two remedies for readout failure: increase width or improve the alignment of the latent state with its future computation. Width savings can therefore be interpreted as additional usable reasoning capacity rather than merely a change in representation geometry.*

5 Why Cumulative Superposition Wins

We now answer three questions. When cumulative needs fewer dimensions? How a Transformer implements its update, and whether another weighting can do better once the full history must remain usable?

5.1 Cumulative needs fewer dimensions

Task regime. Nonlocalization is a state-independent bounded-influence condition, it constrains only the downstream query u and never compares its alignments with a_R and a_F. Since $\|u\|_2 = 1$, a fixed C means that every query draws on at least m/C^2 reached objects, Such readouts arise naturally in set and multiset aggregation, additive objectives over visited objects, and supervision over algorithmic or reasoning traces (Zaheer et al., 2017; Lee et al., 2019; Kool et al., 2019; Veličković et al., 2022; Lightman et al., 2024). Continuous and superposed reasoning provide the corresponding state geometry of multi-step latent search (Hao et al., 2024; Zhu et al., 2025b). On the state side, $k = o(m)$ says that the states activated at the current step grow more slowly than the accumulated computation. It includes chains and fixed-width traces, where $k = O(1)$ and $m = \Theta(c)$ at depth c, as well as grid-like breadth-first search in d dimensions, where $k = \Theta(c^{d-1})$ and $m = \Theta(c^d)$. Fixed C and $k = o(m)$ therefore capture a broad growing-history regime in which queries reuse a constant fraction of an expanding trace while its current frontier occupies a vanishing fraction.

Theorem 5.1 (Sharp separation for nonlocalized query classes)**.** *Fix $F \subsetneq R$, with $|F| = k < m = |R|$, and use the same candidate bank for both states. Let $\mathcal{U} \subseteq \mathcal{A}_R$ be any nonempty compact C-nonlocalized query class. Then*

$$\Gamma(a_R;\mathcal{U}) \geq \frac{1}{C}, \qquad \Gamma(a_F;\mathcal{U}) \leq C\sqrt{\frac{k}{m}}, \tag{23}$$

Consequently, whenever $\Gamma(a_F;\mathcal{U}) > 0$,

$$\frac{D_\star(a_R;\mathcal{U},M,\delta)}{D_\star(a_F;\mathcal{U},M,\delta)} \leq C^4\frac{k}{m}(1+o(1)). \tag{24}$$

For every fixed C, this ratio tends to zero in the growing-history regime $k = o(m)$; thus cumulative needs asymptotically fewer dimensions for the same query class and error. If $\Gamma(a_F;\mathcal{U}) = 0$, frontier instead has infinite critical dimension above chance while cumulative remains finite.

Remark 5.2. *This theorem resolves the paper's main paradox without selecting a special query. For the broad class of nonlocalized computations that reuse an expanding history, storing the full reached set reduces rather than consumes hidden dimension, even though the state represents more objects. A cumulative latent workspace can therefore remain reliable at widths where frontier-only compression fails. For model design, discarding earlier intermediate states is not automatically capacity-efficient when later reasoning may aggregate evidence across the trace.*

The reason is coherent addition. Readout succeeds when

$$\underbrace{\langle a,u\rangle\sqrt{d}}_{\text{useful signal}} \gtrsim \underbrace{\sqrt{2\log M}}_{\text{unrelated competition}}. \tag{25}$$

For cumulative, $\langle a_R,u\rangle = \|u\|_1/\sqrt{m}$: useful historical components add before normalization, while unrelated candidates remain centered fluctuations. The extra stored objects therefore strengthen a history-wide computation instead of acting as independent noise. This is the mathematical source of the apparent reversal.

5.2 Transformer realization and its dimension cost

At depth c, let $N_c = R_{c+1} \setminus R_c$ be the newly reached set and $r_c = |N_c|$. An edge token $e = (s,t)$ stores $[g_s; g_t]$, and the latent token stores the current superposition. A visited-edge mask exposes the current expansion candidates; we assume one active parent per node in N_c so that active destinations are averaged uniformly.

Theorem 5.3 (Transformer realization and routing dimension)**.** *A width-$O(d)$ construction with two attention blocks, a standard residual, and RMSNorm implements the following update with vanishing error:*

$$q_{c+1} = \mathrm{RMSNorm}(\lambda_c q_c + \eta_c o_c), \qquad o_c = \frac{1}{\sqrt{r_c}}\sum_{t\in N_c} g_t + o(1). \tag{26}$$

The first attention block compares the source-membership score $\beta\langle q_c, g_s\rangle$ with a threshold token and writes a gate on each edge; the second averages the destinations of the gated edges. The coefficient choices

$$(\lambda_c,\eta_c) = (0,1) \quad \textit{and} \quad (\lambda_c,\eta_c) = (\sqrt{|R_c|},\sqrt{r_c}) \tag{27}$$

produce frontier and cumulative, respectively, while other nonnegative coefficients produce layer-weighted recursive updates. No FFN is required.

More generally, the same membership head can route from any weighted state sequence $\boldsymbol{a} = (a_c)_{c=0}^{D}$. If every layer contains $\Theta(n)$ independent inactive source keys and $a_{\min,c} := \min_{v\in R_c} a_{c,v}$, its routing critical dimension is

$$d_{\mathrm{route}}(\boldsymbol{a},\delta) = \Theta\left(\max_{c<D}\frac{\log(Dn/\delta)}{a_{\min,c}^2}\right), \tag{28}$$

with infinite dimension when a required coefficient is zero. For a uniform state on s objects this reduces to $\Theta(s\log(Dn/\delta))$.

Remark 5.4. *This result closes the gap between a representation-level capacity comparison and executable continuous reasoning. Frontier, cumulative, and weighted superpositions are not merely hand-designed vectors evaluated by an external decoder: the same two-block Transformer can generate them autoregressively, and the critical-dimension law also governs whether attention can propagate them reliably. Thus, the cumulative advantage applies to an implementable latent workspace, while d_{route} gives a concrete width requirement for preserving and reusing that workspace across reasoning depth.*

5.3 Cumulative is minimax under nonlocalized future use

Consider any weighted sequence readable by the Transformer membership head, $\boldsymbol{a} = (a_c)_{c=0}^{D}$ with $a_c \in \mathcal{A}_{R_c}$. If every reached object may later be queried as an edge source, its smallest coefficient controls the worst membership margin. Theorem 5.3 and Corollary 4.4 show that its actual end-to-end critical dimension is, up to universal constants, the scale

$$\mathfrak{D}(\boldsymbol{a};\mathcal{U}) := \max\left\{\max_{c<D} \frac{\log(Dn/\delta)}{(\min_{v\in R_c} a_{c,v})^2}, \frac{2\log M}{\Gamma(a_D;\mathcal{U})^2}\right\}. \tag{29}$$

This optimizes over the full weighted family $\mathcal{A}_{R_c}$, which contains the recursive updates in Theorem 5.3 as well as their two canonical endpoints.

Theorem 5.5 (Sharp universal and robust optimality of cumulative)**.** *Suppose every reached object must remain eligible for later source-membership attention. Let $\mathcal{U}$ be any nonempty compact C-query class, and let $\boldsymbol{a}_R = (a_{R_c})_{c=0}^{D}$ be the uniform cumulative sequence. Then*

$$\mathfrak{D}(\boldsymbol{a}_R;\mathcal{U}) \le C^2 \inf_{\substack{a_c\in\mathcal{A}_{R_c}\\ 0\le c\le D}} \mathfrak{D}(\boldsymbol{a};\mathcal{U}). \tag{30}$$

The factor C^2 is sharp in the readout-dominated regime. Moreover, cumulative is exactly robust-minimax over the full nonlocalized family:

$$\boldsymbol{a}_R \in \arg\min_{\substack{a_c\in\mathcal{A}_{R_c}\\ 0\le c\le D}} \mathfrak{D}(\boldsymbol{a};\mathcal{U}_C(R_D)). \tag{31}$$

Thus, for fixed $C = O(1)$, cumulative is rate-minimax-optimal for every nonlocalized query class and exactly optimal when nonlocalization is the only prior on future use.

Remark 5.6. *This result turns cumulative from one feasible endpoint into a design rule for latent reasoning. When a model cannot know in advance which part of its trace later computation will reuse, no alternative weighting achieves a smaller worst-case end-to-end width over the full nonlocalized family; for any fixed nonlocalized class, cumulative is still within the sharp C^2 factor. Uniform residual accumulation is therefore a robust default, while learned or query-matched weighting is a specialization that trades away protection against other future uses.*

6 Experiment

6.1 Nonlocalized separation

We test Theorem 5.1 on a grid over $m \in \{64, 128, 256, 512\}$, $k/m \in \{1/32, \dots, 1/2\}$, $C \in \{1, 1.25, 1.5, 2\}$, and $M \in \{2^{10}, 2^{14}\}$, estimating worst-case $D_\star$ by binary search on d with 2000 Monte Carlo trials ($\delta = 0.1$). Each query is drawn independently of G and projected onto the cap $\|u\|_\infty \le C/\sqrt{m}$. *Dense* uses random positive weights on all m coordinates, modeling uneven aggregation over a visited set (Zaheer et al., 2017). *Temporal* centers an exponential bump along the trace, modeling localized process supervision over a CoT (Lightman et al., 2024). *Block* partitions the trace into segments with different scales, modeling staged algorithmic readouts (Veličković et al., 2022). *Union* mixes the three generators. The *sharp witness* is the theorem's example, uniform on $\lceil m/C^2 \rceil$ coordinates that contain F. *Zero-overlap* is uniform on $R \setminus F$, so $\Gamma(a_F;\mathcal{U}) = 0$.

Protocol. We use the ProsQA(Zhong et al., 2022) graph-memory task and hold the graph trace, candidate bank, and Gaussian codebook fixed when comparing state constructions. Queries are drawn before the Gaussian map. The controlled attention experiment uses an identity-regularized source-membership router at $d = 192$ and evaluates five independent Gaussian codebooks.

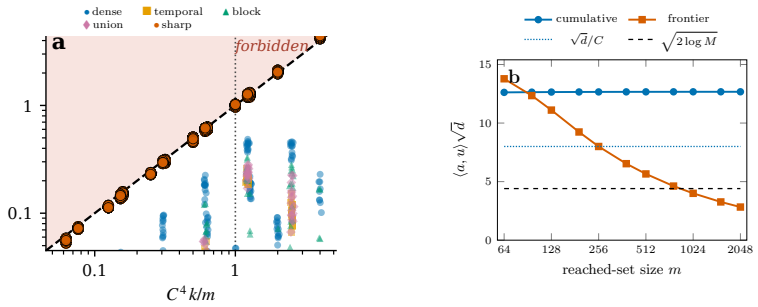

Figure 2: Cumulative superposition is both broadly more dimension-efficient and structurally favored by growing history. (a) Across diverse nonlocalized query families, its critical-dimension ratio remains below the theoretical envelope, while the sharp construction approaches equality. (b) As the reached history grows, cumulative preserves coherent query signal, whereas frontier discards reusable evidence and its signal eventually falls below the fixed competition scale.

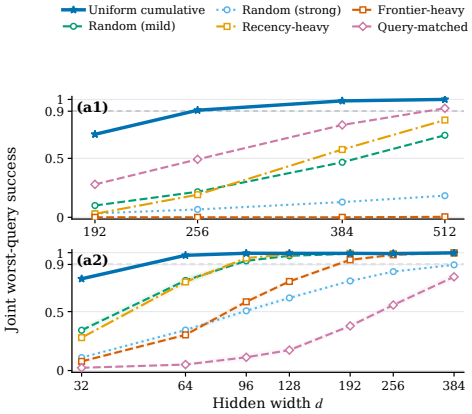

(b) Paired five-seed results

Lightweight two-block model

State	Acc. @ 32	Acc. @ 40	Empirical $d_{0.9}$
Frontier	0.210 ± 0.007	0.322 ± 0.005	110.6 ± 1.6
Cumulative	$\mathbf{0.793 \pm 0.015}$	$\mathbf{0.899 \pm 0.010}$	$\mathbf{40.3 \pm 1.1}$
Ideal frontier	0.216 ± 0.007	0.328 ± 0.006	112
Ideal cumulative	0.947 ± 0.005	0.990 ± 0.002	32

GPT-2 backbone

State	Acc. @ 32	Acc. @ 64	Empirical $d_{0.9}$
Frontier	0.468 ± 0.038	0.900 ± 0.026	66.2 ± 5.9
Cumulative	$\mathbf{0.898 \pm 0.024}$	$\mathbf{0.998 \pm 0.003}$	$\mathbf{34.4 \pm 3.8}$

Critical-width ratio $d^{\mathrm{C}}_{0.9}/d^{\mathrm{F}}_{0.9} \downarrow$	Two-block	GPT-2
Paired estimate	0.365 ± 0.008	0.521 ± 0.067

Figure 3: **One empirical view of dimension separation and weighted optimality.** The shared legend specifies six reusable weighting schemes. (a1) The original regularized two-block diagnostic; (a2) the GPT-2 replication reports the mean with one-standard-deviation bands. Within each panel the trained router is shared by every curve, and uniform cumulative crosses the 90% worst-query criterion first. (b) Paired five-seed results show that cumulative reaches matched held-out accuracy at substantially lower critical width than frontier for both the lightweight two-block model and GPT-2. Entries are means with two-sided 95% Student-t intervals; $d_{0.9}$ is linearly interpolated, and "Ideal" directly supplies the analytic state. Full configurations are deferred to the appendix.

6.2 Cumulative has a lower dimension threshold use

Across the paired five-seed sweeps, cumulative crosses the same 90% held-out success threshold at substantially lower width on both architectures. Its interpolated critical dimension is 40.3 ± 1.1, compared with 110.6 ± 1.6 for frontier, in the lightweight two-block model; with the two-layer GPT-2 backbone it is 34.4 ± 3.8, compared with 66.2 ± 5.9. The corresponding cumulative/frontier ratios are 0.365 ± 0.008 and 0.521 ± 0.067. Thus the separation predicted by Theorem 5.1 persists when the analytic router is replaced by a trained causal Transformer. Figure 3 shows both width sweeps and their paired critical-dimension estimates under matched evaluation protocols.

6.3 Uniform weighting prevents a token-level routing bottleneck

We next test the mechanism behind Theorem 5.3. Figure 4(a) shows the predicted two-term geometry: uniform cumulative is the joint lower-left point after normalizing routing and worst-query readout width by the uniform construction. Panels (b,c) expose what this optimum means inside the router. Uniform weights equalize the coefficient of every reusable memory token; recency and ran-

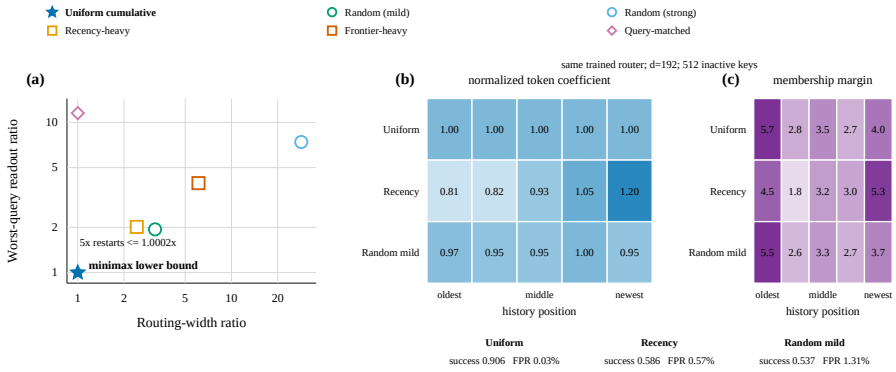

Figure 4: Why uniform cumulative is minimax among reusable weighted states. (a) Pareto geometry of routing and worst-query readout width terms, normalized by uniform cumulative; lower is better. (b) A token-level heatmap of normalized coefficients over reached-history position, averaged over test traces and five Gaussian codebooks. (c) The same router's membership margins. The summary below gives joint routing success and per-inactive-token false-positive rate.

dom weights produce low-coefficient tokens, and those tokens receive lower membership margins. Consequently, joint route success falls from 0.906 to 0.586 and 0.537, while the per-inactive-key false-positive rate rises from 0.03% to 0.57% and 1.31%, respectively.

7 Conclusion

We characterized frontier, cumulative, and general weighted reasoning superpositions through three linked results. First, a sharp Gaussian dimension law reduces reliable readout to state–query alignment; for any fixed nonlocalized query class, cumulative requires asymptotically fewer dimensions than frontier in the growing-history regime $k = o(m)$. Second, two attention blocks, a residual path, and RMSNorm realize both states as endpoints of one weighted Transformer update, attention selects newly reached objects, whereas the residual determines whether previous computation is discarded or retained. The same construction shows that the smallest reusable coefficient controls routing width. Third, optimizing routing and readout jointly over the full weighted family makes uniform cumulative superposition exactly minimax-optimal when reached history remains reusable and future use is exchangeable.

These results explain the apparent paradox that storing more objects can make reasoning easier. cumulative therefore converts reusable history into coherent signal rather than additional interference. This gives "stronger" an operational meaning—fewer dimensions for the same query class, transition, and target error—and identifies cumulative as the capacity-optimal target in this regime. The analysis isolates capacity and representability under Gaussian embeddings and an isotropic candidate bank; Appendix C extends the reduction to correlated banks with positive population margins. Extending the theory to learned anisotropic codebooks and explaining when optimization discovers the required routing and weighting remain important next steps.

REFERENCES

Leonard Bereska, Zoe Tzifa-Kratira, Reza Samavi, and Efstratios Gavves. Superposition as lossy compression: Measure with sparse autoencoders and connect to adversarial vulnerability. *arXiv preprint arXiv:2512.13568*, 2025.

Jonah Brown-Cohen, David Lindner, and Rohin Shah. Quantifying the necessity of chain of thought through opaque serial depth. *arXiv preprint arXiv:2603.09786*, 2026.

Aydar Bulatov, Yuri Kuratov, and Mikhail Burtsev. Recurrent memory transformer. In *Advances in Neural Information Processing Systems*, volume 35, 2022.

Mostafa Dehghani, Stephan Gouws, Oriol Vinyals, Jakob Uszkoreit, and Lukasz Kaiser. Universal transformers. In *International Conference on Learning Representations*, 2019.

Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, et al. Toy models of superposition. *Transformer Circuits Thread*, 2022. URL `https://transformer-circuits.pub/2022/toy_model/index.html`.

Alex Foote. Tackling polysemanticity with neuron embeddings. *arXiv preprint arXiv:2411.08166*, 2024.

Rao Fu, Zixuan Yang, Jiankun Zhang, Jing Ma, Hechang Chen, Yu Li, and Yi Chang. Simply stabilizing the loop via fully looped transformer. *arXiv preprint arXiv:2605.18797*, 2026.

Angeliki Giannou, Shashank Rajput, Jy-Yong Sohn, Kangwook Lee, Jason D. Lee, and Dimitris Papailiopoulos. Looped transformers as programmable computers. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 11398–11442, 2023.

Halil Alperen Gozeten, M. Emrullah Ildiz, Xuechen Zhang, Hrayr Harutyunyan, Ankit Singh Rawat, and Samet Oymak. Continuous chain of thought enables parallel exploration and reasoning. *arXiv preprint arXiv:2505.23648*, 2025.

Manan Gupta and Dhruv Kumar. Disentangling polysemantic neurons with a null-calibrated polysemanticity index and causal patch interventions. *arXiv preprint arXiv:2508.16950*, 2025.

Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian. Training large language models to reason in a continuous latent space. *arXiv preprint arXiv:2412.06769*, 2024.

Lijie Hu, Liang Liu, Shu Yang, Xin Chen, Zhen Tan, Muhammad Asif Ali, Mengdi Li, and Di Wang. Understanding reasoning in chain-of-thought from the hopfieldian view. *arXiv preprint arXiv:2410.03595*, 2024.

Wouter Kool, Herke van Hoof, and Max Welling. Attention, learn to solve routing problems! In *International Conference on Learning Representations*, 2019.

Juho Lee, Yoonho Lee, Jungtaek Kim, Adam Kosiorek, Seungjin Choi, and Yee Whye Teh. Set transformer: A framework for attention-based permutation-invariant neural networks. In *Proceedings of the 36th International Conference on Machine Learning*, pp. 3744–3753, 2019.

Yanyang Li, Shuo Liang, Michael R. Lyu, and Liwei Wang. Making long-context language models better multi-hop reasoners. *arXiv preprint arXiv:2408.03246*, 2024.

Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let's verify step by step. In *International Conference on Learning Representations*, 2024.

Hanyu Lin, Min Cai, Jiawei Wen, and Haodi Zhang. Tyler: Typed latent reasoning for language models – when to think, what to compute, and how much to allocate. *arXiv preprint arXiv:2606.16360*, 2026.

William Merrill and Ashish Sabharwal. The expressive power of transformers with chain of thought. In *International Conference on Learning Representations*, 2024.

Petar Veličković, Adrià Puigdomènech Badia, David Budden, Razvan Pascanu, Andrea Banino, Misha Dashevskiy, Raia Hadsell, and Charles Blundell. The CLRS algorithmic reasoning benchmark. In *Proceedings of the 39th International Conference on Machine Learning*, pp. 22084–22102, 2022.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, volume 35, 2022.

Kevin Xu and Issei Sato. On expressive power of looped transformers: Theoretical analysis and enhancement via timestep encoding. *arXiv preprint arXiv:2410.01405*, 2024.

Ling Yang, Zhaochen Yu, Tianjun Zhang, Shiyi Cao, Minkai Xu, Wentao Zhang, Joseph E. Gonzalez, and Bin Cui. Buffer of thoughts: Thought-augmented reasoning with large language models. *arXiv preprint arXiv:2406.04271*, 2024.

Yuekun Yao, Yupei Du, Dawei Zhu, Michael Hahn, and Alexander Koller. Language models can learn implicit multi-hop reasoning, but only if they have lots of training data. *arXiv preprint arXiv:2505.17923*, 2025.

Manzil Zaheer, Satwik Kottur, Siamak Ravanbakhsh, Barnabas Poczos, Ruslan Salakhutdinov, and Alexander J. Smola. Deep sets. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

Wanjun Zhong, Yifan Gao, Ning Ding, Yujia Qin, Zhiyuan Liu, Ming Zhou, Jiahai Wang, Jian Yin, and Nan Duan. ProQA: Structural prompt-based pre-training for unified question answering. In Marine Carpuat, Marie-Catherine de Marneffe, and Ivan Vladimir Meza Ruiz (eds.), *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 4230–4243, Seattle, United States, July 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.naacl-main.313. URL https://aclanthology.org/2022.naacl-main.313/.

Hanlin Zhu, Shibo Hao, Zhiting Hu, Jiantao Jiao, Stuart Russell, and Yuandong Tian. Emergence of superposition: Unveiling the training dynamics of chain of continuous thought. *arXiv preprint arXiv:2509.23365*, 2025a.

Hanlin Zhu, Shibo Hao, Zhiting Hu, Jiantao Jiao, Stuart Russell, and Yuandong Tian. Reasoning by superposition: A theoretical perspective on chain of continuous thought. In *Advances in Neural Information Processing Systems*, 2025b. arXiv:2505.12514.

A Proofs for the Dimension Limit

Proof of Lemma 4.1. Let $Y := Ga$. Since $\|a\|_2 = 1$, $Y \sim N(0, I_d/d)$ and $\sqrt{d}\,\|Y\| \sim \chi_d$.

If $\mu < 1$, define

$$c := \frac{u_0 - \mu a}{\sqrt{1-\mu^2}}. \tag{32}$$

Then c is unit norm, $c \perp a$, and $u_0 = \mu a + \sqrt{1-\mu^2}c$. If $\mu = 1$, the term involving c below is absent. Gaussianity and $c \perp a$ imply that Gc is independent of Y. Conditioned on Y, rotational invariance gives

$$\sqrt{d}\langle Y/\|Y\|, Gc\rangle \sim N(0,1), \tag{33}$$

and this random variable is independent of $\|Y\|$. Therefore

$$\begin{aligned}\sqrt{d}\langle q_a, Gu_0\rangle &= \mu\sqrt{d}\,\|Y\| + \sqrt{1-\mu^2}\sqrt{d}\langle Y/\|Y\|, Gc\rangle \\ &\overset{d}{=} \mu C_d + \sqrt{1-\mu^2}Z_0.\end{aligned} \tag{34}$$

For each $j \geq 1$, $u_j \perp a$, so Gu_j is independent of Y and

$$\sqrt{d}\langle Y/\left\|Y\right\|, Gu_j\rangle \sim N(0,1). \tag{35}$$

The orthogonality of $u_1, \ldots, u_M$ and their orthogonality to u_0 imply joint independence of all the Gaussian terms. Comparing the scores proves (17).

Let $S_\mu := \mu C_d + \sqrt{1-\mu^2}Z$ and $F_\mu(t) := \mathbb{P}(S_\mu \leq t)$. For $0 < \mu < 1$, put $\sigma^2 := 1-\mu^2$. Differentiating under the integral against the χ_d density gives

$$\frac{\partial F_\mu(t)}{\partial \mu} = -K_{d,\mu,t}\int_0^\infty (c-\mu t)c^{d-1}\exp\left(-\frac{(c-\mu t)^2}{2\sigma^2}\right)dc, \tag{36}$$

where $K_{d,\mu,t} > 0$. Denote the integral by I. Since $(c-\mu t)e^{-(c-\mu t)^2/(2\sigma^2)}$ equals $-\sigma^2$ times the derivative of the exponential, integration by parts yields

$$I = \begin{cases} \sigma^2 \exp\left(-\dfrac{\mu^2 t^2}{2\sigma^2}\right), & d = 1, \\ \sigma^2(d-1)\displaystyle\int_0^\infty c^{d-2}\exp\left(-\frac{(c-\mu t)^2}{2\sigma^2}\right)dc, & d > 1. \end{cases} \tag{37}$$

Thus $I > 0$ and $\partial_\mu F_\mu(t) < 0$ for every t. Continuity handles $\mu = 0, 1$, so S_μ is strictly increasing in first-order stochastic order. Conditional on $S_\mu = s$, success against M independent standard-normal decoys has probability $\Phi(s)^M$, a strictly increasing function of s. Therefore the unconditional success probability is strictly increasing in μ. The critical-dimension conclusion follows immediately. □

Lemma A.1 (Gaussian maximum and chi concentration)**.** *Let $Z_1, \ldots, Z_M$ be independent $N(0,1)$ variables and let $C_d \sim \chi_d$. As $M, d \to \infty$,*

$$\frac{\max_{j\leq M} Z_j}{\sqrt{2\log M}} \to 1, \qquad \frac{C_d}{\sqrt{d}} \to 1 \tag{38}$$

in probability.

Proof. For the upper Gaussian bound, for every $\varepsilon > 0$, the union bound and $\mathbb{P}(Z \geq t) \leq e^{-t^2/2}$ give

$$\begin{aligned} \mathbb{P}\left(\max_{j\leq M} Z_j > (1+\varepsilon)\sqrt{2\log M}\right) &\leq M\exp(-(1+\varepsilon)^2\log M) \\ &= M^{-2\varepsilon-\varepsilon^2} \to 0. \end{aligned} \tag{39}$$

For the lower bound, let $t = (1-\varepsilon)\sqrt{2\log M}$. The Gaussian Mills lower bound gives, for a universal $c > 0$ and all sufficiently large t,

$$\mathbb{P}(Z > t) \geq c\frac{e^{-t^2/2}}{t}. \tag{40}$$

Therefore $M\mathbb{P}(Z > t) \to \infty$, and independence yields

$$\mathbb{P}\left(\max_{j\leq M} Z_j \leq t\right) = (1-\mathbb{P}(Z>t))^M \leq e^{-M\mathbb{P}(Z>t)} \to 0. \tag{41}$$

This proves the first convergence. For the second, C_d^2/d is the average of d independent squared standard Gaussians, so it converges to one in probability by the weak law of large numbers. The square root is continuous on $(0,\infty)$. □

Proof of Theorem 4.2. Divide the two sides of (17) by $\sqrt{2\log M}$. Lemma A.1 gives

$$\frac{\max_{j\leq M} Z_j}{\sqrt{2\log M}} \to 1. \tag{42}$$

If $\lambda = 0$, then

$$\mathbb{E}\left[\left(\frac{\mu_M C_{d_M}}{\sqrt{2\log M}}\right)^2\right] = \frac{d_M\mu_M^2}{2\log M} \to 0, \tag{43}$$

so the aligned term converges to zero in probability without requiring $d_M \to \infty$.

Now suppose $\lambda \in (0, \infty]$. Since $\mu_M \le 1$, the assumed ratio implies $d_M \to \infty$, and The aligned term satisfies

$$\frac{\mu_M C_{d_M}}{\sqrt{2 \log M}} = \sqrt{\frac{d_M \mu_M^2}{2 \log M}} \frac{C_{d_M}}{\sqrt{d_M}} \to \sqrt{\lambda}, \tag{44}$$

where convergence to $+\infty$ is understood when $\lambda = \infty$. Because $|\sqrt{1 - \mu_M^2}| \le 1$ and a standard Gaussian is tight,

$$\frac{\sqrt{1 - \mu_M^2} Z_0}{\sqrt{2 \log M}} \to 0 \tag{45}$$

in probability. If $1 < \lambda < \infty$, choose a constant strictly between 1 and $\sqrt{\lambda}$; the true score lies above that constant and the maximum decoy below it with probability tending to one. If $0 < \lambda < 1$, choose a constant strictly between $\sqrt{\lambda}$ and 1 and reverse the argument. The cases $\lambda = 0$ and $\lambda = \infty$ follow from the two endpoint limits above. This proves (19). □

Proof of Corollary 4.4. For fixed M and $\mu > 0$, the success probability in (17) is nondecreasing in d. To see this, couple $C_{d+1} = \sqrt{C_d^2 + Z^2}$ with an additional independent standard Gaussian Z; then $C_{d+1} \ge C_d$ almost surely, while all other variables in (17) are unchanged.

Fix $\varepsilon > 0$. Theorem 4.2 implies that at dimensions

$$d_- = (1 - \varepsilon) \frac{2 \log M}{\mu_M^2}, \qquad d_+ = (1 + \varepsilon) \frac{2 \log M}{\mu_M^2}, \tag{46}$$

the success probabilities tend to zero and one, respectively. For all sufficiently large M, the critical dimension lies between d_- and d_+. Letting $\varepsilon \downarrow 0$ proves (20). Solving the same transition relation for $\log M$ gives (22). □

B Proofs for Frontier, Cumulative, and Weighted States

Proof of Theorem 5.1. Fix $u \in \mathcal{U}$. Nonnegativity, unit norm, and (14) give

$$1 = \|u\|_2^2 \le \|u\|_\infty \|u\|_1 \le \frac{C}{\sqrt{m}} \|u\|_1 . \tag{47}$$

Therefore

$$\langle a_R, u \rangle = \frac{\|u\|_1}{\sqrt{m}} \ge \frac{1}{C}. \tag{48}$$

For frontier,

$$\langle a_F, u \rangle = \frac{1}{\sqrt{k}} \sum_{v \in F} u_v \le \sqrt{k} \|u\|_\infty \le C \sqrt{\frac{k}{m}}. \tag{49}$$

Taking infima proves (23).

Compactness imply that the worst-query success probability under state a is exactly $p_{d,M}(\Gamma(a; \mathcal{U}))$. If $\Gamma(a_F; \mathcal{U}) > 0$, the uniform critical-dimension law gives

$$\frac{D_\star(a_R; \mathcal{U}, M, \delta)}{D_\star(a_F; \mathcal{U}, M, \delta)} = \frac{\Gamma(a_F; \mathcal{U})^2}{\Gamma(a_R; \mathcal{U})^2} (1 + o(1)) \le C^4 \frac{k}{m} (1 + o(1)). \tag{50}$$

The remainder is uniform over positive alignments by the asymptotic convention. If $\Gamma(a_F; \mathcal{U}) = 0$, compactness supplies a query with zero frontier alignment; its correct answer is exchangeable with the M decoys, giving the stated infinite critical dimension.

It remains to prove sharpness. Suppose first that $s = m/C^2$ is an integer and $k \le s$. Choose $S \subseteq R$ with $|S| = s$ and $F \subseteq S$, and let $u_v = C/\sqrt{m}$ on S and zero elsewhere. Then u is a nonnegative unit vector satisfying (14). For the compact singleton class $\mathcal{U} = \{u\}$,

$$\Gamma(a_R; \mathcal{U}) = \frac{1}{C}, \qquad \Gamma(a_F; \mathcal{U}) = C \sqrt{\frac{k}{m}}, \tag{51}$$

so equality holds in (24). If m/C^2 is not an integer, use $\lfloor m/C^2 \rfloor$ coordinates at the cap and one residual coordinate. The rounding changes the ratio by $1 + o(1)$ whenever $m/C^2 \to \infty$. The condition $k \le m/C^2$ holds eventually in every fixed-C growing-history regime with $k = o(m)$. $\square$

It follows by taking $u = a_R$: then $\langle a_R, u\rangle = 1$ and $\langle a_F, u\rangle = \sqrt{k/m}$, and substitution into Corollary 4.4 gives the two dimensions.

Proof of Theorem 5.3. We use the standard finite-depth convention for this construction: fix D and $\delta \in (0,1)$, let $n \to \infty$, omit empty transitions, and start from the exact intended state q_0. For the reusable weighted sequence, let $I_c \subseteq V \setminus R_c$ be the $\Theta(n)$ independent inactive source keys exposed at layer c, and define

$$\mathcal{E}_{\mathrm{route}}(d) := \bigcap_{c<D} \left\{ \min_{v\in R_c} \langle q_{a_c}, g_v\rangle > \tau_c > \max_{s\in I_c} \langle q_{a_c}, g_s\rangle \right\}. \tag{52}$$

Here $d_{\mathrm{route}}(\boldsymbol{a}, \delta)$ is the least d for which some thresholds $(\tau_c)_{c<D}$ make $\mathbb{P}(\mathcal{E}_{\mathrm{route}}(d)) \ge 1 - \delta$. Frontier uses the identical definition with its required sources F_c in place of R_c; the full-history version above is the one used by Theorem 5.5.

We first prove the routing bounds. Fix a layer and abbreviate $a = a_c$, $\alpha = a_{\min,c}$, and $Y = Ga$. For every required source $v \in R_c$, Gaussian regression with $e_v = a_v a + (e_v - a_v a)$ gives

$$\sqrt{d}\langle q_a, g_v\rangle \overset{d}{=} a_v C_d + \sqrt{1 - a_v^2}\, Z_v, \tag{53}$$

where $C_d = \sqrt{d}\,\|Y\| \sim \chi_d$ and each Z_v is marginally standard normal. The active Z_v need not be mutually independent, which is irrelevant for the union bound. For every inactive $s \in I_c$, $e_s \perp a$, so

$$\sqrt{d}\langle q_a, g_s\rangle \sim N(0,1), \tag{54}$$

and these variables are independent within I_c.

Choose $\tau_c = 3\alpha/8$. On the event

$$\frac{C_d}{\sqrt{d}} \ge \frac{3}{4}, \quad \min_{v\in R_c} Z_v \ge -\frac{\alpha\sqrt{d}}{4}, \quad \max_{s\in I_c} Z_s \le \frac{\alpha\sqrt{d}}{4}, \tag{55}$$

every required score is at least $\alpha/2$ and every inactive score is at most $\alpha/4$. Standard chi-square and Gaussian tails therefore give universal $A, c > 0$ such that

$$\mathbb{P}(\mathcal{E}_{\mathrm{route}}(d)^c) \le A \sum_{c<D} n \exp(-cda_{\min,c}^2). \tag{56}$$

Taking $d \ge C_0 \max_{c<D} \log(Dn/\delta)/a_{\min,c}^2$ in (56), with C_0 sufficiently large, proves the upper bound in (28).

For the lower bound, at each layer select v_c with coefficient $a_{c,v_c} = a_{\min,c}$. Successful common-threshold routing necessarily implies

$$\langle q_{a_c}, g_{v_c}\rangle > \max_{s\in I_c} \langle q_{a_c}, g_s\rangle. \tag{57}$$

This is exactly the readout event of Section 4 with useful direction e_{v_c}, alignment $a_{\min,c}$, and $|I_c| = \Theta(n)$ orthogonal competitors. Corollary 4.4 at fixed δ yields

$$d_{\mathrm{route}} \ge (2 - o(1)) \max_{c<D} \frac{\log n}{a_{\min,c}^2}, \tag{58}$$

For fixed D and δ, $\log(Dn/\delta) = \Theta(\log n)$, so this lower bound and (56) prove (28). If a required coefficient is zero, its score is exchangeable with the inactive scores and its success is at most $1/(|I_c|+1)$; hence the critical dimension is infinite whenever the target success exceeds this chance level, as it eventually does for fixed δ. For a uniform state on s objects, $a_{\min,c} = s^{-1/2}$, giving the final assertion.

We next realize the update. The argument below uses the full-history event (55). For the frontier transition, replace R_c by F_c in that event and in the preceding tail calculation; because its state is

uniform on F_c, the same calculation applies with $a_{\min,c} = |F_c|^{-1/2}$. Augment every token by $O(1)$ constant coordinates. In the first block an edge query $[g_s; 1]$ attends to the latent key $[q_c; 0]$ and threshold key $[0; \tau_c]$; assigning scalar values one and zero writes $h_e = \sigma(\beta(\langle q_c, g_s\rangle - \tau_c))$. On (55) the raw gap is at least $\gamma_c := a_{\min,c}/8$. Thus active gates are $1 + O(e^{-\beta\gamma_c})$ and inactive gates are $O(e^{-\beta\gamma_c})$.

The second block has a constant-query destination head with scalar gate keys and values g_t, together with a self-copy head returning q_c. Denote its raw destination average by $\bar{o}_c$. If its temperature is T, the inactive attention mass is bounded by

$$\sum_{e\in E_-} \alpha_e \le \frac{|E_-|}{|E_+|} \exp\{-T(1 - O(e^{-\beta\gamma_c}))\}. \tag{59}$$

Let $\alpha_\star$ denote the smallest required-source coefficient across the fixed computation (using $|F_c|^{-1/2}$ for the frontier replacement just described). Choose β, T so that $Te^{-\beta\gamma_c} \to 0$ and $T - \log|E| \to \infty$. Since $d \gtrsim \log n$ also implies $\max_v \|g_v\| = O_p(1)$, the temperatures may be increased further so that, uniformly over the fixed number of layers,

$$\left\| \sqrt{r_c}\,\bar{o}_c - \frac{1}{\sqrt{r_c}} X_N \right\| = o_p(\alpha_\star), \qquad X_N := \sum_{t\in N_c} g_t. \tag{60}$$

Thus, after the output scaling, $o_c := \sqrt{r_c}\bar{o}_c$ equals the o_c in (26), with a rate strong enough for both coefficient choices in (27).

The output projection combines the two heads as $(\lambda_c - 1)q_c + \eta_c o_c$; the residual adds back q_c. Directional RMSNorm means the usual positive scalar RMS normalization (with zero stabilization for exact unit normalization); its stabilization and gain do not change the represented direction. Setting $\lambda_c = 0$ leaves the normalized new frontier. For cumulative, let $X_R = \sum_{v\in R_c} g_v$ and use X_N from (60). Since

$$\frac{\|X_R\|}{\sqrt{|R_c|}} = 1 + O_p(d^{-1/2}), \tag{61}$$

the pre-normalization vector satisfies

$$\frac{\left\| \sqrt{|R_c|} q_{R_c} + \sqrt{r_c} o_c - (X_R + X_N) \right\|}{\|X_R + X_N\|} = O_p(d^{-1/2}) + o_p(\alpha_\star). \tag{62}$$

The normalization map obeys $\|x/\|x\| - y/\|y\|\| \le 2\|x - y\|/\|y\|$ whenever $\|x - y\| \le \|y\|/2$; hence RMSNorm returns $q_{R_{c+1}} + o_p(\alpha_\star)$. Indeed, $d\alpha_\star^2 \gtrsim \log n$ makes $d^{-1/2} = o(\alpha_\star)$; hence the state error is smaller than every subsequent membership gap. Induction over fixed D therefore preserves both routing and directional accuracy and proves simultaneous realization. Other non-negative output coefficients give the stated layer-weighted recursion. All operations are attention, a linear output projection, a residual, or RMSNorm, so no FFN is needed. □

Proof of Theorem 5.5. Write $L := \log(Dn/\delta)$ and $B := 2\log M$, and interpret a zero alignment or reusable coefficient as infinite width. For every $c < D$ and every $a_c \in \mathcal{A}_{R_c}$,

$$\left(\min_{v\in R_c} a_{c,v}\right)^2 \le \frac{1}{|R_c|} \sum_{v\in R_c} a_{c,v}^2 = \frac{1}{|R_c|}, \tag{63}$$

with equality only at $a_c = a_{R_c}$. Hence the routing part of $\mathfrak{D}(\boldsymbol{a}_R; \mathcal{U})$ is no larger than that of any $\mathfrak{D}(\boldsymbol{a}; \mathcal{U})$.

For the readout part, Theorem 5.1 gives $\Gamma(a_{R_D}; \mathcal{U}) \ge 1/C$. Cauchy–Schwarz gives $\Gamma(a_D; \mathcal{U}) \le 1$ for every unit a_D, so

$$\frac{B}{\Gamma(a_{R_D}; \mathcal{U})^2} \le C^2 B \le C^2 \frac{B}{\Gamma(a_D; \mathcal{U})^2}. \tag{64}$$

Combining the routing inequality and (64), and using $C \ge 1$, yields

$$\mathfrak{D}(\boldsymbol{a}_R; \mathcal{U}) \le C^2 \mathfrak{D}(\boldsymbol{a}; \mathcal{U}) \tag{65}$$

for every weighted sequence $\boldsymbol{a}$. Taking the infimum proves (30).

The factor is sharp. Suppose $s := |R_D|/C^2$ is an integer, choose $S \subseteq R_D$ with $|S| = s$, and let $u = \mathbf{1}_S/\sqrt{s}$ and $\mathcal{U} = \{u\}$. This class is C-nonlocalized. Cumulative has terminal alignment $1/C$, whereas the weighted terminal state $a_D = u$ has alignment one. Keep all earlier states cumulative. Whenever the readout scale B is at least the cumulative routing scale $L \max_{c<D} |R_c|$, the two objectives are respectively $C^2 B$ and B, so equality holds in (30).

It remains to prove exact robust optimality. Put $m := |R_D|$ and let $u_\star$ minimize $\|u\|_1$ over the nonempty compact set $\mathcal{U}_C(R_D)$. This full class is closed under every coordinate permutation P of R_D. Hence, for any $a_D \in \mathcal{A}_{R_D}$,

$$\begin{aligned}\Gamma(a_D;\mathcal{U}_C(R_D)) &\le \frac{1}{m!}\sum_P \langle a_D, Pu_\star\rangle = \frac{\|a_D\|_1 \|u_\star\|_1}{m} \\ &\le \frac{\|u_\star\|_1}{\sqrt{m}} = \Gamma(a_{R_D};\mathcal{U}_C(R_D)). \end{aligned} \tag{66}$$

The second inequality uses $\|a_D\|_1 \le \sqrt{m}\,\|a_D\|_2 = \sqrt{m}$; the final equality follows because $\langle a_{R_D}, u\rangle = \|u\|_1/\sqrt{m}$. Thus a_{R_D} minimizes the robust readout term. Equation (63) shows that every earlier a_{R_c} minimizes its routing term. The cumulative sequence therefore minimizes their maximum, proving (31). □

C Exact Reduction for an Arbitrary Fixed Bank

The orthogonal bank makes the main dimension law transparent. The following identity shows that the same signal-versus-competition structure persists for an arbitrary fixed correlated bank.

Theorem C.1 (Arbitrary-bank Gaussian reduction)**.** *Fix a unit state a and deterministic candidate directions $u_0, u_1, \dots, u_M$, independent of G. For each $j \ge 1$, let*

$$h_j := u_0 - u_j, \qquad \gamma_j := \langle a, h_j\rangle > 0, \qquad r_j := h_j - \gamma_j a \in a^\perp. \tag{67}$$

Let $g \sim N(0, I_n)$ and $C_d \sim \chi_d$ be independent. Then

$$\mathbb{P}(\widehat{j}(q_a) = 0) = \mathbb{P}\left(\max_{1\le j\le M} \langle g, -r_j/\gamma_j\rangle < C_d\right). \tag{68}$$

Proof. The correct candidate beats candidate j exactly when

$$\begin{aligned} 0 &< \langle Ga, G(u_0 - u_j)\rangle \\ &= \gamma_j \|Ga\|^2 + \langle Ga, Gr_j\rangle. \end{aligned} \tag{69}$$

Since $\gamma_j > 0$, simultaneous success is equivalent to

$$\max_{j\le M} \langle Ga, G(-r_j/\gamma_j)\rangle < \|Ga\|^2. \tag{70}$$

Every r_j is orthogonal to a. Joint Gaussianity therefore makes Ga independent of the family $\{G(-r_j/\gamma_j)\}_{j\le M}$. Conditional on $Y := Ga$, the process on the left of (70) has the same law as

$$\frac{\|Y\|}{\sqrt{d}}\,\{\langle g, -r_j/\gamma_j\rangle\}_{j\le M}. \tag{71}$$

Canceling $\|Y\| > 0$ and using $\sqrt{d}\,\|Y\| \sim \chi_d$ proves (68). □

Theorem C.1 requires no semantic choice of query and no orthogonality. Correlation changes only the maximum of the fixed Gaussian family. In the isotropic model, that maximum is $(1/\mu)\max_{j\le M} Z_j$ up to a common centered Gaussian term, recovering Theorem 4.2. More generally, the critical dimension is the squared typical scale of the Gaussian maximum in (68). The frontier–cumulative comparison remains a comparison under the same bank; the simplified $\log M/\mu^2$ form is obtained when the competition geometry is common to both states.

D Experimental Details and Additional Results

D.1 Dataset and query construction

We use the public graph-4 ProsQA splits released with Zhu et al. (2025b). Table 1 records the full split statistics. The paired width sweep uses every training and test trace and a fixed 120-trace validation subset.

Table 1: ProsQA graph-4 statistics. Graph sizes and solution lengths are averaged over traces.

Split	# traces	$\lvert V\rvert$	$\lvert E\rvert$	solution length
Train	14,785	22.8	36.5	3.5
Validation	257	22.7	36.3	3.5
Test	419	22.7	36.0	3.5

Queries are sampled in coefficient space before the Gaussian embedding. Each query contains the complete terminal frontier but remains strictly partial, uses between 75% and 85.7% of the reached history, and has within-support coefficient ratio at most $\rho = 1.2$. These choices imply the certified cap $C = 1.386$; the largest realized caps are smaller on every split. The same query, Gaussian draw, useful candidate, and decoy bank are reused for frontier and cumulative.

Table 2: Audit of the fixed nonlocalized query bank. "Queries" counts coefficient-space queries before repeated Gaussian test draws.

Split	Used traces	Queries	Support fraction	Realized $C_{\max}$
Train	14,785	118,280	$[0.750, 0.857]$	1.319
Validation	120	480	$[0.750, 0.857]$	1.271
Test	419	3,352	$[0.750, 0.857]$	1.287

D.2 Paired width-sweep configuration

The learned construction is the two-attention router in Theorem 5.3. Its source query and key projections are initialized as $I + 0.08Z/\sqrt{d}$; the membership inverse temperature and destination temperature are initialized to 8 and 10 and trained jointly. Frontier and cumulative share all router parameters and optimizer updates. The best epoch is selected by the sum of their validation state cosines.

Table 3: Configuration of the formal paired experiment.

Component	Configuration
Widths	$16, 24, 32, 40, 48, 56, 64, 80, 96, 112, 128$
Training seeds / embedding seeds	0:4 / 5100:5104
Query-bank size per trace (train/val/test)	$8/4/8$
Decoys (train/val/test)	$256/1{,}024/4{,}096$
Gaussian draws (val/test)	$1/4$
Epochs and optimizer	4; AdamW, learning rate 3×10^{-3}
Regularization	weight decay 10^{-4}; gradient clip 1.0
Gradient accumulation	16 traces
Loss weights (gate/state/query)	$1/2/0.1$
Query seed and certified cap	1701; $C = 1.385640646$
Primary endpoint	first tested d with test success ≥ 0.9

At test time, success is top-1 retrieval of the fixed useful candidate against 4,096 decoys. Each training seed is evaluated on 419 traces, eight queries per trace, and four Gaussian draws. Statistical intervals are therefore computed across the five independent training seeds, rather than by treating the within-seed evaluations as independent replications. During training, each trace uses one query per epoch, cycling deterministically through the pre-generated eight-query bank; hence four epochs expose four query IDs per trace.

D.3 COMPLETE FIXED-QUERY RESULTS

Table 4 gives the complete width sweep underlying Figure 3. Intervals are two-sided 95% Student-t intervals across five training seeds. "Ideal" supplies the analytic state to the same decoder and is not an additional trained model.

Table 4: Held-out top-1 accuracy at every tested width.

d	Learned frontier	Learned cumulative	Ideal frontier	Ideal cumulative
16	0.048 ± 0.001	0.254 ± 0.012	0.052 ± 0.001	0.442 ± 0.003
24	0.114 ± 0.004	0.574 ± 0.008	0.120 ± 0.005	0.790 ± 0.004
32	0.210 ± 0.007	0.793 ± 0.015	0.216 ± 0.007	0.947 ± 0.005
40	0.322 ± 0.005	0.899 ± 0.010	0.328 ± 0.006	0.990 ± 0.002
48	0.430 ± 0.008	0.946 ± 0.004	0.436 ± 0.008	0.9987 ± 0.0002
56	0.533 ± 0.010	0.972 ± 0.005	0.540 ± 0.009	0.9999 ± 0.0002
64	0.622 ± 0.008	0.984 ± 0.003	0.630 ± 0.008	1.0000 ± 0.0001
80	0.758 ± 0.010	0.992 ± 0.002	0.767 ± 0.010	1.000 ± 0.000
96	0.847 ± 0.005	0.997 ± 0.001	0.855 ± 0.006	1.000 ± 0.000
112	0.905 ± 0.006	0.998 ± 0.001	0.913 ± 0.005	1.000 ± 0.000
128	0.941 ± 0.005	0.998 ± 0.001	0.949 ± 0.003	1.000 ± 0.000

The threshold is stable at the individual-seed level (Table 5). The learned cumulative/frontier ratio has mean 0.391 and 95% interval ± 0.067; the fixed-alignment prediction is $(\mu_F/\mu_R)^2 = (0.5011/0.8893)^2 = 0.317$. At $d = 40$, learned-state cosines are 0.966 ± 0.006 for frontier and 0.933 ± 0.006 for cumulative. Thus the cumulative accuracy advantage cannot be explained by a higher state cosine.

Table 5: Empirical critical dimension $d_{0.9}$ for every training seed.

Seed	F	C	C/F	Ideal F	Ideal C	Ideal C/F
0	112	40	0.357	112	32	0.286
1	128	40	0.313	112	32	0.286
2	112	48	0.429	112	32	0.286
3	112	48	0.429	112	32	0.286
4	112	48	0.429	112	32	0.286
Mean	115.2	44.8	0.391	112	32	0.286

D.4 WEIGHTED ROUTING PROTOCOL AND RAW CODEBOOK RESULTS

At $d = 192$, the identity-regularized router is trained for four epochs on 600 traces and six profiles with 64 Gaussian route decoys, AdamW ($\mathrm{lr} = 3 \times 10^{-3}$, weight decay 10^{-4}), clip 1.0, accumulation 16, initialization noise 0.02, and identity penalty 10. We select by mean route success minus 0.02 projection error on 80 validation traces, 256 Gaussian decoys, and one draw. The frozen checkpoint is evaluated without retraining on five codebooks, 120 test traces, two draws, all unreached graph keys, and 512 additional Gaussian decoys per route step.

Table 6: Raw reusable-token routing results at $d = 192$. Each cell reports mean joint route-step success (fraction) / per-inactive-key false-positive rate (%). A route step succeeds iff every currently reached node is accepted and every unreached graph node and all 512 Gaussian decoys are rejected.

Codebook seed	Uniform cumulative	Recency-heavy	Random mild
4903	$0.896/0.032$	$0.576/0.572$	$0.533/1.318$
5004	$0.904/0.028$	$0.594/0.585$	$0.541/1.282$
5105	$0.907/0.028$	$0.593/0.557$	$0.536/1.312$
5206	$0.913/0.026$	$0.587/0.564$	$0.541/1.336$
5307	$0.909/0.029$	$0.578/0.570$	$0.534/1.280$
Mean	$0.906/0.029$	$0.586/0.570$	$0.537/1.306$

The corresponding mean minimum coefficient, normalized by the uniform value, is 1.000, 0.814, and 0.775. The monotone association between this weakest coefficient, joint success, and false positives is the token-level pattern predicted by the $a_{\min}$ mechanism in Theorem 5.5.

D.5 GPT-2 REPLICATION: CONFIGURATION AND COMPLETE RESULTS

The GPT-2 replication uses the same ProsQA graph-4 traces and the same paired comparison, but replaces the lightweight router with a two-layer causal GPT-2 Transformer. Each edge is represented by causal source and destination tokens followed by one continuous-thought token. The two states share the model, optimizer updates, coefficient-space query, Gaussian codebook, useful candidate, and decoy bank; only the frontier versus cumulative residual update differs. Queries are generated before the Gaussian map, use the full reached support ($\alpha = 1$), and draw positive coefficients with imbalance at most $\rho = 1.1$.

Table 7: Complete configuration of the paired GPT-2 width sweep.

Component	Configuration
Backbone	causal GPT-2; 2 layers, 8 heads, FFN width $4d$
Transformer details	GELU, context 512, dropout 0.1, residuals and LayerNorm
Tested widths / maximum codebook width	$32, 64, 96, 128, 192, 256, 384$ / 768
Model seeds / embedding seeds	0:4 / 6100:6104
Used traces (train/validation/test)	$800/120/120$
Queries per trace (train/validation/test)	$8/4/4$
Query construction	full reached support; $\rho = 1.1$; seed 2701
Decoys (train/validation/test)	$256/1{,}024/4{,}096$
Gaussian draws (validation/test)	$1/2$
Optimizer	AdamW; learning rate 10^{-4}; weight decay 10^{-2}
Optimization	clip 1.0; accumulation 16; at most 40 epochs
Loss weights (routing/state/query)	$1/10/0.1$
Stopping rule	after epoch 6: routing $\geq .98$ and both cosines $\geq .80$
Stopping patience / checkpoint selection	2 epochs / best summed validation cosine
Decoder and endpoint	shared top-1 decoder; interpolated $d_{0.9}$

Each seed–width model is evaluated on 120 held-out traces, four fixed queries per trace, and two Gaussian draws, with the useful candidate ranked against 4,096 fixed decoys. Table 8 reports all tested widths; intervals are two-sided 95% Student-t intervals over the five model seeds.

Table 8: Complete GPT-2 held-out top-1 accuracy sweep.

d	Frontier	Cumulative
32	0.468 ± 0.038	0.898 ± 0.024
64	0.900 ± 0.026	0.998 ± 0.003
96	0.978 ± 0.009	1.000 ± 0.000
128	0.993 ± 0.004	1.000 ± 0.000
192	0.998 ± 0.002	1.000 ± 0.000
256	1.000 ± 0.001	1.000 ± 0.000
384	1.000 ± 0.000	1.000 ± 0.000

The seed-level interpolants in Table 9 are the raw values summarized in the main-text figure. Cumulative has a lower critical dimension for every paired seed; the mean ratio is 0.521 ± 0.067.

D.6 WEIGHTED ROUTING PROTOCOL AND RAW CODEBOOK RESULTS

At $d = 192$, the identity-regularized router is trained for four epochs on 600 traces and six profiles with 64 Gaussian route decoys, AdamW ($\mathrm{lr} = 3 \times 10^{-3}$, weight decay 10^{-4}), clip 1.0, accumulation 16, initialization noise 0.02, and identity penalty 10. We select by mean route success minus 0.02 projection error on 80 validation traces, 256 Gaussian decoys, and one draw. The frozen checkpoint

Table 9: GPT-2 interpolated critical dimension $d_{0.9}$ for every model seed.

Seed	Frontier	Cumulative	C/F
0	63.2	32.0	0.507
1	68.6	33.0	0.481
2	73.5	36.2	0.493
3	62.6	32.0	0.511
4	63.0	38.8	0.616
Mean	66.2	34.4	0.521
95% interval	± 5.9	± 3.8	± 0.067

is evaluated without retraining on five codebooks, 120 test traces, two draws, all unreached graph keys, and 512 additional Gaussian decoys per route step.

Table 10: Raw reusable-token routing results at $d = 192$. Each cell reports mean joint route-step success (fraction) / per-inactive-key false-positive rate (%). A route step succeeds iff every currently reached node is accepted and every unreached graph node and all 512 Gaussian decoys are rejected.

Codebook seed	Uniform cumulative	Recency-heavy	Random mild
4903	$0.896/0.032$	$0.576/0.572$	$0.533/1.318$
5004	$0.904/0.028$	$0.594/0.585$	$0.541/1.282$
5105	$0.907/0.028$	$0.593/0.557$	$0.536/1.312$
5206	$0.913/0.026$	$0.587/0.564$	$0.541/1.336$
5307	$0.909/0.029$	$0.578/0.570$	$0.534/1.280$
Mean	$0.906/0.029$	$0.586/0.570$	$0.537/1.306$

The corresponding mean minimum coefficient, normalized by the uniform value, is 1.000, 0.814, and 0.775. The monotone association between this weakest coefficient, joint success, and false positives is the token-level pattern predicted by the $a_{\min}$ mechanism in Theorem 5.5.

D.7 Software and computing resources

We use Python 3.10, PyTorch 2.5.1, CUDA 12.4, NumPy 2.1.3, and Transformers 4.46.2 on NVIDIA A100 GPUs. The lightweight sweep contains 55 trained models and records 37.6 summed GPU-hours; the GPT-2 replication contains 35 single-GPU training runs. The four-width weighted run records 0.17 GPU-hours and its token analysis reuses the $d = 192$ checkpoint.