arXiv:2609.27279v1 [cs.AI] 23 Sep 2026

ENSIMEM: ENTITY-STRUCTURED INDEXING FOR LONG-TERM AGENT MEMORY

Xuanyu Meng[1], **Xing Fan**[2], **Xinyi Fan**[1], **Chenlei Guo**[2], **Yixuan Xie**[1], **Jiawei Han**[1]
[1]University of Illinois Urbana-Champaign
[2]Amazon
{xuanyum2, xfan31, xie39, hanj}@illinois.edu
{fanxing, guochenl}@amazon.com

ABSTRACT

An agent that interacts with users over long periods must recall facts, preferences, events, and changes from a continuously growing interaction history. Existing memory systems often compress interactions into generic summaries or retrieve anonymous text chunks, making it difficult for an agent to identify the correct entity, property, and supporting evidence. We present EnSIMem, an entity-structured long-term memory architecture for an agent. During offline construction, the system organizes interactions into theme-coherent episodes and builds dialogue-grounded index entries of the form [*entity*][*entity_type*][*property* : *value*]. Each entry preserves its source turns, temporal information, and available multimodal fields. During online interaction, the agent's request is decomposed into evidence requirements whose properties are aligned with the memory index. Entity-property lookup and adaptive retrieval then collect the evidence needed for point, temporal, compositional, and aggregation reasoning. The agent generates its response from the preserved source evidence rather than from lossy memory summaries. On long-term agent-memory benchmarks, EnSIMem achieves high answer accuracy while maintaining compact contexts and favorable online efficiency. These results show that entity-structured indexing and episode-level provenance provide a reliable foundation for long-term memory in agents. The code of our model is available at https://github.com/RamonMeng/EnSIMem.

1 INTRODUCTION

An agent deployed over a *long term* must accumulate and reuse information from its past interactions. This information may include user preferences, personal facts, plans, commitments, events, feedback, and decisions that influence future actions. For an agent, memory is therefore not simply a larger context window: it is a persistent interface between past experience and current reasoning, as reflected by memory-stream, tiered-memory, and evolving-memory architectures (Park et al., 2023; Packer et al., 2023; Xu et al., 2025; Kang et al., 2025). The memory system must help the agent recover the right evidence at the right time while preserving enough context for the agent to interpret that evidence correctly.

Long-term agent memory is particularly challenging because relevant information is distributed across many interactions and may be expressed in different ways. This challenge is central to long-term conversational-memory benchmarks and interactive memory evaluations (Maharana et al., 2024; Wu et al., 2024). A user may describe an event with a concrete action while a later request refers to the same event using a more general expression. Important information may also be introduced through pronouns, temporal references, correlations, or multimodal observations. For example, an interaction may state that a person traveled to Chicago by helicopter, while a later request asks how that person traveled to Chicago. A useful memory system must connect these expressions without collapsing the underlying event into an uninformative category such as "activity" or "travel." It must also retain the original evidence so that the agent can distinguish an explicit fact from an inference.

Existing approaches expose a tension between scalability and specificity. Providing the complete interaction history gives the agent access to all evidence, but causes context growth, higher latency, and distraction effects (Liu et al., 2024; Du et al., 2025; Wang et al., 2024; Bai et al., 2025). Summarization reduces the context size but can remove details, temporal qualifiers, and provenance. Conventional retrieval-augmented generation systems retrieve anonymous text chunks or vector-nearest memories (Lewis et al., 2020; Karpukhin et al., 2020; Izacard et al., 2022; Khattab & Zaharia, 2020); more recent systems use iterative or structured retrieval to improve multi-step reasoning (Asai et al., 2024; Trivedi et al., 2023; Edge et al., 2024). These representations are compact, but they do not explicitly indicate which entity, property, or event makes a memory relevant. EnSIMem addresses this tension by using structure to identify the right evidence and the original dialogue to interpret it: a large interaction history is converted into a small, query-specific, evidence-rich buffer instead of being passed wholesale to the agent.

We present EnSIMem, an entity-structure indexing-based long-term memory architecture for agentic systems, building on entity-structured retrieval and structure-augmented reasoning ideas (Meng et al., 2026; Parekh et al., 2025). Figure 1 summarizes the overall EnSIMem architecture. The central design principle is to use structure as an address system for memory, rather than as a lossy replacement for the original interaction. During offline construction, EnSIMem organizes conversations into theme-coherent episodes and builds dialogue-grounded index records of the form $[entity][entity_type][property : value]$. The entity may be a person, object, or salient event. The property is selected at an intermediate level of granularity: broad enough to connect paraphrases, but specific enough to preserve the identity of the underlying relation or event. When available, a finer-grained value records the concrete realization of that property. Every record remains linked to its source turns, timestamps, and multimodal fields.

Upon interaction, the query is decomposed into explicit evidence requirements. This follows the broader use of iterative query decomposition and agentic query rewriting for complex retrieval tasks (Trivedi et al., 2023; Shankar et al., 2024). The query planner extracts entities and properties at the same granularity as the memory index and identifies whether the query requires point evidence, temporal comparison, compositional reasoning, or aggregation. Structured entity-property lookup is combined with dense fallback retrieval when the wording differs from the indexed property. Retrieval then proceeds adaptively until the requirements are sufficiently covered. This allows the system to retrieve a small candidate set for a point query while preserving exhaustive coverage for a query that requires counting or comparing multiple events.

The final response is generated from the retrieved source evidence rather than from index entries alone, consistent with evidence-grounded graph and structure-aware generation approaches (He et al., 2024; Hu et al., 2024). This separation gives the agent both efficient access and evidentiary grounding: the index identifies where to look, while the original episode provides the context needed for interpretation and answer synthesis. It also makes the memory pipeline inspectable. Each answer can be traced from a query requirement, through an entity-property access path, to the dialogue turns and multimodal observations that support it.

EnSIMem separates query-independent memory construction from query-dependent recall. The offline representation can therefore be reused across many future queries, while the online stage spends computation only on the evidence required by the current query. On long-term conversational memory benchmarks, EnSIMem achieves high answer accuracy while maintaining compact retrieval contexts and favorable online efficiency. The results suggest that theme-coherent episodes, dialogue-grounded surface-property indexing, and requirement-aware retrieval provide a practical foundation for reliable long-term memory in agents.

Our contributions are:

1. **Theme-coherent episodic memory construction.** We organize long interactions into theme-coherent, contiguous episodes that preserve local context, temporal order, and provenance.

2. **Dialogue-grounded entity-property indexing.** We index originally mentioned entities and salient events with structured records linked to their original dialogue turns and multimodal observations.

3. **Granularity-controlled property alignment.** We align corpus and query properties at an intermediate level that supports paraphrase matching without erasing event-specific distinctions.

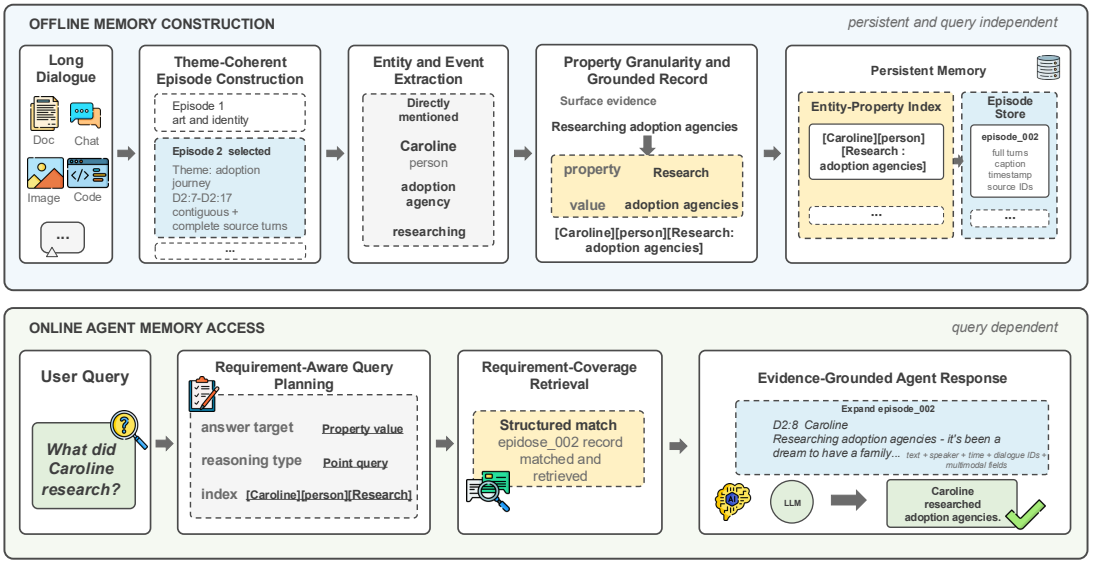

Figure 1: Overview of EnSIMem. Offline memory construction preserves multimodal dialogue, partitions it into theme-coherent episodes, extracts dialogue-grounded entity-property records, and stores an entity-property index linked to the original episodes. Online, the agent decomposes the request, performs requirement-aware retrieval with structured matching, expands the selected episode into source evidence, and generates a grounded response.

4. **Requirement-aware evidence localization.** We decompose agent requests into explicit evidence requirements and retrieve memories through structured entity-property matching.
5. **Query-type-aware adaptive retrieval.** We distinguish point queries from temporal, compositional, and aggregation queries and allocate evidence budgets appropriate to their reasoning requirements.
6. **An evidence-rich short-context buffer for agent response generation.** We convert long interaction histories into compact, query-specific reasoning buffers that retain relevant source evidence while filtering unrelated noise, without additional training or supervision. This methodology leads to state-of-the-art results, as shown in our experiments.

2 Related Work

Memory architectures for long-running agents. Prior agents use memory streams and reflection (Park et al., 2023), tiered context management (Packer et al., 2023), continual personalization (Zhong et al., 2023), verbal experience (Shinn et al., 2023), or reusable skills (Wang et al., 2023). Recent systems add temporal graphs, consolidated memories, evolving notes, and hierarchical stores (Rasmussen et al., 2025; Chhikara et al., 2025; Xu et al., 2025; Kang et al., 2025). EnSIMem complements these approaches with source-grounded entity-property access for factual and relational recall.

Entity-structured retrieval for long documents. Dense retrieval and graph-aware methods support general-purpose and multi-hop access (Lewis et al., 2020; Karpukhin et al., 2020; Khattab & Zaharia, 2020; Edge et al., 2024; Jiménez Gutiérrez et al., 2024). EnSI-RAG introduces entity-structure indexing for long-document question answering (Meng et al., 2026); EnSIMem extends this idea to persistent agent memory by linking dialogue-grounded entity-property records to theme-coherent episodes and query-aware evidence budgets.

Structured agent memory. Memora separates rich values from lightweight abstractions and cue anchors, while EnSIMem retains the original source turns and uses explicit entity-property alignment to localize evidence (Xia et al., 2026). Detailed benchmark definitions and baseline provenance are deferred to Appendix A.

3 ENSIMEM

3.1 SYSTEM OBJECTIVE AND DESIGN PRINCIPLE

EnSIMem is an index-structured long-term memory layer for an agent that can interact with systems/users over an extended period of time. Given an interaction corpus $\mathcal{D}$ and a new agent request q, the system must identify the relevant past evidence and provide it to the agent for reasoning and response generation.

We separate memory construction from memory access:

$$\mathcal{M} = F_{\text{memory}}(\mathcal{D}), \qquad \hat{y} = G(q, R(q, \mathcal{M})), \tag{1}$$

where $\mathcal{M}$ is a persistent, query-independent memory representation of corpus $\mathcal{D}$ via its memory construction process F_{memory}, R is the query-dependent retrieval process upon query q, and G is the agent's response-generation process, which generates $\hat{y}$.

The central design principle is to use an index structure as an address system for evidence, rather than a lossy replacement for the original interaction. Indexed structured records make it possible to identify which entity, event, and property are relevant, while the original dialogue remains available for interpretation and answer synthesis.

3.2 THEME-COHERENT EPISODIC MEMORY

The interaction corpus is an ordered collection of sessions containing text, speakers, timestamps, and optional multimodal observations. EnSIMem groups neighboring turns into contiguous, theme-coherent episodes rather than treating every turn or an entire session as a retrieval unit. Because the partition is query-independent and unsupervised, we use a fixed semantic policy: turns remain together when they share an entity, event, goal, or unresolved conversational reference, and a boundary is introduced only when a new topic or task persists. A speaker change alone does not trigger a split, and multimodal fields remain attached to the turn that introduced them.

At a high level, the query-independent partition can be written as

$$\Pi(S_n) = (e_{n,1}, \ldots, e_{n,m_n}), \qquad \bigcup_{j=1}^{m_n} e_{n,j} = S_n, \quad e_{n,j} \cap e_{n,k} = \varnothing \ (j \neq k), \tag{2}$$

where each $e_{n,j}$ is a contiguous, theme-coherent portion of session S_n. The equation emphasizes coverage and non-overlap, while the semantic policy determines the useful granularity. A theme is a short description of what the turns are jointly about, rather than a fixed ontology or a query-specific label. The policy keeps follow-up turns and references with the material they depend on, but separates a sustained change of topic, event, or goal.

This construction is a middle ground between individual turns and complete sessions: per-turn units can lose anaphoric context, whereas whole-session units can mix unrelated topics and add noise. Because segmentation is unsupervised, we do not claim a unique correct boundary for every conversation; nearby choices at the same semantic scale should preserve the same underlying evidence and have similar behavior. The episode-granularity ablation in Section 4 tests this trade-off by comparing theme-coherent, per-turn, and per-session units.

3.3 DIALOGUE-GROUNDED ENTITY–PROPERTY INDEX

For each episode e, EnSIMem extracts explicitly mentioned entities and salient events, including people, objects, places, activities, and named events. It then associates each entity with properties or relations supported by the local dialogue.

The extractor produces records of the form

$$r = \langle x, t, p, v, e, \pi \rangle, \tag{3}$$

where x is an entity, t its type, p a property or relation, v an optional concrete value, e the containing episode, and π a provenance link to the supporting evidence. The record is an address for finding the episode, not a replacement for it.

The retrieval-facing index is therefore written as:

$$[x][t][p : v] \longrightarrow \{episode_id\}. \tag{4}$$

The value is retained when the episode states a concrete realization of the property and is empty when no value is available. For example, a discussion of Caroline researching adoption agencies yields the index handle

$$[\text{Caroline}][\text{person}][\text{research : adoption agencies}] \longrightarrow \{e\}. \tag{5}$$

At query time, this handle identifies a relevant episode; the agent then reads the preserved source context to recover the answer and its qualifications. The index remains compact and reusable while the source retains temporal, conversational, and multimodal information.

3.4 Granularity-Controlled Property Alignment

A central requirement of EnSIMem is that properties extracted from the corpus and properties extracted from an agent request must have *compatible granularity*.

A property should be broad enough to support paraphrase or similarity match, but specific enough to preserve the identity of the underlying event or relation. An overly narrow property reduces recall because semantically equivalent expressions fail to match. But an overly broad property fails to support discriminative selection. For example, mapping *research*, *attending a class*, and *playing football* to a generic property *activity* makes it impossible to identify which activity it is about.

EnSIMem treats property extraction as a controlled abstraction process. The property captures the semantic relation needed for retrieval, while the value preserves the concrete realization when it is present. The index does not rely on a fixed, dataset-specific ontology. Instead, corpus properties and request properties are generated under the principle of *compatible semantic granularity* and compared through compatibility rather than exact string match.

This alignment allows expressions such as "*took a helicopter to Chicago*" and "*traveled to Chicago*" to share a travel-related retrieval property while preserving the specific transportation detail in the value. At the same time, unrelated actions remain distinguishable because they are not collapsed into a universal property (e.g., *activity*).

3.5 Requirement-Aware Query Planning

For a request q, the query planner extracts a complete set of atomic retrieval requirements from the question. We denote this requirement set by $H(q)$:

$$H(q) = \{h_1, \ldots, h_K\}, \qquad h_i = (x_i, t_i, p_i, v_i, c_i, \rho_i), \tag{6}$$

where each requirement identifies a target entity x_i, its type t_i, an intermediate-granularity property p_i, an optional value v_i or condition c_i, and the role ρ_i that the evidence will play in answering the request. The planner also classifies the reasoning type of the request. EnSIMem distinguishes, for example:

- **point questions**, which require one or a small number of relevant episodes;
- **temporal questions**, which require evidence ordered or compared by time;
- **compositional questions**, which require multiple related properties or entities; and
- **aggregation and frequency questions**, which require collecting all relevant episodes before counting or comparing them.

This classification determines the evidence budget before retrieval begins. A frequency question such as "How often did Sam get health checkups?" cannot be treated as an ordinary top-k question, because the answer may require counting multiple separate episodes even when no single episode states the final frequency.

3.6 REQUIREMENT-COVERAGE RETRIEVAL

For each requirement, structured matching first searches the entity-property index for records compatible in entity, type, property, value, and conditions. Matching records identify candidate episodes. Let $\mathcal{E}(q)$ denote the deduplicated set of candidate episodes returned for query q. If I denotes the structured index and T its textual representation, then

$$\mathcal{E}(q) = \mathrm{Dedup}(\mathrm{Structured}(I, H(q)) \cup \mathrm{Dense}(T, H(q))) . \tag{7}$$

The structured path is deliberately used as an address system: it can identify an episode even when the useful answer is stated across several nearby turns. When the structured path is incomplete or the wording differs from the indexed property, dense retrieval over the textual index representation supplies additional candidates. The two paths are complementary: the index provides a precise address, while dense retrieval protects recall.

The resulting candidates are merged and deduplicated by episode identity, and the system tracks which requirements are supported by the evidence collected so far. Point questions can stop after a small set covers the answer. Temporal and compositional questions continue until their comparisons or reasoning bridges are supported. Aggregation and frequency questions use an exhaustive stopping condition: all necessary episodes are retrieved before generation. Stopping is therefore determined by the query's evidence requirements rather than by an arbitrary numeric similarity threshold.

3.7 EVIDENCE-GROUNDED AGENT RESPONSE

After retrieval, selected episodes are expanded back into source evidence, including turn text, speaker, time, and available multimodal fields. If $\Gamma(\mathcal{E}(q))$ denotes this expansion, the agent response is conceptually

$$\hat{y} = G(q, H(q), \Gamma(\mathcal{E}(q))) . \tag{8}$$

Here G is the response-generation function: it receives the request, its retrieval requirements, and the expanded source evidence, and produces the grounded answer $\hat{y}$. The agent receives this compact evidence buffer together with the request and its requirements, and generates the answer from the source rather than from index records alone. Explicit facts are answered directly, uncertainty is acknowledged when evidence is incomplete, temporal wording preserves the precision of the source, and aggregation is computed across the retrieved episodes.

For example, for the request "What did Caroline research?" the planner identifies the entity Caroline and the property *research*. The index points to the episode in which she says that she is researching adoption agencies, and the agent expands that episode and answers "adoption agencies." The same process preserves an image caption, image-associated query, or image URL when it contains information not present in the dialogue text.

3.8 SEPARATION OF OFFLINE MEMORY CONSTRUCTION AND ONLINE AGENT REQUEST PROCESSING

Our design separates offline persistent construction from online request-dependent access. Offline processing turns the agent interaction corpus into reusable, provenance-linked indexed memory. For each request, the online stage selects a requirement-covered set of episodes and expands it into a compact evidence buffer for agent reasoning. This separation allows long-context history to be reused across requests while keeping each online reasoning context focused and grounded in the original agent interaction.

4 EXPERIMENTS

4.1 EVALUATION ON THE LOCOMO BENCHMARK DATASET

We report evaluation-model accuracy using the Memora-compatible protocol (Xia et al., 2026). Because answers are open-ended, BLEU and F1 which measure lexical overlap rather than semantic correctness and do not represent a convincing measure; we therefore use an LLM-based, fixed binary evaluation protocol. LLM model-based evaluation is a scalable proxy for human assessment, while

Table 1: Evaluation: LLM-model-based accuracy on LoCoMo. Published baseline and MEMORA scores are transcribed from the Memora benchmark table; asterisks mark baselines reported there from prior work. We report the same EnSIMem answers under GPT-4o-mini and Qwen3-32B evaluation models.

Method	Multi-hop	Temporal	Open-domain	Single-hop	Overall
Full Context	76.6%	81.9%	50.0%	88.5%	82.5%
RAG	55.7%	54.8%	45.8%	71.0%	63.3%
HippoRAG	39.0%	22.4%	51.0%	58.7%	47.1%
Zep*	53.7%	60.2%	43.8%	66.9%	61.6%
Mem0	62.4%	66.0%	50.0%	67.7%	65.3%
LangMem*	71.0%	50.8%	59.0%	84.5%	73.4%
Nemori*	75.1%	77.6%	51.0%	84.9%	79.4%
MEMORA (S)	78.4%	85.1%	59.4%	90.0%	84.9%
MEMORA (P)	78.7%	86.6%	59.4%	91.8%	86.3%
EnSIMem (evaluated by GPT-4o-mini)	**88.7%**	89.0%	70.8%	94.2%	**90.6%**
EnSIMem (evaluated by Qwen3-32B)	83.3%	**89.6%**	**72.9%**	**94.5%**	90.0%

Table 2: Evaluation: LLM-model-based accuracy on LongMemEval. "SS" denotes single-session. Published comparison rows are transcribed from Memora; EnSIMem reports the results from both evaluation models separately. Bold denotes the highest value in each column; underlining denotes the second-highest value.

Method	SS-pref.	SS-assist.	Temporal	Multi-session	Knowledge-update	SS-user	Average
Full Context	16.7%	**98.2%**	60.2%	51.1%	76.9%	85.7%	65.6%
Nemori	**86.7%**	92.9%	72.2%	55.6%	79.5%	90.0%	74.6%
MEMORA (S)	76.7%	76.8%	84.2%	73.7%	96.2%	97.1%	83.8%
MEMORA (P)	83.3%	78.6%	**89.5%**	78.2%	**97.4%**	**98.6%**	87.4%
EnSIMem (evaluated by GPT-4o-mini)	83.3%	91.7%	89.3%	**93.0%**	96.2%	**98.6%**	**92.8%**
EnSIMem (evaluated by Qwen3-32B)	**86.7%**	94.6%	89.1%	91.4%	96.1%	**98.6%**	**92.8%**

known evaluator biases motivate reporting two evaluation models separately (Zheng et al., 2023). The model used for answer generation in both offline and online stages is GPT-4.1-mini. We report the GPT-4o-mini and Qwen3-32B evaluation-model results separately rather than averaging them. Benchmark definitions and category mappings are provided in Appendix A.

EnSIMem is strongest on single-hop and temporal questions, while open-domain questions remain the most challenging slice. Overall, GPT-4o-mini evaluates the answer and gives an overall accuracy of 90.6%, and Qwen3-32B gives an overall accuracy of 90.0%. We report these evaluation-model results separately; the GPT-4o-mini row remains the directly comparable result against the published Memora numbers.

4.2 Evaluation on the LongMemEval Benchmark Dataset

We report evaluation-model accuracy using the same two-model protocol. Comparison-row provenance is provided in Appendix A.

GPT-4o-mini reaches 92.8% overall on LongMemEval, improving over MEMORA (P) by 5.41 points; Qwen3-32B also reaches 92.8%, improving by 5.44 points. The largest gains are on SS-assistant and multi-session (+13.1/+14.8 points for GPT-4o-mini and +16.0/+13.2 for Qwen3-32B), while the remaining categories differ by at most 1.30 points except SS-pref., which is 3.4 points higher. We therefore report the two evaluation-model results separately rather than averaging them.

4.3 Efficiency

We instrument Steps 4–6 on the current LoCoMo run, recording wall-clock latency, context-token counts, and retrieval search steps. The instrumentation does not change ranking, stopping rules, or generated answers.

Table 3: Online efficiency of EnSIMem on LoCoMo.

Stage	Mean s	P50 s	P95 s	Mean tokens
Step 4 planning	10.3	9.68	16.2	3,740
Step 5 retrieval	1.59	1.12	4.07	7,470
Step 6 generation	2.57	2.31	4.93	7,470
End to end	14.4	13.6	22.4	7,470

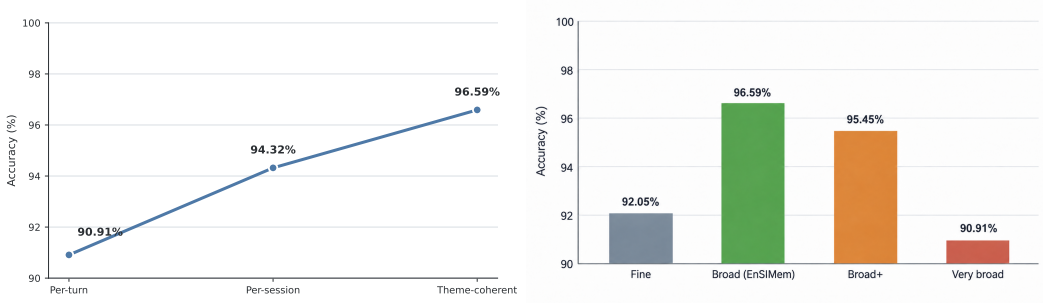

Figure 2: Ablation results on LoCoMo. The left panel compares episode granularity, showing that theme-coherent episodes outperform per-turn and per-session units. The right panel compares four property granularities: Fine, Broad (EnSIMem setting), Broad+, and Very broad. Broad+ is a slightly broader variant that merges additional action-like predicates into the shared property *activity*, while still preserving intermediate property distinctions. Higher accuracy is better.

The system averages 1.92 retrieval steps per query. Planning accounts for 71.1% of mean end-to-end latency, compared with 11.0% for retrieval and 17.9% for generation. This overhead constructs a more complete short reasoning buffer by identifying the property granularity, reasoning type, and evidence budget. Memora reports 5.70 seconds mean end-to-end latency and 4.61 seconds mean search latency (Xia et al., 2026). However, this cannot be treated as a direct comparison with our EnSIMem, as we use different processing engines and execution settings.

4.4 Ablation Study

We conduct controlled LoCoMo ablations, changing one design choice while fixing the answer model, evaluation model, and evaluation protocol. Each comparison isolates one source of the final gain rather than retuning the full pipeline.

Episode granularity. We compare theme-coherent episodes with per-turn and per-session units. This tests whether coherent middle-sized units improve the relevance–noise trade-off for an agent's reasoning buffer. The results are shown in Figure 2. EnSIMem obtains 96.59% accuracy, compared with 94.32% for per-session episodes and 90.91% for per-turn episodes. Thus, theme-coherent episodes improve accuracy by 2.27 and 5.68 percentage points over the two alternatives, respectively. This isolates memory-unit granularity: per-turn units can lose local references, whereas per-session units can admit unrelated context.

Property granularity. Using the same episodes, we compare four extraction granularities. Fine preserves surface predicates; Broad+ is slightly broader than EnSIMem's Broad prompt and merges additional action-like predicates into *activity*; Very broad maps more properties to *fact* or broad concepts. Accuracy is 92.05%, 96.59%, 95.45%, and 90.91% for Fine, Broad (EnSIMem), Broad+, and Very broad, respectively, showing that moderate normalization helps while excessive coarsening hurts retrieval. These results suggest that moderate normalization improves property alignment, while overly coarse properties remove distinctions needed for retrieval.

Structured matching versus dense retrieval. Dense retrieval ranks complete episodes by semantic similarity, whereas EnSIMem first matches entity, type, and property fields and links records to source episodes. With the answer and evaluation models fixed, EnSIMem reaches 96.59% versus

Table 4: Retrieval-route ablation on LoCoMo. Both variants use the same answer model and GPT-4o-mini evaluation model.

Retrieval variant	Accuracy
Dense retrieval only	86.36%
EnSIMem: structured entity/type/property retrieval	**96.59%**

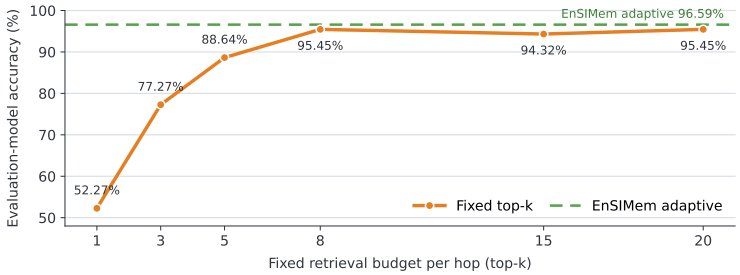

Figure 3: Fixed top-k budgets versus EnSIMem's adaptive budget on LoCoMo. The dashed line shows the adaptive result. Higher accuracy is better.

86.36% for dense retrieval (Table 4), showing the benefit of the structured access path. It therefore tests whether explicit index structure contributes beyond dense semantic similarity alone.

Retrieval Episode Strategy. We compare the adaptive budget with fixed top-k policies for $k \in \{1, 3, 5, 8, 15, 20\}$. Fixed retrieval rises from 52.27% at $k = 1$ to 95.45% at $k = 8$ before saturating or declining, while EnSIMem reaches 96.59% (Figure 3). This isolates query-type-aware allocation: fixed budgets trade recall and noise uniformly, whereas the adaptive policy can spend evidence where the question requires coverage.

5 Discussion, Limitations, and Future Work

Our study focuses on episodic conversational memory and does not yet cover procedural or semantic memory. The pipeline relies on LLM-based episode segmentation, property extraction, and query planning; although our ablations show robust trends, different models or prompts may produce different intermediate representations, motivating future work on confidence estimation and cross-model validation.

LoCoMo and LongMemEval primarily evaluate conversational recall and reasoning, rather than skill execution or persistent factual updates. In addition, evaluation-model preferences may affect absolute scores, while runtime comparisons depend on the underlying processing engine and deployment configuration.

A natural extension is to integrate EnSIMem with procedural and semantic memory. Procedural records could represent reusable skills with goals, preconditions, actions, outcomes, and provenance, while semantic records could maintain entity–property facts with validity intervals and supporting episodes. An agent harness could coordinate these stores through shared entity identifiers and provenance links, selecting the appropriate memory type for each request. Future evaluation should therefore include skill reuse, fact consolidation, temporal updates, and cross-memory reasoning.

6 Conclusion

EnSIMem reframes long-term conversational memory as a short-context reasoning problem. Theme-coherent episodes organize local context, while entity–property records provide precise handles that connect each request to the relevant source evidence. Requirement-aware planning and adaptive retrieval select only the evidence needed for point, temporal, compositional, and aggrega-

tion questions. Unlike generalized summaries or anonymous chunks, the reasoning buffer retains the original text, provenance, and available multimodal evidence, reducing irrelevant context while preserving the details required for grounded answers. This evidence-preserving design provides a foundation for extending structured memory access beyond episodic memory.

AI USE DISCLOSURE

In this work, we used generative AI tools to polish the writing, verify grammar errors, and implement and debug software. We also used GPT-4.1-mini as the answer-generation model in our benchmark experiments and GPT-4o-mini and Qwen3-32B as evaluation models for assessing generated answers. These model-based evaluations were conducted as part of the experimental protocol and were not used to create synthetic training data or alter the ground-truth annotations.

We did not use generative AI tools to generate synthetic datasets, provide proofs, or establish mathematical claims without author verification. All AI-assisted code was reviewed, executed, and tested by the authors. The authors independently checked the manuscript text, citations, mathematical notation, figures, experimental configurations, and reported results. We take full responsibility for the final content of this paper.

REFERENCES

Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations*, 2024.

Yushi Bai, Shangqing Tu, Jiajie Zhang, Hao Peng, Xiaozhi Wang, Xin Lv, Shulin Cao, Jiazheng Xu, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. LongBench v2: Towards deeper understanding and reasoning on realistic long-context multitasks. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2025. doi: 10.18653/v1/2025.acl-long.183.

Prateek Chhikara, Dev Khant, Saket Aryan, Taranjeet Singh, and Deshraj Yadav. Mem0: Building production-ready ai agents with scalable long-term memory. *arXiv preprint arXiv:2504.19413*, 2025.

Yufeng Du, Minyang Tian, Srikanth Ronanki, Subendhu Rongali, Sravan Bodapati, Aram Galstyan, Azton Wells, Roy Schwartz, Eliu A Huerta, and Hao Peng. Context length alone hurts llm performance despite perfect retrieval. *arXiv preprint arXiv:2510.05381*, 2025.

Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, Dasha Metropolitansky, Robert Osazuwa Ness, and Jonathan Larson. From local to global: A graph RAG approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*, 2024. doi: 10.48550/arXiv.2404.16130.

Xiaoxin He, Yijun Tian, Yifei Sun, Nitesh V. Chawla, Thomas Laurent, Yann LeCun, Xavier Bresson, and Bryan Hooi. G-Retriever: Retrieval-augmented generation for textual graph understanding and question answering. *arXiv preprint arXiv:2402.07630*, 2024. doi: 10.48550/arXiv.2402.07630.

Yuntong Hu, Zhihan Lei, Zheng Zhang, Bo Pan, Chen Ling, and Liang Zhao. GRAG: Graph retrieval-augmented generation. *arXiv preprint arXiv:2405.16506*, 2024. doi: 10.48550/arXiv.2405.16506.

Gautier Izacard, Mathilde Caron, Lucas Hosseini, Sebastian Riedel, Piotr Bojanowski, Armand Joulin, and Edouard Grave. Unsupervised dense information retrieval with contrastive learning. *Transactions on Machine Learning Research*, 2022.

Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. HippoRAG: Neurobiologically inspired long-term memory for large language models. *arXiv preprint arXiv:2405.14831*, 2024. doi: 10.48550/arXiv.2405.14831.

Jiazheng Kang, Mingming Ji, Zhe Zhao, and Ting Bai. Memory OS of AI agent. *arXiv preprint arXiv:2506.06326*, 2025. doi: 10.48550/arXiv.2506.06326.

Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pp. 6769–6781, 2020. doi: 10.18653/v1/2020.emnlp-main.550.

Omar Khattab and Matei Zaharia. ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. In *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 39–48, 2020. doi: 10.1145/3397271.3401075.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pp. 9459–9474, 2020.

Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024. doi: 10.1162/tacl_a_00638.

Adyasha Maharana, Dong-Ho Lee, Sergey Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Fang. Evaluating very long-term conversational memory of LLM agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 13829–13849, 2024. doi: 10.18653/v1/2024.acl-long.747.

Xuanyu Meng, Jiashuo Sun, Jash Rajesh Parekh, and Jiawei Han. EnSI-RAG: Entity-structure-indexed retrieval-augmented generation for long-document question answering. *arXiv preprint arXiv:2608.21252*, 2026. doi: 10.48550/arXiv.2608.21252.

Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Memgpt: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*, 2023. doi: 10.48550/arXiv.2310.08560.

Jash Rajesh Parekh, Pengcheng Jiang, and Jiawei Han. Structure-augmented reasoning generation. *arXiv preprint arXiv:2506.08364*, 2025.

Joon Sung Park, Joseph C. O'Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023. doi: 10.1145/3586183.3606763.

Preston Rasmussen, Pavlo Paliychuk, Travis Beauvais, Jack Ryan, and Daniel Chalef. Zep: A temporal knowledge graph architecture for agent memory. *arXiv preprint arXiv:2501.13956*, 2025. doi: 10.48550/arXiv.2501.13956.

Shreya Shankar, Tristan Chambers, Tarak Shah, Aditya G. Parameswaran, and Eugene Wu. DocETL: Agentic query rewriting and evaluation for complex document processing. *arXiv preprint arXiv:2410.12189*, 2024. doi: 10.48550/arXiv.2410.12189.

Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, 2023.

Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 10014–10037, 2023. doi: 10.18653/v1/2023.acl-long.557.

Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023. doi: 10.48550/arXiv.2305.16291.

Minzheng Wang, Longze Chen, Cheng Fu, Shengyi Liao, Xinghua Zhang, Bingli Wu, Haiyang Yu, Nan Xu, Lei Zhang, Run Luo, Yunshui Li, Min Yang, Fei Huang, and Yongbin Li. Leave no document behind: Benchmarking long-context LLMs with extended multi-document QA. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024. doi: 10.18653/v1/2024.emnlp-main.322.

Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. Longmemeval: Benchmarking chat assistants on long-term interactive memory. *arXiv preprint arXiv:2410.10813*, 2024. doi: 10.48550/arXiv.2410.10813.

Menglin Xia, Xuchao Zhang, Shantanu Dixit, Paramaguru Harimurugan, Rujia Wang, Victor Ruhle, Robert Sim, Chetan Bansal, and Saravan Rajmohan. Memora: A harmonic memory representation balancing abstraction and specificity. *arXiv preprint arXiv:2602.03315*, 2026. doi: 10.48550/arXiv.2602.03315.

Wujiang Xu, Zujie Liang, Kai Mei, Hang Gao, Juntao Tan, and Yongfeng Zhang. A-mem: Agentic memory for LLM agents. *arXiv preprint arXiv:2502.12110*, 2025. doi: 10.48550/arXiv.2502.12110.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging llm-as-a-judge with mt-bench and chatbot arena. In *Advances in Neural Information Processing Systems*, 2023. URL https://arxiv.org/abs/2306.05685.

Wanjun Zhong, Lianghong Guo, Qiqi Gao, He Ye, and Yanlin Wang. Memorybank: Enhancing large language models with long-term memory. *arXiv preprint arXiv:2305.10250*, 2023.

A APPENDIX

A.1 BENCHMARK OVERVIEW

Table 5. Benchmark scope, question families, and evaluation protocol used in this paper.

Benchmark	Scope and question families	Protocol and role in this paper
LoCoMo	**Scope:** 10 long, multi-session conversations with speaker turns, timestamps, and optional image observations. **Question families:** multi-hop, temporal, open-domain, and single-hop questions testing entity tracking, paraphrase alignment, temporal grounding, multimodal evidence use, and aggregation.	**Protocol:** Evaluation-model scoring on the same generated answers using GPT-4o-mini and Qwen3-32B, reported separately. **Result:** 90.6% and 90.0% overall, respectively. Published comparison rows are transcribed from Memora (Maharana et al., 2024; Xia et al., 2026).
LongMemEval	**Scope:** interactive long-term memory tasks pairing a conversation history with a later user request. **Question families:** single-session preference, single-session assistant, temporal, multi-session, knowledge update, and single-session user questions.	**Protocol:** the same two evaluation-model procedure, with GPT-4o-mini and Qwen3-32B reported separately. **Result:** 92.8% overall for each evaluation model after rounding. Full Context, Nemori, and MEMORA values follow the published Memora table (Wu et al., 2024; Xia et al., 2026).

A.2 BASELINE AND METHOD OVERVIEW

Table 6. Baseline and method comparison used in the benchmark tables.

Method	Representation and access	Comparison role
Full Context	Complete interaction history is passed to the answer model; there is no persistent selection structure.	Upper reference that minimizes retrieval misses, but exposes the model to irrelevant context and high input cost.
RAG	Anonymous text chunks retrieved by dense semantic similarity.	Standard retrieval baseline; it does not explicitly address the relevant entity, property, provenance, or reasoning type (Lewis et al., 2020; Karpukhin et al., 2020).
HippoRAG	Graph-inspired associative memory with propagation across related information.	Tests whether graph retrieval alone matches explicit entity-property addressing for multi-hop recall (Jiménez Gutiérrez et al., 2024).
Zep	Temporally aware knowledge graph with evolving facts and relations.	Tests explicit temporal organization; it does not use EnSIMem's source episodes or requirement-specific evidence budgets (Rasmussen et al., 2025).
Mem0	Consolidated facts and preferences, with graph variants for relations.	Production-oriented persistent memory; consolidation can remove local dialogue context (Chhikara et al., 2025).
Nemori	Persistent conversational-memory representation.	Established memory baseline; we report the published comparison values rather than reimplementing or retuning it (Xia et al., 2026).
MEMORA (S)	Harmonic entries combine an abstraction, a richer value, and cue indices.	Sparse-retrieval baseline testing the abstraction-specificity trade-off (Xia et al., 2026).
MEMORA (P)	The same harmonic representation with an adaptive policy retriever.	Strongest published comparator and closest policy-driven baseline (Xia et al., 2026).

B CASE STUDIES

We present three representative LoCoMo cases in which the evaluation model labels EnSIMem correct while labeling Memora incorrect. The cases are selected to expose three distinct mechanisms: cross-episode evidence coverage, temporal grounding, and fine-grained entity-property localization. Each table records the question, the required operation, the memory representation, the retrieved evidence, the generated responses, and the specific failure mode of the competing system.

Table 7. Case study 1: cross-episode evidence coverage for a multi-fact question.

Case component	Details
Question	*What do Melanie's kids like?* (query b0008; category 1). The question refers to the children as a group and asks for their salient interests, not for one isolated activity from a single conversation.
Required operation	The system must resolve the subject *Melanie's kids*, collect answer-bearing facts across sessions, and avoid replacing the requested interests with a generic list of family activities. The gold answer is *dinosaurs, nature*.
Memora representation	Memora exposes family-related memories as broad activity summaries. These memories contain camping, painting, beach, hiking, and pottery, but they do not provide an explicit address for the entity *Melanie's kids* together with the property *likes* and its answer-bearing values.
Memora response	"Melanie's kids like camping, painting, beach, hiking, pottery." **Incorrect**. The response is topically related, but it omits both benchmark-targeted facts: the children love nature and are interested in dinosaurs.
Why Memora fails	Memora's broad summary conflates nearby family activities with the specific values requested by the question. The retrieval result is not empty; rather, it fails to preserve the right evidence with sufficient priority. This is a coverage-and-specificity failure caused by semantic aggregation without a structured child-interest address.
EnSIMem query structure	The query is represented as an entity-property requirement over Melanie's children and their interests. Because the answer is aggregative, the planner keeps expanding evidence until the relevant child-interest facts are covered instead of stopping after one semantically similar episode.
EnSIMem representation	Theme-coherent episodes preserve the local dialogue and keep each extracted entity-property record linked to its source turns. This prevents facts about Melanie's own activities from being confused with facts about her children.
EnSIMem retrieval	EnSIMem retrieves `conv-26::session_4::episode_002` (D4:8), where Melanie says that the two younger children love nature, and `conv-26::session_6::episode_003` (D6:6), where the children are excited about a dinosaur exhibit. The two episodes jointly cover the gold answer.
Evidence trace	D4:8 states that the children "love nature" after a family camping trip. D6:6 states that they were excited about a dinosaur exhibit and enjoyed learning about animals. These are direct, dialogue-grounded supports rather than inferred preferences.
EnSIMem response	EnSIMem identifies the children's interest in nature and dinosaurs and retains the supporting episode references. **Correct**.
Why EnSIMem succeeds	Structured matching first identifies the correct entity and property, while adaptive expansion supplies the second episode needed for complete coverage. The answer buffer therefore contains the answer-bearing facts instead of only a broad collection of related activities.
Key difference	Memora returns a generic semantic activity summary. EnSIMem uses an addressable entity-property index and query-type-aware evidence coverage to aggregate the right facts across multiple episodes before generation.

Table 8. Case study 2: temporal grounding from a relative-time expression.

Case component	Details
Question	*When did Melanie read the book "Nothing is Impossible"?* (query b0039; category 2). The question requires resolving a relative temporal expression rather than retrieving a date stated as an absolute year.
Required operation	Link the book-reading event to its temporal qualifier, interpret "last year" relative to the dated conversation, and return the normalized year. The gold answer is *2022*.
Memora representation	Memora's memory store contains many book and reading memories, but it does not maintain a directly addressable relation between Melanie, the reading event, the book reference, and the relative-time phrase in the source episode.
Memora response	"The memories do not mention Melanie reading the book 'nothing is impossible'." **Incorrect**. The source statement is present in the conversation, but it is not surfaced as an answerable memory.
Why Memora fails	The failure is relational rather than purely lexical. The question names the book explicitly, while the source says "this book" and places the event inside a dated conversation. A generic memory lookup therefore treats the title as absent and does not connect the event to its relative-time qualifier.
EnSIMem query structure	The planner creates a temporal requirement for Melanie's reading activity and the referenced book. It marks the relative-time expression as answer-bearing, so retrieval must preserve both the source timestamp and the phrase "last year."
EnSIMem representation	The entity-property record links Melanie to the reading event and retains the value-level temporal qualifier. The original episode remains available for interpreting the reference "this book" and for checking the date used in normalization.
EnSIMem retrieval	EnSIMem retrieves `conv-26::session_7::episode_002`, specifically D7:8: "This book I read last year reminds me to always pursue my dreams." The episode is dated 12 July 2023.
Evidence trace	The phrase "last year" is directly grounded in D7:8. Applying the episode date as the reference point gives 2022, which is the benchmark answer.
EnSIMem response	EnSIMem answers that Melanie read the book "last year" and normalizes the relative expression to 2022. **Correct**.
Why EnSIMem succeeds	The index preserves the event, entity, temporal qualifier, and source timestamp as a connected retrieval path. Query-specific planning then performs the small amount of temporal normalization needed after retrieval.
Key difference	Memora requires a directly matching memory statement and reports no evidence. EnSIMem preserves temporal relations as searchable structure and uses the source episode to resolve relative time without losing provenance.

Table 9. Case study 3: entity-property specificity for multimodal art evidence.

Case component	Details
Question	*What kind of art does Caroline make?* (query b0017; category 1). The question asks for the relevant artistic form, not merely whether Caroline engages in visual art.
Required operation	Identify Caroline as the entity, distinguish the artistic-form property from the broad activity of making art, and use both textual and image-grounded evidence. The gold answer is *abstract art*.
Memora representation	Memora stores broad activity-level descriptions such as painting and drawing. These summaries do not expose a sufficiently specific property for the artistic form, and the associated image descriptions are not retained as an explicit answer handle.
Memora response	"Painting and drawing." **Incorrect**. The response is related to Caroline's activities but is too coarse to answer the requested subtype.
Why Memora fails	The representation collapses several distinct records into one generic art activity. It therefore cannot distinguish the requested style from the broader fact that Caroline paints or draws, and it does not use the abstract-landscape image description as a retrieval cue.
EnSIMem query structure	The planner addresses the entity Caroline and requests the property describing the kind or form of her art. The requirement is specific enough to prefer an abstract-art record over generic records about painting, drawing, or art as emotional expression.
EnSIMem representation	EnSIMem indexes Caroline's art-related records at an intermediate property granularity and links each record to its source episode, original dialogue, timestamp, and multimodal fields. This creates a direct path from Caroline to the artistic-form property.
EnSIMem retrieval	The system retrieves `conv-26::session_9::episode_003` (D9:14), whose image caption describes an abstract landscape, and `conv-26::session_11::episode_003` (D11:8 and D11:12), which preserve Caroline's paintings about inclusivity, identity, and self-acceptance.
Evidence trace	D9:14 supplies the image-grounded abstract-landscape cue. D11:8 and D11:12 confirm that the artwork belongs to Caroline and connect the visual artifacts to her stated artistic practice rather than to Melanie's art activities.
EnSIMem response	The retrieved evidence supports the more specific answer *abstract art*. The evaluation model labels the EnSIMem answer **Correct**.
Why EnSIMem succeeds	The entity-property index retains the distinction between a broad activity and a requested subtype, while provenance-preserving episode expansion supplies the surrounding dialogue and image evidence needed to interpret the artwork.
Key difference	Memora retrieves a generic activity category. EnSIMem preserves the entity, property granularity, source episode, and image-grounded evidence together, allowing the agent to answer at the specificity required by the benchmark.

C Evaluation Model Rationale

Our evaluation follows the answer-scoring protocol used by Memora and its underlying Mem0 evaluation implementation (Xia et al., 2026). This choice is appropriate for long-term agent memory because the benchmark answers are open-ended: a correct answer may paraphrase the reference, use a relative temporal expression, or include additional explanation without matching the reference string word for word.

LoCoMo decision rule. For every query, the evaluation program passes the question, the gold answer, and the generated answer to an evaluation model. The Memora/Mem0-compatible prompt asks whether the generated answer is semantically consistent with the gold answer and requests a binary `CORRECT` or `WRONG` label in JSON. The implementation maps `CORRECT` to one and `WRONG` to zero, then aggregates these decisions overall and by LoCoMo category. Category 5 is excluded according to the Memora-compatible scoring policy. Thus, the reported LoCoMo number measures answer-level semantic correctness rather than token overlap.

LongMemEval decision rule. LongMemEval uses the corresponding Memora-style task-specific checker. The evaluation program again separates answer generation from evaluation, but selects a prompt according to the question type: preference questions use a personalization rubric, knowledge-update questions check the updated answer, temporal questions include their temporal tolerance, and unanswerable items check whether the model correctly abstains. The evaluator receives the question, reference answer or rubric, and generated response, and returns `yes` or `no`; the implementation maps `yes` to correct and aggregates accuracy by question type and overall. This preserves the benchmark's intended semantics while keeping the evaluation procedure independent of the retrieval pipeline.

Why this is preferable to lexical metrics. BLEU and F1 primarily reward shared surface tokens. They can penalize a valid paraphrase, fail to recognize equivalent date formats, and provide little evidence that the answer identifies the requested entity, property, or event. The evaluation-model prompt instead exposes the complete question–gold–prediction triple and explicitly allows semantically equivalent wording and time expressions. This makes the protocol better aligned with the actual objective of an agent-memory system: recovering the right fact or event for the user's request.

Reproducibility and model dependence. LoCoMo evaluation uses temperature zero, a fixed seed of 42, and a fixed Memora-compatible prompt; the LongMemEval checker uses temperature zero and its fixed task-specific prompts. The same protocol and aggregation code are used across methods within each benchmark. Because model-based evaluation can still reflect evaluator-specific preferences, we report GPT-4o-mini and Qwen3-32B results separately rather than treating their arithmetic mean as a new ground truth. Agreement in the qualitative trends across the two evaluation models provides a more transparent robustness check while preserving the identity of each evaluation model.

Relation to the benchmark comparison. Using the Memora-compatible protocol keeps our LoCoMo comparison on the same answer-level scale as the published Memora rows. It also makes the source of every score explicit: the answer model generates the response, the evaluation model makes the binary semantic decision, and the reported accuracy is the fraction of correct decisions. This separation avoids conflating answer generation with evaluation and allows future work to replace or add evaluation models without changing the stored answers or retrieval pipeline.

Table 10: Consistency of GPT-4o-mini and Qwen3-32B evaluation-model scores. Δ is defined as Qwen3-32B minus GPT-4o-mini; all difference columns are in percentage points. The correlation is computed across benchmark categories, excluding the overall column.

Benchmark	m	$\overline{\Delta}$	MAE	SD(Δ)	max $\lvert\Delta\rvert$	$\lvert\Delta_{\text{overall}}\rvert$	r
LoCoMo	4	-0.620	2.11	2.88	5.46	0.610	0.946
LongMemEval	6	$+0.750$	1.37	1.79	3.37	0.030	0.936

D Evaluation-Model Consistency

We additionally examine whether the reported scores are stable across the two evaluation models. For each benchmark category, let g_c and q_c denote the accuracy obtained with GPT-4o-mini and Qwen3-32B, respectively, and define the signed difference $\Delta_c = q_c - g_c$ in percentage points (pp). We report the mean signed difference, mean absolute difference (MAE), the population standard deviation of the category differences, the largest absolute category difference, the gap between the two overall scores, and the Pearson correlation across category accuracies. The analysis uses the unrounded category-level results underlying the main tables before their three-significant-digit display; it does not treat the arithmetic mean of the two evaluation models as a new score.

The two evaluation models therefore give highly similar aggregate conclusions. On LoCoMo, the overall scores differ by only 0.610 pp (90.6% versus 90.0%), and the category-level MAE is 2.11 pp. The larger LoCoMo discrepancy is concentrated in multi-hop questions, where Qwen3-32B is 5.46 pp lower; the remaining category gaps are at most 2.09 pp. On LongMemEval, the overall gap is only 0.030 pp (92.8% versus 92.8% in three significant figures), with a smaller category-level MAE of 1.37 pp. The largest LongMemEval difference is 3.37 pp on single-session preference questions, while temporal reasoning, knowledge update, and single-session user questions differ by at most 0.250 pp. The high category-level correlations ($r = 0.946$ for LoCoMo and $r = 0.936$ for LongMemEval), together with the small absolute gaps, indicate that the headline findings are consistent across the two evaluation models even though individual categories can show model-specific sensitivity.

This is an aggregate stability analysis rather than a claim of perfect per-query agreement. A paired reliability analysis would require retaining both binary labels for every identical query. With those labels, one could additionally report exact agreement, Cohen's κ, a paired McNemar test, and bootstrap confidence intervals. We do not infer those quantities from category totals alone; instead, we report the two evaluation-model results separately and use the agreement of their aggregate trends as the robustness check.

E PROMPTS

This appendix records the core prompts used by EnSIMem. The text inside each box is the prompt template passed to the corresponding LLM stage; braces denote runtime substitutions such as a conversation, an episode, a query, or a retrieval plan. We show the prompts that determine the memory representation, query decomposition, evidence-grounded answering, and evaluation. Repair and validation prompts are invoked only when a model returns malformed JSON and follow the same schemas shown here. Structured matching, dense fallback, evidence-budget expansion, and score aggregation are deterministic operations and do not invoke an LLM. The corresponding source files are `LoCoMo/prompts.py`, `LoCoMo/08_memora_llm_judge.py`, and `LongMemEval/memora_evaluation.py`; the runtime substitutions shown below are made by the pipeline scripts.

E.1 OFFLINE MEMORY CONSTRUCTION

Theme-coherent episode partitioning. The partition prompt creates contiguous evidence units before any entity-property records are extracted. It preserves dialogue order and keeps image metadata attached to the turn that introduced it.

Theme-Coherent Episode Partitioning: System Prompt

```
You partition a complete dialogue session into contiguous theme-coherent memory episodes.
    These boundaries define immutable evidence units. Return one JSON object only.
```

Theme-Coherent Episode Partitioning: User Prompt

```
Partition this session.

A memory episode is one contiguous span mainly focused on one coherent entity, event, goal,
     or bounded topic. Use theme_type entity, event, or topic. Split only on a genuine
    semantic shift; keep greetings, acknowledgements, clarification questions, and short
    follow-ups with the material they support.

Structural rules:
- Cover every supplied dialogue turn exactly once and in order, without gaps or overlaps.
- Do not split merely because the speaker changes.
- Avoid one-turn fragments unless a turn clearly introduces a separate durable theme.
- Keep text, image captions, image-retrieval descriptions, and image URLs attached to their
     DIA.
- Copy start_dia_id and end_dia_id exactly.

Return JSON only:
{"segments": [{"start_dia_id": "D1:1", "end_dia_id": "D1:5", "theme": "concise canonical
    theme", "theme_type": "entity|event|topic", "theme_description": "one-sentence scope",
     "boundary_reason": "session_start or semantic shift"}]}

Conversation: {conversation_id}
Session: {session_id}
Observed at: {observed_at}
Turns:
{session_text}
```

Entity-property index extraction. This prompt is query-independent: it extracts atomic records once and keeps the source episode as the only answer evidence. The property policy is deliberately intermediate-grained, while values and source DIA identifiers retain concrete detail.

Entity-Property Indexing: System Prompt

```
You are the high-recall, precision-preserving information-extraction stage of an episodic
    memory system. Build a query-independent entity-structured index from one complete
    theme episode. The episode remains the only answer evidence; records are navigation
    handles. Follow the minimum-sufficient predicate policy and return one valid JSON
    object only.
```

Entity-Property Indexing: User Prompt

```
Extract an exhaustive set of atomic records from the episode below.

Read every dialogue turn and every displayed metadata field. Image captions and image-
    retrieval descriptions are auxiliary textual evidence; an image URL is opaque
    provenance and must never be used for outside lookup.

Extraction requirements:
- Scan each DIA independently and emit a record for every explicit factual clause.
- Resolve I/my/me using the speaker on that turn and cite exact evidence_dia_ids.
- Preserve the speaker-to-entity relation for explicit first-person actions and emit
    explicit inverse relations only when entailed.
- Use one entity and one minimum-sufficient reusable property. Keep distinctive predicates
    such as research, travel, attend, read, paint, checkup, and support; do not collapse
    them all into activity or concatenate the subject into the property.
- Put concrete source detail in value. Preserve relative time, modality, polarity,
    conditions, and image metadata.
- Do not infer facts from world knowledge, episode adjacency, or an image URL.

Return one object only:
{"records": [{"entity": "...", "entity_type": "person|organization|place|object|event|
    activity|concept|other", "property": "short predicate", "value": "explicit value or
    empty", "condition_property": "... or empty", "condition_value": "... or empty", "
    property_kind": "aspect|relation", "modality": "observed|planned|desired|hypothetical|
    negated|uncertain", "valid_time": "... or empty", "source": "speaker", "
    evidence_dia_ids": ["D1:1"], "projection_kind": "direct|inverse|state_projection|
    event_access", "confidence": 0.0}]}

Episode metadata:
Episode ID: {episode_id}
Conversation: {conversation_id}
Session: {session_id}
Observed at: {observed_at}
Episode theme: {episode_theme} ({episode_theme_type})

Episode turns:
{episode_text}
```

E.2 Online retrieval and response generation

Requirement-aware query planning. The planner converts a question into explicit requirements and an evidence-seeking hop graph. It uses the same schema as the index, marks aggregation questions as `all_matching`, and leaves unknown answer values empty.

Requirement-Aware Query Planning: System Prompt

```
You are the exhaustive query-decomposition stage of an episodic memory system. Convert the
    question into searchable requirements and an evidence-seeking hop graph using the same
     [entity][entity_type][property:value][condition:value] schema as the memory index.
    Preserve every meaningful discriminator, but never guess an answer. Do not answer the
    question. Return exactly one valid JSON object.
```

Requirement-Aware Query Planning: User Prompt

```
Build an exhaustive retrieval plan for the question below.

Planning rules:
- Include the named entity, the relationship that disambiguates the target, the requested
    property, and all explicit time, quantity, and location constraints.
- Use minimum-sufficient properties aligned with the index; do not invent subject-prefixed
    or compound labels.
- Keep unknown answer values empty. Use only information stated in the question or a
    declared bridge.
- Add separate hops for independently useful discriminators. For example, a daughter's
    birthday requires relationship=daughter, property=birthday, and property=time.
- Use reasoning_type=aggregation or comparison and retrieval_scope=all_matching for counts,
     exhaustive lists, frequencies, intervals, and first/second questions. Otherwise use
    point retrieval.
- Preserve relative time and explicit conditions. Do not answer the question in this stage.

Return one object only:
{"answer_target": {"type": "time|value|entity|location|count|list|boolean|likelihood|
    explanation", "description": "..."}, "reasoning_type": "direct|aggregation|comparison|
```

```
    temporal|causal|counterfactual|commonsense_inference|multi_hop|unanswerable", "
    retrieval_scope": "point|all_matching", "required_properties": [{"entity": "...", "
    entity_type": "...", "broad_property": "...", "property_text": "...", "value": "...",
    "role": "entity|relation|answer_property|time|constraint"}], "hops": [{"hop_id": "h1",
     "purpose": "...", "anchor": {"entity": "", "entity_type": "", "property": "", "value
    ": "", "condition_property": "", "condition_value": ""}, "depends_on": [], "
    bridge_request": ""}]}

Observed index property vocabulary:
{index_property_vocabulary}

Question: {question}
```

Evidence-grounded answer generation. The answer model receives complete retrieved episodes, not only the structured records. The structured records and hop plan are navigation aids; the final answer must be supported by the original dialogue and attached image metadata.

Evidence-Grounded Answer Generation: System Prompt

```
Answer only from the supplied complete original conversation episodes. The structured
    indexes and hop plan are navigation aids, not evidence. Read evidence from every hop,
    combine facts when the plan is multi-hop, and keep entities correctly bound. Do not
    use outside knowledge. Treat captions and retrieval descriptions as textual evidence
    attached to their exact DIA; treat image URLs as opaque provenance.

For list, count, frequency, interval, and first/second questions, inventory every
    qualifying event across all supplied episodes before answering. For temporal questions
    , bind the requested event first and preserve approximate source wording. If the
    evidence does not establish the answer, output exactly: Unknown. Return only the
    shortest sufficient answer without explanation.
```

Evidence-Grounded Answer Generation: User Prompt

```
Question: {question}

Retrieval plan:
{plan}

Complete original episodes retrieved across all hops:
{episodes}

Answer:
```

E.3 Evaluation prompts

The benchmark evaluation is separate from answer generation. The evaluator receives the question, the benchmark gold answer (or rubric), and the generated answer. It does not receive the retrieval plan or hidden retrieval diagnostics. LoCoMo uses the Memora-compatible binary prompt below; LongMemEval selects the corresponding task-specific variant according to the question type.

LoCoMo Evaluation Model: Single User Prompt

```
Your task is to label an answer to a question as CORRECT or WRONG. You will be given (1) a
    question, (2) a gold answer, and (3) a generated answer. Be generous with grading: a
    longer answer is correct when it touches the same topic as the gold answer, and a time
     answer is correct when it refers to the same date or period even if the format
    differs.

Question: {question}
Gold answer: {gold_answer}
Generated answer: {generated_answer}

Return only a JSON object with the key "label" and the value CORRECT or WRONG.
```

LongMemEval Evaluation Model: Task-Specific User Prompt

```
I will give you a question, a reference answer or rubric, and a model response. Answer yes
    if the response is semantically correct and no otherwise. The response may use
```

```
    equivalent wording. For preference questions, check whether it satisfies the
    personalization rubric. For knowledge-update questions, accept the updated answer even
     if older information is also present. For temporal-reasoning questions, do not
    penalize an off-by-one error when the benchmark asks for a number of days, weeks, or
    months. For unanswerable questions, answer yes only when the response correctly
    identifies that the requested information is not given.

Question: {question}
Reference answer or rubric: {answer}
Model response: {hypothesis}

Answer yes or no only.
```

The LoCoMo implementation maps `CORRECT` to one and `WRONG` to zero; the LongMemEval implementation maps `yes` to one and `no` to zero. Both protocols use temperature zero. The LoCoMo evaluator also fixes the seed to 42, while the LongMemEval checker uses top-$p = 1$. Category-wise and overall accuracies are computed from these binary decisions.