Learning from How Users Interpret: Personalizing Recommendation Explanation Order

Deo Munduku
Paris-Dauphine University, PSL Research University
UMR CNRS 7243, LAMSADE
Paris, France
deo.munduku@lamsade.dauphine.fr

Elsa Negre
Paris-Dauphine University, PSL Research University
UMR CNRS 7243, LAMSADE
Paris, France
elsa.negre@lamsade.dauphine.fr

arXiv:2611.08785v1 [cs.HC] 5 Nov 2026

Abstract

Personalized explanations typically adapt their content or presentation based on preferences, interaction histories, or other user characteristics. We examine whether regularities inferred from users' prior textual interpretations can also support personalization of explanation organization. In a controlled within-participant study with 122 participants, prior interpretations were represented along three dimensions: information selection, relative salience, and relational organization, with the relational dimension partially operationalized through co-mobilization. These representations were aggregated into fixed historical profiles and used to organize otherwise identical explanation sentences. Participants encountered four organization procedures: their own historical profile (OWN), another user's historical profile (DONOR), a leave-one-participant-out population profile (POP), and a target-aware language-model procedure using the user's historical profile (PROFILE-LLM). Post-exposure objective comprehension was the primary outcome. OWN achieved 86.5% correct responses, compared with 66.0% for DONOR, 63.9% for POP, and 73.0% for PROFILE-LLM. Mixed-effects analyses showed higher comprehension with OWN than with each alternative procedure. No reliable overall condition effect was detected for recommendation acceptance. These findings indicate that historical profiles derived from prior interpretations can provide an actionable signal for personalizing explanation order while keeping explanatory content and wording fixed.

CCS Concepts

• **Human-centered computing** → **User models**; **Empirical studies in HCI**; • **Information systems** → *Users and interactive retrieval*.

Keywords

personalized explanations, explanation organization, objective comprehension, historical user profiles, sentence ordering, human-information interaction

1 Introduction

Personalized explanations aim to adapt the information presented to the characteristics of each user. Existing approaches rely, among other signals, on preferences, interaction histories, past behaviors, or other individual characteristics to adapt the content, form, or presentation of an explanation [10, 17, 22]. These signals can characterize what a user prefers or values, but they do not necessarily describe how that user organizes information when interpreting it. Two users with similar preferences may, for instance, differ in the elements they mobilize, in the relative importance they assign to them, or in the pieces of information they consider jointly [8, 9]. Previous interpretations constitute a complementary source of information about the user. When a person expresses an interpretation of a situation, the resulting response may reveal certain properties of how information is mobilized: which elements are selected, their relative salience, and which elements are associated in the expressed interpretation [3]. A single interpretation may, however, depend strongly on the particular situation in which it is produced. When some of these properties recur across multiple interpretations from the same user, they may constitute interpretive regularities derived from that user's history. The value of these regularities lies in their potential for reuse. They may provide a signal complementary to preferences for personalizing the presentation of new information. In the context of explanations, this opens a particular possibility: personalizing not only the information that is presented, but also its organization according to how the user has previously mobilized information. Prior work on personalized explanations has examined adaptations of content, form, and presentation [10, 17, 22]. Other approaches also exploit user history to adapt generation or model behavior [11, 21, 23]. More directly, work on information ordering has shown that explanation organization can itself serve as a dimension of personalization when guided by priorities explicitly stated by the user [15].

To our knowledge, interpretive regularities derived from prior textual interpretations have not yet been exploited as a signal for personalizing explanation organization. We therefore ask the following central research question:

> **RQ.** Does organizing an explanation from interpretive regularities derived from the user's prior interpretations improve objective comprehension of the presented information compared with alternative organization procedures?

To investigate this question, we represent prior interpretations along three dimensions: information selection, relative salience, and relational organization, with the relational dimension partially operationalized through co-mobilization. These dimensions allow us to construct, for each user, a historical profile reflecting regularities observed across their prior interpretations and to use this profile to personalize the organization of new explanations. We evaluate this approach in a controlled study in which objective comprehension is the primary outcome, considered separately from recommendation acceptance and subjective judgments of the explanation [16, 24]. Our contributions are threefold:

(1) We show, in a controlled study with 122 participants, that organizing an explanation from interpretive regularities

derived from the user's prior interpretations improves objective comprehension compared with three alternative organization procedures.
(2) We use a fixed historical profile as an actionable signal for ordering explanation sentences while holding explanation content and wording constant, and compare the resulting procedure with another user's historical profile, a leave-one-participant-out population profile, and a target-aware LLM reference procedure.
(3) We evaluate recommendation acceptance separately and detect no reliable overall condition effect in the present study, thereby distinguishing the observed comprehension advantage from the acceptance outcome.

2 Related Work

2.1 Personalized Explanations: Content, Form, and Organization

Research on personalized explanations distinguishes several targets of adaptation. Schneider and Handali [22] identify complexity, decision information, and presentation, and distinguish these targets from the user information used to personalize them. This separates what a system knows about a user from the explanation properties that this knowledge may guide. In recommender systems, personalization can concern the arguments used to explain an item's relevance. Kouki et al. [10] generate personalized explanations from multiple sources, including user interactions and social connections, and compare explanation styles, combinations of styles, and textual or visual formats. Earlier work proposed considering users' interpretive schemas as a basis for contextualizing recommendation explanations [12, 14]. This conceptual perspective motivates attention to how users interpret explanatory information. Munduku and Negre [13] also examine user interpretive schemas in recommender system explanations. Their work on time-constrained explanations [15] isolates organization more directly by varying information order while holding explanation content and wording constant. Personalized ordering follows the importance explicitly assigned to information by the user, and its effects depend on the time available to consult the explanation. These contributions motivate attention to interpretation and information organization in personalized explanations. However, the dimension being adapted must be distinguished from the evidence used to adapt it. Explicitly stated priorities and regularities derived from earlier textual interpretations provide different kinds of information about the user. Evidence supporting organization based on the former does not establish the usefulness of organization based on the latter. Our study examines whether historical interpretations provide a useful basis for personalizing sentence order in a subsequent explanation.

2.2 From User History to Explanation Adaptation

Historical data can support personalization through different mechanisms. LaMP retrieves relevant entries from user profiles to condition classification and generation tasks [21], whereas OPPU combines user-specific parameter-efficient adaptation with retrieval and textual profiles [23]. For recommendation explanations, PEPLER uses personalized prompt learning associated with users and items [11]. Logic-Scaffolding selects relevant historical items, extracts their aspects, and generates explanations through guided reasoning [18]. ReasoningRec constructs profiles of preferences and dislikes and uses synthetic explanations to train a model to recommend and explain [2]. Other user properties can also guide adaptation: in human–robot interaction, Gebellí et al. [4] use knowledge–memory models to adjust explanation detail and refer to previously encountered concepts. The relevant distinction is therefore not between explicit profiles and LLM-based approaches, but between what a user representation captures and how it guides an explanation. The studies above rely on properties such as preferences, item aspects, and acquired knowledge. We instead examine information selection, relative salience, and relational organization in historical textual interpretations as inputs to an ordering procedure, using co-mobilization as a partial empirical operationalization of the relational dimension. We additionally compare this rule-based procedure with an LLM that receives the same structured historical profile, together with target-specific context, and reorders the same predefined explanation sentences.

2.3 Evaluating Personalized Explanations: Comprehension and Decisions

Explanation evaluations address objectives that are not interchangeable. Tintarev and Masthoff [24] distinguish, among others, effectiveness, satisfaction, transparency, persuasiveness, and efficiency. Their studies show that personalizing feature-based explanations can increase satisfaction while reducing measured decision effectiveness. Alignment with a user profile therefore does not guarantee benefits across these objectives. Presentation choices also involve trade-offs. Gedikli et al. [5] find that content-based tag-cloud explanations can improve perceived transparency and satisfaction while requiring greater cognitive effort. Nimmo et al. [16] examine actual understanding of a system providing explanations and find that few of the individual characteristics studied are associated with this outcome in their experimental setting. These findings motivate direct assessment of whether the properties selected for personalization provide the intended benefit to users. Generated-explanation studies provide complementary evidence: PEPLER evaluates textual quality and feature coverage [11], while Logic-Scaffolding includes human judgments of relevance, readability, and factuality [18]. Such measures characterize explanation outputs but do not directly establish their effects on users' comprehension. We consequently evaluate objective comprehension separately from effects on decisions. A change in choice or recommendation acceptance is not, by itself, evidence that an explanation has been better understood.

Positioning. The Interpretive Schema Manifestation Approach (ISMA) [1] characterizes expressed interpretations along three dimensions: information selection, relative salience, and relational organization. We use this representation to construct a fixed historical profile from prior interpretations. The present study does not introduce this framework, nor does it re-evaluate the existence, person-specificity, or cross-situational informativeness of these interpretive regularities. It instead addresses a distinct downstream question: whether a historical profile based on this representation

can serve as an actionable signal for organizing explanatory information. Specifically, we use the user's historical profile to determine the order of otherwise identical explanation sentences and test whether this organization improves objective comprehension relative to another user's historical profile, a leave-one-participant-out population profile, and a target-aware LLM procedure. Recommendation acceptance is examined separately.

3 History-Derived Explanation Organization

We use ISMA-derived representations as input to a history-based explanation-organization procedure. Historical interpretations are first represented and aggregated into a fixed historical profile before explanation evaluation. The procedure examined in the present study then maps this historical profile to the sentence order of a new explanation. For each target situation, the explanation consists of six predefined sentences. Their content, wording, and typography remain unchanged: only their order varies.

3.1 Historical Trace Representation

Historical traces are represented along the three ISMA dimensions of expressed interpretation: information selection, relative salience, and relational organization. Selection concerns which information is mobilized in a response; salience characterizes its relative prominence within that response; relational organization concerns how information units are associated. Each historical trace is a written response produced in a source situation presenting six information units. These units are aligned across situations through six functional roles: cost or price, access or travel, core attributes of the offering, experience conditions, availability, and usage or contractual conditions. Core attributes describe what the offering includes, provides, or enables, such as an activity's content or an event's programme. The role labels are adapted to the materials of the present study and are intended to preserve functional alignment rather than exact lexical identity with earlier operationalizations. This functional alignment provides common coordinates without assuming semantic equivalence across domains. Throughout this section, u denotes a user and t a source episode. The indices $i, j \in \{1, \ldots, 6\}$, with $i \neq j$, identify distinct functional roles and their corresponding information units within an episode. A hat denotes a quantity estimated from textual responses. For each historical trace, the extraction procedure produces an empirical representation

$$\widehat{\Gamma}_{u,t} = \left(\widehat{\alpha}_{u,t}, \widehat{\omega}_{u,t}, \widehat{\rho}_{u,t}\right).$$

These indicators characterize properties of expressed interpretations; they neither directly observe a latent interpretive schema nor fully reconstruct the meaning attributed by the user. Responses undergo Unicode NFKC normalization, whitespace normalization, and segmentation at sentence-ending punctuation and line breaks. Segments containing fewer than four words are discarded. Retained segments and presented information units are encoded using `paraphrase-multilingual-MiniLM-L12-v2` within Sentence Transformers [19, 20].

Selection. An information unit is identified as mobilized when its cosine similarity with at least one retained segment reaches the fixed threshold of 0.46. This threshold was determined before the present study in a separate development calibration [1] against human annotations, using cross-validation grouped by participant. The resulting configuration is applied here without recalibration on the present dataset. Failure to detect an information unit does not establish that the user did not consider it.

Relative salience. For each selected unit, we retain its maximum similarity to the retained segments. These values are divided by their sum over selected units; unselected units receive zero. This indicator is a proxy for textual salience, rather than explicitly declared importance, and depends on the selection procedure.

Co-mobilization. We use Qwen3-8B [25] to identify the information units mobilized in each retained segment. The model receives the segment and the six reference information units and returns the identifiers of the units detected. For two selected units in a historical trace containing retained segments, the empirical co-mobilization coordinate is the proportion of segments in which both units are detected:

$$\widehat{\rho}_{u,t,ij} = \frac{\left|\left\{s \in S_{u,t} : i, j \in U(s)\right\}\right|}{|S_{u,t}|}. \tag{1}$$

Here, $S_{u,t}$ contains the retained segment instances, and $U(s)$ contains the information-unit identifiers detected in segment s. The numerator counts segments in which both units are detected; the denominator counts all retained segments. This expression applies when both units are selected and $|S_{u,t}| > 0$. If at least one of the two units is unselected, the corresponding relational coordinate is left undefined. If the coordinate is defined and equals zero, both units were selected, but they were never jointly detected in any retained segment. The representation contains six selection coordinates, six salience coordinates, and fifteen possible relational coordinates, some of which may be undefined. The co-mobilization indicator is symmetric and untyped: it captures joint detection without specifying relation direction or semantic type, such as causality, opposition, or compensation. It therefore provides a partial empirical operationalization of relational organization. The extraction configuration was fixed before explanation evaluation and is not adjusted using comprehension, acceptance, or other evaluation outcomes. The present study evaluates the explanation-organization procedure constructed from the resulting historical profiles; it does not provide a new human-annotation validation of extraction accuracy on this corpus. Model versions, prompts, calibration details, extraction settings, and the handling of edge cases, including responses without retained segments and invalid model outputs, are reported in the supplementary material.

3.2 Historical Profile Construction

For each user, the four source episodes yield four episode-level representations. These representations are aggregated into a single historical profile, comprising selection, relative salience, and relational components, denoted by $\widehat{h}_u^{\alpha}$, $\widehat{h}_u^{\omega}$, and $\widehat{h}_u^{\rho}$, respectively.

Historical selection and salience. For selection and salience, each episode-level vector is first normalized by its Euclidean norm. The four normalized vectors are then averaged, and the resulting vector is normalized again:

$$\widehat{h}_u^z = \mathcal{N}_2\left(\frac{1}{4}\sum_{t\in T_u^{\text{hist}}} \mathcal{N}_2(\widehat{z}_{u,t})\right), \qquad z \in \{\alpha, \omega\}. \tag{2}$$

Here, T_u^{hist} contains the four source-episode indices for user u, and $\widehat{z}_{u,t}$ denotes the corresponding six-dimensional selection or salience vector. Normalization is defined as

$$\mathcal{N}_2(v) = \begin{cases} v/\|v\|_2, & \|v\|_2 > 0, \\ \mathbf{0}, & \|v\|_2 = 0. \end{cases} \tag{3}$$

Zero vectors are retained in the four-episode average. These historical components summarize the relative structure of the extracted information; their coordinates are not probabilities or explicitly reported importance ratings.

Historical co-mobilization. Each empirical relational coordinate is averaged over the source episodes in which it is defined:

$$\widehat{h}^{\rho}_{u,ij} = \frac{1}{|T^{\text{def}}_{u,ij}|}\sum_{t\in T^{\text{def}}_{u,ij}} \widehat{\rho}_{u,t,ij}, \qquad |T^{\text{def}}_{u,ij}| > 0. \tag{4}$$

The set $T^{\text{def}}_{u,ij}$ contains the source episodes in which the relational coordinate for roles i and j is defined. Each contributing episode has equal weight. Defined zero values contribute to the average, whereas undefined values are omitted. A coordinate that is undefined in all four source episodes remains undefined in the historical profile. No additional normalization is applied to the relational component. For each relational coordinate, we also record $m_{u,ij} = |T^{\text{def}}_{u,ij}|$, the number of episodes contributing a defined value, and $m^{+}_{u,ij}$, the number of these episodes with positive co-mobilization. These counts document historical support without changing the coordinate's value or its eligibility for the ordering procedure. A positive historical coordinate may therefore be supported by only one episode. The resulting historical profile summarizes regularities in the four observed interpretations; it is not assumed to represent an immutable psychological trait. Historical profiles are computed and fixed before explanation construction. Outcomes from the subsequent evaluation are not used to construct or update them.

3.3 From Historical Profiles to Explanation Order

Each sentence in the target explanation is associated with one of the six functional roles. The adaptation rule uses this correspondence to map the historical profile to a sentence order. All sentences are retained and appear exactly once, without added boldface, color, or boxes.

Initial ordering. Sentences are sorted by decreasing historical salience. Ties are broken by decreasing historical selection, followed by increasing role identifier, fixed in advance. Selection therefore acts as a secondary criterion; it does not have a role equivalent to salience in determining the order.

Pair selection. Among pairs with defined and strictly positive historical co-mobilization, we select the pair with the highest value. For each pair, the first member is the sentence appearing earlier in the initial order, and the second member is the sentence appearing later. Tied pairs are ranked by the initial position of their first member, with earlier positions preferred. If these positions are identical, the initial positions of their second members are compared in the same way.

Table 1: Illustrative transformation from the initial order to the final explanation order.

Stage	Sentence order
Initial order	Price → Core attributes → Experience conditions → Availability → Usage conditions → Access
Selected pair	Price–Access
Final order	Price → Access → Core attributes → Experience conditions → Availability → Usage conditions

Final ordering. The second member of the selected pair is placed immediately after the first. The relative order of the two members and that of the remaining four sentences are preserved. If the pair is already consecutive, no displacement occurs. If no eligible pair is available, the initial order is retained. Table 1 illustrates this transformation. The labels denote sentence roles; the sentences themselves are not rewritten.

In this example, access information moves from the sixth to the second position. Bringing the pair together therefore gives local priority to proximity between two information units, at the expense of modifying the order based on their individual salience. This choice constitutes a design hypothesis: presenting two historically co-mobilized information units consecutively may facilitate their joint consideration. It does not follow necessarily from ISMA, and its benefit remains to be evaluated. In particular, the selected pair may rely on a single positive observation; placing its members consecutively establishes neither a repeated association nor a particular semantic relation. For a given historical profile and the same six functional roles, the rule produces the same role order across target situations. It thus transfers an organization derived from historical traces without reassessing the semantic requirements of each target situation. The evaluation presented in Section 4 concerns this complete explanation-organization procedure and its effects on objective comprehension and recommendation acceptance. It does not, by itself, attribute any observed benefit to individual historical-profile components.

4 Experimental Study

This study investigates whether a historical profile derived from users' prior interpretations can guide explanation organization in new situations. The evaluation focuses on post-exposure objective comprehension and stated recommendation acceptance while holding explanation content constant and varying only sentence order. We distinguish *source situations*, used to collect written interpretations and construct historical profiles, from *target situations*, reserved for the evaluation of explanations organized from those profiles. Target situations were distinct from those used during profile construction.

4.1 Participants and Profile Construction

Profile-construction data were collected specifically for the present study between August 30 and September 21, 2026. A total of 135 participants began this phase. Of these, 122 provided the four written interpretations required to construct a historical profile, yielding 488 source episodes. Among these participants, 80 completed the tasks in French and 42 in English. 65% identified as men, 30% as women, and 5% as non-binary. Participants were recruited in university lecture halls, through social media, and within a doctoral-student discussion group. Participation was voluntary, unpaid, and unrelated to course assessment or academic grades. All participants were at least 18 years old and provided informed consent before taking part. No formal ethics-committee approval was obtained. Participants were identified only through pseudonymous study codes, and no directly identifying personal information was collected. Participants were informed about the purpose and procedure of the study before providing their responses. Individual written responses were not published, and the study codes were used only to link responses across the different phases of the study. The participant sample and data collection were specific to the present study and did not overlap with the participant sample used in the separate study underlying the ISMA reference [1]. Each participant completed four interpretation tasks in a counterbalanced order. Within each source situation, the presentation order of the six information units was also counterbalanced across participants. The situations concerned an outdoor activity, a meal-kit delivery service, a mobile-phone plan, and a home-cleaning service. For each task, participants read a context accompanied by six information units and then wrote an open-ended interpretation.

Participants were not asked to select information units, assign importance scores, or explicitly identify relationships between units. The resulting responses therefore constitute historical textual traces of expressed interpretation rather than direct reports of the structural dimensions later inferred from them. Each information unit corresponded to one of the six functional roles defined in Section 3.1. This alignment provides common coordinates across situations without assuming semantic equivalence across domains. Information selection and relative salience were estimated using the fixed extraction procedure described in Section 3.1; relational organization was partially operationalized through the extracted co-mobilization indicator. Processing the 488 responses involved 2,153 textual segments, with no parsing failures in the extraction outputs. The four source responses of each participant were aggregated according to Section 3.2, yielding 122 historical profiles. All profiles had nonzero selection and salience components. Historical profiles were computed and fixed before explanation construction. Evaluation outcomes, including comprehension responses, subjective judgments, and recommendation-acceptance responses, were not used to construct or update them. All 122 participants subsequently completed the explanation-evaluation phase between September 23 and September 30, 2026. The same study code linked each evaluation session to the corresponding historical profile.

Table 2: Target situations used for explanation evaluation.

ID	Decision context	Recommendation
T1	Watch a film in a new city	Screening of *Horizon 47*
T2	Choose a cultural activity	*Cartographies de lumière* exhibition
T3	Participate in a cooking workshop	*Saveurs du Levant*
T4	Attend a musical performance	Ensemble Nova concert

Table 3: Explanation-organization conditions.

Condition	Profile source	Ordering procedure
OWN	User	Fixed rule from Section 3.3
DONOR	Another participant, same language	Same rule as OWN
POP	Other 121 participants, aggregated	Same rule as OWN
PROFILE-LLM	User	Mistral with target context, recommendation name, and six sentences

4.2 Explanation Construction and Organization Conditions

Target situations and explanation content. Four target situations involving leisure and cultural activities were used for evaluation. They were distinct from the source situations used for profile construction and are summarized in Table 2.

Each target explanation contained six predefined sentences, one per functional role. Historical-profile coordinates were mapped to target sentences through these roles, relying on functional alignment rather than semantic identity across domains. Explanations were available in French and English and were presented in the language used by the participant during profile construction. For a given participant–target-situation combination, all four conditions used exactly the same six sentences. No information was added, removed, or reformulated; only sentence order varied. Typography, layout, and visual presentation were otherwise identical. Sentence lengths ranged from 9 to 16 words in French (median 13) and from 7 to 17 words in English (median 12). Detailed sentence lengths and the full bilingual materials will be provided in the supplementary material.

Organization conditions. The four conditions differed in the profile used and in the procedure determining sentence order, as summarized in Table 3.

OWN: user's historical profile. The deterministic ordering rule defined in Section 3.3 was applied to the user's historical profile. The resulting role order was then applied to the four target situations for that participant.

DONOR: another participant's historical profile. The same ordering procedure was applied using the historical profile of another participant from the same language group. Assignment was bijective and excluded self-assignment, so that each profile served

exactly once as a donor within its language group. Donor assignments were fixed across target situations and before evaluation outcomes were collected. Full assignment details and the random seed will be documented in the supplementary material.

POP: population profile. For each user, a leave-one-participant-out population profile was constructed from the other 121 participants. The historical-profile selection and salience components of the other 121 participants were averaged and L2-normalized; relational coordinates were averaged over participants for whom they were defined, retaining defined zeros and excluding undefined values. The OWN ordering rule was then applied to this population profile. POP therefore provided a population-level historical reference, whereas DONOR used another participant's historical profile.

PROFILE-LLM: target-aware ordering from the user's historical profile. Mistral-7B-Instruct-v0.3 [7] received the same structured historical profile used in OWN, together with the target context, recommendation name, and six predefined sentences. The prompt distinguished undefined relational coordinates from coordinates defined with a zero value. The model did not generate or rewrite sentence content; it returned only a permutation of identifiers 1–6, from which the explanation was assembled. All 488 PROFILE-LLM outputs were valid permutations. The complete prompt, model revision, and generation parameters will be provided in the supplementary material. Unlike OWN, DONOR, and POP, which produced a fixed role order for a given historical profile, PROFILE-LLM could produce target-specific orders because it additionally received target context and sentence content. The comparison therefore concerns complete explanation-organization procedures rather than an isolated effect of the language model.

Formative pretest. Twelve pretest participants, distinct from the 122 main-study participants, completed a formative pretest, with six participants per language. Each examined four explanations and rated clarity, coherence, and naturalness on three-point scales while also providing free-text comments. Mean ratings were 3.00 for clarity and 2.21 for both coherence and naturalness in each language. The pretest informed clarification of the third functional-role label and correction of one English sentence, after which PROFILE-LLM orders were regenerated using the finalized materials. Because the pretest concerned an earlier material version, it was treated as formative rather than as validation of the final explanation sequences. Detailed pretest materials will be provided in the supplementary material.

Checks on constructed explanations. A total of 1,952 candidate explanations were constructed (122 participants $\times$ 4 target situations $\times$ 4 conditions), of which 488 were actually presented during evaluation. All candidates were valid permutations and preserved identical sentence content across conditions for a given participant–target-situation combination. The four candidate orders were distinct in 478 of the 488 combinations; the remaining ten comprised five PROFILE-LLM–OWN and five PROFILE-LLM–DONOR matches, all retained. Across OWN profiles, the selection tie-breaker changed none of the 122 initial orders. The relational grouping step modified 28 of the 122 orders; in the remaining 94, the selected pair was already adjacent. Thus, historical salience accounted for most observed sentence-order structure, while the relational step modified a subset of cases. Detailed construction diagnostics will be provided in the supplementary material.

4.3 Experimental Design and Procedure

All 122 participants in the profile-construction sample completed the explanation-evaluation session between September 23 and September 30, 2026. The evaluation followed a within-participant design with four conditions: OWN, DONOR, POP, and PROFILE-LLM. Each participant encountered each target situation and each organization condition exactly once, yielding four trials per participant and one explanation version per target situation.

Assignment and counterbalancing. Four condition–situation assignments defined by a Latin square were crossed with four cyclic rotations of T1–T2–T3–T4, producing 16 experimental schedules. This distributed conditions across target situations and trial positions, although it did not balance every possible carryover sequence. Because each target situation used the same comprehension questions across organization conditions, this counterbalancing also distributed the fixed comprehension questions across conditions. Participant codes were randomly assigned to schedules within language groups subject to predefined allocation counts. For the 80 French-speaking participants, each schedule received five participants; among the 42 English-speaking participants, ten schedules received three participants and six received two. The allocation and randomization were fixed before evaluation responses were collected. Participants were not shown condition labels. Complete assignment schedules and randomization details will be documented in the supplementary material.

Procedure. Source-data collection, historical-profile construction, donor assignment, population-profile construction, and candidate-explanation generation were completed before evaluation responses were collected. Evaluation outcomes were not used to modify these objects. The explanation-evaluation phase took place between September 23 and September 30, 2026, in a separate session using the same study code to link responses to the corresponding historical profile. For each trial, participants read the target context, recommendation, and explanation in the same language as their source interpretation tasks. Reading was self-paced, with no imposed time limit. Explanation-display duration was automatically recorded from explanation onset to the click that opened the evaluation questions; it therefore measures exposure duration rather than necessarily active reading time. After the questions were opened, the explanation was removed and could not be revisited. Participants first reported recommendation acceptance, then answered the factual and integration comprehension questions, and finally completed the three subjective judgments described in Section 4.4. The procedure was repeated across the four trials, with no correctness feedback provided during the session.

4.4 Measures

Each trial collected six responses: one recommendation-acceptance response, two objective-comprehension responses, and three subjective judgments, yielding 24 questionnaire responses per participant across four trials. Explanation-display duration was recorded automatically and was not an additional questionnaire item.

Recommendation acceptance. Immediately after the explanation was removed, participants answered: "In the situation presented, would you accept this recommendation?" Response options were "No," "Uncertain," and "Yes," with uncertainty retained as a separate category. Acceptance was collected before the comprehension questions so that those questions would not first direct attention toward particular information. It represents stated acceptance within the experimental scenario and is not, by itself, interpreted as evidence of a better decision.

Post-exposure objective comprehension. Each target situation included two multiple-choice questions. The first assessed a factual detail relevant to the recommendation; the second required combining at least two pieces of information to identify a practical consequence or applicable condition. Each question had one correct answer and plausible distractors. For a given situation and language, questions and response options were identical across organization conditions. Response options were randomly permuted except when a meaningful numerical or temporal progression required a fixed order. Correct responses were coded as 1 and incorrect responses as 0; missing responses remained distinct from incorrect responses. A descriptive 0–2 trial score was also computed when both questions were answered. The primary objective outcome was question-level correctness across the eight objective-comprehension questions. Differences between factual and integration questions were examined in the secondary analysis described in Section 4.5. Because participants responded after the explanation had been removed, performance reflects both comprehension and short-term retention and does not identify a unique cognitive mechanism. The eight questions were piloted with the same twelve formative pretest participants, who were distinct from the main sample. The pilot assessed clarity, uniqueness of the correct answer, and plausibility of distractors. Because these participants had already encountered the explanations, it was not used to estimate first-exposure question difficulty. Final questions, response options, and answer keys were fixed before the main evaluation.

Subjective judgments. After the objective questions, participants responded in a fixed order to three items:

(1) **Perceived comprehension:** "I understood the information presented in this explanation."
(2) **Mental effort:** "Understanding this explanation required a lot of mental effort."
(3) **Perceived naturalness:** "The sequence of information in this explanation felt natural."

Each item used a three-point agreement scale:

1 = disagree, 2 = neither agree nor disagree, 3 = agree.

The three items were analyzed separately and treated as single-item subjective indicators rather than as a composite or validated psychometric scale.

Exposure duration and language versions. Explanation-display duration was treated as a complementary descriptive measure rather than a direct measure of attention or active reading. French and English formulations were examined by the six pretest participants in each corresponding language, and identified discrepancies were resolved before the main evaluation. Complete bilingual questionnaire materials, response options, and answer keys will be provided in the supplementary material.

4.5 Statistical Analysis

Question-level correctness on the objective-comprehension questions was the primary outcome. Recommendation acceptance and the three subjective judgments were secondary outcomes, while exposure duration was treated descriptively.

Primary analysis. Question-level correctness was analyzed using a logistic mixed-effects model with organization condition, question identity, language, and trial position as fixed effects. Trial position was treated as categorical, and question identity accounted for differences in difficulty among the eight objective questions. The initial model included random intercepts for participants and for trials nested within participants. The trial-level variance was estimated at the boundary ($< 10^{-12}$), resulting in a singular fit. The trial-level random intercept was therefore removed. The final model retained a participant random intercept and was non-singular. Three contrasts of primary interest compared OWN with DONOR, POP, and PROFILE-LLM. Tests were two-sided and Holm-adjusted [6] across these three contrasts to control the familywise error rate at 0.05. Results are reported as odds ratios, 95% confidence intervals, Holm-adjusted p-values, and estimated marginal probabilities on the response scale, averaged over question identity, language, and trial position. The descriptive 0–2 trial score did not replace the question-level analysis.

Question-type analysis. A secondary analysis tested whether condition effects differed between factual and integration questions by adding a condition-by-question-type interaction while retaining question identity. A separate main effect of question type was not interpreted because question type is determined by item identity. Differences between condition effects were evaluated through the interaction rather than by comparing significance patterns across question types. The interaction was evaluated using a likelihood-ratio test.

Recommendation acceptance. The three-category acceptance outcome ("No," "Uncertain," "Yes") was analyzed using a multinomial mixed-effects framework with organization condition, target situation, language, and trial position as fixed effects and participant as a random intercept. The overall effect of organization condition was evaluated using a joint Wald test of the six condition coefficients across the two non-reference outcome categories.

Subjective outcomes. Perceived comprehension, mental effort, and perceived naturalness were analyzed separately using cumulative-logit mixed-effects models with organization condition, target situation, language, and trial position as fixed effects and participant

as a random intercept. For each outcome, the overall effect of organization condition was evaluated using a likelihood-ratio test comparing models with and without the condition term.

Multiplicity for secondary inference. One omnibus condition-by-question-type interaction test and four omnibus condition tests for recommendation acceptance and the three subjective outcomes formed a family of five secondary tests. Holm's method [6] controlled the familywise error rate across this family. Additional pairwise contrasts or subgroup analyses were treated as exploratory.

Exposure duration. Exposure duration was summarized by condition using medians and interquartile ranges. It was not included as a covariate in the primary model because explanation organization could itself affect duration; any additional duration analyses were therefore treated as exploratory.

Incomplete responses and exclusions. Available valid objective responses were retained, and missing responses were not recoded as incorrect. Exclusions were limited to documented technical failures preventing reliable delivery or response recording, records that could not be linked to a valid participant and trial, and duplicate records. For duplicate submissions to the same question, the earliest valid response was retained. Low correctness, uncertain acceptance, and unusually short or long exposure duration were not exclusion criteria, and objective questions were not treated as attention checks. Missingness, if present, was described by condition, language, and trial position. Because all 122 participants completed all four evaluation trials, a separate participant-level complete-case analysis did not define a different sample.

Sample size. The evaluation sample comprised the 122 participants with complete historical profiles and was not determined by a prospective power analysis. Inference therefore relies on estimated effect sizes, 95% confidence intervals, and the mixed-effects analyses above rather than post-hoc observed power.

Scope of inference. The analyses concern the complete explanation-organization procedures rather than the independent causal contribution of information selection, relative salience, or relational organization. Comparisons involving PROFILE-LLM likewise concern its complete procedure, including access to target-context information, rather than an isolated language-model effect. Differences between factual and integration performance may inform interpretation but do not identify a unique mechanism through which personalization operates. Conclusions are therefore limited to the participants, situations, questions, languages, and organization procedures represented in this study.

5 Results

All 122 participants completed the four evaluation trials, yielding 488 trials and 976 objective-comprehension responses. No responses were missing, and no trial was excluded under the exclusion criteria described above.

RQ: Does organizing an explanation from interpretive regularities derived from the user's prior interpretations improve objective comprehension of the presented information compared with alternative organization procedures?

Table 4: Primary objective-comprehension results. Odds ratios compare OWN with each alternative organization procedure.

Condition / Contrast	Accuracy	OR [95% CI]	p_{Holm}
OWN	86.5%	–	–
DONOR	66.0%	–	–
POP	63.9%	–	–
PROFILE-LLM	73.0%	–	–
OWN vs. DONOR	–	3.82 [2.36, 6.18]	< 0.001
OWN vs. POP	–	4.23 [2.61, 6.84]	< 0.001
OWN vs. PROFILE-LLM	–	2.67 [1.64, 4.36]	< 0.001

Table 5: Objective-comprehension accuracy by question type and organization condition. Values are raw proportions.

Condition	Factual	Integration
OWN	89.3%	83.6%
DONOR	72.1%	59.8%
POP	68.9%	59.0%
PROFILE-LLM	81.1%	64.8%

We report the primary objective-comprehension results first, followed by question-type differences, secondary outcomes, exposure duration, and exploratory consistency across contexts.

5.1 Objective Comprehension

Objective-comprehension accuracy was highest in the OWN condition. Participants answered 211 of 244 questions correctly in OWN (86.5%), compared with 161 of 244 in DONOR (66.0%), 156 of 244 in POP (63.9%), and 178 of 244 in PROFILE-LLM (73.0%). The primary mixed-effects logistic model supported these comparisons (Table 4). Compared with DONOR, OWN was associated with 3.82 times higher odds of a correct response (95% CI [2.36, 6.18], $p_{\text{Holm}} < 0.001$). The corresponding odds ratios were 4.23 relative to POP (95% CI [2.61, 6.84], $p_{\text{Holm}} < 0.001$) and 2.67 relative to PROFILE-LLM (95% CI [1.64, 4.36], $p_{\text{Holm}} < 0.001$).

Figure 1 shows the corresponding model-predicted probabilities of a correct response with 95% confidence intervals. The predicted probability was highest for OWN (89.1%), followed by PROFILE-LLM (75.3%), DONOR (68.2%), and POP (65.9%).

Across all three primary contrasts, organizing the explanation from the user's historical profile resulted in higher objective comprehension of the presented information than each of the alternative organization procedures tested.

5.2 Factual and Integration Questions

OWN showed the highest descriptive accuracy for both question types. For factual questions, accuracy was 89.3% in OWN, compared with 72.1% in DONOR, 68.9% in POP, and 81.1% in PROFILE-LLM. For integration questions, accuracy was 83.6% in OWN, compared with 59.8% in DONOR, 59.0% in POP, and 64.8% in PROFILE-LLM (Table 5).

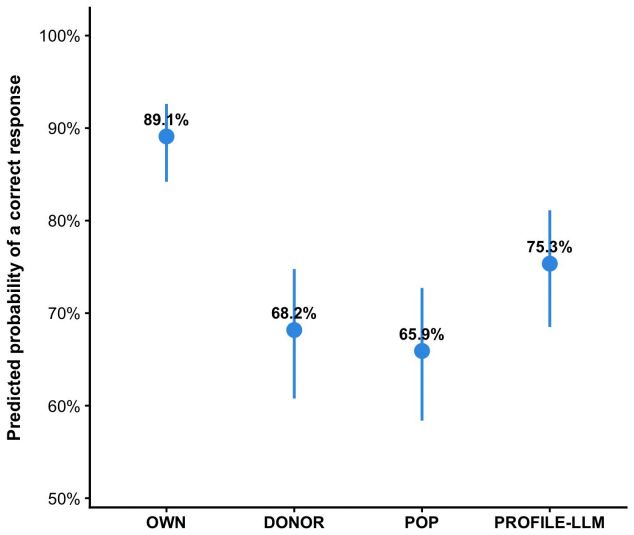

Figure 1: Estimated marginal probability of a correct objective-comprehension response by organization condition from the primary mixed-effects logistic model, averaged over question identity, language, and trial position. Vertical lines indicate 95% confidence intervals.

A likelihood-ratio test found no evidence that the effect of organization condition differed by question type, $\chi^2(3) = 1.285$, $p = 0.733$, $p_{\text{Holm}} = 1.000$. Thus, although OWN had the highest descriptive accuracy for both factual and integration questions, the data did not indicate a differential condition effect across the two question types.

5.3 Subjective Judgments and Recommendation Acceptance

Subjective judgments differed descriptively across organization conditions (Table 6). For perceived comprehension, 63.1% of participants selected the highest response category in OWN, compared with 32.8% in DONOR, 34.4% in POP, and 44.3% in PROFILE-LLM. A likelihood-ratio test indicated an overall condition effect, $\chi^2(3) = 29.27$, $p < 0.001$, $p_{\text{Holm}} < 0.001$. For perceived naturalness, the corresponding percentages were 54.9% in OWN, 27.9% in DONOR, 36.1% in POP, and 37.7% in PROFILE-LLM. A likelihood-ratio test also indicated an overall condition effect, $\chi^2(3) = 29.29$, $p < 0.001$, $p_{\text{Holm}} < 0.001$. For mental effort, the highest response category was selected by 23.0% of participants in OWN and 36.1% in each of DONOR, POP, and PROFILE-LLM. No reliable overall condition effect was established after multiplicity correction (likelihood-ratio test, $\chi^2(3) = 7.41$, $p = 0.060$, $p_{\text{Holm}} = 0.180$).

Recommendation acceptance showed a different descriptive pattern: 39.3% of responses were Yes in OWN, 41.0% in DONOR, 45.9% in POP, and 45.9% in PROFILE-LLM. A joint Wald test found no evidence of an overall condition effect on recommendation acceptance, $\chi^2(6) = 3.11$, $p = 0.795$, $p_{\text{Holm}} = 1.000$. Overall, OWN showed the highest descriptive percentages for perceived comprehension and perceived naturalness, for which omnibus condition effects were

Table 6: Secondary outcomes by organization condition. Subjective values report the percentage selecting response category 3; acceptance reports the percentage selecting Yes.

Condition	Compr.	Effort	Natural.	Accept. Yes
OWN	63.1%	23.0%	54.9%	39.3%
DONOR	32.8%	36.1%	27.9%	41.0%
POP	34.4%	36.1%	36.1%	45.9%
PROFILE-LLM	44.3%	36.1%	37.7%	45.9%

Note. Compr. = perceived comprehension; Natural. = perceived naturalness; Accept. Yes = percentage of Yes responses.

significant. No reliable overall condition effect was established for mental effort after multiplicity correction or for recommendation acceptance.

5.4 Exposure Duration

Median exposure duration ranged from 38.2 to 42.2 s across conditions. It was 38.2 s in OWN (IQR [28.1, 50.3]), 39.8 s in DONOR (IQR [29.2, 51.7]), 42.2 s in POP (IQR [30.8, 51.9]), and 40.6 s in PROFILE-LLM (IQR [28.8, 53.8]). Exposure duration is reported descriptively only and was not included as a primary covariate because it may itself be affected by explanation organization.

5.5 Exploratory Consistency Across Contexts

We descriptively examined whether the overall pattern observed in the primary analysis was also visible across languages, target situations, and trial positions (Table 7). Across languages, OWN had the highest objective-comprehension accuracy in both English (85.7%) and French (86.9%). The corresponding PROFILE-LLM accuracies were 70.2% and 74.4%, respectively. Across target situations, OWN had the highest accuracy in T1 (91.7%), T2 (89.7%), and T3 (91.9%). The pattern was not uniform: in T4, PROFILE-LLM had the highest accuracy (78.3%), compared with 73.4% for OWN. Across trial positions, OWN had the highest accuracy in positions 1 (84.4%), 2 (87.9%), and 4 (91.9%). At position 3, PROFILE-LLM was slightly higher than OWN (82.3% vs. 81.7%).

Overall, the exploratory analyses were broadly consistent with the primary pattern while also revealing context-specific exceptions.

6 Discussion

Historical profiles as an organizing signal. Explanations organized from the user's own historical profile yielded higher objective comprehension than the three alternative organization procedures. Because explanation content and wording were held constant, the observed advantage concerns information organization rather than the provision of additional explanatory content. The findings therefore support the use of historical profiles derived from prior interpretations as an actionable signal for organizing subsequent explanations.

What the procedure comparisons establish. The comparisons with DONOR and POP indicate an advantage for using the user's own historical profile rather than another participant's profile or a leave-one-participant-out population profile. OWN also outperformed the

Table 7: Exploratory objective-comprehension accuracy across languages, target situations, and trial positions. Values are raw proportions.

Language		
Condition	**EN**	**FR**
OWN	85.7%	86.9%
DONOR	64.3%	66.9%
POP	64.3%	63.7%
PROFILE-LLM	70.2%	74.4%

Target situation				
Condition	**T1**	**T2**	**T3**	**T4**
OWN	91.7%	89.7%	91.9%	73.4%
DONOR	65.6%	71.7%	67.2%	59.7%
POP	62.9%	56.2%	75.0%	62.1%
PROFILE-LLM	74.1%	67.7%	71.9%	78.3%

Trial position				
Condition	**1**	**2**	**3**	**4**
OWN	84.4%	87.9%	81.7%	91.9%
DONOR	69.4%	57.8%	65.5%	71.7%
POP	65.0%	59.7%	54.7%	77.6%
PROFILE-LLM	75.9%	65.0%	82.3%	68.8%

target-aware PROFILE-LLM procedure tested here, although both used the user's historical profile. This should not be interpreted as a general superiority of deterministic rules over language models, since the two are complete organization procedures that use profile and target information differently. Likewise, the design does not isolate the independent contributions of information selection, relative salience, and relational organization. Because historical salience accounted for most of the observed sentence-order structure, the comprehension advantage should not be interpreted as evidence that the three profile components contributed equally.

Objective comprehension, subjective experience, and acceptance. Perceived comprehension and perceived naturalness showed reliable condition effects, whereas no reliable overall condition effect was established for mental effort after multiplicity correction or for recommendation acceptance. The absence of a condition-by-question-type interaction also provided no evidence that the comprehension advantage differed between factual and integration questions. These secondary outcomes delimit the primary result and do not identify a specific cognitive mechanism.

Contextual variation. The OWN advantage was not uniform across all exploratory subsets: PROFILE-LLM was descriptively higher in T4 and slightly higher at trial position 3. Because the study was not designed to identify the causes of these reversals, a post-hoc mechanistic explanation would be unwarranted. These exceptions instead motivate examining when historical-profile-based organization is most beneficial and when target-adaptive procedures may be preferable.

Relation to prior work. Prior work has personalized explanatory content, presentation, and organization [10, 22]. Explanation ordering under time constraints has also been guided by explicitly elicited user importance [15]. The present study differs by deriving the organizing signal from prior open-ended interpretations and by evaluating objective comprehension rather than recommendation acceptance as the primary outcome.

Implications. The findings suggest that user representations may capture not only what information a person prefers or selects, but also regularities in how information is emphasized and organized across prior interpretations. Such representations can inform presentation while leaving the underlying content unchanged. Explanation organization may therefore serve as a personalization dimension under the controlled procedures studied here, although its generalization to richer adaptive interfaces and longer-term use remains to be established.

6.1 Limitations

Several limitations delimit the scope of these findings. Participants were recruited through convenience channels, and the study was limited to French and English, four controlled target situations, six predefined functional roles, and fixed sentence-level explanations. Of the 135 participants who began profile construction, 122 provided the four interpretations required for evaluation, so the analytical sample also reflects completion of the profile-construction phase. Each historical profile was derived from only four prior interpretations, and the automatic extraction procedure, although fixed before evaluation, was not independently revalidated on the present corpus. OWN should also be interpreted as a complete explanation-organization procedure: the design does not isolate the independent contributions of information selection, relative salience, and relational organization. Moreover, historical salience accounted for most of the observed sentence-order structure, while the relational grouping step modified only a subset of orders. The observed advantage therefore should not be interpreted as evidence that all profile components contributed equally. Relational organization was only partially operationalized through co-mobilization. This indicator is symmetric and untyped and does not identify semantic, causal, directional, or otherwise typed relations. A positive historical co-mobilization coordinate may also be supported by only one source episode. PROFILE-LLM represents one specific model and prompting procedure and also conditions its ordering on target-specific information. Its comparison with OWN therefore does not establish a general advantage over language-model-based personalization. The study also did not include a direct preference-based personalization condition; therefore, it does not establish whether historical profiles provide benefits beyond personalization based on explicitly elicited preferences. Finally, objective comprehension was assessed through eight multiple-choice questions after the explanation had been removed, so performance reflects both comprehension and short-term retention and does not identify a unique cognitive mechanism. The subjective outcomes were single-item, three-point task-specific indicators rather than validated psychometric scales. Broader and longer-term studies are needed to assess generalization to less controlled information-interaction settings.

7 Conclusion

This study examined whether historical profiles derived from prior interpretations can guide the organization of personalized explanations. Explanations organized from the user's own historical profile yielded higher post-exposure objective comprehension than DONOR, POP, and PROFILE-LLM while explanation content and wording remained fixed. Because comprehension was assessed after the explanation had been removed, performance also reflects short-term retention. These findings indicate that historical profiles can provide an actionable signal for explanation organization. Higher objective comprehension did not coincide with a detectable overall condition effect on recommendation acceptance. More broadly, the results suggest that personalization can operate not only through content selection but also through information organization. Future work should examine this approach in broader populations, richer explanation settings, and longer-term deployments.

References

[1] Anonymous. 2027. Reference withheld for double-blind review.

[2] Millennium Bismay, Xiangjue Dong, and James Caverlee. 2025. ReasoningRec: Bridging Personalized Recommendations and Human-Interpretable Explanations through LLM Reasoning. In *Findings of the Association for Computational Linguistics: NAACL 2025*, Luis Chiruzzo, Alan Ritter, and Lu Wang (Eds.). Association for Computational Linguistics, Albuquerque, New Mexico, 8147–8163. doi:10.18653/v1/2025.findings-naacl.454

[3] Kathleen Carley and Michael Palmquist. 1992. Extracting, Representing, and Analyzing Mental Models. *Social Forces* 70, 3 (1992), 601–636. doi:10.1093/sf/70.3.601

[4] Ferran Gebellí, Anaís Garrell, Jan-Gerrit Habekost, Séverin Lemaignan, Stefan Wermter, and Raquel Ros. 2025. Personalised Explanations in Long-term Human-Robot Interactions. In *2025 34th IEEE International Conference on Robot and Human Interactive Communication (RO-MAN)*. IEEE, Eindhoven, Netherlands, 775–782. doi:10.1109/RO-MAN63969.2025.11217900

[5] Fatih Gedikli, Dietmar Jannach, and Mouzhi Ge. 2014. How should I explain? A comparison of different explanation types for recommender systems. *International Journal of Human-Computer Studies* 72, 4 (2014), 367–382. doi:10.1016/j.ijhcs.2013.12.007

[6] Sture Holm. 1979. A Simple Sequentially Rejective Multiple Test Procedure. *Scandinavian Journal of Statistics* 6, 2 (1979), 65–70.

[7] Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. 2023. Mistral 7B. *CoRR* abs/2310.06825 (2023). doi:10.48550/arXiv.2310.06825

[8] Philip N. Johnson-Laird. 2010. Mental models and human reasoning. *Proceedings of the National Academy of Sciences* 107, 43 (2010), 18243–18250. doi:10.1073/pnas.1012933107

[9] Walter Kintsch. 1988. The Role of Knowledge in Discourse Comprehension: A Construction-Integration Model. *Psychological Review* 95, 2 (1988), 163–182. doi:10.1037/0033-295X.95.2.163

[10] Pigi Kouki, James Schaffer, Jay Pujara, John O'Donovan, and Lise Getoor. 2019. Personalized Explanations for Hybrid Recommender Systems. In *Proceedings of the 24th International Conference on Intelligent User Interfaces* (Marina del Rey, California, USA) *(IUI '19)*. Association for Computing Machinery, New York, NY, USA, 379–390. doi:10.1145/3301275.3302306

[11] Lei Li, Yongfeng Zhang, and Li Chen. 2023. Personalized Prompt Learning for Explainable Recommendation. *ACM Transactions on Information Systems* 41, 4, Article 103 (2023), 26 pages.

[12] Deo Munduku. 2025. Contextualisation des explications dans les systèmes de recommandation : une approche basée sur les schémas d'interprétation des utilisateurs. Congrès INFORSID 2025 – Forum Jeunes Chercheurs Jeunes Chercheuses (JCJC). Communication, Pau, France.

[13] Deo Munduku and Elsa Negre. 2026. Bridging the Interpretive Gap in Recommender System Explanations through User Interpretive Schemas. *Procedia Computer Science* (2026). https://www.sciencedirect.com/science/article/pii/S1877050926037865

[14] Deo Munduku and Elsa Negre. 2026. Contextualisation des explications dans les systèmes de recommandation : une approche basée sur les schémas d'interprétation des utilisateurs. *Revue ouverte d'ingénierie des systèmes d'information* 6 (2026), 112–122. Numéro spécial INFORSID 2025. doi:10.21494/ISTE.OP.2026.1451

[15] Deo Munduku and Elsa Negre. 2026. When Attention Is Bounded, Structure Matters: Personalizing Recommendation Explanations Under Time Constraints. In *Proceedings of the 20th ACM Conference on Recommender Systems (RecSys '26)*. Association for Computing Machinery. doi:10.1145/3773078.3831820

[16] Robert Nimmo, Marios Constantinides, Ke Zhou, Daniele Quercia, and Simone Stumpf. 2024. User Characteristics in Explainable AI: The Rabbit Hole of Personalization?. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. Association for Computing Machinery. doi:10.1145/3613904.3642352

[17] Ingrid Nunes and Dietmar Jannach. 2017. A Systematic Review and Taxonomy of Explanations in Decision Support and Recommender Systems. *User Modeling and User-Adapted Interaction* 27, 3–5 (2017), 393–444. doi:10.1007/s11257-017-9195-0

[18] Behnam Rahdari, Hao Ding, Ziwei Fan, Yifei Ma, Zhuotong Chen, Anoop Deoras, and Branislav Kveton. 2024. Logic-Scaffolding: Personalized Aspect-Instructed Recommendation Explanation Generation using LLMs. In *Proceedings of the 17th ACM International Conference on Web Search and Data Mining (WSDM '24)*. Association for Computing Machinery, New York, NY, USA, 1078–1081. doi:10.1145/3616855.3635689

[19] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. Association for Computational Linguistics, Hong Kong, China, 3982–3992. doi:10.18653/v1/D19-1410

[20] Nils Reimers and Iryna Gurevych. 2020. Making Monolingual Sentence Embeddings Multilingual using Knowledge Distillation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 4512–4525. doi:10.18653/v1/2020.emnlp-main.365

[21] Alireza Salemi, Sheshera Mysore, Michael Bendersky, and Hamed Zamani. 2024. LaMP: When Large Language Models Meet Personalization. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 7370–7392. doi:10.18653/v1/2024.acl-long.399

[22] Johannes Schneider and Joshua Peter Handali. 2019. Personalized Explanation for Machine Learning: A Conceptualization. In *Proceedings of the 27th European Conference on Information Systems (ECIS 2019)*. Stockholm and Uppsala, Sweden.

[23] Zhaoxuan Tan, Qingkai Zeng, Yijun Tian, Zheyuan Liu, Bing Yin, and Meng Jiang. 2024. Democratizing Large Language Models via Personalized Parameter-Efficient Fine-tuning. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (Eds.). Association for Computational Linguistics, Miami, Florida, USA, 6476–6491. doi:10.18653/v1/2024.emnlp-main.372

[24] Nava Tintarev and Judith Masthoff. 2012. Evaluating the Effectiveness of Explanations for Recommender Systems: Methodological Issues and Empirical Studies on the Impact of Personalization. *User Modeling and User-Adapted Interaction* 22, 4-5 (2012).

[25] An Yang et al. 2025. Qwen3 Technical Report. *arXiv preprint arXiv:2505.09388* (2025).