arXiv:2611.08946v1 [cs.AI] 6 Nov 2026

Flex-RSI

Technical Report

Flex-RSI: Benchmarking Recursive Self-Improvement of Frontier Agents

Flex-RSI Team

Texas A&M University · University of Cambridge · National University of Singapore · NVIDIA · Massachusetts Institute of Technology · MIT-IBM Watson AI Lab · The University of Texas at Austin · Carnegie Mellon University · University of Pennsylvania · University of Southern California · Purdue University · Adobe · University of Toronto

Frontier language models are increasingly deployed as agents, and how well an agent performs depends not only on its model but also on its harness: the prompts, tools, memory, and control loop through which the model acts. When an agent is applied to a new problem, a natural question is whether it can use the feedback of the task to improve the way it works, and not merely its answer. Most existing benchmarks keep the agent's harness fixed and score the solution it produces, or target large-scale model training. We introduce Flex-RSI, a benchmark for the recursive self-improvement of agents on vision and robotics tasks under a fixed budget. In each round, an executor improves the task solution, which is kept only if it scores higher on held-out validation data, and an improver then revises the harness, including the executor and itself. Every run starts from the same general-purpose harness, and the final solution is scored on hidden test data. Evaluating closed-source and open-weight models on tasks spanning perception, reasoning, robot action, 3D reconstruction, visual creation, and infrastructure, we find that self-improvement helps most when the agent builds its own solution, that feedback recording only success or failure holds back open-weight models, and that open-weight models fall furthest behind where outputs must pass strict checks of physical stability or exact correctness. Paired runs further show that editing the harness itself contributes to the gains.

Project page: https://flex-rsi.github.io
Authors: The full list of authors and their contributions is in Section 7

1 Introduction

Frontier large language models now routinely act as agents that write code, conduct experiments, and read the results over multiple steps (Anthropic, 2026a; OpenAI, 2026b; Google DeepMind, 2026; Qwen Team, 2026b; DeepSeek-AI, 2026; GLM-5 Team, 2026). An agent's performance depends not only on the underlying model but also on its *harness*: the prompts, tools, memory, and control loop through which the model acts. This raises a practical question for anyone who deploys an agent on a new problem: given only the task, can the agent use the task's feedback to improve the way it works under a fixed budget, and not merely its answer?

Most existing benchmarks run each model inside a fixed harness, often the vendor's own, and score only the solution or artifact it produces, mostly on machine-learning research tasks with a continuous metric (Chan et al., 2025; Wijk et al., 2024); those that also let the agent revise its harness (RSI-Exam Team, 2026) or that center on upstream model training at scale (OpenRSI Foundation, 2026) remain within machine-learning research. We introduce Flex-RSI to measure how well a frontier agent improves both its solution and its own harness on general vision and robotics tasks, a setting in which agents are increasingly deployed, each attempt consumes real compute, and the goal is a better solution within hours to days.

In Flex-RSI, the harness consists of two components that alternate over rounds (Figure 1a). The *executor* improves the task solution: it modifies the solution, runs experiments, and evaluates them on the labeled development data, and its candidate is kept only if it also scores higher on held-out validation data. The *improver* then reviews the round and modifies the harness itself, including the executor's prompts, tools, and

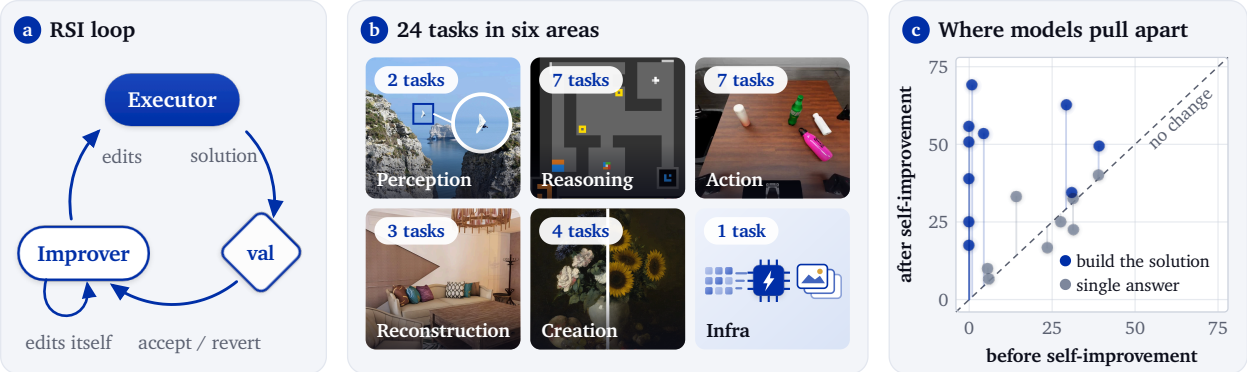

Figure 1 **Flex-RSI at a glance. (a)** The RSI loop. In each round, the executor edits the task solution, which is kept or reverted according to its score on held-out validation data; the improver then edits the harness, including the executor and itself. **(b)** The 24 tasks fall into six areas. Each tile shows an example task from one area, and its badge gives the number of tasks in that area. **(c)** Each dot is a task, placed by the score difference between the best and the worst model before (horizontal) and after (vertical) self-improvement; its height above the diagonal shows how much further apart the models end up. Blue: tasks in which the agent builds the solution. Gray: tasks whose solution is a single answer.

workflow as well as its own prompts, settings, and lessons. Every run starts from the same general-purpose harness and, on each task, the same starter solution.

The tasks differ in what the agent must produce (a single answer or a solution it builds), how the evaluator scores it (on a scale or only as success or failure), and how strictly the output is checked. From experiments on 24 tasks spanning perception, reasoning, robot action, 3D reconstruction, visual creation, and infrastructure (Figure 1b), with 18 closed-source and open-weight models, we find:

1. **Tasks in which the agent builds the solution leave more room for self-improvement:** tasks whose solution is a single answer to each question leave little room to improve, while tasks in which the agent builds its own solution leave the most, and there models pull apart (Figure 1c).
2. **The granularity of feedback matters:** feedback that scores each episode on a scale helps all models, whereas feedback that records only success or failure leaves open-weight models barely improving on the solutions they build.
3. **Strict checks of physical stability and exact correctness pose the greatest challenge:** outputs that look plausible often fail checks of physical stability or exact correctness, and open-weight models fall furthest behind on them.

Paired runs further show that letting the improver edit the harness raises the test score in six of seven pairs of otherwise identical runs (Section 5.4).

2 Related work

2.1 Benchmarks for self-improving agents

Benchmarks for machine-learning engineering and scientific research measure whether agents can turn a task description into a working result. MLE-bench evaluates machine-learning engineering on competition tasks (Chan et al., 2025); RE-Bench compares agents with human experts on research-engineering problems (Wijk et al., 2024); and PaperBench assesses the replication of machine-learning research (Starace et al., 2025). Related evaluations examine the design and execution of machine-learning experiments (Huang et al., 2023b) and the generation of executable programs for data-driven scientific discovery (Chen et al., 2025). Their primary evaluation targets are the models, code, and research results that an agent produces.

Evaluations framed around recursive self-improvement cover several different objects of improvement. RSI-Exam supports executable research on both task methods and harnesses for frozen models, followed by hidden-test evaluation (RSI-Exam Team, 2026). OpenRSI Index spans foundation-model development, including large-scale training and lighter-weight inference, agent, and systems tasks (OpenRSI Foundation, 2026). Autonomous post-training is evaluated through controlled training runs and assessments of the resulting models (Rank et al., 2026; Huang et al., 2026c), while algorithm-design tasks ask agents to improve training procedures that are subsequently evaluated by retraining (Chi et al., 2026). Sustained machine-learning research with repeated validation feedback and a separate hidden test is evaluated in AutoResearchExam (Ramesh et al., 2026). Harness engineering also appears alongside model-training tasks in Vals RSI Index (Chen et al., 2026b). TraceTrial-Bench evaluates the reliability of agent harnesses themselves rather than their improvement (Cao et al., 2026). These settings share an interest in autonomous experimentation, but differ in whether the optimized artifact is data, a training algorithm, model weights, or the harness around a model.

Complementary studies examine how experience becomes reusable improvement. Data-centric research under a fixed post-training stack is studied in RSIBench-Data (Meng et al., 2026). Procedural-ability transfer and the conversion of episodic experience into reusable skills are evaluated in EvoAgentBench (Gao et al., 2026b) and SkillEvolBench (Lei et al., 2026), respectively. EvolveTool-Bench checks whether evolved tools satisfy independent conformance tests beyond completing the current task (Kaliyev and Maryanskyy, 2026). The contribution of persistent experience to later personal-agent behavior is isolated through matched persistence-on/off comparisons in PAST-Bench (Xue et al., 2026). Whether self-evolving language-model agents learn new tasks without forgetting earlier ones is evaluated in a related study (Chen et al., 2026e). EmbodiedCLBench evaluates whether embodied agents can reuse the experience gained on basic tasks to solve unseen compositions of them (Huang et al., 2026b), and Embodied Skill Bench measures how packaged embodied skills help agents on physical tasks (Han et al., 2026c). Together, these studies motivate evaluating both immediate task gains and improvements that remain useful on subsequent tasks.

Flex-RSI studies task-solution and harness improvement together on vision, robotics, and related downstream tasks. For each task, every model starts from the same general-purpose harness and starter solution, then alternates execution and harness revision under a fixed budget. The improver can edit the executor's prompts, tools, and workflow as well as its own prompts, settings, and lessons. This shared protocol separates three questions: how much the task solution improves, what additional gains come from editing the harness, and whether a harness built by one model helps another model on the same task. A review of the evidence behind recursive self-improvement claims finds that few studies compare improvement procedures from identical starting artifacts at a matched budget (Chen et al., 2026g); the paired runs with and without the improver (Section 5.4) make this comparison.

2.2 Frontier agents

Language-model agents connect reasoning to actions and environmental feedback through an execution loop. ReAct interleaves reasoning and tool use (Yao et al., 2023), Reflexion carries verbal feedback into subsequent trials through episodic memory (Shinn et al., 2023), and Voyager accumulates executable skills through open-ended interaction in Minecraft (Wang et al., 2023). Coding interfaces such as Claude Code and Codex CLI expose file editing and command execution (Anthropic, 2026c; OpenAI, 2026d). OpenHands provides infrastructure for software-development agents (Wang et al., 2025b), while SWE-agent studies how an agent–computer interface supports repository navigation, editing, and testing (Yang et al., 2024). Research agents extend these capabilities to experimentation: AIDE searches over candidate solution code (Jiang et al., 2025), and The AI Scientist combines research ideation, experiments, paper writing, and review in an automated workflow (Lu et al., 2024). Frontier models have also been evaluated as robot agents that act in the control loop (Huang et al., 2026a; Zhou et al., 2026b; Yao et al., 2026c) or write reusable robot tools (Ge et al., 2026), across many embodiments (Yang et al., 2026b) and against learned policies (Zhang et al., 2026f), and coding agents have been asked to reconstruct dynamic scenes as executable graphics programs (Shen et al., 2026a).

Automated agent design makes parts of this execution machinery optimization variables. Promptbreeder evolves both task prompts and the mutation prompts that revise them (Fernando et al., 2023), while GEPA uses

reflection on execution traces to propose and select prompt updates (Agrawal et al., 2026). ADAS searches over agent programs using an archive of prior designs (Hu et al., 2025), and AFlow searches over code-represented workflows (Zhang et al., 2025b). A builder model can also learn principles for constructing the execution environment of a fixed target model (Qian et al., 2026). AlphaEvolve applies evolutionary program search to scientific and algorithmic discovery, with executable evaluators selecting candidate task programs (Novikov et al., 2025). These approaches optimize objects ranging from prompts to complete programs, with different boundaries around the optimization procedure itself.

Self-modifying agents extend this search to their own implementations. STOP applies its code improver to the improver itself (Zelikman et al., 2024); Gödel Agent and SICA likewise revise agent logic and orchestration code (Yin et al., 2025; Robeyns et al., 2025). The Darwin Gödel Machine retains an archive of empirically evaluated, self-modified agents (Zhang et al., 2025a), while the Huxley–Gödel Machine also considers a candidate's potential to produce better descendants (Wang et al., 2026d). Meta-Harness searches harness code using previous candidates, scores, and execution traces (Lee et al., 2026b). RHI represents the harness as a prompt-level specification of the agent loop and refines it using pairwise feedback over its own revision history (Lee et al., 2026a). Self-Harness uses execution failures to guide revisions to the model's own harness (Zhang et al., 2026a), and AIDE^2 recursively revises research-agent code (Srikanth et al., 2026).

Several systems explicitly distinguish task improvement from improvement of the process that produces it. Hyperagents combine a task agent and a meta agent in a shared editable program, allowing the self-modification mechanism itself to evolve and studying transfer of meta-level improvements (Zhang et al., 2026b). HSI separates a task harness, an evolver, and a meta-evolver under a fixed outer anchor (Zhou, 2026). In robotics, RegenHarness revises harness configuration from execution records, with regression checks and release gates governing acceptance (Wang et al., 2026a). Harness evolution around frozen models now spans video understanding (Luo et al., 2026; Liang et al., 2026; Lin et al., 2026), image generation with a frozen generator (Han et al., 2026a), 3D reconstruction (Yao et al., 2026a), and embodied agents, including robot manipulation, where the harness is revised from diagnosed execution failures (Tian et al., 2026a; Feng et al., 2026; Yu et al., 2026b; Du et al., 2026), experience is organized into knowledge or executable stacks shared across tasks (Bao et al., 2026; Chen et al., 2026f; Wen et al., 2026), the system decides where to spend robot trials (Song et al., 2026), and changes are tested in simulation before they reach the robot (Hu et al., 2026). Studies of these loops find that the number of rollouts the optimizer sees in each round and the selection rule determine whether an evolved robot harness generalizes (Wei et al., 2026b), that accepting any edit that raises an aggregate validation score admits regressions and favors noise (Wang et al., 2026e), that proxy improvements need not hold where the update is deployed (Wang et al., 2026b), and that an exact validation score after every attempt opens the validation split to adaptive selection (Yao et al., 2026b). RRSI addresses the risk that harness evolution overfits its training tasks, with gains that shrink or vanish on out-of-distribution benchmarks, by constraining how edits are proposed and selected (Xia et al., 2026). Other studies ask which failures harness evolution can repair and when weight training should take over (Tian et al., 2026b; Wang et al., 2026c), argue that rewriting the scaffold around frozen weights avoids diminishing returns but not an eventual ceiling (Bobadilla-Suarez et al., 2026), and report that weak evaluators and early failures in long skill chains stall improvement over long robot runs (Wang, 2026). Surveys organize these developments by the components that evolve and the feedback that drives their updates (Gao et al., 2026a; Fang et al., 2025; Chen et al., 2026a), and recent surveys and position papers cover recursive self-improvement in physical AI (Han et al., 2026b; Zhou et al., 2026a), agent harnesses for video generation (Fan et al., 2026), and the use of frontier models for 3D modeling, computational design, and robotics (Dou et al., 2026).

These systems propose particular ways to accumulate experience, search for agent designs, or revise an improvement procedure. Flex-RSI evaluates the ability to make such revisions through a shared executor–improver harness (Section 3.2). Each frontier model operates both components from the same initial harness and, on each task, the same starter solution and budget. The executor changes the task solution; the improver changes the executor and its own instructions and lessons. The evaluator, data splits, and budget accounting remain fixed. This setting makes the model's capacity to improve its operating machinery the object of

comparison across vision, robotics, and related tasks, while keeping changes to the solution distinct from changes to the harness.

2.3 Flex-RSI tasks

Flex-RSI draws its 24 tasks from six areas, combining established benchmarks, source datasets, and our own curated creation tasks. Table 1 summarizes their sources and metrics. Each task provides a starter solution, development data, hidden validation and test splits, and a fixed improvement budget, so that these resources support the evaluation of changes to both task solutions and agent harnesses.

Action. The manipulation tasks use simulated environments from RoboDojo (Chen et al., 2026d) and RoboCasa365 (Nasiriany et al., 2026), which extends RoboCasa's household simulation platform (Nasiriany et al., 2024). Active Search uses H*Bench's panoramic environments for visual search (Yu et al., 2026a), while RoboLab uses the RoboLab-120 suite to evaluate a generalist tabletop manipulation policy (Yang et al., 2026a). The LIBERO-based tasks (Liu et al., 2023) examine robustness to changes in initial conditions and task specifications through LIBERO-PRO (Zhou et al., 2025), and adaptation to changes during execution through LIBERO-MAX (Zhang et al., 2026e). Embodied-SWE draws on EmbodiedSWE-Bench, where agents write robot-control programs and receive scores based on physical task progress (Shen et al., 2026b). Related simulation benchmarks extend such evaluation to whole-body humanoid loco-manipulation, as in SIMPLE (Wei et al., 2026a).

Reasoning. Ego-Exo Memory uses the Ego-Exo4D portion of EgoExoMem for memory reasoning across synchronized first- and third-person videos (Liu et al., 2026b; Grauman et al., 2024). Kinematics uses QuantiPhy to infer physical quantities such as sizes, velocities, and accelerations from video (Puyin et al., 2026); Mass adapts product photographs and metadata from Amazon Berkeley Objects for weight estimation (Collins et al., 2022). Abstract reasoning spans grid transformations in ARC-AGI-1 and ARC-AGI-2 (Chollet, 2019; Chollet et al., 2025) and interactive exploration in ARC-AGI-3 (ARC Prize Foundation, 2026). Spatial-VSI uses VSI-Bench's spatial questions about indoor videos (Yang et al., 2025).

Creation. Image Editing uses our own curated images and editing instructions, and Text-to-Image uses original Chinese and English prompts. Multi-turn Editing follows the five-turn sessions of MSE-Bench, introduced with VINCIE, in which each instruction acts on the previous result (Qu et al., 2026). Video Generation uses the compositional prompts and evaluation criteria of T2V-CompBench (Sun et al., 2025).

Reconstruction. 3D Tracking adapts IGGT's instance-tracking setting (Li et al., 2026) on ScanNet and ScanNet++ (Dai et al., 2017; Yeshwanth et al., 2023). Scene from Photo uses HARMONY300, introduced with HARMONY for reconstructing furnished rooms from single images (Sun et al., 2026); its synthetic scenes derive from 3D-FRONT (Fu et al., 2021). Scene from Video uses the indoor scenes released with HoloScene (Xia et al., 2025), drawn from Replica and ScanNet++ (Straub et al., 2019; Yeshwanth et al., 2023), together with CoDimRecon's geometry, rendering, and physical-stability evaluation setting (Xie et al., 2026).

Perception. Visual Search combines high-resolution questions from V*Bench and HR-Bench 8K (Wu and Xie, 2024; Wang et al., 2025a). Video Quality combines subjective quality scores for YouTube-UGC videos (Wang et al., 2019; Yim et al., 2020) with human preferences on generated-video pairs from Rapidata's Veo 3.1 comparison dataset (Rapidata, 2025).

Infrastructure. VAE Decode uses the variational autoencoder released with MiniMax H3 (MiniMax, 2026b) and asks the agent to accelerate its decoder while preserving the reference outputs on unseen inputs.

3 Flex-RSI

Flex-RSI evaluates whether an agent, given a task and a fixed budget, can improve not only its solution to the task but also the harness through which it operates. We first define the task setting (Section 3.1), then the self-improving harness (Section 3.2), and finally how runs are evaluated (Section 3.3).

Table 1 The task suite by area. Source benchmark or dataset, primary metric, and size of the hidden test split of each task. SPL: success weighted by path length; MRA: mean relative accuracy; T-mIoU: temporal mean IoU of the instance tracks; RHAE: the official ARC-AGI-3 score of action efficiency relative to humans; SRCC: Spearman rank correlation; VLM: vision-language model. Ours: data collected for Flex-RSI (Section A).

Area	Task	Source	Metric	Test split
Action	RoboDojo	RoboDojo	Task score	10 tasks × 5 layouts
	RoboCasa	RoboCasa365	Success rate	8 tasks × 5 episodes
	Active Search	H*Bench	SPL	47 episodes
	RoboLab	RoboLab-120	RoboLab score	24 tasks × 2 episodes
	LIBERO-MAX	LIBERO-MAX	Dynamic success rate	24 episode pairs
	LIBERO-PRO	LIBERO-PRO	Success rate	30 episodes
	Embodied-SWE	EmbodiedSWE-Bench	Rubric score	4 seeds × 8 environments
Reasoning	Ego-Exo Memory	EgoExoMem (Ego-Exo4D)	Accuracy	120 questions
	Kinematics	QuantiPhy	MRA	67 questions
	Mass	Amazon Berkeley Objects	MRA	600 items
	ARC-AGI-1	ARC-AGI-1 evaluation set	pass@2	400 puzzles
	ARC-AGI-2	ARC-AGI-2 evaluation set	pass@2	120 puzzles
	ARC-AGI-3	ARC-AGI-3 public games	RHAE	10 games
	Spatial-VSI	VSI-Bench	VSI-Bench score	320 questions
Creation	Image Editing	Ours	VLM judge + checks	80 images
	Multi-turn Editing	MSE-Bench	VLM judge	40 sessions
	Video Generation	T2V-CompBench	Composition score	20 prompts
	Text-to-Image	Ours	VLM judge	80 prompts
Reconstruction	3D Tracking	IGGT (ScanNet, ScanNet++)	T-mIoU	10 clips
	Scene from Photo	HARMONY300	Composite score	60 photos
	Scene from Video	HoloScene	Composite score	6 scenes
Perception	Visual Search	V*Bench, HR-Bench 8K	Accuracy	140 questions
	Video Quality	YouTube-UGC, Veo 3.1 pairs	(SRCC + pairwise acc.) / 2	40 videos, 20 pairs
Infrastructure	VAE Decode	MiniMax H3 VAE decoder	$\log_2$ speedup	12 scenes

3.1 Problem setup

Each task provides the agent with a starter solution, compute resources, a budget, and data partitioned into three disjoint splits. The *development split*, including its labels, is available to the agent for experimentation and evaluation. The *validation split* is withheld from the agent and determines which candidate solutions are retained. The *test split* is likewise withheld and is used only once, after the run terminates, to evaluate the final solution. The objective of the agent is to maximize the test score of its final solution within the budget.

On a given task, every run starts from the same starter solution and is subject to the same budget, which limits wall-clock time, GPU hours, model usage, the number of validation evaluations, and the number of rounds. The 24 tasks are listed in Table 1 and described in Section A.

3.2 The self-improving harness

A harness is the program that mediates between a language model and a task. Through its prompts, tools, control loop, and memory, it determines the context presented to the model and the actions available to it. Since the harness consists of code and text, the agent can modify it in the same manner as any other program, while the model weights remain fixed. In Flex-RSI, the harness $H = (E, I)$ comprises two components that query the same model. The *executor* E operates on the task solution: it modifies the solution, conducts experiments, and evaluates the results on the development split. The *improver* I operates on the harness: it analyses the executor's trajectories and results, diagnoses failures, and revises the prompts, tools, and workflow of the executor as well as its own instructions and accumulated lessons.

Let s_0 denote the starter solution, $\mathcal{D}_{\text{dev}}$, $\mathcal{D}_{\text{val}}$, and $\mathcal{D}_{\text{test}}$ the development, validation, and test splits, and $S(s; \mathcal{D})$ the score that the task evaluator assigns to a solution s on a split $\mathcal{D}$. A run starts from the harness $H_0 = (E_0, I_0)$

and proceeds over rounds $t = 1, 2, \ldots$. In round t, the executor of the current harness modifies the accepted solution s_{t-1}, using only the development split, and submits a candidate

$$\tilde{s}_t = E_{t-1}\big(s_{t-1};\ \mathcal{D}_{\text{dev}}\big). \tag{1}$$

The candidate is then evaluated on the validation split and retained only if it improves on the accepted solution by more than a noise margin δ, which the starting harness sets to twice the standard deviation of the starter's validation score and which the improver may revise (Section B):

$$s_t = \begin{cases} \tilde{s}_t & \text{if } S(\tilde{s}_t; \mathcal{D}_{\text{val}}) > S(s_{t-1}; \mathcal{D}_{\text{val}}) + \delta, \\ s_{t-1} & \text{otherwise.} \end{cases} \tag{2}$$

Finally, the improver examines the records $\tau_{1:t}$ of all rounds so far, which comprise the executor's trajectories, experiments, and scores, and produces the harness for the next round:

$$(E_t,\ I_t) = I_{t-1}\big(E_{t-1},\ I_{t-1};\ \tau_{1:t}\big). \tag{3}$$

The run terminates at the round T in which the budget is exhausted or progress stops (Section B), and its result is the final solution s_T.

Equation (3) makes the recursion explicit: the improver is both the operator that produces the next harness and one of its arguments, so that a modification of I changes how all subsequent modifications are made. The decomposition also separates the two kinds of improvement. A change to the solution takes effect only after it has been confirmed on data unseen by the agent (Equation (2)), whereas a change to the harness is reflected in the candidates of subsequent rounds (Equation (1)).

All runs start from the same general-purpose harness H_0, which contains no task- or model-specific components. The evaluator S, the data splits, and the accounting of the budget are fixed and cannot be modified by the agent. Differences between runs can therefore be attributed to the agents themselves, and runs of different models on the same task are directly comparable.

3.3 Evaluation

We report the test score of the final solution, $S(s_T; \mathcal{D}_{\text{test}})$, and its improvement over the starter solution, $S(s_T; \mathcal{D}_{\text{test}}) - S(s_0; \mathcal{D}_{\text{test}})$. Agent code is executed in a sandbox without access to the validation or test data, so no run can read the held-out data.

4 Experiments

This section reports what the agents achieve, in the order in which a reader would ask about it: how well each model does over all tasks and within each area (Section 4.2), how much of that is due to self-improvement (Section 4.3), in which rounds the gains arrive (Section 4.4), and what they cost (Section 4.5). It reports what happens; Section 5 examines why.

4.1 Models and setup

We evaluate 18 models, comprising 11 closed-source models, namely Claude Opus 5.5 (Anthropic, 2026a), Claude Sonnet 5.5 (Anthropic, 2026b), GPT-6 Astra (OpenAI, 2026b), GPT-6.1 Sol (OpenAI, 2026c), GPT-5.6 Sol (OpenAI, 2026a), Gemini 4 Argon (Google, 2026b), Gemini 3.8 Flash (Google DeepMind, 2026), Gemini 3.7 Flash (Google, 2026a), Grok 4.6 (xAI, 2026), Muse Spark 1.3 (Meta, 2026a), and Muse Wonder 1.0 (Meta, 2026b), and 7 open-weight models, namely Qwen3.8-27B and Qwen3.8-Flash-Next (Qwen Team, 2026b), Qwen3.6-35B-A3B (Qwen Team, 2026a), Gemma-4-31B-it (Gemma Team, 2026), MiniMax-M3 (MiniMax, 2026a), GLM-5.3 Flash (GLM-5 Team, 2026), and DeepSeek-V4.1 Flash (DeepSeek-AI, 2026).[1] Closed-source models are accessed through their providers' interfaces, and open-weight models are served on our own GPUs. Every model operates within the same starting harness (Section 3.2) and, on a given task, under the same budget.

[1]We evaluate the 1017 version of DeepSeek-V4.1 Flash; it is referred to as DeepSeek-V4.1 Flash throughout.

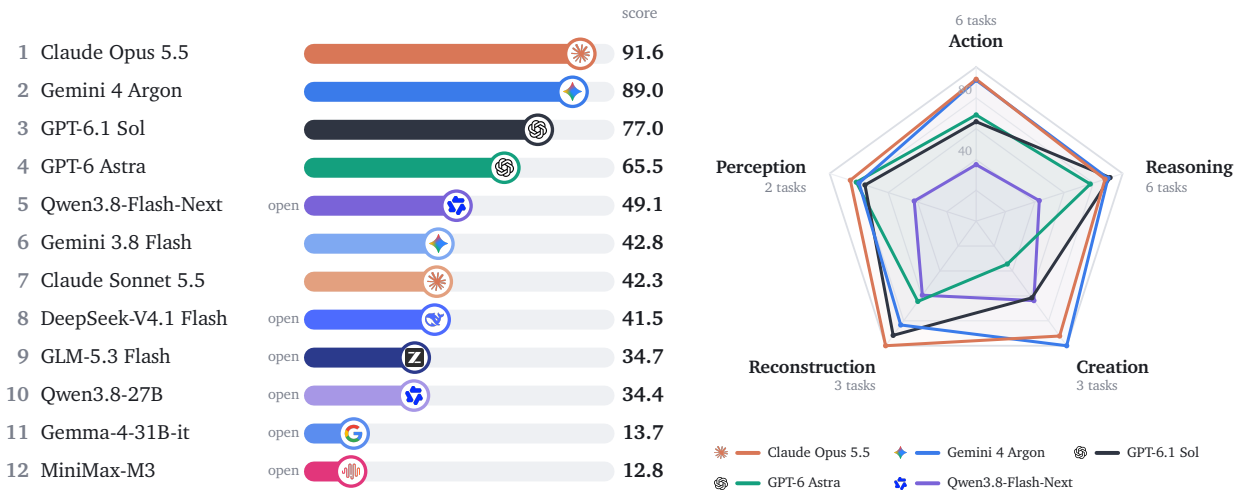

(a) Score of every model run on at least five tasks.

(b) Per-area score of the five leading models.

Figure 2 **Leaderboard.** On each task a model scores the share of the other models it beats (100 = best). (b) gives a model's mean over the tasks of each area it was run on, further out is stronger; (a) gives the mean of these five area scores.

How results are reported. Each model is run once on each task it is evaluated on; where a run was repeated, the latest run is reported. RoboDojo and RoboCasa are suites of 10 and 8 robot tasks, each run separately for 10 rounds and aggregated into one suite score. A run's result is the test score of its final solution on the task's primary metric, on a 0–100 scale. Because the metrics differ across tasks, models are compared through ranks: on each task, a model scores the share of the other models evaluated on the task that it beats, from 0 (last) to 100 (first). Its score in an area is the mean over the tasks of the area it was run on, and its overall score is the mean of its five area scores. To measure self-improvement itself, we also report the validation score of the starter solution and the best validation score reached during the run. The main results cover the 20 tasks evaluated with at least six models; the four tasks evaluated with two models so far are reported in Table 12.

4.2 Main results

Figure 2a gives one score per model and Figure 2b breaks it down by area. Claude Opus 5.5 has the highest score (91.6): it is one of two models run on all 20 tasks and has the best test score on 11 of them. Gemini 4 Argon follows closely (89.0) over the nine tasks it was run on, then GPT-6.1 Sol (77.0) and GPT-6 Astra (65.5). The best open-weight model, Qwen3.8-Flash-Next (49.1), places fifth, ahead of Gemini 3.8 Flash and Claude Sonnet 5.5. The leaders differ by area: Sol matches Opus and Argon on reasoning but trails them on action and creation, and Astra is weakest on creation. The rest of this section takes the areas in turn; each table gives the test score of every model on every task of the area, and Section C adds the starter's score, the validation scores and the rounds of every run (and the average rank of each model, Table 11).

Action. The model's solution drives an agent: an active searcher in 360° scenes (Active Search, SPL), one generalist tabletop policy (RoboLab), a Franka arm in two LIBERO variants, a dual-arm ARX X5 robot (RoboDojo, 10 tasks), and a mobile manipulator in a simulated kitchen (RoboCasa, 8 tasks). Opus has the best test score on five of the six tasks, and Argon on Active Search (79.7). The gap between closed-source and open-weight models is widest here: on RoboDojo and RoboCasa the best open-weight models reach 17.8 and 7.5, against 55.1 and 57.5 for Opus. These two suites start from nothing, the starter scoring 0 on validation in 4 of 8 RoboDojo runs and 7 of 10 RoboCasa runs, so a model must first write a working policy before a round can improve the score. Tables 13 and 14 give the results of the two suites task by task.

Reasoning. The model answers memory questions over synchronized egocentric and exocentric video (Ego-Exo Memory), estimates physical quantities from video and images (Kinematics and Mass, mean relative accuracy), and solves ARC-AGI-1 and -2 grid puzzles and the interactive games of ARC-AGI-3. Closed models

Table 2 **Action.** Test score on each task (0–100); models ordered by average rank over these tasks.

Model	Avg. rank	Active Search	RoboLab	LIBERO-MAX	LIBERO-PRO	RoboDojo	RoboCasa
Claude Opus 5.5	**1.5**	74.2	**71.9**	**45.8**	**63.3**	**55.1**	**57.5**
Gemini 4 Argon	1.7	**79.7**	–	–	–	46.8	40.0
GPT-6.1 Sol	2.8	74.4	62.5	41.7	53.3	41.5	32.5
GPT-6 Astra	3.5	78.2	33.0	–	–	28.6	37.5
Qwen3.6-35B-A3B	3.5	–	–	25.0	53.3	–	–
GPT-5.6 Sol	4.0	–	37.8	–	–	–	–
Qwen3.8-Flash-Next	5.2	59.3	17.4	**45.8**	53.3	8.4	5.0
Qwen3.8-27B	6.0	69.0	31.6	25.0	50.0	–	0.0
Claude Sonnet 5.5	6.0	55.8	41.0	–	–	–	–
GLM-5.3 Flash	6.0	60.5	–	–	–	–	7.5
Gemini 3.8 Flash	6.0	65.5	–	–	–	–	–
Muse Wonder 1.0	6.0	–	–	–	–	10.6	–
Gemma-4-31B-it	6.4	52.2	5.2	29.2	53.3	–	2.5
DeepSeek-V4.1 Flash	8.0	47.1	–	–	–	17.8	0.0
MiniMax-M3	9.0	32.5	8.3	–	–	0.0	0.0

Table 3 **Reasoning.** Test score on each task (0–100); models ordered by average rank over these tasks.

Model	Avg. rank	Ego-Exo	Kinematics	Mass	ARC-1	ARC-2	ARC-3
GPT-6.1 Sol	**1.8**	67.5	–	–	98.8	**95.8**	**44.2**
Claude Opus 5.5	1.8	**72.5**	54.6	**63.5**	98.0	95.0	40.4
Gemini 4 Argon	2.0	70.0	–	–	–	–	–
Claude Sonnet 5.5	2.5	–	**57.2**	62.5	–	–	–
GPT-5.6 Sol	2.5	–	50.0	63.0	–	–	–
GPT-6 Astra	2.8	65.8	–	–	**99.8**	93.3	36.5
Qwen3.8-Flash-Next	5.2	59.2	49.2	60.4	80.8	22.5	0.9
Gemini 3.7 Flash	5.5	–	46.3	60.6	–	–	–
Qwen3.8-27B	5.7	60.8	43.4	62.7	77.2	19.2	0.3
DeepSeek-V4.1 Flash	6.0	50.8	–	–	91.0	46.7	0.6
Gemini 3.8 Flash	6.0	63.3	48.2	57.5	–	–	–
Grok 4.6	7.5	–	37.2	59.3	–	–	–
GLM-5.3 Flash	7.7	56.7	–	–	63.5	14.2	–
Gemma-4-31B-it	7.8	59.2	–	–	54.0	1.7	0.0
MiniMax-M3	9.0	51.7	–	–	–	4.2	–

lead every task: Opus on Ego-Exo Memory (72.5) and Mass (63.5), Sonnet on Kinematics (57.2), Astra on ARC-AGI-1 (99.8), and Sol on ARC-AGI-2 (95.8) and ARC-AGI-3 (44.2). Open-weight models are close on the physical quantities (Qwen3.8-27B scores 62.7 on Mass) but fall away on ARC: the best of them reaches 46.7 on ARC-AGI-2 and 0.9 on ARC-AGI-3. ARC-AGI-3 is also one of the tasks where self-improvement does least, with a median validation score of 3.8 for the starter and 4.1 for the best solution.

Creation. The model edits images from instructions (Image Editing), edits one photo over five turns (Multi-turn Editing), and builds text-to-video workflows around a frozen generator (Video Generation, composition score). Scores are compressed: the 12 models span 60.6 to 74.2 on Image Editing, and the top six on Multi-turn Editing lie within 2.8 points. Open-weight models are most competitive here: Qwen3.8-Flash-Next ranks third in the area and scores 91.4 on Video Generation, against 93.9 for Opus.

Reconstruction. The model tracks object instances in 3D (3D Tracking, temporal mIoU) and rebuilds a furnished room from a photo or an indoor scene from a video (scene score). Opus leads all three tasks (79.6, 66.5 and 70.9), 13 to 18 points ahead of the next model, GPT-6 Astra, on the two scene tasks. Reconstruction gains the most from self-improvement: the median validation score rises from 19.4 to 62.2 on 3D Tracking and from 7.7 to 38.0 on Scene from Photo.

Perception. The model finds small targets in high-resolution images (Visual Search, accuracy) and rates the perceptual quality of real user videos (Video Quality). The starters already score high, with a median

Table 4 Creation. Test score on each task (0–100); models ordered by average rank over these tasks. Image, Multi-turn and Video: Image Editing, Multi-turn Editing and Video Generation.

Model	Avg. rank	Image	Multi-turn	Video
Gemini 4 Argon	**1.0**	**74.2**	**71.1**	–
Claude Opus 5.5	1.7	73.8	70.3	**93.9**
Qwen3.8-Flash-Next	3.7	70.0	69.2	91.4
GPT-6.1 Sol	4.0	69.9	68.3	91.5
Qwen3.8-27B	4.7	69.7	70.2	81.8
GPT-6 Astra	6.5	69.7	67.8	–
Gemma-4-31B-it	6.7	68.4	66.3	87.1
Claude Sonnet 5.5	7.0	66.0	69.8	–
MiniMax-M3	7.5	64.3	–	87.3
DeepSeek-V4.1 Flash	8.0	67.6	–	–
Gemini 3.8 Flash	9.0	67.2	–	–
GLM-5.3 Flash	9.5	60.6	–	79.6

Table 5 Reconstruction. Test score on each task (0–100); models ordered by average rank over these tasks. Photo and Video: Scene from Photo and Scene from Video.

Model	Avg. rank	3D Track.	Photo	Video
Claude Opus 5.5	**1.0**	**79.6**	**66.5**	**70.9**
GPT-6.1 Sol	2.0	77.3	–	–
Gemini 4 Argon	3.0	73.8	–	–
GLM-5.3 Flash	4.0	61.2	–	–
Qwen3.8-Flash-Next	4.0	58.9	40.0	34.5
GPT-6 Astra	4.3	52.8	53.2	53.4
DeepSeek-V4.1 Flash	5.0	60.2	–	–
Qwen3.8-27B	6.3	52.4	29.0	26.7
Muse Wonder 1.0	7.0	57.9	–	–
MiniMax-M3	7.0	27.9	33.3	26.7
Gemini 3.8 Flash	8.0	53.7	–	–
Gemma-4-31B-it	8.0	32.0	15.4	17.5
Claude Sonnet 5.5	11.0	47.9	–	–

Table 6 Perception. Test score on each task (0–100); models ordered by average rank over these tasks.

Model	Avg. rank	Visual Search	Video Quality
Claude Opus 5.5	**2.0**	96.4	82.0
GPT-6 Astra	3.0	93.6	**82.1**
GPT-6.1 Sol	3.0	96.4	80.8
Gemini 4 Argon	3.0	**98.6**	80.6
Gemini 3.8 Flash	5.5	95.0	77.5
DeepSeek-V4.1 Flash	6.5	77.1	82.0
Qwen3.8-27B	6.5	85.7	78.0
Qwen3.8-Flash-Next	7.0	88.6	76.6
GLM-5.3 Flash	8.0	85.7	76.2
Claude Sonnet 5.5	9.0	80.7	–
Gemma-4-31B-it	10.5	72.9	74.3
MiniMax-M3	11.5	70.7	69.8

validation score of 80.0 and 78.2, and the leaders are a few points apart: Gemini 4 Argon scores 98.6 on Visual Search, and on Video Quality GPT-6 Astra (82.1), Claude Opus 5.5 and DeepSeek-V4.1 Flash (82.0) are within 0.1 of each other. With little headroom left, the best open-weight model is 10 points behind on Visual Search and level on Video Quality.

4.3 How much self-improvement adds

The scores above combine what a model can do from the start with what it gains by improving. Figure 3 separates the two by comparing the starter's validation score with the best validation score reached during the run. Self-improvement raises the validation score in 150 of the 177 runs (85%), by a median of 10.4 points.

The gain depends on how much the task leaves to build. It is largest where the agent builds most of its solution from a weak starter: on 3D Tracking, RoboLab and Active Search the median validation score rises from 19.4, 8.3 and 23.0 to 62.2, 45.8 and 52.7 (Figure 3a). It is smallest where the starter already solves most of the task, as on ARC-AGI-1 and ARC-AGI-2 (median starter 91.9 and 82.5, a median gain of 0 points on both). Two tasks do not fit this pattern: ARC-AGI-3 and Scene from Video start low (3.8 and 17.3) and still barely improve (to 4.1 and 18.5).

The gain also depends on the model. Closed-source models improve more often and by more: they raise the validation score in 94% of their 85 runs, by a median of 14.9 points, against 76% of 92 runs and 8.5 points for open-weight models. Per model (Figure 3b), the mean gain ranges from 26.7 points for Claude Opus 5.5 and

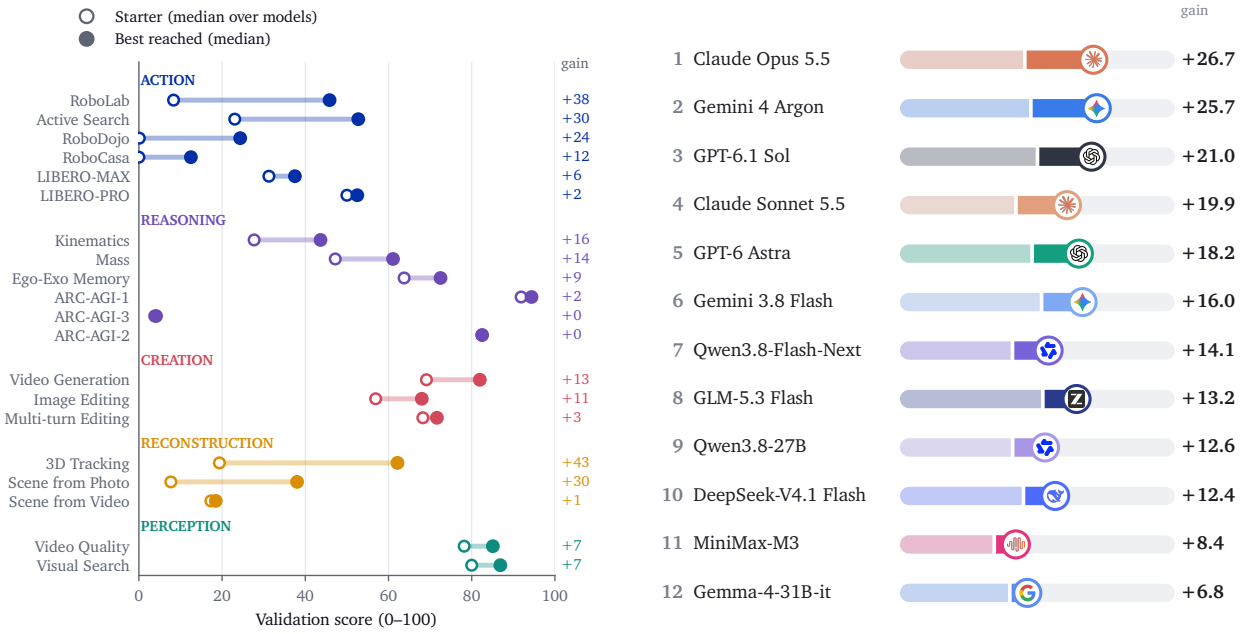

(a) By task: median over its models. **(b)** By model: mean over its runs.

Figure 3 Where self-improvement helps. Validation score of the starter (open circle; light bar) and best validation score reached during the run (filled circle; dark bar), on a 0–100 scale. (a) One row per task, grouped by area and sorted by gain; the number on the right is the difference between the two medians. (b) One row per model run on at least five tasks, sorted by the mean gain.

25.7 for Gemini 4 Argon to 6.8 for Gemma-4-31B-it, and the four models with the highest scores in Figure 2a have four of the five largest gains.

4.4 Improvement over rounds

Figure 4a shows in which rounds the gains arrive. For each run that improves, we express its best validation score after every round as a share of the run's whole gain, and take the median over the runs of each area. Most of the gain comes early: the median run reaches 90% of its gain by round 3 in Perception and Reasoning and by round 4 in Creation. Action and Reconstruction take longer, reaching 90% only by rounds 6 and 8. These are the two areas with the weakest starters (median validation score 16.7 and 19.4, against 60.0 to 78.8 in the other three), where the agent likely has to build a working solution before it can refine it.

4.5 Cost

Runs differ widely in what they spend. A run on one task, counting each robot task of RoboDojo and RoboCasa separately, uses a median of 49.7 million tokens, and the middle 80% of runs use between 8.6 and 93.4 million. We compare tokens rather than prices, since the open-weight models are served on our own GPUs. To compare models on the same footing, Figure 4b expresses each model's token use relative to the median model on the same task. Spending more does not buy a higher score: across the 12 models, the rank correlation between relative token use and score is −0.05. Claude Opus 5.5, GPT-6.1 Sol and GPT-6 Astra, three of the four highest-scoring models, use 0.35 to 0.61 times the median, whereas Gemini 4 Argon reaches a score close to that of Opus with 1.66 times. Five of the six open-weight models use at least the median (1.0 to 1.75 times), and all six score below 50.

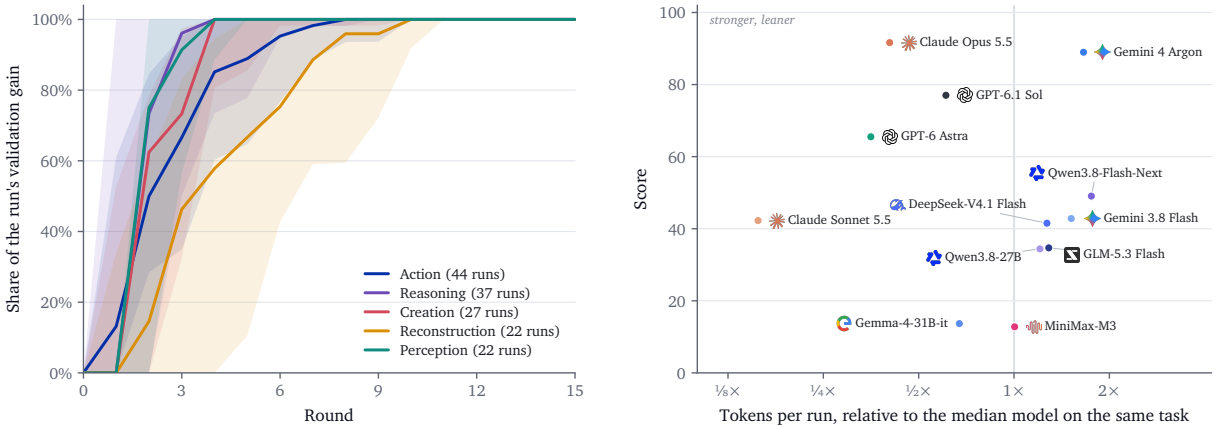

(a) Share of the gain reached by each round.

(b) Token use against score.

Figure 4 **When the gains arrive and what they cost.** (a) For each run with a validation gain, the best validation score after each round as a share of the run's whole gain; the line is the median over the runs of an area and the band the middle half. Runs that stop earlier keep their final value, and runs longer than 15 rounds are shown up to round 15. (b) Tokens per run relative to the median model on the same task (geometric mean over the tasks with at least three models with usage records), against the score of Figure 2a.

4.6 Qualitative examples

Figures 5 to 9 compare the behavior of solutions before and after self-improvement on robot manipulation, active search, abstract reasoning, image editing, and 3D reconstruction from a photo. Further examples on these and the remaining tasks are given in Section D.

Figure 5 Robot manipulation. Each row shows one episode of a robot task played by the solution of an earlier round (Before) and by that of a later round (After) of the same run, three frames each, cropped to the part of the scene where the two differ. The rows give the benchmark, the model of the run, the round of each solution and the outcome of the episode.

Object search. Find the ad 'BROADWAY SHOWS'.

Path search. If you want to enter the red house, where should you move?

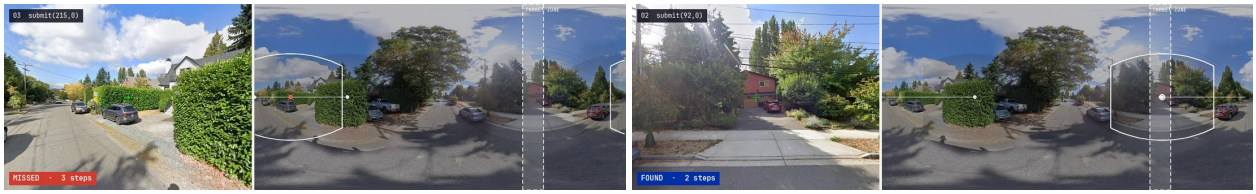

Path search. You want to travel to Frankfurt, your flight number is 'LH496', where'd you go?

Figure 6 **Active Search.** Episodes of the starter solution (left of each row) and of the final solution of the run of Claude Opus 5.5 (right). Each image shows the last view of the episode (left part) and the panorama with the path of the view and the target zone (right part); the label gives the outcome and the number of actions, and the top line the last action. Object search asks for a named object, path search for the direction in which to move.

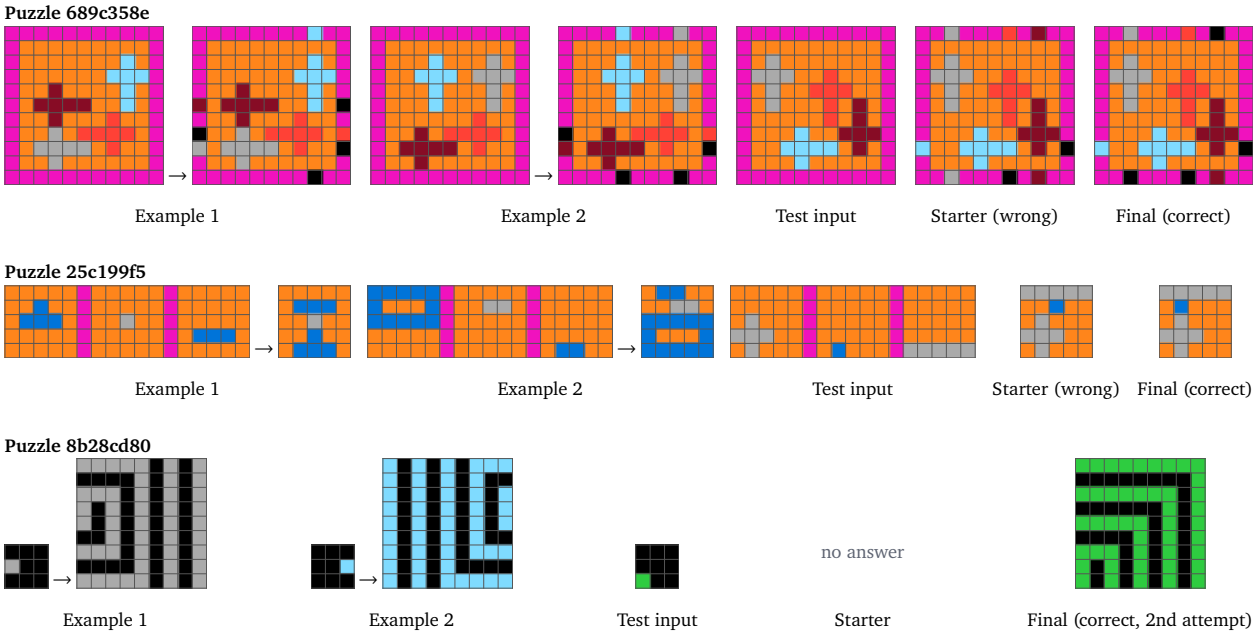

Figure 7 **ARC-AGI-2.** Three validation puzzles: two of the training examples (input → output), the test input, and the answers of the starter solution and of the final solution of the run of Claude Opus 5.5. A puzzle counts as solved if one of two attempts is correct.

Add object, exact count. *Add one more bee-eater of the same species on the branch so there are three in total*

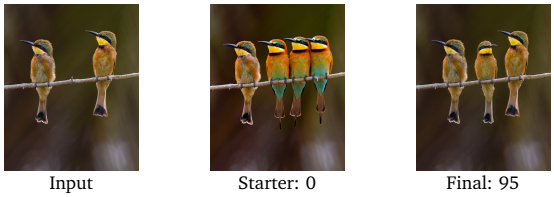

Input Starter: 0 Final: 95

Solve, then write. *Write the answer to this expression after the equals sign on the chalkboard*

Input Starter: 0 Final: 89

What happens over time. *Show what these bananas look like a week later*

Input Starter: 0 Final: 77

Clean up and fill in. *Clean up the litter on the reedy shore and fill it in naturally*

Input Starter: 0 Final: 83

Replace one thing, keep the rest. *Replace the flowers in the painting with sunflowers; keep the vase, the peach and the grapes unchanged*

Input Starter: 63 Final: 91

Soft-effect removal, haze. *Remove the haze over the city so the sky is clear and the buildings look crisp with natural colors; keep the buildings unchanged*

Input Starter: 51 Final: 75

Figure 8 **Image Editing.** For each edit instruction, the input image and the outputs of the starter solution and of the final solution of the run of Claude Opus 5.5, with the score that the task evaluator assigns to each output (multiplied by 100).

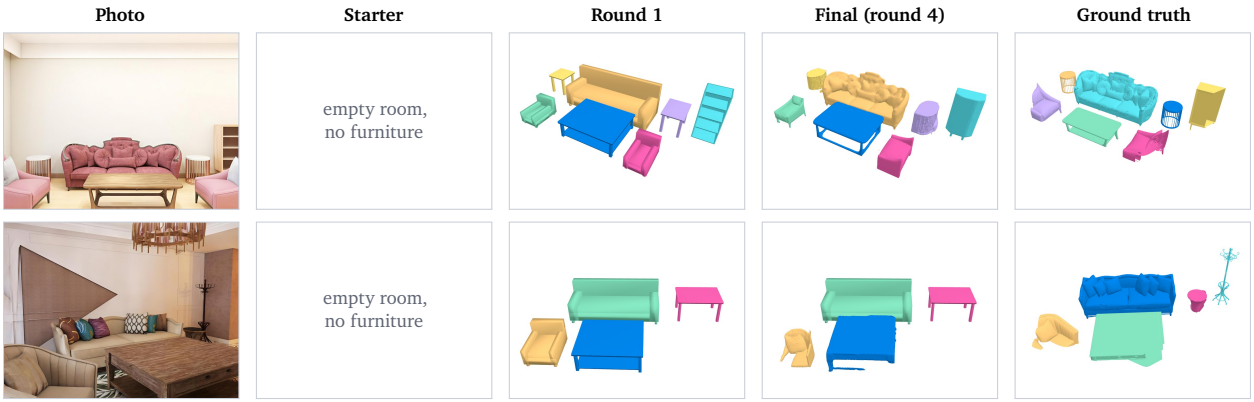

Figure 9 **Scene from Photo.** Development photos and the furniture reconstructed by the starter solution, by the solution after round 1, and by the final solution of the run of Claude Opus 5.5, next to the annotated furniture, seen from above the photo's camera; each color is one piece, and the room box is not drawn. The starter returns an empty room box; the solution of round 1 places a box template for each piece; the final solution generates the shape of each piece.

5 Analysis and takeaways

Table 7 lists, for every task evaluated with at least six models, the quantities on which the findings of Section 1 rest. The tasks fall into two groups according to what the agent has to produce. On seven tasks, the solution gives a single answer to each question, and the starter solution obtains it by passing the question to a model: the model being evaluated, or on Kinematics and Mass a fixed open-weight model. On the other thirteen, the agent has to build the solution itself, namely a program, a robot policy, or a processing pipeline that produces the output. These thirteen tasks are divided further according to how the task evaluator scores an episode: on nine it assigns a score on a scale, and on four it records only whether the episode or level succeeded. Sections 5.1 to 5.3 examine the three findings in turn, and Section 5.4 examines whether editing the harness contributes to the improvement.

Table 7 Evidence for the findings. Starter solution gives the median validation score of the starter solution over the runs of a task, and Largest gain the largest increase of the validation score over the starter solution in any run. Best minus worst model gives the difference between the highest and the lowest validation score among the models, at the start of the run and at the best validation score reached. Median gain gives the median increase of the validation score over the runs of closed-source and of open-weight models, and Best test score the highest test score of a closed-source and of an open-weight model. The four tasks evaluated with two models are omitted. All scores are on the 0–100 scale of the result tables.

			Best minus worst model		Median gain		Best test score	
Task	Starter solution	Largest gain	Start	End	Closed-source	Open-weight	Closed-source	Open-weight
The solution gives a single answer to each question								
Visual Search	80.0	20.0	31.3	22.5	6.2	11.9	98.6	88.6
Video Quality	78.2	14.8	23.5	16.7	6.7	12.7	82.1	82.0
Ego-Exo Memory	63.7	13.8	27.5	25.0	7.5	5.6	72.5	60.8
Kinematics	27.7	33.5	14.2	33.1	8.8	28.5	57.2	49.2
Mass	47.2	16.6	6.0	5.8	13.9	11.2	63.5	62.7
ARC-AGI-1	91.9	11.2	31.2	32.5	1.2	0.0	99.8	91.0
ARC-AGI-2	82.5	3.8	38.8	40.0	1.2	0.0	95.8	46.7
The agent builds the solution; the evaluator scores each episode on a scale								
RoboDojo	0.1	68.4	0.9	68.9	46.8	9.9	55.1	17.8
Active Search	23.0	35.4	30.8	34.4	28.4	29.3	79.7	69.0
RoboLab	8.3	54.2	29.2	62.5	39.6	18.7	71.9	31.6
Image Editing	56.9	22.3	4.6	13.9	15.1	9.3	74.2	70.0
Multi-turn Editing	68.3	4.2	0.4	2.3	3.7	2.1	71.1	70.2
Video Generation	69.1	23.9	1.0	20.2	21.6	8.7	93.9	91.4
3D Tracking	19.4	65.4	0.0	55.6	42.8	39.8	79.6	61.2
Scene from Photo	7.7	60.1	0.0	50.6	53.6	23.0	66.5	40.0
Scene from Video	17.3	38.8	0.0	38.8	24.7	0.0	70.9	34.5
The agent builds the solution; the evaluator records only success or failure								
RoboCasa	0.0	55.0	5.0	58.8	51.2	5.6	57.5	7.5
LIBERO-MAX	31.2	18.8	0.0	18.8	12.5	3.1	45.8	45.8
LIBERO-PRO	50.0	25.0	0.0	25.0	17.5	0.0	63.3	53.3
ARC-AGI-3	3.8	18.4	39.0	49.3	16.4	0.2	44.2	0.9

5.1 Tasks in which the agent builds the solution leave more room for self-improvement

> **Finding 1.** When the agent builds the solution, self-improvement raises the score the most and separates the models; when the solution is a single answer, the starter solution already scores high and the models end about as far apart as they began. On 3D Tracking, every run starts from 19.4 and the best run gains 65.4 points.

On tasks whose solution is a single answer, the starter solution already passes each question to the model, so it scores high from the start and its score depends on the model: its median validation score over these seven tasks is 78.2, and the models are 27.5 points apart before the first round (medians over tasks). The

largest gain that any run achieves is 14.8 points at the median, and at the end of the runs the models are 25.0 points apart, about as far as at the start. On the tasks in which the agent builds the solution, every model starts from the same program, although its score can depend on the model that it calls. On the nine of them whose evaluator scores each episode on a scale, this program scores 19.4 at the median, the largest gain is 38.8 points at the median, and the models, 0.9 points apart at the start, end 38.8 points apart; on the four whose evaluator records only success or failure, they move from 2.5 to 37.2 points apart. The same pattern appears in Figure 1c. The exceptions are Kinematics, whose solution is a single answer but whose starter solution scores only 27.7 and on which a run gains 33.5 points, and Multi-turn Editing, on which the agent builds the solution but no run gains more than 4.2 points.

5.2 The granularity of feedback matters

> **Finding 2.** When the evaluator scores each episode on a scale, open-weight models also improve the solutions they build; when it records only success or failure, they barely improve, while closed-source models still do. On RoboCasa, the median open-weight run gains 5.6 points and the median closed-source run 51.2.

On the nine tasks in which the agent builds the solution and the evaluator scores each episode on a scale, the median gain over the open-weight runs of a task is 9.9 points at the median over tasks, against 28.4 for the closed-source runs, and on 3D Tracking and Active Search it is close to that of the closed-source runs (39.8 against 42.8, and 29.3 against 28.4). On the four tasks whose evaluator records only whether an episode or a level succeeded, namely RoboCasa, LIBERO-MAX, LIBERO-PRO, and ARC-AGI-3, the median open-weight gain falls to 1.6 points, while the closed-source runs still gain 16.9. The median open-weight run does not improve at all on LIBERO-PRO and gains 0.2 points on ARC-AGI-3. Since a candidate is retained only if it raises the validation score (Equation (2)), a score on a scale rewards a change that improves part of an episode, whereas a score that records only success or failure rewards a change only once it turns failed episodes into successful ones.

5.3 Strict checks of physical stability and exact correctness pose the greatest challenge

> **Finding 3.** Outputs that look correct often fail checks of physical stability or exact correctness, and tasks scored by such checks are where open-weight models fall furthest behind. On Scene from Video, no object placed by an open-weight model remains stable in simulation, against 70.7% for Claude Opus 5.5.

The difference between the best closed-source and the best open-weight model is largest on tasks whose evaluator applies such a check: whether a robot task succeeds in simulation (RoboCasa, 50.0 points; RoboLab, 40.3; RoboDojo, 37.3), whether the output grid is exactly right (ARC-AGI-2, 49.2) or a level is cleared (ARC-AGI-3, 43.3; on ARC-AGI-1, whose puzzles the starter solutions already nearly solve, it is 8.8), and whether a reconstructed scene is physically and geometrically consistent (Scene from Video, 36.4; Scene from Photo, 26.5). It is smallest on tasks scored by a judge or with a tolerance: Video Quality (0.1), Mass (0.8), Multi-turn Editing (1.0), Video Generation (2.5), and Image Editing (4.2). The component scores of the two scene tasks show how plausible outputs fail these checks (Tables 29 and 30). On Scene from Photo, the renderings of the open-weight scenes resemble the photo, with an Appearance score of up to 81.8, but the furniture is not recovered as complete and correctly placed objects: their Objects score is at most 16.5, against 62.5 for both closed-source models. On Scene from Video, the Physics score, the fraction of objects that remain stable in simulation, is 0.0 for all four open-weight models, against 70.7 for Claude Opus 5.5 and 37.7 for GPT-6 Astra. The two LIBERO tasks are the exceptions among the robot tasks: the best open-weight model matches the best closed-source model on LIBERO-MAX (45.8) and trails it by 10.0 points on LIBERO-PRO.

5.4 Does editing the harness help?

Finding 4. Letting the improver edit the harness raises the test score in six of seven pairs of otherwise identical runs, and most where the unchanged starting harness makes little progress. On Active Search, Qwen3.8-Flash-Next reaches 21.0 without the improver and 59.3 with it.

Table 8 Frozen and self-improving harness. Each row compares two runs of the same model on the same task whose settings (task version, budget, protocol, and seed) are identical except for the improver. With the frozen harness, the improver does not run and the starting harness is never modified, so only the executor works on the solution; with the self-improving harness, the improver edits the executor and itself after every round, as in all other runs of this report. Validation is the best validation score reached during the run, Test the test score of the final solution, and Δ the test score with the self-improving harness minus that with the frozen one. The robot tasks are single RoboCasa tasks with 5 episodes per split. Claude Opus 5.5 with the frozen harness stopped after round 8, when its validation score reached the maximum; a second self-improving run of Claude Opus 5.5 on Turn On Sink Faucet with the same settings reached a test score of 40.0. In the RoboCasa runs of DeepSeek-V4.1 Flash and Qwen3.8-Flash-Next, the model also reasons while acting as the policy, so the Open Drawer run of DeepSeek-V4.1 Flash differs from the one in Table 14. All scores are multiplied by 100.

		Frozen harness			Self-improving harness			
Task	**Model**	**Validation**	**Test**	**Rounds**	**Validation**	**Test**	**Rounds**	**Δ Test**
RoboCasa: Open Drawer	DeepSeek-V4.1 Flash	0.0	0.0	5	40.0	40.0	10	+40.0
RoboCasa: Turn On Sink Faucet	Qwen3.8-Flash-Next	20.0	0.0	9	40.0	20.0	10	+20.0
	Claude Opus 5.5	100.0	80.0	8	80.0	80.0	10	0.0
Active Search	Qwen3.8-27B	41.2	50.7	10	51.7	69.0	15	+18.3
	Qwen3.8-Flash-Next	16.5	21.0	5	50.3	59.3	12	+38.3
3D Tracking	Qwen3.8-27B	48.9	45.9	11	53.9	52.4	9	+6.5
	Qwen3.8-Flash-Next	63.4	56.9	9	65.7	58.9	12	+2.0

Table 9 Starting a weaker model from a stronger model's harness. On two RoboCasa tasks, Active Search, and 3D Tracking, a weaker model is run from the harness that a stronger model had at the end of its own run on the task (the source run, last row of each block). Executor gives the executor the run starts from: the starting harness of every run, or the final executor of the source run. Improver gives the improver that runs between rounds: the model's own, which edits the executor and itself; none, in which case the executor stays frozen; the final improver of the source run, kept fixed, which edits only the executor; or the model's own improver starting from the final improver of the source run, which edits the executor and itself. The first row of each block is the self-improving run of Table 8. Validation is the best validation score reached during the run and Test the test score of the final solution. Runs from a source run's harness use the code and protocol of the source run; the RoboCasa runs use 5 episodes per split. All scores are multiplied by 100.

Task	**Model**	**Executor**	**Improver**	**Validation**	**Test**	**Rounds**
RoboCasa: Open Drawer	DeepSeek-V4.1 Flash	Starting harness	Own	40.0	40.0	10
	DeepSeek-V4.1 Flash	From Gemini 4 Argon	None (executor frozen)	40.0	20.0	10
	DeepSeek-V4.1 Flash	From Gemini 4 Argon	Fixed, from Gemini 4 Argon	80.0	60.0	10
	Gemini 4 Argon (source run)	Starting harness	Own	60.0	80.0	10
RoboCasa: Turn On Sink Faucet	Qwen3.8-Flash-Next	Starting harness	Own	40.0	20.0	10
	Qwen3.8-Flash-Next	From Claude Opus 5.5	None (executor frozen)	20.0	0.0	10
	Qwen3.8-Flash-Next	From Claude Opus 5.5	Fixed, from Claude Opus 5.5	0.0	0.0	10
	Claude Opus 5.5 (source run)	Starting harness	Own	80.0	80.0	10
Active Search	Qwen3.8-27B	Starting harness	Own	51.7	69.0	15
	Qwen3.8-27B	From Gemini 4 Argon	Own, from Gemini 4 Argon	39.5	49.2	9
	Gemini 4 Argon (source run)	Starting harness	Own	67.3	79.7	12
3D Tracking	Qwen3.8-27B	Starting harness	Own	53.9	52.4	9
	Qwen3.8-27B	From Claude Opus 5.5	Own, from Claude Opus 5.5	59.0	53.8	15
	Claude Opus 5.5 (source run)	Starting harness	Own	84.7	79.6	15

Table 8 compares pairs of runs that differ only in whether the improver runs. In the run without the improver, labeled frozen harness in the table, the starting harness is never modified; in the run with it, labeled self-improving harness, the improver edits the executor and itself after every round, as in all other runs of this report. The test score with the improver is higher in six of the seven pairs, by 2.0 to 40.0 points; in the seventh, Claude Opus 5.5 on Turn On Sink Faucet reaches 80.0 in both runs, and its run without the improver reached the maximum validation score in round 8. The difference is largest where the unchanged starting harness makes little progress: DeepSeek-V4.1 Flash remains at 0.0 on Open Drawer without the improver and reaches 40.0 with it, and Qwen3.8-Flash-Next gains 20.0 on Turn On Sink Faucet and 38.3 on Active Search. On 3D Tracking, the executor alone already raises the two Qwen models from the 19.4 of the starter solution to validation scores of 48.9 and 63.4, and the improver adds 6.5 and 2.0 points on the test split. The runs without the improver also end after fewer rounds in six of the seven pairs.

Table 9 asks whether a harness built by a stronger model helps a weaker one. On Open Drawer, DeepSeek-V4.1 Flash started from the final executor of Gemini 4 Argon reaches 20.0 when the executor is not modified further and 60.0 when Argon's final improver continues to edit it, against 40.0 when it improves its own harness from the start and 80.0 for Argon. On Turn On Sink Faucet, Qwen3.8-Flash-Next started from the harness of Claude Opus 5.5 reaches 0.0 in both settings, below the 20.0 it reaches when it improves its own harness and the 80.0 of Opus. On Active Search and 3D Tracking, Qwen3.8-27B started from the final harness of the source run and improved it further with its own improver. On Active Search, it reaches 49.2 from the harness of Gemini 4 Argon, against 69.0 when it improves its own harness from the start and 79.7 for Argon. On 3D Tracking, it reaches 53.8 from the harness of Claude Opus 5.5, against 52.4 from the starting harness and 79.6 for Opus. A harness built by a stronger model thus raised the score of the weaker model on two of the four tasks, by 20.0 points on Open Drawer and 1.4 points on 3D Tracking, and in neither case did the weaker model approach the score of the stronger one.

6 Conclusion

We introduced Flex-RSI, a benchmark that measures whether a frontier agent can improve both its task solution and its own harness within a fixed budget, with every run starting from the same general-purpose harness and every final solution scored on a hidden test split. Across 24 tasks and 18 models, three findings emerge. Self-improvement raises the score most on tasks in which the agent builds its own solution, and on these tasks models that start from the same program end far apart, whereas on tasks whose solution is a single answer the starter solution already scores high and the models end about as far apart as they began. Open-weight models improve when the evaluator scores each episode on a scale, but barely improve when it records only success or failure. The largest gaps between closed-source and open-weight models arise where the evaluator checks physical stability or exact correctness, even when the outputs look plausible. Paired runs further show that editing the harness contributes to the improvement, most where the unchanged starting harness makes little progress, whereas a harness built by a stronger model raised the score of a weaker model on only two of four tasks.

The improvement measured here is mostly at the level of the harness. The evaluated model's weights never change, and on most tasks no model is trained; on some, the solution may train or fine-tune helper models on the development split. Flex-RSI therefore does not measure a model that modifies its own parameters. Extending the benchmark to weight-level self-improvement, and studying how it interacts with changes to the harness, is left to future work. The leaderboard is maintained on the project page, `https://flex-rsi.github.io`.

7 Contributors

Flex-RSI is a community effort. Organizers lead the project, contributors bring its tasks, code and experiments, and advisors guide its direction. The organizers are listed first; the other contributors and the advisors follow in alphabetical order of surname, and the contributions of each person are given below.

Organizers

Jingdong Zhang* Texas A&M University

Jionghao Wang*	Texas A&M University
Zhengming Yu	Texas A&M University

Contributors (alphabetical order)

Yancheng Cai	University of Cambridge
Wei Chow	National University of Singapore
Chris Choy	NVIDIA
Zhiyang Dou	Massachusetts Institute of Technology
Hang Hua	MIT-IBM Watson AI Lab
Jiachen Li	The University of Texas at Austin
Zhixuan Liu	Carnegie Mellon University
Chen Wang	University of Pennsylvania
Songlin Wei	University of Southern California
Siyuan Yang	Texas A&M University
Yu Yuan	Purdue University
Xiaohang Zhan	Adobe
Yucheng Zhang	Purdue University
Shuhong Zheng	University of Toronto

Advisors (alphabetical order)

Xin Li	Texas A&M University
Lingjie Liu	University of Pennsylvania
Wojciech Matusik	Massachusetts Institute of Technology
Zhengzhong Tu	Texas A&M University
Wenping Wang	Texas A&M University

*Organizers, equal contribution.

Contributions.

- **Project leads:** Jingdong Zhang, Jionghao Wang, Zhengming Yu (benchmark protocol, organization, project page, and paper).
- **Action tasks:** Jionghao Wang (RoboDojo, RoboCasa), Zhiyang Dou (RoboDojo, RoboCasa), Yucheng Zhang (RoboDojo, RoboCasa), Songlin Wei (RoboDojo, RoboCasa), Jingdong Zhang (Active Search, LIBERO-MAX, LIBERO-PRO), Xiaohang Zhan (Active Search), Chris Choy (RoboLab), Jiachen Li (LIBERO-MAX, LIBERO-PRO), Hang Hua (Embodied-SWE).
- **Reasoning tasks:** Jingdong Zhang (Ego-Exo Memory), Yu Yuan (Kinematics, Mass), Zhengming Yu (ARC-AGI-1, ARC-AGI-2, ARC-AGI-3), Hang Hua (Spatial-VSI).
- **Creation tasks:** Xiaohang Zhan (Image Editing), Zhixuan Liu (Multi-turn Editing), Shuhong Zheng (Video Generation), Jingdong Zhang (Video Generation), Wei Chow (Text-to-Image).
- **Reconstruction tasks:** Jingdong Zhang (3D Tracking), Xiaohang Zhan (3D Tracking), Chen Wang (Scene from Photo, Scene from Video).
- **Perception tasks:** Jingdong Zhang (Visual Search, Video Quality), Xiaohang Zhan (Visual Search, Video Quality), Yancheng Cai (Video Quality).
- **Infrastructure task:** Siyuan Yang (VAE Decode).
- **Harness and codebase:** Jingdong Zhang, Jionghao Wang, Zhengming Yu, Chris Choy, Jiachen Li, Shuhong Zheng, Wei Chow, Zhixuan Liu, Zhiyang Dou, Hang Hua, Xiaohang Zhan, Chen Wang.
- **Experiments and ablations:** Jingdong Zhang, Jionghao Wang, Zhengming Yu, Zhiyang Dou, Yucheng Zhang.

References

Lakshya A Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziems, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J Ryan, Meng Jiang, et al. GEPA: Reflective prompt evolution can outperform reinforcement learning. In *International Conference on Learning Representations*, volume 2026, pages 8479–8565, 2026.

Ekin Akyürek, Mehul Damani, Adam Zweiger, Linlu Qiu, Han Guo, Jyothish Pari, Yoon Kim, and Jacob Andreas. The surprising effectiveness of test-time training for few-shot learning. *arXiv preprint arXiv:2411.07279*, 2024.

Codruta O. Ancuti, Cosmin Ancuti, and Radu Timofte. NH-HAZE: An image dehazing benchmark with non-homogeneous hazy and haze-free images. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 444–445, 2020.

Peter Anderson, Angel Chang, Devendra Singh Chaplot, Alexey Dosovitskiy, Saurabh Gupta, Vladlen Koltun, Jana Kosecka, Jitendra Malik, Roozbeh Mottaghi, Manolis Savva, and Amir R. Zamir. On evaluation of embodied navigation agents. *arXiv preprint arXiv:1807.06757*, 2018.

Anthropic. System card: Claude Opus 5.5. https://www-cdn.anthropic.com/fc1b44717c85dc068bc6ba5024219938094694bd/Claude%20Opus%205.5%20System%20Card.pdf, September 2026a.

Anthropic. System card: Claude Sonnet 5.5. https://www.anthropic.com/claude-sonnet-5-5-system-card, September 2026b.

Anthropic. Claude Code: Overview. Official documentation, 2026c. URL https://code.claude.com/docs/en/overview. Accessed October 29, 2026.

ARC Prize Foundation. ARC-AGI-3: A new challenge for frontier agentic intelligence. *arXiv preprint arXiv:2603.24621*, 2026.

Andreea Ardelean, Mert Özer, and Bernhard Egger. Gen3DSR: Generalizable 3D scene reconstruction via divide and conquer from a single view. In *International Conference on 3D Vision (3DV)*, pages 616–626, 2025.

Shuai Bai, Yuxuan Cai, et al. Qwen3-VL technical report. *arXiv preprint arXiv:2511.21631*, 2025a.

Shuai Bai et al. Qwen2.5-VL technical report. *arXiv preprint arXiv:2502.13923*, 2025b.

Shifeng Bao, Fanding Huang, Yihan Lin, Youhe Feng, Guanlin Li, Chen Zhao, Yang Li, Jiawei He, Cheng Chi, and Jing Zhang. RoboHarn-Evo: Evolving hierarchical physical knowledge for self-improving robotic manipulation. *arXiv preprint arXiv:2609.37583*, 2026.

Gilad Baruch, Zhuoyuan Chen, Afshin Dehghan, Tal Dimry, Yuri Feigin, Peter Fu, Thomas Gebauer, Brandon Joffe, Daniel Kurz, Arik Schwartz, and Elad Shulman. ARKitScenes: A diverse real-world dataset for 3D indoor scene understanding using mobile RGB-D data. *arXiv preprint arXiv:2111.08897*, 2021.

Sebastian Bobadilla-Suarez, Bob Suh, and Ryan Fortin. Audit the scaffold, not the checkpoint: A stationarity dichotomy for recursive self-improvement in agentic coding. *arXiv preprint arXiv:2609.34924*, 2026.

Ellis Brown, Jihan Yang, Shusheng Yang, Rob Fergus, and Saining Xie. Benchmark designers should "train on the test set" to expose exploitable non-visual shortcuts. *arXiv preprint arXiv:2511.04655*, 2025.

Zhiyong Cao, Hengjie Zhu, Yongxin Wang, Xuchen Li, Jintao Li, Shuxia Liu, Dunqiang Liu, Bo Wang, Kairan Zhang, Cong Lin, JiangDi Yu, Siyuan Liu, et al. Trial by trace: Benchmarking AI agent harnesses for reliability and failure attribution. *arXiv preprint arXiv:2610.32285*, 2026.

Jun Shern Chan, Neil Chowdhury, Oliver Jaffe, et al. MLE-bench: Evaluating machine learning agents on machine learning engineering. In *International Conference on Learning Representations*, 2025.

Mingguang Chen, Licheng Wang, and Bo Qu. Recursive self-improvement in AI: From bounded self-refinement to autonomous research loops. *arXiv preprint arXiv:2607.07663*, 2026a.

Oliver Chen, Geby Jaff, Langston Nashold, and Rayan Krishnan. Vals RSI Index. Vals AI, August 2026b. URL https://vals.ai/benchmarks/rsi_index.

Siyi Chen, Hugo Hadfield, Alex Zook, Mikaela Angelina Uy, Chan Hee Song, Erwin Coumans, Xuning Yang, Faisal Ladhak, Qing Qu, Stan Birchfield, Jonathan Tremblay, and Valts Blukis. VoLo: A physical orchestrator for open-vocabulary long-horizon manipulation. *arXiv preprint arXiv:2606.07723*, 2026c.

Tianxing Chen, Yue Chen, Zixuan Li, Junyuan Tang, Kailun Su, Haoran Lu, Weijie Wan, Baijun Chen, Songling Liu, Haowen Yan, et al. RoboDojo: A unified sim-and-real benchmark for comprehensive evaluation of generalist robot manipulation policies. *arXiv preprint arXiv:2607.04434*, 2026d.

Xiaoyin Chen, Yizheng Huang, Yinglun Xu, Ziyun Zhao, Emmett McQuinn, Julie Liu, and Hope Yao. Are self-evolving agents continual learners? *arXiv preprint arXiv:2610.25858*, 2026e.

Yangzhi Chen, Feiran Liu, Zishuo Wang, Zijian Jin, Xin Zhu, Lecheng Ruan, and Wenjuan Han. Learning to compose: Recursive self-improvement of shared executable stacks for zero-shot long-horizon robot manipulation. *arXiv preprint arXiv:2610.24464*, 2026f.

Zihan Chen, Di Zhu, and Lei Nico Zheng. Auditing the product, not the recursion: The evaluation evidence behind recursive self-improvement claims. *arXiv preprint arXiv:2610.27610*, 2026g.

Ziru Chen, Shijie Chen, Yuting Ning, Qianheng Zhang, Boshi Wang, Botao Yu, Yifei Li, Zeyi Liao, Chen Wei, Zitong Lu, et al. ScienceAgentBench: Toward rigorous assessment of language agents for data-driven scientific discovery. In *International Conference on Learning Representations*, volume 2025, pages 96934–96990, 2025.

Bowen Cheng, Ishan Misra, Alexander G. Schwing, Alexander Kirillov, and Rohit Girdhar. Masked-attention mask transformer for universal image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1290–1299, 2022.

Yizhe Chi, Wenyi Li, Deyao Hong, Xiaoqiu Wang, Mingju Gao, Kaisen Yang, Bingxiang He, Youjie Zheng, Calvin Xiao, and Qinhuai Na. AI4AI-Bench: Benchmarking LLM agents in algorithmic design for recursive self-improvement. *arXiv preprint arXiv:2608.20318*, 2026.

François Chollet. On the measure of intelligence. *arXiv preprint arXiv:1911.01547*, 2019.

Francois Chollet, Mike Knoop, Gregory Kamradt, and Bryan Landers. ARC prize 2024: Technical report. *arXiv preprint arXiv:2412.04604*, 2024.

Francois Chollet, Mike Knoop, Gregory Kamradt, Bryan Landers, and Henry Pinkard. ARC-AGI-2: A new challenge for frontier AI reasoning systems. *arXiv preprint arXiv:2505.11831*, 2025.

Jasmine Collins, Shubham Goel, Kenan Deng, Achleshwar Luthra, Leon Xu, Erhan Gundogdu, Xi Zhang, Tomas F Yago Vicente, Thomas Dideriksen, Himanshu Arora, et al. ABO: Dataset and benchmarks for real-world 3D object understanding. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 21094–21104. IEEE, 2022.

Angela Dai, Angel X Chang, Manolis Savva, Maciej Halber, Thomas Funkhouser, and Matthias Nießner. ScanNet: Richly-annotated 3D reconstructions of indoor scenes. In *2017 IEEE conference on computer vision and pattern recognition (CVPR)*, pages 2432–2443. IEEE, 2017.

DeepSeek-AI. DeepSeek-V4: Towards highly efficient million-token context intelligence. *arXiv preprint arXiv:2606.19348*, 2026.

Zhiyang Dou, Jamison Meindl, Akihisa Watanabe, Anna Deng, Tianyu Huang, Igor Sadalski, Harrison Liang, Minghao Guo, Benjamin Tod Jones, and Wojciech Matusik. On the opportunities and risks of frontier models for 3D modeling, computational design and robotics. Computational Design and Fabrication Group (CDFG), MIT CSAIL, September 2026. URL https://mit-cdfg.github.io/Survey-AI-for-3D-modeling-Robotics/. Living survey.

Tingwen Du, Haoyu Guo, Hanxue Zhang, Xuanxu Wang, Junting Dong, Chao Gao, Mingfei Han, Tao Lu, Mulin Yu, Chunhua Shen, Xiaojun Chang, and Zhihui Li. From VLM planners to robotic agents: Self-evolving harnesses for history-dependent manipulation. *arXiv preprint arXiv:2610.29560*, 2026.

Liaoyuan Fan, Yuhao Zheng, Dongxing Mao, Hangyu Ran, Siyuan Hu, Weijia Mao, Rui Zhao, Weijia Wu, Guian Fang, Yanzhe Chen, Zeyu Zhu, Yuxin Jiang, et al. Video generation meets agent harness: Agentic video generation, agentic world modeling and recursive self-improvement. *arXiv preprint arXiv:2610.40191*, 2026.

Jinyuan Fang, Yanwen Peng, Xi Zhang, Yingxu Wang, Xinhao Yi, Guibin Zhang, Yi Xu, Bin Wu, Siwei Liu, Zihao Li, et al. A comprehensive survey of self-evolving ai agents: A new paradigm bridging foundation models and lifelong agentic systems. *arXiv preprint arXiv:2508.07407*, 2025.

Senyu Fei, Siyin Wang, Junhao Shi, Zihao Dai, Jikun Cai, Pengfang Qian, Li Ji, Xinzhe He, Shiduo Zhang, Zhaoye Fei, Jinlan Fu, Jingjing Gong, and Xipeng Qiu. LIBERO-Plus: In-depth robustness analysis of vision-language-action models. *arXiv preprint arXiv:2510.13626*, 2025.

ZhiYuan Feng, Zhuoshi Pan, Zhaolu Kang, Zhongxia Chen, Shuo Yang, Sen Ye, Heng Wang, and Han Li. EnactOpt: Diagnosis-guided harness evolution for embodied agents. *arXiv preprint arXiv:2610.14977*, 2026.

Chrisantha Fernando, Dylan Banarse, Henryk Michalewski, Simon Osindero, and Tim Rocktäschel. Promptbreeder: Self-referential self-improvement via prompt evolution. *arXiv preprint arXiv:2309.16797*, 2023.

Martin A. Fischler and Robert C. Bolles. Random sample consensus: A paradigm for model fitting with applications to image analysis and automated cartography. *Communications of the ACM*, 24(6):381–395, 1981.

Huan Fu, Bowen Cai, Lin Gao, Ling-Xiao Zhang, Jiaming Wang, Cao Li, Qixun Zeng, Chengyue Sun, Rongfei Jia, Binqiang Zhao, et al. 3D-FRONT: 3D furnished rooms with layouts and semantics. In *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 10913–10922. IEEE, 2021.

Galbot Team, Xuchuan Chen, Xiaoqian Cheng, Yu Deng, Lihe Ding, Shaocong Dong, Xiangjun Gao, Haozhe Jia, Zekai Li, Zhoujian Li, et al. Systematically exploring the capabilities of GPT-6 Astra as embodied policies. *arXiv preprint arXiv:2609.38537*, 2026.

Huan-ang Gao, Jiayi Geng, Wenyue Hua, Mengkang Hu, Xinzhe Juan, Hongzhang Liu, Shilong Liu, Jiahao Qiu, Xuan Qi, Yiran Wu, et al. A survey of self-evolving agents: What, when, how, and where to evolve on the path to artificial super intelligence. *Transactions on Machine Learning Research*, 2026a.

Xingze Gao, Chuanrui Hu, Hongda Chen, Pengfei Yao, Zhao Wang, Yi Bai, Zhengwei Wu, Yunyun Han, Xiaofeng Cong, Jie Gui, Yafeng Deng, and Teng Li. EvoAgentBench: Benchmarking agent self-evolution via ability transfer. In *Findings of the Association for Computational Linguistics: EMNLP 2026*, 2026b.

Shijia Ge, Alex Zhou, Jianshu Zeng, Yexing Wan, Di Wu, Zelin Zheng, Yazhe Wang, Zhiqi Jia, Xuan Shangguan, Jay Zhu, Yijun Liu, Lingyu He, et al. Make code as policy great again: Frontier agents write, call, and evolve robot tools. *arXiv preprint arXiv:2609.39018*, 2026.

Gemma Team. Gemma 4 technical report. *arXiv preprint arXiv:2607.02770*, 2026.

Dhruba Ghosh, Hannaneh Hajishirzi, and Ludwig Schmidt. GenEval: An object-focused framework for evaluating text-to-image alignment. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.

GLM-5 Team. GLM-5: from vibe coding to agentic engineering. *arXiv preprint arXiv:2602.15763*, 2026.

GLM-V Team. GLM-4.5V and GLM-4.1V-Thinking: Towards versatile multimodal reasoning with scalable reinforcement learning. *arXiv preprint arXiv:2507.01006*, 2025.

Google. Gemini 3.7 Flash. Gemini API documentation, https://ai.google.dev/gemini-api/docs/models/gemini-3.7-flash, 2026a.

Google. Gemini 4 Argon. Gemini API documentation, https://ai.google.dev/gemini-api/docs/models/gemini-4-argon, 2026b.

Google DeepMind. Gemini 3.8 Flash model card. https://deepmind.google/models/model-cards/gemini-3-8-flash/, September 2026.

Kristen Grauman, Andrew Westbury, Lorenzo Torresani, Kris Kitani, Jitendra Malik, Triantafyllos Afouras, Kumar Ashutosh, Vijay Baiyya, Siddhant Bansal, Bikram Boote, et al. Ego-Exo4D: Understanding skilled human activity from first- and third-person perspectives. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 19383–19400, 2024.

Shuo Han, Zhilin Wang, Peng Ling, Wenjia Gong, Yang Liu, Guanjun Jiang, Yukun Cao, and Xike Xie. ImageRSI: Verification-driven recursive self-improvement for agentic image generation. *arXiv preprint arXiv:2610.32120*, 2026a.

Xiaofeng Han, Kaize Luo, Zenghuang Fu, Zichong Zhu, Zhe Feng, Changwei Wang, Rongtao Xu, Weiliang Meng, and Xiaopeng Zhang. Recursive self-improvement in physical AI: Mechanisms, evidence, and evaluation. *arXiv preprint arXiv:2610.38266*, 2026b.

Zhi Han, Zihan Guo, Chenxi Zeng, Zicai Cui, Yue Xin, Ming Zhou, Dongzhan Zhou, Weinan Zhang, and Yang Li. Embodied Skill Bench: Benchmarking how well agents with embodied skills perform on physical tasks. *arXiv preprint arXiv:2610.16238*, 2026c.

Jiaheng Hu, Roberto Martin-Martin, Peter Stone, Rocky Duan, Zhenyu Jiang, and Guanya Shi. SimEX: Simulation-integrated robotics AutoResearch. *arXiv preprint arXiv:2609.38982*, 2026.

Shengran Hu, Cong Lu, and Jeff Clune. Automated design of agentic systems. In *International Conference on Learning Representations*, volume 2025, pages 21344–21377, 2025.

Da Huang, Jichao Wang, Junhui Wu, Tong Ran, Linting Guo, Siying Zhang, Chuang Men, Sukai Wang, Yue Zhou, Maoqing Yao, and Mo Wu. Frontier LLMs enter the world of robotics: A cross-model evaluation on RoboColiseum. *arXiv preprint arXiv:2610.27865*, 2026a.

Jie Huang, Yanan Chen, Ruixun Liu, Kaichen He, Xinyi He, Rui Huang, Zilong Zheng, Zhenliang Zhang, and Yiwu Zhong. EmbodiedCLBench: Evaluating continual learning for self-evolving embodied agents. *arXiv preprint arXiv:2610.13021*, 2026b.

Kaiyi Huang, Kaiyue Sun, Enze Xie, Zhenguo Li, and Xihui Liu. T2I-CompBench: A comprehensive benchmark for open-world compositional text-to-image generation. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023a.

Qian Huang, Jian Vora, Percy Liang, and Jure Leskovec. MLAgentBench: Evaluating language agents on machine learning experimentation. *arXiv preprint arXiv:2310.03302*, 2023b.

Yue Huang, Yanbo Wang, Lichi Li, Yuanyuan Chen, Andreas Haupt, Xiaonan Luo, Zheyuan Liu, Jihye Choi, Felix Juefei-Xu, Xiangliang Zhang, Radha Poovendran, Misha Sra, Ying Liu, Ben Burtenshaw, Thomas Wolf, Alex Pentland, Zichen Chen, and Zhangchen Xu. RSI Arena: Towards recursive self-improvement with human feedback, 2026c. URL https://arxiv.org/abs/2610.33624.

Zhengyao Jiang, Dominik Schmidt, Dhruv Srikanth, Dixing Xu, Ian Kaplan, Deniss Jacenko, and Yuxiang Wu. AIDE: AI-driven exploration in the space of code. *arXiv preprint arXiv:2502.13138*, 2025.

Alibek Kaliyev and Artem Maryanskyy. Beyond task completion: A verification-vs.-conformance gap in tool-evolving agents, 2026. URL https://arxiv.org/abs/2604.00392. KDD 2026 Workshop on Agentic AI Evaluation and Trustworthiness.

Yifan Kang, Zihan Wang, Zhiwen Fan, and Bangya Liu. Encore: Few-shot agentic discovery of manipulation strategies. *arXiv preprint arXiv:2609.37359*, 2026.

Max Ku, Dongfu Jiang, Cong Wei, Xiang Yue, and Wenhu Chen. VIEScore: Towards explainable metrics for conditional image synthesis evaluation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12268–12290, 2024.

So Kuroki, Kento Kawaharazuka, and Yujin Tang. Learning and transferring closed-loop robot software. *arXiv preprint arXiv:2609.19906*, 2026.

Hyunin Lee, Jinglue Xu, Jeffrey Seely, Donghyun Lee, Matei Zaharia, and Yujin Tang. Recursive harness self-improvement. *arXiv preprint arXiv:2607.15524*, 2026a.

Yoonho Lee, Roshen Nair, Qizheng Zhang, Kangwook Lee, Omar Khattab, and Chelsea Finn. Meta-harness: End-to-end optimization of model harnesses. *arXiv preprint arXiv:2603.28052*, 2026b.

Yingtie Lei, Zhongwei Wan, Jiankun Zhang, Samiul Alam, Zixuan Zhong, Peizhou Huang, Xin Wang, Jingxuan Zhang, Donghao Zhou, Yunta Hsieh, et al. SkillEvolBench: Benchmarking the evolution from episodic experience to procedural skills. *arXiv preprint arXiv:2605.24117*, 2026.

Hao Li, Zhengyu Zou, Fangfu Liu, Xuanyang Zhang, Fangzhou Hong, Yukang Cao, Yushi Lan, Manyuan Zhang, Gang Yu, Dingwen Zhang, et al. IGGT: Instance-grounded geometry transformer for semantic 3D reconstruction. In *International Conference on Learning Representations*, 2026.

Susan Liang, Jianmin Wu, and Daxiang Dong. VidHarness: Evolving agent harnesses for cost-efficient long video understanding. *arXiv preprint arXiv:2609.38413*, 2026.

Zihao Lin, Haoyu Dong, Haobo Xu, Chi-Chih Chang, Xuying Ning, Yinglong Xia, Shuang Yang, Zhuokai Zhao, Xinying Song, Dongfang Liu, Lifu Huang, Yang Song, et al. VideoHarness-SRSI: Structured recursive self-improvement for video understanding harness. *arXiv preprint arXiv:2610.19424*, 2026.

Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. LIBERO: Benchmarking knowledge transfer for lifelong robot learning. *Advances in Neural Information Processing Systems*, 36:44776–44791, 2023.

Fangcheng Liu, Yeqing Shen, Anda Cheng, Weishi Mi, Chao Tang, Chenyuan Liu, Yushun Xiang, Tingguang Li, Yong-Lu Li, and Yehui Tang. RoboICL: Embodied in-context learning with GPT-6 Astra. *arXiv preprint arXiv:2609.34261*, 2026a.

Haotian Liu, Chunyuan Li, Yuheng Li, Bo Li, Yuanhan Zhang, Sheng Shen, and Yong Jae Lee. LLaVA-NeXT: Improved reasoning, OCR, and world knowledge. Blog post, 2024a. URL https://llava-vl.github.io/blog/2024-01-30-llava-next/.

Ruiping Liu, Junwei Zheng, Yufan Chen, Di Wen, Shaofang Quan, Chengzhi Wu, Jiaming Zhang, Kailun Yang, Kunyu Peng, and Rainer Stiefelhagen. EgoExoMem: Cross-view memory reasoning over synchronized egocentric and exocentric videos. *arXiv preprint arXiv:2605.18734*, 2026b.

Shilong Liu et al. Grounding DINO: Marrying DINO with grounded pre-training for open-set object detection. In *European Conference on Computer Vision (ECCV)*, 2024b.

Shiyu Liu, Yucheng Han, et al. Step1X-Edit: A practical framework for general image editing. *arXiv preprint arXiv:2504.17761*, 2025.

David G. Lowe. Distinctive image features from scale-invariant keypoints. *International Journal of Computer Vision*, 60(2): 91–110, 2004.

Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI Scientist: Towards fully automated open-ended scientific discovery. *arXiv preprint arXiv:2408.06292*, 2024.

Bingjun Luo, Jialin Guo, and Siqi Li. Video-RSI: Recursive self-improvement of video understanding agents via harness evolution. *arXiv preprint arXiv:2609.37950*, 2026.

Fanqing Meng, Lingxiao Du, Qiguang Chen, Ziqi Zhao, Haocheng Lu, Mengkang Hu, and Michael Qizhe Shieh. RSIBench-Data: Benchmarking data-centric research for recursive self-improvement. *arXiv preprint arXiv:2607.25886*, 2026.

Meta. Introducing Muse Spark 1.3. https://research.meta.ai/blog/introducing-muse-spark-1-3, September 2026a.

Meta. Introducing Muse Wonder and Muse Wonder Full: Frontier models for agents. https://about.fb.com/news/2026/10/introducing-muse-wonder-models/, October 2026b.

MiniMax. MiniMax-M3. Model card, https://huggingface.co/MiniMaxAI/MiniMax-M3, 2026a.

MiniMax. MiniMax H3. Hugging Face model card, 2026b. URL https://huggingface.co/MiniMaxAI/MiniMax-H3.

Soroush Nasiriany, Abhiram Maddukuri, Lance Zhang, Adeet Parikh, Aaron Lo, Abhishek Joshi, Ajay Mandlekar, and Yuke Zhu. RoboCasa: Large-scale simulation of everyday tasks for generalist robots. In *Robotics: Science and Systems*, 2024.

Soroush Nasiriany, Sepehr Nasiriany, Abhiram Maddukuri, and Yuke Zhu. RoboCasa365: A large-scale simulation framework for training and benchmarking generalist robots. In *International Conference on Learning Representations*, volume 2026, pages 98643–98667, 2026.

Alexander Novikov, Ngân Vũ, Marvin Eisenberger, Emilien Dupont, Po-Sen Huang, Adam Zsolt Wagner, Sergey Shirobokov, Borislav Kozlovskii, Francisco JR Ruiz, Abbas Mehrabian, et al. AlphaEvolve: A coding agent for scientific and algorithmic discovery. *arXiv preprint arXiv:2506.13131*, 2025.

OpenAI. GPT-5.6 system card. https://deploymentsafety.openai.com/gpt-5-6, July 2026a.

OpenAI. GPT-6 Astra system card. https://deploymentsafety.openai.com/gpt-6-astra, September 2026b.

OpenAI. Addendum: GPT-6.1 Sol system card. https://cdn.openai.com/pdf/38e3efcf-545e-44cd-99ec-2b7eb395f4cc/oai_GPT_6_1_Sol.pdf, September 2026c.

OpenAI. Codex CLI. Official documentation, 2026d. URL https://developers.openai.com/codex/cli/. Accessed October 29, 2026.

OpenRSI Foundation. OpenRSI Index. https://index.openrsi.foundation, 2026.

Federico Perazzi, Jordi Pont-Tuset, Brian McWilliams, Luc Van Gool, Markus Gross, and Alexander Sorkine-Hornung. A benchmark dataset and evaluation methodology for video object segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 724–732, 2016.

Li Puyin, Tiange Xiang, Ella Mao, Shirley Wei, Xinye Chen, Adnan Masood, Li Fei-Fei, and Ehsan Adeli. QuantiPhy: A quantitative benchmark evaluating physical reasoning abilities of vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 33174–33184, 2026.

Cheng Qian, Kunlun Zhu, Beibin Li, Zhenhailong Wang, and Heng Ji. Learning meta-skills for agent harness design in test-time AI4AI. *arXiv preprint arXiv:2609.38143*, 2026.

Leigang Qu, Feng Cheng, Ziyan Yang, Qi Zhao, Shanchuan Lin, Yichun Shi, Yicong Li, Wenjie Wang, Tat-Seng Chua, and Lu Jiang. VINCIE: Unlocking in-context image editing from video. In *International Conference on Learning Representations*, volume 2026, pages 114873–114918, 2026.

Qwen Team. Qwen3.6-35B-A3B: Agentic coding power, now open to all. https://qwen.ai/blog?id=qwen3.6-35b-a3b, 2026a.

Qwen Team. Qwen3.8-Max: A new bar for coding and cowork. https://qwen.ai/blog?id=qwen3.8, August 2026b.

Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, volume 139 of *Proceedings of Machine Learning Research*, pages 8748–8763, 2021.

Anirudha Ramesh, Siddartha Devic, Shivank Garg, Advait Parulekar, Drish Mahnot, Shreyas Pimpalgaonkar, Vishnu Suresh, Alex Dimakis, and Maheswaran Sathiamoorthy. AutoResearchExam, September 2026. URL https://benchmarks.bespokelabs.ai/autoresearchexam/. Published September 9, 2026.

Ben Rank, Hardik Bhatnagar, Ameya Prabhu, Shira Eisenberg, Karina Nguyen, Matthias Bethge, and Maksym Andriushchenko. PostTrainBench: Can LLM agents automate LLM post-training? *arXiv preprint arXiv:2603.08640*, 2026.

Rapidata. Rapidata Video Generation Veo 3.1 Human Preference. Hugging Face dataset card, 2025. URL https://huggingface.co/datasets/Rapidata/text-2-video-human-preferences-veo3.1.

Nikhila Ravi, Valentin Gabeur, Yuan-Ting Hu, Ronghang Hu, Chaitanya Ryali, Tengyu Ma, Haitham Khedr, et al. SAM 2: Segment anything in images and videos. *arXiv preprint arXiv:2408.00714*, 2024.

Maxime Robeyns, Martin Szummer, and Laurence Aitchison. A self-improving coding agent. *arXiv preprint arXiv:2504.15228*, 2025.

RSI-Exam Team. RSI-Exam: Benchmarking recursive self-improvement through executable research. https://github.com/aiming-lab/RSI-Exam, 2026.

SAM 3D Team, Xingyu Chen, Fu-Jen Chu, Pierre Gleize, Kevin J. Liang, Alexander Sax, et al. SAM 3D: 3Dfy anything in images. *arXiv preprint arXiv:2511.16624*, 2025.

Bokui Shen, Fei Xia, Chengshu Li, Roberto Martín-Martín, Linxi Fan, Guanzhi Wang, Claudia Pérez-D'Arpino, Shyamal Buch, Sanjana Srivastava, Lyne P. Tchapmi, Micael E. Tchapmi, Kent Vainio, Josiah Wong, Li Fei-Fei, and Silvio Savarese. iGibson 1.0: A simulation environment for interactive tasks in large realistic scenes. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2021.

Haozhan Shen, Kangjia Zhao, Tiancheng Zhao, Ruochen Xu, Zilun Zhang, Mingwei Zhu, and Jianwei Yin. ZoomEye: Enhancing multimodal LLMs with human-like zooming capabilities through tree-based image exploration. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025.

Ruihong Shen, Žiga Kovačič, Peter Kulits, Xingrui Wang, Zizhang Li, Joshua B. Tenenbaum, Alan Yuille, Jieneng Chen, and Jiajun Wu. 4DCodeBench: Benchmarking agents on inverse graphics of dynamic scenes. *arXiv preprint arXiv:2610.03715*, 2026a.

Zeyu Shen, Haoxiang You, Yilang Liu, Zhicheng Zheng, Lihan Zha, Kashu Yamazaki, Mingtong Zhang, Suning Huang, Jiankai Sun, Qianzhong Chen, et al. EmbodiedSWE: Coding agents for long horizon dexterous robotics. *arXiv preprint arXiv:2609.27308*, 2026b.

Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *Advances in neural information processing systems*, 36:8634–8652, 2023.

Python Song, Zhixuan Liang, Kelsey Fu, Mengdi Wang, Junfeng Yang, and Shilong Liu. EmbodiedRSI: Active continual robot learning through hypothesis-guided co-evolution. *arXiv preprint arXiv:2610.10498*, 2026.

Charles Spearman. The proof and measurement of association between two things. *The American Journal of Psychology*, 15 (1):72–101, 1904.

Dhruv Srikanth, Bingchen Zhao, Dixing Xu, Yuxiang Wu, and Zhengyao Jiang. Recursive self-improvement of AI research agents. *arXiv preprint arXiv:2609.26457*, 2026.

Trevor Standley, Ozan Sener, Dawn Chen, and Silvio Savarese. image2mass: Estimating the mass of an object from its image. In *Proceedings of the 1st Annual Conference on Robot Learning*, volume 78 of *Proceedings of Machine Learning Research*, pages 324–333, 2017.

Giulio Starace, Oliver Jaffe, Dane Sherburn, et al. PaperBench: Evaluating AI's ability to replicate AI research. *arXiv preprint arXiv:2504.01848*, 2025.

Julian Straub, Thomas Whelan, Lingni Ma, Yufan Chen, Erik Wijmans, Simon Green, Jakob J Engel, Raul Mur-Artal, Carl Ren, Shobhit Verma, et al. The replica dataset: A digital replica of indoor spaces. *arXiv preprint arXiv:1906.05797*, 2019.

Kaiyue Sun, Kaiyi Huang, Xian Liu, Yue Wu, Zihan Xu, Zhenguo Li, and Xihui Liu. T2V-CompBench: A comprehensive benchmark for compositional text-to-video generation. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 8406–8416. IEEE, 2025.

Shufan Sun, Chen Wang, Enxin Song, Jiatao Gu, and Lingjie Liu. HARMONY: Hierarchical agentic reasoning for MONocular image-to-scene synthesis. *arXiv preprint arXiv:2609.26793*, 2026.

Tencent Hunyuan Foundation Model Team. HunyuanImage 3.0 technical report. *arXiv preprint arXiv:2509.23951*, 2025.

Wenqing Tian, Zeyu Zhang, Zhaocheng Liu, Fengwei Liu, Qiang Liu, and Liang Wang. PhysEvo: Astra can act, let it. *arXiv preprint arXiv:2610.08995*, 2026a.

Yuan Tian, Bing Hu, Hao Wang, Binghang Lu, and Fang Wu. Harness evolution hits a ceiling: When weight training should begin. *arXiv preprint arXiv:2610.11655*, 2026b.

Emanuel Todorov, Tom Erez, and Yuval Tassa. MuJoCo: A physics engine for model-based control. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 5026–5033, 2012.

Zhengzhong Tu, Yilin Wang, Neil Birkbeck, Balu Adsumilli, and Alan C. Bovik. UGC-VQA: Benchmarking blind video quality assessment for user generated content. *IEEE Transactions on Image Processing*, 30:4449–4464, 2021.

Florin-Alexandru Vasluianu, Tim Seizinger, and Radu Timofte. WSRD: A novel benchmark for high resolution image shadow removal. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 1826–1835, 2023.

Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023.

Jiaming Wang. What stops recursive self-improvement in robotics? lessons from 123 rounds of agentic skill discovery. *arXiv preprint arXiv:2609.31760*, 2026.

Kailin Wang, Haoxiang Jie, Yaoyuan Yan, Zhiyou Heng, and Zhaosong Li. RegenHarness: A robot agent harness with evidence-gated recursive self-improvement. *arXiv preprint arXiv:2609.27612*, 2026a.

Ke Wang, Zijie Zhao, Zhiyi Yuan, and Changlun Li. When better gets worse: Improvement fidelity for self-improving agents in adaptive worlds. *arXiv preprint arXiv:2609.32677*, 2026b.

Peidong Wang, Ying Chang, Zhiming Ma, Xiaocui Yang, Shi Feng, Daqi Liu, Yuzhe Yang, Xu Cao, Chuang Wang, Jian Su, and Daling Wang. AdaRSI: On-demand recursive self-improvement for embodied agents. *arXiv preprint arXiv:2610.20474*, 2026c.

Wenbin Wang, Liang Ding, Minyan Zeng, Xiabin Zhou, Li Shen, Yong Luo, Wei Yu, and Dacheng Tao. Divide, conquer and combine: A training-free framework for high-resolution image perception in multimodal large language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 7907–7915, 2025a.

Wenyi Wang, Piotr Piękos, Li Nanbo, Firas Laakom, Yimeng Chen, Mateusz Ostaszewski, Mingchen Zhuge, and Jürgen Schmidhuber. Huxley-Gödel Machine: Human-level coding agent development by an approximation of the optimal self-improving machine. In *International Conference on Learning Representations*, volume 2026, pages 80356–80386, 2026d.

Xingyao Wang, Boxuan Li, Yufan Song, Frank F Xu, Xiangru Tang, Mingchen Zhuge, Jiayi Pan, Yueqi Song, Bowen Li, Jaskirat Singh, et al. OpenHands: An open platform for AI software developers as generalist agents. In *International Conference on Learning Representations*, volume 2025, pages 65882–65919, 2025b.

Yihao Wang, Linhan Xia, Rui Liu, Zhaofeng Zhang, Hongyu Wu, Yang Yang, Jinglu He, Yu Guo, and Kai Lei. SAGE: A statistical acceptance gate for self-evolving agents. *arXiv preprint arXiv:2609.36043*, 2026e.

Yilin Wang, Sasi Inguva, and Balu Adsumilli. YouTube UGC dataset for video compression research. In *2019 IEEE 21st international workshop on multimedia signal processing (MMSP)*, pages 1–5. IEEE, 2019.

Zhou Wang, Alan C. Bovik, Hamid R. Sheikh, and Eero P. Simoncelli. Image quality assessment: From error visibility to structural similarity. *IEEE Transactions on Image Processing*, 13(4):600–612, 2004.

Songlin Wei, Zhenhao Ni, Jie Liu, Zhenyu Zhao, Junjie Ye, Hongyi Jing, Junkai Xia, Xiawei Liu, Michael Leong, Liang Heng, Di Huang, and Yue Wang. SIMPLE: Simulation-based policy learning and evaluation for humanoid loco-manipulation. *arXiv preprint arXiv:2606.08278*, 2026a.

Xinyue Wei, Minghua Liu, Zhan Ling, and Hao Su. Approximate convex decomposition for 3D meshes with collision-aware concavity and tree search. *ACM Transactions on Graphics*, 41(4), 2022. Article 42.

Zhijie Wei, Ferris Tan, Jinghui Wang, and Hao Zhou. Scale and selection: What makes automatic harness evolution work for visual-interface robot agents. *arXiv preprint arXiv:2609.39304*, 2026b.

Zimo Wen, Yijin Chen, Yuxuan Cao, Wendi Chen, Yanwen Zou, Wenye Yu, Fuhang Kuang, Han Xue, Jun Lv, Chuan Wen, and Cewu Lu. RoboRSI: Stable, efficient, and reusable robot self-evolution in complex real-world environments. *arXiv preprint arXiv:2610.12424*, 2026.

Hjalmar Wijk, Tao Lin, Joel Becker, et al. RE-Bench: Evaluating frontier AI R&D capabilities of language model agents against human experts. *arXiv preprint arXiv:2411.15114*, 2024.

Chenfei Wu et al. Qwen-Image technical report. *arXiv preprint arXiv:2508.02324*, 2025a.

Diankun Wu, Fangfu Liu, Yi-Hsin Hung, and Yueqi Duan. Spatial-MLLM: Boosting MLLM capabilities in visual-based spatial intelligence. In *Advances in Neural Information Processing Systems*, 2025b.

Haoning Wu, Erli Zhang, Liang Liao, Chaofeng Chen, Jingwen Hou, Annan Wang, Wenxiu Sun, Qiong Yan, and Weisi Lin. Exploring video quality assessment on user generated contents from aesthetic and technical perspectives. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 20144–20154, 2023.

Penghao Wu and Saining Xie. V*: Guided visual search as a core mechanism in multimodal LLMs. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 13084–13094. IEEE, 2024.

xAI. Model card: Grok 4.6. https://media.x.ai/v1/website/card-4p6-4cd2dc57.pdf, August 2026.

Hongchi Xia, Chih-Hao Lin, Hao-Yu Hsu, Quentin Leboutet, Katelyn Gao, Michael Paulitsch, Benjamin Ummenhofer, and Shenlong Wang. HoloScene: Simulation-ready interactive 3D worlds from a single video. *Advances in Neural Information Processing Systems*, 38:32501–32524, 2025.

Peng Xia, Rujun Han, Zifeng Wang, Yanfei Chen, Yufan Zhuang, Yoonho Lee, Chengsong Huang, Han Yu, Zhongying CuiZhu, Yifei Ming, Huaxiu Yao, Burak Gokturk, Tomas Pfister, and Chen-Yu Lee. RRSI: Regularized recursive self-improvement of agent harnesses. *arXiv preprint arXiv:2609.24972*, 2026.

Shuzhao Xie, Lelin Wang, Guying Lin, Zhi Wang, and Minchen Li. CoDimRecon: Agentic reconstruction of sim-ready 3D scenes with deformable curves, surfaces, and volumes. *arXiv preprint arXiv:2609.36024*, 2026.

Shuhan Xue, Zixin Ding, Yichen Shen, Yinjie Wang, Zhenfei Yin, Yingcheng Wu, Yuxin Chen, Mengdi Wang, and Ling Yang. PAST-Bench: Benchmarking the foundations of recursive self-improvement in personal agents. *arXiv preprint arXiv:2608.04003*, 2026.

Jihan Yang, Shusheng Yang, Anjali W Gupta, Rilyn Han, Li Fei-Fei, and Saining Xie. Thinking in space: How multimodal large language models see, remember, and recall spaces. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10632–10643. IEEE, 2025.

John Yang, Carlos Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. SWE-agent: Agent-computer interfaces enable automated software engineering. *Advances in Neural Information Processing Systems*, 37:50528–50652, 2024.

Xuning Yang, Rishit Dagli, Alex Zook, Hugo Hadfield, Ankit Goyal, Stan Birchfield, Fabio Ramos, and Jonathan Tremblay. RoboLab: A high-fidelity simulation benchmark for analysis of task generalist policies. In *Robotics: Science and Systems*, 2026a.

Zhiqin Yang, Chenxin Li, Xiaomeng Hu, Yibin Liu, Weidong Huang, Jiankai Sun, Haitao Li, Zijian Wu, Yuzhi Huang, Fanding Huang, Hanwen Sun, Jiashun Liu, et al. RobotWorld: Benchmarking multimodal agents for robot use across diverse tasks and embodiments. *arXiv preprint arXiv:2610.10409*, 2026b.

Haoyu Yao, Chenguo Lin, Ruize Du, and Yadong Mu. Learning to code the world: Harness evolution for compositional 3D reconstruction. *arXiv preprint arXiv:2610.37046*, 2026a.

Kaixin Yao, Longwen Zhang, Xinhao Yan, Yan Zeng, Qixuan Zhang, Lan Xu, Wei Yang, Jiayuan Gu, and Jingyi Yu. CAST: Component-aligned 3D scene reconstruction from an RGB image. *ACM Transactions on Graphics*, 44(4), 2025. doi: 10.1145/3730841.

Kehui Yao, Chuanbo Hu, Jiao Chen, Ramin Giahi, Reza Yousefi Maragheh, Lalitesh Morishetti, Ziwei Hong, Yiyun Su, Jianpeng Xu, Evren Korpeoglu, and Kannan Achan. An explicit feedback contract for iterative ML-Agent harnesses. *arXiv preprint arXiv:2610.28798*, 2026b.

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023.

Yifei Yao, Yude Zou, Guanbo Jiao, Xing Gao, Ning Gao, Yukai Wang, Jiantong Chen, Zihou Zhu, Xinyu Li, Ningyu Yan, Shuai Wang, Jianming Ma, et al. Profiling a frontier agent with EBench: From direct control to harness with policy. *arXiv preprint arXiv:2610.35832*, 2026c.

Chandan Yeshwanth, Yueh-Cheng Liu, Matthias Nießner, and Angela Dai. ScanNet++: A high-fidelity dataset of 3D indoor scenes. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 12–22. IEEE, 2023.

Joong Gon Yim, Yilin Wang, Neil Birkbeck, and Balu Adsumilli. Subjective quality assessment for youtube ugc dataset. In *2020 IEEE International Conference on Image Processing (ICIP)*, pages 131–135. IEEE, 2020.

Xunjian Yin, Xinyi Wang, Liangming Pan, Li Lin, Xiaojun Wan, and William Yang Wang. Gödel Agent: A self-referential agent framework for recursive self-improvement. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 27890–27913, 2025.

Heyang Yu, Yinan Han, Xiangyu Zhang, Baiqiao Yin, Bowen Chang, Xiangyu Han, Xinhao Liu, Jing Zhang, Marco Pavone, Chen Feng, Saining Xie, and Yiming Li. Thinking in 360°: Humanoid visual search in the wild. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026a. URL https://arxiv.org/abs/2511.20351.

Heyang Yu, Haomin Li, Sihang Li, Jixin Yan, Shiyun Peng, Wanqi Zhong, Tianyu Xin, Yuan Yao, Xiaoyu Tian, and Yiming Li. RoboDoctor: Training-free recursive self-improvement for robot policies. *arXiv preprint arXiv:2610.35604*, 2026b.

Eric Zelikman, Eliana Lorch, Lester Mackey, and Adam Tauman Kalai. Self-taught optimizer (STOP): Recursively self-improving code generation. In *Conference on Language Modeling*, 2024.

Hangfan Zhang, Shao Zhang, Kangcong Li, Chen Zhang, Yang Chen, Yiqun Zhang, Lei Bai, and Shuyue Hu. Self-harness: Harnesses that improve themselves. *arXiv preprint arXiv:2606.09498*, 2026a.

Jenny Zhang, Shengran Hu, Cong Lu, Robert Lange, and Jeff Clune. Darwin Gödel Machine: Open-ended evolution of self-improving agents, 2025a. URL https://arxiv.org/abs/2505.22954.

Jenny Zhang, Bingchen Zhao, Wannan Yang, Jakob Foerster, Jeff Clune, Minqi Jiang, Sam Devlin, and Tatiana Shavrina. Hyperagents. *arXiv preprint arXiv:2603.19461*, 2026b.

Jiayi Zhang, Jinyu Xiang, Zhaoyang Yu, Fengwei Teng, Xionghui Chen, Jiaqi Chen, Mingchen Zhuge, Xin Cheng, Sirui Hong, Jinlin Wang, et al. AFlow: Automating agentic workflow generation. In *International Conference on Learning Representations*, volume 2025, pages 34040–34077, 2025b.

Jingdong Zhang, Xin Li, Jan Kautz, Wenping Wang, and Chris Choy. S4VY: Segment anything in feed-forward 4D visual geometry. *arXiv preprint arXiv:2609.36875*, 2026c.

Jingdong Zhang, Yizhou Wang, Zhengzhong Tu, Xin Li, Wenping Wang, and Xiaohang Zhan. Beyond thinking: Imagining in 360° for humanoid visual search. *arXiv preprint arXiv:2605.09146*, 2026d.

Jingdong Zhang et al. UniSER: A foundation model for unified soft effects removal. *arXiv preprint arXiv:2511.14183*, 2025c.

Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 586–595, 2018.

Yunbei Zhang, Zijian Jin, Yuanzhe Liu, Janet Wang, Xilun Zhang, Yuyou Zhang, Zhenyu Zhang, Daoan Zhang, Shuaicheng Niu, Gen Li, et al. LIBERO-MAX: Do robot policies adapt when the world changes? *arXiv preprint arXiv:2609.36518*, 2026e.

Zexi Zhang, Zecheng Zhu, Zidong Chen, Zulkhuu Tuya, and Stephen James. BiGym 2.0: Benchmarking learned and agent-developed policies for humanoid household manipulation. *arXiv preprint arXiv:2610.07594*, 2026f.

Jinliang Zheng, Jianxiong Li, Zhihao Wang, Dongxiu Liu, Xirui Kang, Yuchun Feng, Yinan Zheng, Jiayin Zou, Yilun Chen, Jia Zeng, Ya-Qin Zhang, Jiangmiao Pang, Jingjing Liu, Tai Wang, and Xianyuan Zhan. X-VLA: Soft-prompted transformer as scalable cross-embodiment vision-language-action model. *arXiv preprint arXiv:2510.10274*, 2025.

Lianmin Zheng et al. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Ziwei Zheng, Michael Yang, Jack Hong, Chenxiao Zhao, Guohai Xu, Le Yang, Chao Shen, and Xing Yu. DeepEyes: Incentivizing "thinking with images" via reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2026.

Hai Zhou, Fan Yang, Zeyi Chen, Lejun Zhang, Chenyang Fu, Jianping Li, Jiayi Zhu, Xianyang Liu, Yingxuan Yang, Lili Meng, Lu Wang, Xinyang Han, et al. Who improves the robot? toward physical recursive self-improvement. *arXiv preprint arXiv:2610.18842*, 2026a.

Tailin Zhou. Hierarchical self-improvement: A framework for task-specific evolvable agent harnesses. *arXiv preprint arXiv:2608.08466*, 2026.

Xueyang Zhou, Yangming Xu, Guiyao Tie, Yongchao Chen, Guowen Zhang, Duanfeng Chu, Pan Zhou, and Lichao Sun. LIBERO-PRO: Towards robust and fair evaluation of vision-language-action models beyond memorization. *arXiv preprint arXiv:2510.03827*, 2025.

Youwei Zhou, Bohan Hou, Yankai Fu, Yaoxu Lyu, Yifan Ye, Shilin Shan, Chen Guo, Jiaming Liu, Ran He, Caifeng Shan, Shanghang Zhang, Jianfei Yang, et al. Large multimodal models for direct robot control: GPT-6 as a beginning. *arXiv preprint arXiv:2610.28259*, 2026b.

Zhengyu Zou, Hao Li, Kuixuan Jiao, Liu Liu, Tingyang Xiao, Xiaolin Zhou, Fangzhou Hong, Zhizhong Su, Dingwen Zhang, and Ziwei Liu. IGGT4D: Streaming 4D instance-grounded geometry transformer. *arXiv preprint arXiv:2607.19228*, 2026.

A Task cards

This appendix describes each of the 24 tasks of Table 1: what it measures, how its data were built from its source and divided into splits, what the agent has to produce and with which resources, the starter solution, the metric and where it comes from, and the prior work on the source benchmark. Every task follows the setting of Section 3.1: the development split is available to the agent with its labels, while the validation and test splits are withheld. Unless stated otherwise, the solution may call the model being evaluated within the stated limits, and it may not use the annotations of the source benchmark.

A.1 Action

RoboDojo

Purpose. Whether the agent can turn a general multimodal model into a closed-loop policy for a dual-arm robot. No policy is trained: the model decides the motions at evaluation time, and the agent improves the program that queries it.

Source and construction. Ten simulation tasks of RoboDojo (Chen et al., 2026d). Full credit (success) requires: four bottles in a dustbin, using a handover between the hands; a garment folded with sleeves in and hem aligned to the shoulders; digit tiles placed in the order forming the largest number; three tubes inserted upright into a rack; all five own pieces placed in tic-tac-toe against a randomly playing opponent arm; a T-shaped block pushed onto its pad without lifting; eight xylophone keys struck left to right with a mallet; the conveyor object shown on a board placed in a lifted basket; four eggs in a holder with its lid closed; and five nesting dolls lined up by size. Several tasks also require a return to the start pose. All layouts come from RoboDojo's evaluation seed 0: per task, validation takes the first five, test the next five, and development the rest, of which a fixed panel of five is evaluated. Validation layouts of three tasks are those of Galbot Team et al. (2026). Tasks with a randomized variant (fold clothes, largest number, Push-T, nesting dolls) use two standard and three randomized layouts per split. The agent sees full development results, the validation score only, and no test results.

Setting. An ARX X5 dual-arm robot in Isaac Sim at 25 Hz, run through the pinned single-layout RoboDojo client of RoboICL (Liu et al., 2026a), 500 to 1100 steps per episode, starting in an elbow-up pose slightly in front of RoboDojo's. The policy receives three RGB cameras, proprioception, and the head-camera calibration, never object poses or rewards, and returns chunks of at most 64 end-effector or joint targets, limited to 3 cm and 0.15 rad per step. It may call the run's model at most 80 times per episode and no other model; scripted motions and its own inverse kinematics are allowed.

Starter solution. One model call per chunk of up to ten end-effector motions, given the head-camera image at 320×240, the robot state, and its notes from earlier calls.

Metric. Each episode receives the progress score $s \in [0, 100]$ of its RoboDojo rubric (for example 10, 25, 40, 100 for one to four bottles), as defined in the RoboDojo task documentation (Chen et al., 2026d), 100 only on success, scored as the scene stands when it ends. The task score is $S_t = \frac{1}{5} \sum_e s_{t,e}$ over the test layouts and the suite score $\frac{1}{10} \sum_t S_t$, a task without a run counting as zero. The success rate is also reported.

Prior work. RoboDojo has 42 simulation tasks and evaluates 30 trained manipulation policies (Chen et al., 2026d). Frontier models have also been used as the policy: Galbot Team et al. (2026) reach 48% success on a RoboDojo subset with GPT-6 Astra combined with $\pi_{0.5}$; RoboICL (Liu et al., 2026a) conditions the model on demonstrations and interaction memory (overall score 50.64 on 30 tasks against 33.68 for the strongest baseline); and Encore (Kang et al., 2026) writes policy programs from a few demonstrations.

RoboCasa

Purpose. Whether the agent can obtain a model-driven policy for a mobile manipulator that carries over from the kitchens and objects used in development to kitchens and objects it has never seen.

Source and construction. Eight tasks of RoboCasa365 (Nasiriany et al., 2026), an extension of RoboCasa (Nasiriany et al., 2024), with RoboCasa's own success checks. Six are atomic tasks from RoboCasa365's seen-task target set: turn on the microwave (on, gripper away from the button), turn on the sink faucet (handle turned

over 0.4 rad), open the named cabinet (every door at least 90% open), open the named drawer (at least 95% open), place a mug under the coffee dispenser, and turn off the named burner. Two are composite: portion hot dogs (from its unseen composite set; one bun and one sausage on each of two plates) and kettle boiling (kettle within 15 cm of a burner that is on). An episode is fixed by a kitchen (layout and style), object instances, and a seed from which RoboCasa places the robot and objects and words the instruction. Development has five fixed episodes in RoboCasa's pretraining kitchens with pretraining objects; validation and test have five episodes each in the same five target kitchens (the first five of RoboCasa's ten evaluation kitchens) with target objects, differing by seed; validation and test were reduced from ten to five episodes each, and every episode was checked to build and reset. The agent sees full development results, the validation success rate only, and no test results; neither it nor the policy may load target kitchens or objects.

Setting. The PandaOmron robot (a Franka Panda arm on a mobile base with a torso lift) in MuJoCo at 20 Hz, with the per-task horizons of RoboCasa's dataset registry (450 to 2250 steps). The policy receives three 512×512 RGB cameras, proprioception, and calibration, never object states, and returns chunks of at most 64 rows of RoboCasa's 12-dimensional action. It may call the run's model at most 80 times per episode and no other model.

Starter solution. One model call per chunk of up to ten actions, given one base-camera image at 256×256, its calibration, the robot state, and its notes from earlier calls.

Metric. An episode succeeds, and ends, as soon as the task's success check of RoboCasa365 (Nasiriany et al., 2026) holds after a step; otherwise, including when the policy stops early, it counts as a failure. The task score is the test success rate $\mathrm{SR}_t = \frac{1}{5}\sum_e \mathbb{1}[\mathrm{success}_{t,e}]$ and the suite score $\frac{1}{8}\sum_t \mathrm{SR}_t$, a task without a run counting as zero.

Prior work. RoboCasa365 trains and benchmarks generalist policies, evaluating in target kitchens separate from the pretraining ones (Nasiriany et al., 2026). Galbot Team et al. (2026) evaluate GPT-6 Astra as a direct and hybrid policy on it, and Kuroki et al. (2026) optimize closed-loop control code, raising success on four RoboCasa source tasks from 28.3% to 64.2% and reaching 57.0% on nine transfer tasks.

Active Search

Purpose. Active visual search: a humanoid robot that sees part of a 360-degree scene turns its head to find a named object (object search) or the direction to move in (path search), with as few movements as possible; most targets start out of view.

Source and construction. H*Bench (Yu et al., 2026a) annotates search tasks on real panoramas of transport hubs, shops and streets: an instruction, a target yaw range (and, for objects, a pitch range), and starting directions and difficulty levels; its public benchmark split has 2,400 object-search and 1,600 path-search episodes. Development and validation come from it: each task becomes one episode, with one start drawn by a fixed seed from its annotated starts; each panorama yields at most one episode, so the splits share no scene; both are stratified by search kind and difficulty in proportion to H*Bench. Development has 80 episodes (40 object, 40 path) and validation 100 (50 and 50). The test split uses the non-public in-the-wild panoramas of Imagining in 360° (Zhang et al., 2026d), eight from Google Street View and seven generated by image models. A target was annotated by clicking its center (objects, widened by ±30° in yaw and ±20° in pitch) or the direction to move in (±10°), with starts at multiples of 90°; one start per task gives 47 episodes (33 object, 14 path) in 15 scenes. The agent sees the development episodes with their annotations and panoramas.

Setting. The solution receives the instruction and a 100-degree view (960×720 pixels), never the panorama. With at most ten actions it rotates the head by a relative angle or submits an absolute direction, which ends the episode; no submission fails. As in H*Bench, success requires the submitted direction to lie in the target ranges widened to at least the tolerances above. Up to 20 model calls per episode on average and open models on one GPU are allowed; other H*Bench data are not.

Starter solution. The H*Bench baseline with its prompts: one model call per step on the whole conversation (every view with the environment's feedback), answered with one rotation or submission.

Metric. Success weighted by path length (SPL) (Anderson et al., 2018), counting actions: episode e scores $s_e\,\ell_e/\max(\ell_e, p_e)$, with $s_e \in \{0, 1\}$ success, p_e the actions taken (including the submission), and ℓ_e the fewest

that can succeed (1 if the target center lies within 50° of the starting yaw, else 2). Scores are averaged within each search kind, and the two averages are averaged.

Prior work. H*Bench reports about 30% success for the best proprietary models; post-training Qwen2.5-VL-3B (HVS-3B) with supervised fine-tuning and multi-turn reinforcement learning raises success from 14.83% to 47.38% (objects) and from 6.44% to 24.94% (paths) (Yu et al., 2026a). Imagining in 360° adds an Imaginator that predicts the layout of seen and unseen regions for a frozen actor, raising HVS-3B from 48.04% to 62.75% and from 24.12% to 39.38% on H*Bench (Zhang et al., 2026d).

RoboLab

Purpose. Whether the agent can obtain one tabletop manipulation policy that follows instructions for tasks it has never seen; the splits contain different tasks, so the score measures generalization across tasks.

Source and construction. The pool is every RoboLab-120 (Yang et al., 2026a) task of difficulty simple or moderate with a horizon of at most 60 s (73 tasks). Each task is given one group of RoboLab's taxonomy: procedural (affordance, reorientation, sorting, stacking), else relational (conjunction, counting, spatial), else visual (color, semantics, size). Tasks are shuffled within each (group, difficulty) stratum with a fixed seed and dealt by systematic sampling to development (16 tasks), validation (12), test (24), and an unused remainder, so every split has each stratum in proportion; test has 4 procedural, 9 relational, 10 visual, and 1 unattributed task. Development tasks include stacking bowls, packing cans into a bin, standing a jug upright, and selecting a hammer by color. Each task keeps its RoboLab scene, default instruction, horizon, success predicate, and subtask tracker unchanged. Start states are deterministic, so a task's episodes differ only through the policy; development runs one episode per task, validation and test two (16, 24, and 48 episodes). The agent sees development tasks and results, validation aggregates only, and no held-out task names.

Setting. The DROID setup, a Franka Panda arm with a Robotiq 2F-85 gripper, in Isaac Lab at 15 Hz. The policy receives the instruction, over-the-shoulder and wrist RGB images, over-the-shoulder depth, calibration, and proprioception, never the task name or object poses, and returns chunks of absolute end-effector poses and gripper commands, limited to 5 cm and 0.25 rad per step. It may call the run's model at most 40 times within 1200 s of thinking per episode, no other model, and may not key behavior on task identity.

Starter solution. One model call per gripper waypoint, given both images, the pose, and calibration, followed by a straight-line move and an open or close.

Metric. RoboLab's total score (Yang et al., 2026a), Score $= \frac{100}{N}\sum_{e=1}^{N}\left(\mathbb{1}[\text{success}_e] + (1-\mathbb{1}[\text{success}_e])\,q_e\right)$ over the $N = 48$ test episodes, where $q_e \in [0,1]$ is the fraction of subtasks completed when the episode ends. The success rate is also reported.

Prior work. RoboLab analyses task-generalist policies along visual, procedural, and relational competencies (Yang et al., 2026a). Chen et al. (2026c) evaluate agents without a trained policy on its 120 tasks with vague instructions: code-as-policy (CaP-X) reaches 14.4% success, VLM-based task and motion planning (TiPToP) 18.9%, and their orchestrator without a vision–language–action policy 17.8%.

LIBERO-MAX

Purpose. Whether the agent can build a robot controller that adapts when the world changes during execution, not only before it starts.

Source and construction. LIBERO-MAX (Zhang et al., 2026e) has 8,000 paired cases on LIBERO (Liu et al., 2023) scenes, with starting configurations from LIBERO-Plus (Fei et al., 2025) and LIBERO-PRO (Zhou et al., 2025). A case pairs a Base episode with a Dynamic episode that has the same start and policy seed and receives one online event: target or receptacle relocation, camera shift, sensor-noise onset, illumination or visual-theme switch, a burst of distractors, or an inserted obstacle. In the selected cases the event fires when the robot comes within 0.18 m of a designated object, before the grasp. Cases are selected from official metadata, never from outcomes. All splits use three LIBERO tasks (a bowl from a drawer onto a plate, a sauce bottle into a basket, a plate pushed to the front of a stove) and the same number of cases per event type: 8, 16 and 24 pairs for development, validation and test. No initial physical state appears in two splits, and states used in earlier experiments are excluded from validation and test. The agent sees the development case definitions.

Setting. A simulated Franka Panda (MuJoCo). Each step, the solution receives RGB and metric depth from an external and a wrist camera, their calibration, the robot state and the instruction, and returns an absolute end-effector pose and gripper command. It may query a frozen X-VLA policy (Zheng et al., 2025) for 30-step action chunks, possibly on transformed images, and call the model up to 16 times per episode; no other model, no training, no simulator state. The Dynamic episode replays the Base episode exactly until the event.
Starter solution. Executes the X-VLA chunks unchanged, without calling the model.
Metric. Dynamic success rate $\frac{1}{P}\sum_{p=1}^{P} s_p^{\mathrm{dyn}}$ over all P planned pairs, with $s_p^{\mathrm{dyn}} \in \{0,1\}$ the LIBERO success check of the Dynamic episode; an episode that reaches the horizon or its 20-minute limit unsolved fails. The Base success rate and the paired difference, as in LIBERO-MAX, are also reported.
Prior work. Across fourteen vision–language–action, hybrid and world-action policies, LIBERO-MAX events lower success by 11.0 to 25.7 percentage points, mostly through geometry and observation changes (Zhang et al., 2026e).

LIBERO-PRO

Purpose. Whether the agent can build a robot controller that stays successful when a familiar task is perturbed before the episode starts, rather than relying on the conditions the underlying policy was trained under.
Source and construction. LIBERO-PRO (Zhou et al., 2025) perturbs LIBERO (Liu et al., 2023) tasks in their objects, initial states, instructions and environments, and releases the perturbed task definitions and initial states. Ten of its categories are used: object appearance, object shape, object texture, object positions, paraphrased instruction, changed task goal, camera view angle, initial robot pose, visual noise or glare, and view occlusion. Each case is one perturbed initial state of one of the three LIBERO tasks of LIBERO-MAX. Cases are selected from official metadata, with equal numbers per category in each split, spread over the three tasks: 10, 20 and 30 episodes for development, validation and test (ten per task in test). No initial physical state appears in two splits or in LIBERO-MAX, and states used in earlier experiments are excluded from validation and test. The agent sees the development case definitions.
Setting. As in LIBERO-MAX: a controller that may query the frozen X-VLA policy (Zheng et al., 2025) and call the model up to 16 times per episode, with a 20-minute limit per episode and no other model.
Starter solution. Executes the X-VLA chunks unchanged, without calling the model.
Metric. Success rate $\frac{1}{N}\sum_{i=1}^{N} s_i$ over all N planned episodes, with s_i the native LIBERO success check; an episode that reaches the horizon or the time limit unsolved fails.
Prior work. LIBERO-PRO reports that models above 90% success under the standard LIBERO evaluation fall to 0.0% under its generalized setting, a sign of memorized training conditions (Zhou et al., 2025). LIBERO-Plus (Fei et al., 2025) analyses robustness along seven dimensions: object layout, camera viewpoint, initial robot state, language, lighting, background texture and sensor noise.

Embodied-SWE

Purpose. Whether the agent can write, offline with a simulator, a control program for a long-horizon, contact-rich tool-use task that then runs without any model.
Source and construction. The spatula task of EmbodiedSWE-Bench (Shen et al., 2026b): a Franka arm under operational-space control in Isaac Lab wedges a spatula under a slice of bread in a pan, flips it, reloads it onto the blade and serves it flat onto a plate, touching the bread only with the tool. The environment is built with the benchmark's tools in its setting without agent tools, and the grader is withheld. Each evaluation runs eight parallel environments with random bread size and pose, plate position and spatula pose. Development uses the build seed (8 trajectories); validation 2 and test 4 hidden random seeds (16 and 32 trajectories). The agent receives the benchmark's task statement, its rule for batched solutions and its staged-planning guide.
Setting. A program that commands the robot only through the environment's step function and may read object poses. Writing simulator state, applying external forces, changing physics or the scene, or altering the grader scores zero. A rollout has 30 minutes of wall time, and no model is called during evaluation. During development the agent has full simulator access on one GPU.

Starter solution. Holds the arm still for ten steps and scores zero.

Metric. The EmbodiedSWE-Bench rubric score (Shen et al., 2026b): each trajectory receives the credit of the furthest stage it reaches, namely spatula lifted (0.2), bread on the blade (0.4), bread flipped and flat (0.6), bread carried on the blade (0.8) and bread on the plate (1.0). The score is the mean over all trajectories of the split; a crashed rollout scores zero, and a timed-out one is scored as it stands. The final-state success rate is also reported.

Prior work. EmbodiedSWE-Bench (Shen et al., 2026b) evaluates coding agents that write such programs in one four-hour run per task, and uses their solutions to generate demonstrations for vision–language–action models. On this task, single runs of frontier coding agents reach rubric scores from 0.2 to 1.0, with one run disqualified for reward hacking.

A.2 Reasoning

Ego-Exo Memory

Purpose. Memory reasoning across two synchronized views of one activity. The evidence for an answer is visible only at particular moments, often in one view, and the model sees few frames per call, so the solution must choose which moments and views to show.

Source and construction. The questions are the Ego-Exo4D portion of EgoExoMem (Liu et al., 2026b): 764 human-verified four-option questions on 125 recordings of cooking and bicycle repair from Ego-Exo4D (Grauman et al., 2024). Each recording has an egocentric view from the glasses of the person performing the activity (448×448 pixels) and an exocentric view from a fixed camera (796×448 pixels), synchronized at 30 frames per second and lasting from 20 seconds to 23 minutes. The questions concern habitual location, instantaneous position, resulting location, egocentric direction, object state, allocentric relation, and temporal order. All recordings of one session (the same room and people) are assigned to one split, stratified by activity, with about 35%, 25%, and 40% of the questions in the development, validation, and test splits. Each split then keeps a fixed-seed sample stratified by activity and question type and spread over its videos: 80 questions on 38 videos, 80 on 33, and 120 on 47. The answer options are shuffled once and the videos are anonymized. With each development question the agent also receives the annotated moments at which the answer is visible and the view that shows it.

Setting. Given the question, both videos, and frames of both views at one frame per second, the solution returns an option letter. It may call the model three times per question on average over a split, with at most $16 \times 448 \times 448$ pixels per call; further calls are refused. It may run open-weight models, for example for frame retrieval, on one GPU, has no network access during evaluation, and may not use the annotations of EgoExoMem or Ego-Exo4D.

Starter solution. Shows the model eight evenly spaced frames of each view, each labeled with its view and time and downscaled so that all sixteen fit the pixel limit, followed by the question, in a single call.

Metric. Accuracy over the questions of a split; a question without a valid letter counts as wrong. Chance is 0.25.

Prior work. EgoExoMem evaluates multimodal language models on 2.6K cross-view memory questions of eight types. On the full benchmark, the best existing model reaches 55.3%, and the training-free frame selection method E^2-Select, which allocates the frame budget by relevance and samples frames per view with a k-DPP, reaches 58.2% (Liu et al., 2026b).

Kinematics

Purpose. Quantitative physical reasoning from video: building a measurement procedure around a fixed vision–language model, so that quantities follow from pixels, time, and a known scale.

Source and construction. QuantiPhy (Puyin et al., 2026) contains 3,355 questions on 569 videos from physics simulations, laboratory recordings, internet clips, and segmented objects. The task uses its public labeled validation set, 158 questions on 24 videos. Each question asks for the size, distance, velocity, or acceleration of an object, possibly at a stated time, in stated units, and gives one prior in physical units: a known length,

speed, or acceleration, and for some videos depths. The videos are assigned to fixed lists, so that no video appears in two splits and every split mixes the four video sources and the static and dynamic priors and targets: 39 questions on 6 videos for development, 52 on 8 for validation, and 67 on 10 for testing.

Setting. The solution returns a number in the unit of the question. It does not call the model being evaluated: the only language or vision–language model it may use is Qwen3.8-27B with open weights, served locally with at most 32 images ($32 \times 640 \times 360$ pixels in total) per call and eight calls per question on average over a split. It may add non-language vision models and classical code on one GPU, such as point trackers, optical flow, detectors, segmenters, and depth or pose estimators, but no other language model, neither in the solution nor to label data, and no network access during evaluation.

Starter solution. Shows the fixed model four evenly spaced frames, the question, and the prior in a single call with reasoning disabled, and reads the first number of the reply.

Metric. The mean relative accuracy (MRA). For a prediction $\hat{y}_i$ of the true value y_i,

$$\mathrm{MRA} = \frac{1}{N}\sum_{i=1}^{N}\frac{1}{|C|}\sum_{\theta\in C}\mathbf{1}\left[\frac{|\hat{y}_i - y_i|}{|y_i|} < 1-\theta\right], \qquad C = \{0.1, 0.2, \ldots, 0.9, 0.95\},$$

and a question without a numeric answer scores zero. MRA was introduced by VSI-Bench (Yang et al., 2025) with thresholds from 0.5 to 0.95; the threshold set above is that of QuantiPhy (Puyin et al., 2026).

Prior work. QuantiPhy evaluates 21 vision–language models that answer directly from the video. The best reaches an MRA of 53.1, below the average human score of 55.6; the models rely on pretrained world knowledge more than on the given video and prior, and a four-step chain of thought (measure the source property in pixels, estimate the scale, measure the target in pixels, convert) lowers the score of 19 of the 21 models (Puyin et al., 2026).

Mass

Purpose. Estimating a property that is not directly visible: weight depends on material, on whether an object is solid or hollow, and on how its volume is filled.

Source and construction. The items are product listings of Amazon Berkeley Objects (Collins et al., 2022) that state a weight, all three dimensions, and a main photograph. Consumables (food, supplements, cleaning and medical supplies) and multipacks or sets, recognized by name, are removed, as are items outside 10 g to 150 kg and items whose weight divided by the volume of their bounding box lies outside 5 to 5000 kg/m^3, which removes unit and listing errors. Variants of one product (the same brand and the same first six words of the name) form one group, and one item per group is drawn with a fixed seed. The groups are split into 300 development, 300 validation, and 600 test items, with at most 3% of a split from one product type, giving about 130 product types per split. Each item keeps its main photograph and up to two others, with the longer side at most 1024 pixels.

Setting. Given the photographs, a coarse product type, and the listed dimensions in centimetres, but not the name, brand, or description, the solution states the weight in kilograms. As in Kinematics, the only language or vision–language model is the fixed Qwen3.8-27B, with at most eight images ($8 \times 1024 \times 1024$ pixels) per call and four calls per item on average. Non-language models such as segmenters, depth estimators, image encoders, and regressors trained on the development split may run on one GPU; other language models, text encoders applied to other data, and any external source of the weights are not allowed.

Starter solution. Shows the fixed model the main photograph, the product type, and the dimensions in a single call with reasoning disabled, and reads the first number of the reply.

Metric. The mean relative accuracy with the thresholds of Kinematics, averaged over items; an item without a numeric answer scores zero.

Prior work. Image2mass (Standley et al., 2017) collected about 150,000 online product listings with images, sizes, and weights, and estimates mass from an image together with the object's dimensions, comparing learned models with human estimates. Amazon Berkeley Objects (Collins et al., 2022) provides catalogue images, metadata, and 3D models of real products for object understanding.

ARC-AGI-1 and ARC-AGI-2

Purpose. Abstract reasoning from few examples: inferring a grid transformation from a few demonstration pairs and applying it to new inputs, on puzzles designed to require little prior knowledge.

Source and construction. The puzzles come from the official public repositories at fixed versions: 400 training and 400 evaluation puzzles for ARC-AGI-1 (Chollet, 2019), and 1000 training and 120 evaluation puzzles for ARC-AGI-2 (Chollet et al., 2025). Each puzzle has demonstration pairs and one or more test inputs; grids are at most 30×30 with ten colors. Training puzzles identical to an evaluation puzzle are removed (one for ARC-AGI-1). The remaining training puzzles are ordered by a seeded hash: the first 80 form the validation split, and the rest form the development pool, which the agent receives with answers (319 puzzles for ARC-AGI-1); a fixed panel of 32 development puzzles is scored in development evaluations. The test split consists of all official evaluation puzzles, 400 for ARC-AGI-1 and 120 for ARC-AGI-2, so validation and test puzzles are disjoint.

Setting. Given the demonstration pairs and the test inputs, the solution returns at most two candidate grids per test input. It may call the model up to 2000 times per evaluation, with four puzzles solved in parallel in independent sessions, and runs in a sandbox without a GPU or network access. It may use program search, hand-written rules, or models trained on the development puzzles, but may not obtain or store held-out answers. Invalid predictions score zero.

Starter solution. Asks the model for two candidate outputs as JSON in a single call per puzzle; an unparseable reply counts as no answer.

Metric. Pass@2, in percent: with g_{ij}^1, g_{ij}^2 the two candidates for test input j of puzzle i,

$$\text{pass@2} = \frac{100}{N} \sum_{i=1}^{N} \prod_{j} \mathbf{1}\left[y_{ij} \in \{g_{ij}^1, g_{ij}^2\} \right],$$

so a puzzle counts only when every test output is matched exactly. The two attempts per test input follow the ARC Prize protocol (Chollet et al., 2024).

Prior work. ARC (Chollet, 2019) was proposed to measure the efficiency of skill acquisition on novel problems. In ARC Prize 2024 the best private-set score rose from 33% to 55.5%, mainly through deep-learning-guided program synthesis and test-time training (Chollet et al., 2024); test-time training brings an 8B language model to 53.0% on the public validation set, and 61.9% when ensembled with program synthesis (Akyürek et al., 2024). ARC-AGI-2 raises the difficulty: on its semi-private set, o3 (medium) scores 3.0%, against 53.0% on ARC-AGI-1 (Chollet et al., 2025).

ARC-AGI-3

Purpose. Learning an interactive environment without instructions: finding out by acting what the controls do and what completes a level, then completing levels with few actions.

Source and construction. ARC-AGI-3 (ARC Prize Foundation, 2026) consists of turn-based games on grids of up to 64×64 cells with 16 colors, each with several levels and a human baseline action count per level. Its public set has 25 games, and official scores use 110 further private games. The task uses the 25 public games at a fixed version of the official toolkit, split by game with a seeded hash into 10 development, 5 validation, and 10 test games. The agent sees only the identifiers of the development games and plays them through the interface; the game code, the human baselines, and the validation and test games are not accessible to it.

Setting. The solution observes the current frame, the available actions, and the number of completed levels, and submits actions: directions, a game-specific action, a click at a grid position, and undo, whose effects must be inferred; a reset is allowed only after the game is lost. Each game allows at most 1000 actions, 120 model calls, and 1200 seconds, and each evaluation at most 1200 model calls. Games are played one at a time, each from a fresh start with a fixed seed, and no memory or files carry over between games. The solution may not contain game solutions or game code.

Starter solution. Sends the model the current frame as an image and as rows of color values, a short memory, and the recent actions, and executes up to eight actions returned per call along with an updated memory.

Metric. Relative human action efficiency (RHAE), computed by the official toolkit. With h_ℓ the human baseline and a_ℓ the agent's actions on level ℓ, a completed level scores $S_\ell = \min(1.15,\ h_\ell/a_\ell)^2$ and an uncompleted level zero; with weights $w_\ell = \ell$ over the n levels of a game of which k are completed,

$$E = \min\left(\frac{\sum_{\ell=1}^{k} w_\ell}{\sum_{\ell=1}^{n} w_\ell},\ \frac{\sum_{\ell=1}^{n} w_\ell S_\ell}{\sum_{\ell=1}^{n} w_\ell}\right),$$

and the score is the mean of E over games, multiplied by 100 (ARC Prize Foundation, 2026). A game stopped by a limit keeps its partial score.

Prior work. ARC-AGI-3 (ARC Prize Foundation, 2026) reports that humans solve all of its environments, while frontier AI systems scored below 1% as of March 2026; in its preview competition, the best entry, which trains a convolutional network with reinforcement learning to predict which actions change the frame, scored 12.58%.

Spatial-VSI

Purpose. Spatial reasoning about indoor scenes from a moving camera: counts, metric sizes and distances, relative directions, routes, and the order of appearance.

Source and construction. VSI-Bench (Yang et al., 2025) contains over 5000 questions on 288 videos of real rooms from ScanNet, ScanNet++, and ARKitScenes (Dai et al., 2017; Yeshwanth et al., 2023; Baruch et al., 2021). It has ten question types in eight groups: object counting, absolute distance, object size, and room size, answered by numbers, and relative distance, relative direction (easy, medium, hard), route planning, and appearance order, answered by option letters. The scenes of a fixed dataset version are shuffled with a fixed seed and dealt to the test, validation, and development splits in proportions 40%, 30%, and 30%; each split then keeps questions stratified so that the eight groups are weighted as in the official score. The splits hold 80 questions from 46 scenes for development, 480 from 84 scenes for validation, and 320 from 100 scenes for testing.

Setting. Given the video (640 by 480 pixels, one to three minutes), the question, and the options, the solution returns an answer. It may call the model twice per question on average over a split and uses no other model and no downloaded weights; its own code, such as geometry and classical vision with OpenCV, is allowed. It may not read the annotations of the source datasets or access the network, has no GPU, and has four hours per evaluation.

Starter solution. Shows the model 16 evenly spaced frames, 512 pixels wide, with the official VSI-Bench prompt, in one call per question.

Metric. The overall score of VSI-Bench (Yang et al., 2025). A multiple-choice answer is correct when its first word is the option letter; a numerical answer receives the mean relative accuracy with thresholds $\{0.50, 0.55, \ldots, 0.95\}$ and the condition $|\hat{y} - y|/y \le 1 - \theta$, and zero if it is not a number. Scores are averaged within each group, the three direction levels are first averaged into one, and the eight groups are averaged, giving a value between 0 and 1.

Prior work. VSI-Bench finds that the best model it tests, Gemini-1.5 Pro, averages 45.4%, against about 79% for humans on a subset; chain-of-thought, self-consistency, and tree-of-thoughts prompting do not help, whereas asking the model to draw a cognitive map raises relative-distance accuracy from 46.0% to 56.0% (Yang et al., 2025). VSI-Bench-Debiased removes questions that can be answered without the video (Brown et al., 2025), and Spatial-MLLM adds a 3D spatial encoder initialized from a visual geometry model to a multimodal language model (Wu et al., 2025b).

A.3 Creation

Image Editing

Purpose. Whether the agent can build an instruction-based image editor that performs the edit requested in English or Chinese while leaving the rest of the photograph intact, across local, global, text, retouching and restoration, reasoning, and compound edits. The editing model is not prescribed: choosing, combining, and adapting open models is part of the task.

Source and construction. The evaluation set was built for this task in two versions. The first has 100 images; the second, used here, adds 86 images and 14 restoration photographs, for 200 images. The 186 photographs, paintings, and illustrations are CC0 or public-domain works from Wikimedia Commons, each checked for its licence and excluding identifiable people, minors, trademarks, and third-party works; two images of puzzles (an arithmetic problem on a blackboard and a number sequence) were made for the task, and four images carry a reproducible synthetic degradation (blur, color cast, or low contrast) to be undone. The 14 restoration photographs show haze, fog, smoke, lens flare, reflections on glass, or cast shadows; all come from the data of UniSER (Zhang et al., 2025c), including ten of its 39 in-the-wild evaluation images and three of its paired images from NH-HAZE (Ancuti et al., 2020), FlareReal600, and WSRD+ (Vasluianu et al., 2023). Images are resized to at most 1536×1024 pixels with sides that are multiples of 16, and span outdoor (85), product (34), indoor (32), low-quality (21), illustration (13), text (11), and document (4) images. Each image has one instruction, written in Chinese (106) or English (94) together with a translation into the other language, and is assigned to one of 23 kinds of edit in six families: local semantic edits, namely adding, removing, or replacing an object, changing its attributes, changing a pose, setting a count, or moving objects (71 images); global edits of style, background, or lighting, weather, and season (30); text editing, namely replacing, writing, or editing Chinese text (22); retouching and restoration, namely tone, old or degraded photographs and soft effects, removal of clutter, and outpainting to a new aspect ratio (34); edits that require inference, such as candles after three hours of burning, the farthest of several glasses, leaves in late autumn, or a completed sudoku (25); and compound or constrained instructions (18). Each sample is labeled with a difficulty (44, 114, and 42 at three levels) and the size of the edited region (small, medium, large, or global), and is annotated for scoring: a mask of the region allowed to change (92 images), objects that must appear (43) or disappear (32), a target count (11), text that must be readable (24) or gone (14), a target aspect ratio (5), content that must be kept (74), the parts of a compound instruction (13), and a description of the intended result of an inference edit (31). The images are divided into 50 development, 70 validation, and 80 test images, stratified by kind of edit. The development samples come with their annotations, and their evaluations return per-sample judge rationales and side-by-side comparisons; validation and test return only the aggregate score.

Setting. The solution is an editing program that receives the image, the instruction, and its language, and writes an image at the input size. It may use any open-weight editing or generation model, several models at once, vision models, classical image processing, and its own fine-tuning on the development split; Qwen-Image-2.1, Qwen-Image-Edit-2511 with distilled few-step adapters (Wu et al., 2025a), and HunyuanImage-3.0-Instruct (Tencent Hunyuan Foundation Model Team, 2025) are provided. It may not use the judge's model family, remote services, or the annotations. Evaluation runs offline on four 80 GB GPUs with, on average, 150 s per image per GPU; the solution does not call the model being evaluated.

Starter solution. Qwen-Image-Edit-2509 with an eight-step distilled adapter, given the instruction unchanged, producing one output per image.

Metric. A frozen judge, GLM-4.5V (GLM-V Team, 2025) at temperature 0, sees the source and the output and rates instruction compliance SC and preservation PR, and, from the output alone, perceptual quality PQ, each from 0 to 10. The preservation rubric depends on the kind of edit (for example, only content and composition for style transfer, only the kept subjects for a background change), and the judge is told the annotated parts, the kept content, and, for inference edits, the intended result. Global deformation is measured rather than judged: a SIFT (Lowe, 2004) and RANSAC (Fischler and Bolles, 1981) homography from input to output gives the 90th-percentile displacement p_{90} in pixels at a 512-pixel short side, $g = \exp(-\max(0, p_{90} - 0.5)/6)$, and $\mathrm{PR}' = \max(0, \mathrm{PR} - 5(1-g))$. The judge score is $J = \frac{1}{10}\sqrt{\mathrm{SC} \cdot \mathrm{PQ}}\,(0.5 + 0.05\,\mathrm{PR}')$. The programmatic score P averages the checks that apply: object presence and absence and the exact count, answered by the judge; for text, one minus the lowest character error rate over two OCR readings and the judge's transcription, and the absence of text to be removed; the aspect ratio of outpainted images, $\max(0, 1 - 5|r/r^* - 1|)$; preservation, $\mathrm{clip}(1 - (\mathrm{MAD} - 0.02)/0.10)$ from the mean absolute color difference outside the dilated edit region after alignment; and detail, $\mathrm{clip}((cs - 0.55)/0.40)$ from the contrast–structure term of SSIM (Wang et al., 2004) at a 1024-pixel short side. Without a mask, the edit region is formed from the judge's boxes of the edited objects. A sample scores $0.5J + 0.5P$ (J if no check applies), and 0 if $\mathrm{SC} < 3$ or no image is produced. The score is the mean of the six family means, in $[0, 1]$.

Prior work. GEdit-Bench, introduced with the Step1X-Edit editor (Liu et al., 2025), evaluates instruction-based editing on real user requests with a vision–language judge of semantic consistency and perceptual quality, which VIEScore (Ku et al., 2024) combines as $\sqrt{\text{SC} \cdot \text{PQ}}$; the judge score above takes this form and adds the preservation term and the programmatic checks. UniSER (Zhang et al., 2025c) removes soft effects such as haze, flare, shadows, and reflections with a single diffusion model and evaluates its in-the-wild images, which lack ground truth, with reference-free metrics and a vision–language judge.

Multi-turn Editing

Purpose. Whether the agent can improve a multi-turn editing system, in which each instruction is applied to the previous result and becomes known only at its turn, so that errors in content and image quality accumulate over the session.

Source and construction. All 100 sessions of MSE-Bench, the multi-turn editing benchmark introduced with VINCIE (Qu et al., 2026), at a fixed revision. Each session is a real photograph, used unchanged at its original size (many are small), and five instructions that add, remove, or replace objects, change attributes, backgrounds, posture, expression, orientation, position, actions, or interactions, restyle the image, or move the camera (dolly, pan, tilt, orbit); MSE-Bench labels 20 operation types. The sessions are split into 30 development, 30 validation, and 40 test sessions by a seeded search that balances the per-turn counts of the 20 operation types, the sessions with a camera operation, and the images under 512 pixels on the short side. For the development sessions the agent also sees the MSE-Bench annotations (operation types and descriptions of the intended results); the solution never receives them.

Setting. The solution is a session object that is reset with the source image and returns one image per instruction; it may keep earlier images and instructions. The editing model is fixed to the original Qwen-Image-Edit (August 2025) (Wu et al., 2025a) with its four- and eight-step distilled adapters, since the later 2509 and 2511 releases degrade into blocky noise when edited repeatedly. Prompt rewriting, sampler settings, candidate selection, masking and compositing, and Qwen2.5-VL-7B (Bai et al., 2025b) as a helper are allowed; training the editing model, network access, and any copy of MSE-Bench or of VINCIE's outputs are not. Each of four GPUs runs one process with 300 s of start-up and, on average, 60 s per turn. A failed turn, an aspect-ratio change above 5%, or a timeout scores 0 for that turn and the rest of the session. The solution does not call the model being evaluated.

Starter solution. Edits the previous image with each turn's instruction, using the eight-step adapter and a fixed seed.

Metric. Each turn k is judged by a frozen GLM-4.5V (GLM-V Team, 2025) at about one megapixel. Given the instructions of turns $1..k$ and the images up to turn k, MSE-Bench's own evaluation prompt rates prompt following PF and consistency CONS with the previous image; a strict prompt rates the perceptual quality PQ of the turn's image alone (all 0–10). The turn score is $\frac{1}{10}\sqrt{\text{PF} \cdot \text{PQ}}\,(0.5 + 0.05\,\text{CONS})$, or 0 if $\text{PF} < 3$; the score averages over the five turns and then over sessions. MSE-Bench's cumulative success rate, the fraction of sessions in which every turn up to k has $\text{PF} > 6$ and $\text{CONS} > 6$, is reported for $k = 1, 3, 5$.

Prior work. VINCIE (Qu et al., 2026) learns in-context image editing from video and introduced MSE-Bench to measure multi-turn editing by the success rate at each turn, judged by GPT-4o. Here the editing model is fixed and the agent improves the system around it.

Video Generation

Purpose. Whether the agent can build an executable text-to-video workflow around a frozen video generator, deciding how to plan prompts, spend a fixed generation budget, inspect frames, and select candidates, so that the videos follow compositional prompts.

Source and construction. The prompts come from T2V-CompBench (Sun et al., 2025), at a fixed revision of its prompt lists, restricted to four of its seven categories: consistent attribute binding, action binding, generative numeracy, and spatial relationships (two-dimensional only). Each split has 20 prompts, five per category. Near-duplicate prompts were grouped into families (connected components of normalized token Jaccard similarity of at least 0.65); the validation and test prompts were drawn in a salted hash order, excluding the

families of development prompts and taking at most one prompt per family. The agent sees the development prompts and their scores; validation and test prompts are private.

Setting. The generator is MiniMax H3 (MiniMax, 2026b) with frozen weights, producing 124 frames at 864×480 and 24 frames per second. For each prompt the workflow may make at most three generation requests totalling 50 denoising steps, choosing prompts, seeds, step counts, and flow shift; sample frames; call the model being evaluated at most twice (40 calls per evaluation), with images; and process videos locally. It returns one video. It has no network, no direct GPU access, and no access to the evaluators, and may not add text or cues aimed at the evaluators. Each evaluation runs on one 80 GB H100 within 150 minutes; prompts unfinished when the workflow time runs out score 0.

Starter solution. A single generation request with the prompt unchanged and the default seed, 50 steps, and flow shift.

Metric. The original T2V-CompBench evaluators score each video against the original prompt with their default prompts, thresholds, and seeds: attribute binding and action binding with the multimodal-LLM evaluation using LLaVA-v1.6-34B (Liu et al., 2024a) on sampled frames, and numeracy and spatial relationships with the detection-based evaluation using Grounding DINO (Liu et al., 2024b) (16 sampled frames, six of them in a grid for the language model). Each video scores in $[0, 1]$, a missing or corrupt video scores 0, and the score is the mean of the four category means. Because of the subset and the fixed resolution, the scores are not comparable with the T2V-CompBench leaderboard.

Prior work. T2V-CompBench (Sun et al., 2025) benchmarks text-to-video models on 1,400 prompts in seven compositional categories, with multimodal-LLM-based, detection-based, and tracking-based metrics validated against human judgments. Here the generator is fixed, and the agent improves the workflow around it.

Text-to-Image

Purpose. Whether the agent can make a text-to-image program satisfy every constraint of new prompts in Chinese and English while keeping image quality, either with a fixed generator or with any open model.

Source and construction. The 200 prompts and their rubrics were written for this task: 100 in Chinese and 100 in English (not translations of each other), ten in each of 20 subcategories of six categories: subjects and attributes (identity, color, material), counting and binding (single and multiple classes, attribute binding, excluded content), spatial relations and interaction (position, depth and containment, actions and roles), text and layout (English, Chinese, multiple text blocks), style and presentation (medium, lighting and season, camera and composition), and reasoning and composition (world knowledge, causality and time, logical quantity, multiple constraints). Each subcategory has two easy, five medium, and three hard prompts by the author's judgment. Each prompt has a fixed seed and size (1024×1024 for 145 prompts, 1280×768 or 768×1280 for 55) and a rubric of atomic criteria (954 in total) of four types: semantic, count, absence, and exact text, the latter copied verbatim from the prompt; each requirement is credited once. After authoring, every prompt and rubric was reviewed for physical consistency, count logic, text literals, coverage, and duplicate credit, which revised 32 prompts, 51 rubrics, and 34 sizes; scripts check quotas, uniqueness, field consistency, and split isolation, and flagged cross-split pairs of similar wording were reviewed by hand. The prompts are split into 50 development, 70 validation, and 80 test prompts by semantic group (178 groups of prompts sharing a central scene), each split half Chinese and half English with 20/50/30% easy, medium, and hard prompts. Development prompts come with rubrics and per-image feedback; validation returns only the score.

Setting. The solution receives a prompt, its language, size, and seed, and writes one image. With the fixed generator, all images come from the frozen Qwen-Image-2.1 (Wu et al., 2025a); prompt rewriting, layout planning, sampling settings, several candidates with selection, Qwen2.5-VL-7B (Bai et al., 2025b) as a helper, and ordinary image tools are allowed, but no training or adapters. With any model, other open-weight generators, routing, and training on the development split are also allowed. Neither may use the judge's model family, the rubrics, the network, or hosted models at evaluation, which runs on four GPUs with at most 600 s per image.

Starter solution. One image from Qwen-Image-2.1 with the prompt unchanged and 40 sampling steps.

Metric. A frozen GLM-4.5V (GLM-V Team, 2025) judges each image: a semantic criterion scores 0, 0.25, 0.5, 0.75, or 1; a count, absence, or exact-text criterion scores 0 or 1; and the quality $Q \in [0, 1]$ rates rendering defects, excluding prompt compliance and intended style. With criterion scores c_i and weights w_i (all 1), $C = \sum_i w_i c_i / \sum_i w_i$ and the image scores $C\,(0.5 + 0.5\,Q)$; a missing or wrongly sized image scores 0. Scores are averaged within each category and then over the six categories. The reported score is the mean of the test scores of the two settings.

Prior work. GenEval (Ghosh et al., 2023) and T2I-CompBench (Huang et al., 2023a) evaluate compositional text-to-image generation with detectors and other vision models. The prompts here are new and these benchmarks' scores are not reproduced.

A.4 Reconstruction

3D Tracking

Purpose. Whether the agent can assemble a system that segments every object in a few widely separated views of an indoor scene and gives each object the same identity in every view, from RGB images alone. Without camera poses or depth, the views must be associated through appearance and implicit geometry: the cross-view instance tracking that IGGT (Li et al., 2026) introduced as an evaluation of instance-grounded 3D reconstruction.

Source and construction. The validation and test clips are the 3D tracking benchmark released with IGGT, which samples a few views per scene from ScanNet (Dai et al., 2017) and ScanNet++ (Yeshwanth et al., 2023) for spatial coverage with sufficient overlap and annotates a subset of objects by hand with identities consistent across views. Each clip has six to nine RGB views, resized to 592×448 pixels, typically with four to eight annotated objects such as furniture, lamps, and decorations; walls, floors, and ceilings are not annotated. The 8 ScanNet++ clips form the validation split and the 10 ScanNet clips the test split. The development split consists of 100 clips built from ScanNet++ training scenes that share no scene with the validation clips: eight views are drawn from a random stretch covering 30 to 80% of a capture, and up to eight object-like instances are kept (visible in at least two views, at most a quarter of any view, and at least 0.1% of the clip's pixels); a clip with fewer than four is drawn again. The agent sees the development views with their instance maps.

Setting. The solution is a program that receives the views of one clip and returns one instance map per view, with 0 for the background and identities consistent across views. It may use any publicly released pretrained model or library and needs no training; it does not call the model being evaluated. It runs on one GPU without network access, and a whole split must finish within 45 minutes.

Starter solution. Clusters the colors of each view into six clusters with k-means, splits them into connected regions covering 0.2 to 20% of the image, and links each region to the existing track with the nearest mean color, or starts a new track.

Metric. Evaluation is class-agnostic. Predicted and annotated tracks are matched one to one, greedily in decreasing order of their temporal IoU, which accumulates intersection and union over the views: for an annotated object with masks G_f and its matched prediction P_f in views $f = 1, \ldots, S$,

$$\text{T-IoU} = \frac{\sum_f |P_f \cap G_f|}{\sum_f |P_f \cup G_f|}.$$

T-mIoU is the mean of T-IoU over the annotated objects of a clip, averaged over clips (0 to 1). An annotated object without a match scores 0, and predicted tracks without a match are not penalised. T-SR is the share of annotated objects whose per-view IoU exceeds 0.5 in every view in which they appear. The region similarity $\mathcal{J}$ (mean per-view IoU), the boundary measure $\mathcal{F}$ (tolerance 0.8% of the image diagonal), and $\mathcal{J}\&\mathcal{F}$ of DAVIS (Perazzi et al., 2016) are also reported. IGGT named T-mIoU and T-SR but defined them per view, averaging the IoU over views and counting an object as tracked when its prediction is non-empty in every view; the accumulated IoU, the IoU threshold of T-SR, and the greedy matching follow the evaluation of S4VY (Zhang et al., 2026c).

Prior work. IGGT reconstructs geometry and predicts 3D-consistent instance features with one transformer, trained with a 3D-consistent contrastive objective; IGGT4D (Zou et al., 2026) extends it to streaming 4D input;

S4VY (Zhang et al., 2026c) adds persistent object queries to a feed-forward 4D geometry model. On the ten ScanNet clips under this metric, S4VY reports a T-mIoU of 0.803, against 0.781 for IGGT4D, 0.635 for SAM 2 (Ravi et al., 2024) with incremental object discovery, and 0.517 for IGGT.

Scene from Photo

Purpose. Whether the agent can build a single-image 3D scene reconstruction system: from one photograph of a furnished room, recover the camera, the room, and every piece of furniture as a complete object with plausible unseen sides, with physically based materials and lighting, so that a path tracer reproduces the photograph.

Source and construction. The photographs are those of HARMONY300 (Sun et al., 2026): 100 renders of 3D-FRONT rooms (Fu et al., 2021), which have ground-truth geometry, and 200 real photographs, half of moderate and half of high complexity. Each difficulty is split separately with a fixed seed (the renders stratified by their number of annotated objects) into 10, 10, and 20 photographs for development, validation, and test, giving 30, 30, and 60; the remaining 179 photographs are held out of every split, and one render whose image differs from its ground truth is excluded. On the renders the annotated furniture and its meshes come from the 3D-FRONT ground truth. On the real photographs up to eight pieces of furniture per photograph are annotated by Mask2Former (Cheng et al., 2022) panoptic segments of furniture classes (score at least 0.5, at least 0.5% of the image); on the renders this procedure recovers 85% of the annotated objects. The agent sees the development photographs with all labels.

Setting. The solution writes a Blender scene per photograph within 120 s per photograph; the evaluator renders it with Cycles from the scene's own camera under fixed rules (only meshes, physically based shader nodes, non-emissive materials, at most 64 lights). It may use any public pretrained model on one GPU, but not HARMONY300, 3D-FRONT, 3D-FUTURE, or HARMONY's code, and does not call the model being evaluated.

Starter solution. An empty room box in flat colors matched to the photograph, one ceiling light, and the same guessed camera with a 60° field of view for every photograph.

Metric. $s = (A + O + G)/3$ in $[0, 1]$, with A and O averaged over all photographs and G over the renders. Appearance $A = \frac{1}{3}[\max(0, 1 - \text{NCLIP}/0.25) + \max(0, 1 - \text{PL}/0.1) + 1 - \text{LPIPS}]$, where N-CLIP (one minus the cosine similarity of CLIP ViT-B/32 embeddings (Radford et al., 2021)) and the photometric loss PL are HARMONY's metrics and LPIPS is that of Zhang et al. (2018). Objects $O = \text{mean}_k\, r_k \min(1, \text{IoU}_k/0.5)$, where IoU_k compares the meshes assigned to object k with its mask and the realness $r_k = \text{clip}((\min_j P_j - 0.5)/0.4, 0, 1)$; P_j is the probability that Qwen3-VL-8B (Bai et al., 2025a), shown untextured views from four directions, judges the reconstruction a more real object than shortcut j (its visible surface as a sheet, that sheet extruded to a prism, or its bounding box). Geometry $G = (F_{0.1} + F_{0.01} + F^{\text{obj}} + F^{\text{obj}}_{\text{hidden}})/400$, with HARMONY's normalized scene F-scores in the protocol of Gen3DSR (Ardelean et al., 2025) (furniture only, both point sets normalized by their bounding box and aligned by ICP, thresholds 0.1 and 0.01) and per-object F-scores at 5 cm, in percent, after scaling by the median ratio of true to rendered depth; $F^{\text{obj}}_{\text{hidden}}$ counts recall only on the parts of the object the photograph does not show. A missing or unloadable scene, or one with more than 4 million triangles, renders black and has no geometry. For calibration, the true meshes textured with the photograph score 0.879 on the development renders, and HARMONY's released reconstructions score 0.422 on the test split.

Prior work. HARMONY reconstructs such scenes by hierarchical agentic reasoning with a vision-language model and was compared with Gen3DSR, CAST (Yao et al., 2025), and SAM 3D (SAM 3D Team et al., 2025), among others.

Scene from Video

Purpose. Whether the agent can build a video-to-simulation pipeline: from the frames of an indoor video alone, estimate the cameras and reconstruct a scene of separate textured objects on a fixed background that has the true shape, reproduces the frames, and remains at rest under physical simulation.

Source and construction. The scenes are the nine released with HoloScene (Xia et al., 2025), each with frames, ground-truth cameras, per-object instance masks, a support graph, and a ground-truth surface. The test split is the six scenes evaluated by CoDimRecon (Xie et al., 2026): three Replica rooms (Straub et al., 2019) and

three ScanNet++ captures (Yeshwanth et al., 2023). Validation is one iGibson house (Shen et al., 2021); development is another iGibson house and a game room filmed with a phone (undistorted). The iGibson and phone sequences keep every third frame. Each face of the ground-truth surface is labeled by the masks that see it, and each object is marked dynamic if it rests on the floor or on another object and static if it hangs from a wall or the ceiling. The agent sees the development frames with all labels and the scoring code.

Setting. The input is only the RGB frames, without cameras, masks, or depth, as in CoDimRecon. The solution writes a glTF scene (at most 4 million triangles) and a camera for every frame, within 1200 s per scene, on one GPU and 16 CPU cores. The scene is placed in the ground-truth frame by the similarity transform that aligns its cameras with the true ones, rendered unlit from the true cameras, and each mesh node is assigned to the object or background whose masks it covers most. Any public pretrained model may be used, but not the HoloScene release or its source datasets; the model being evaluated is not called.

Starter solution. An empty room box in the mean color of the frames, with every camera at one point.

Metric. Per scene, $s = (G + A + P)/3$ with $G = (\mathrm{F1}_{5\,\mathrm{cm}} + \mathrm{NC})/200$ on 500,000 points of the whole scene, $A = (\mathrm{SSIM} + 1 - \mathrm{LPIPS})/2$ over all input frames (Wang et al., 2004; Zhang et al., 2018), and P the share of dynamic objects that are reconstructed (mask IoU at least 0.5) and stay in place; s is averaged over scenes. Stability follows CoDimRecon: each dynamic object, decomposed into convex hulls by CoACD (Wei et al., 2022), is released alone in MuJoCo (Todorov et al., 2012) for 3 s with all else fixed, and falls if it tilts more than 30° or moves more than $\max(5\,\mathrm{cm}, 0.2\,d)$, with d its bounding-box diagonal. A missing or unloadable scene scores as empty. The ground-truth surface colored from the frames scores 0.936 on the test split.

Prior work. HoloScene reconstructs simulation-ready interactive scenes from a video given its cameras and instance masks. CoDimRecon reconstructs such scenes agentically from RGB views alone, representing objects as curves, surfaces, and volumes, and compares with HoloScene and an agent that writes the scene in Blender.

A.5 Perception

Visual Search

Purpose. Whether the agent can build a visual search procedure around a model with a limited field of view: questions concern small objects in photographs of 2K to 8K resolution, while each model call may contain at most one megapixel, so the solution must decide where to look.

Source and construction. The questions come from V*Bench (Wu and Xie, 2024) (191 multiple-choice questions on about 2K photographs: 115 on object attributes with four options, 76 on relative position with two) and HR-Bench 8K (Wang et al., 2025a) (200 questions with four options on 151 photographs of about 8K resolution, half on a single instance and half across instances; of its four rotated listings of each question one is drawn with a fixed seed). V*Bench is split by category and HR-Bench by photograph, so no photograph appears in two splits. The test split has 140 questions (70 from each benchmark). Development and validation are stratified subsets of 70 questions (33 and 37) and 80 questions (40 and 40). The agent sees the development photographs, questions, answers, and, for V*Bench, the boxes of the queried objects.

Setting. The solution answers each question with an option letter. It may call the model being evaluated on average at most five times per question, with at most 1024×1024 pixels per call in total; it may crop, resize, and run open-weight models on one GPU, and may not use either benchmark beyond the development split.

Starter solution. Shows the model the whole photograph reduced to one megapixel with the question, in a single call.

Metric. Accuracy over the questions of the split; a question without a valid letter counts as wrong. Accuracies on V*Bench and HR-Bench 8K are reported separately. Chance accuracy on the validation split is 0.303.

Prior work. V*Bench was introduced with SEAL, which guides a multimodal model to search the image for the queried region, and HR-Bench with a training-free method that divides the image into patches and combines their answers. Later methods explore zoomed regions as a tree (ZoomEye (Shen et al., 2025)) or learn to crop by reinforcement learning (DeepEyes (Zheng et al., 2026)).

Video Quality

Purpose. Whether the agent can turn a general multimodal model into a judge of perceptual video quality that agrees with people, both for real user videos and for generated videos.

Source and construction. The rated videos are 100 videos of YouTube-UGC (Wang et al., 2019), chosen from its 1,380 videos with mean opinion scores (MOS) excluding 4K: seven per content category for 14 categories (eight for sports and gaming), evenly spaced in MOS within each category. Each lasts 16.6 to 20 s at 360p to 1080p, and its MOS (1 to 5) averages 105 to 154 crowdsourced ratings from the subjective study of Yim et al. (2020). The pairs are 50 pairs from a public human-preference study of Veo 3.1 (Rapidata, 2025), each pitting Veo 3.1 against one of seven other text-to-video models on the same prompt (5 to 6 s clips). Each pair has 15 reliability-weighted crowd votes on three questions: prompt alignment, fewer visual artifacts, and overall preference, the last two judged without the prompt. A clip with more than 55% of the votes wins; otherwise the question is a tie. The videos are dealt in order of category and MOS, and the pairs in order of opposing model, into 30, 30, and 40 videos and 15, 15, and 20 pairs; the clip order in each pair is randomized. The agent sees the development videos, frames at two per second, and labels with vote shares and rating spreads.

Setting. The solution scores each video and answers the three questions for each pair. It may call the model being evaluated on average at most three times per item, with at most eight 1024×576 frames' worth of pixels per call, and may run open-weight models on one GPU; it may not use the source datasets.

Starter solution. Asks for a rating from 1 to 5 given eight evenly spaced frames, and answers all three questions about a pair given six frames of each clip.

Metric. $s = (\text{SRCC} + \text{Acc})/2$, where SRCC is Spearman's rank correlation (Spearman, 1904) between the solution's scores and the MOS over the split's videos (a video without a score receives the median of the others), and Acc is the share of non-tied pair questions answered correctly, pooled over the three criteria (a missing answer is wrong). PLCC and per-criterion accuracies are also reported. Chance is 0.25.

Prior work. YouTube-UGC was collected for video compression research, and its MOS serve to evaluate blind quality models such as VIDEVAL (Tu et al., 2021) and DOVER (Wu et al., 2023).

A.6 Infrastructure

VAE Decode

Purpose. Whether the agent can lower the latency of a fixed numerical computation, the video decoder of a generative model, without changing what it computes.

Source and construction. The computation is the decoder of the MiniMax H3 autoencoder (MiniMax, 2026b), a 36-layer vision transformer of about 2.4 billion parameters with FP32 weights run under FP16 autocast. It decodes a window of seven latent frames of 24 channels into 28 RGB frames, in 256×256 tiles with at least 64 pixels of overlap and linear blending; the reference implementation and weights are pinned by hash. Inputs are latent windows of encoded videos at two sizes, 544×960 and 480×832 pixels, split by scene: 6 development, 6 validation, and 12 test scenes. Each request is either a real window with fresh noise ($+0.05\,\mathcal{N}(0, 1)$) or a synthetic stress window (Gaussian latents scaled per channel, clipped at ± 5). Per size and repetition there are four real and one stress request on development (two repetitions) and validation (six), and eight and two on test (six), so 20, 60, and 120 requests; no input repeats. The development windows are visible.

Setting. The solution is a decoder that receives latents in host memory and writes the full output into a host buffer. Custom Triton or CUDA kernels built from source, fusions, CUDA graphs, compilation, batching of tiles, and overlapping transfers are allowed. Lower precision, approximate or sparse attention, skipped tokens, tiles, or frames, reuse across requests, special treatment of stress windows, more than one GPU, shipped binaries, and existing optimized implementations of this decoder are not; accepted versions are reviewed for these rules. One H200, 90 s of setup, 30 s per request; the solution does not call the model being evaluated.

Starter solution. A direct eager PyTorch translation of the reference, bit-exact with it.

Metric. Every frame must be finite with a per-frame RMSE below 6.3×10^{-4} (1.9×10^{-3} for stress windows) and a maximum absolute error below 0.12 relative to the reference; these bounds are twice the largest error of legal reorderings, while FP16 accumulation, BF16, FP8, or dropped tokens exceed them. Fresh processes of the

solution and of the starter serve the same requests side by side, timed from input to copied output. With cost $C = \bar{t} + t_{\text{setup}}/1024$ ($\bar{t}$ the mean latency), a repetition scores $0.75 \log_2(C_{\text{starter}}/C)_{544\times960} + 0.25 \log_2(C_{\text{starter}}/C)_{480\times832}$; s is the mean over repetitions without the highest and lowest. One gate failure, crash, or timeout, or a median cost or 95th-percentile latency above 1.5 times the starter's, gives $s = -1$. The reported speedup is 2^s.

Prior work. The serving frameworks vLLM-Omni and SGLang (Zheng et al., 2024) include optimized implementations of the H3 decoder; these and code derived from them may not be used.

B Protocol and harness

This appendix specifies the parts of the protocol that Sections 3.2 and 3.3 describe only in outline: the structure of a round, the capabilities of the executor and the improver, the acceptance rule, the budget and the conditions under which a run ends, and the measures that keep the evaluation intact. The rules are identical for all models and tasks; the budget and two settings of the executor phase are fixed per task and listed in Table 10.

B.1 Round structure

A run begins by scoring the starter solution on the validation split, which gives the score against which the first candidate is compared. Each round then proceeds in five steps: an executor phase that produces a candidate solution; a check that reverts every change the executor made outside the files it may modify; the evaluation of the candidate on the validation split; the acceptance decision of Equation (2); and an improver phase that revises the harness.

The executor phase normally ends when the executor declares its work finished, and the solution in its final state becomes the candidate. It ends earlier when one of the executor's evaluations on the development split exceeds the development score of the accepted solution by a fixed share of the distance between the starter's validation score and the validation score of a strong known solution (Table 10); the solution as it was in that evaluation is then submitted. If the phase is interrupted before the executor finishes, because it reached its call limit, its time limit, or the budget, the best state that the executor's development evaluations showed to exceed the accepted solution is submitted, and nothing is submitted if there is none. A round that leaves the solution unchanged does not use a validation evaluation.

B.2 Executor

The executor works in a copy of the run's repository. Its tools allow it to run shell commands; to read, write, and edit files; to view images; to run experiments on the compute allocated to the task, including GPUs where the task provides them; to score the solution, or several variants of it at once, on the development split; and to query the remaining budget. It may write only to the solution, to a notes file that persists across rounds and is shared with the improver, and to a scratch directory that is emptied before each round. It observes development scores directly and the validation scores of earlier rounds through the run records, but it has no access to the validation or test data or to the evaluator. The starting harness allows 60 model calls per executor phase. Independently of the harness, the protocol stops a phase after the number of model calls given in Table 10 or after 180 minutes (600 minutes on the ARC tasks), and an experiment after 240 minutes.

B.3 Improver

After each round, the improver receives the executor's transcript, its experiments and notes, and the scores of all rounds so far. It may modify the executor's prompts, settings, control loop, and tools, as well as its own prompts, settings, acceptance rule, and a file of lessons that it maintains across rounds. A single improver phase may change at most 400 lines of the harness; a larger change is reverted in full. Every four rounds, the improver additionally reviews its own process: which kinds of harness changes were followed by accepted candidates and which by stalls, whether the acceptance rule is too strict or too permissive, and which lessons to record. Changes to the harness take effect from the following round. In the frozen-harness runs of Table 8, the improver is not run.

B.4 Acceptance rule

The margin δ in Equation (2) is determined by the acceptance rule, which is part of the improver's harness. In the starting harness, a candidate is accepted if its validation score exceeds that of the accepted solution by more than 2σ, where σ is the standard deviation of the starter's validation score over repeated evaluations (Table 10). On tasks for which σ was not measured, it is set to zero, and any improvement on the validation split is accepted. Because the rule belongs to the harness, the improver may revise it. If the rule fails to return a decision, and in every run with a frozen harness, the starting rule $\delta = 2\sigma$ is applied.

B.5 Budget and termination

A run ends at the first of the following events: the wall-clock time, GPU hours, tokens (input, including cached input, and output), or cost in dollars reach the task's limit; the number of validation evaluations reaches its limit; the round cap is reached; or five consecutive rounds produce neither a candidate nor a change to the harness, a rule that is not applied on LIBERO-MAX and LIBERO-PRO. The budget is checked before each model call, experiment, and development evaluation and is charged when each of them ends, so that calls running concurrently can exceed a limit by at most their combined cost. A submitted candidate is evaluated on the validation split even when the budget is exhausted. Model calls that the evaluator makes on the validation and test splits, for tasks whose solution calls the model, are recorded but not charged. Periods during which the model provider is unavailable are not charged either, and a phase interrupted by such an outage is repeated. After the run, the final solution and every accepted solution are scored on the test split (Section 3.3).

Table 10 gives the limits of each task. On a given task, all models have the same limits, with the following exceptions. Gemini 3.8 Flash and Gemini 4 Argon are stopped after 80 model calls per executor phase on Active Search, Ego-Exo Memory, 3D Tracking, Visual Search, and Video Quality. On Active Search, the dollar limit for open-weight models is 250. On RoboLab, the runs of Claude Sonnet 5.5, GPT-5.6 Sol, and GPT-6.1 Sol have 240 GPU hours. On Image Editing, the runs of Claude Opus 5.5, Claude Sonnet 5.5, GPT-6.1 Sol, and both Gemini models have 60 wall-clock hours and a dollar limit of 150. On Kinematics and Mass, the wall-clock limit of every model other than Claude Opus 5.5 is the time used by the run of Claude Opus 5.5 (11.3 and 11.0 hours). On Embodied-SWE and Spatial-VSI, Qwen3.8-Flash-Next is limited to the tokens used by the run of Claude Opus 5.5 (22M and 21M). On 3D Tracking, Claude Sonnet 5.5 has a dollar limit of 50. On RoboCasa, Muse Spark 1.3 has 48 GPU hours and 80M tokens. On RoboDojo and RoboCasa, runs that continued past round 10 are scored with the solution they had after round 10.

B.6 Isolation and integrity

After every phase, any change outside the paths that the phase may modify, including changes made through shell commands, is reverted and recorded as a violation; the evaluator, the data splits, and the protocol therefore cannot be altered by the agent. Files larger than 5 MB are excluded from the run's repository. If a phase crashes in two consecutive rounds, the harness is restored to the last version that completed a round without a crash. When a phase ends, all processes it started are terminated, so that no process left running can modify the solution afterwards. All model calls, experiments, and development evaluations are routed through the benchmark, which charges them to the budget and records them.

Table 10 Budget and run settings of each task. The limits of one run, as recorded by the runs of Section 4.2 and Table 12; for RoboDojo and RoboCasa they apply to each of their tasks. Wall: wall-clock hours. GPU: GPU hours. Tokens: input tokens, including cached input, and output tokens, in millions. Dollars: cost limit, which applies only to models with a price. Val.: validation evaluations. Rounds: the round cap. Calls: model calls after which an executor phase is stopped. Submit: the gain on the development split, as a share of the distance between the starter's validation score and that of a strong known solution, at which an executor phase ends and submits its candidate (–: only when the executor finishes). σ: the standard deviation of the starter's validation score, on the scale of the result tables (VAE Decode: $\log_2$ speedup; Text-to-Image: fixed-generator setting, then any-model setting); 0 where it was not measured. Settings that differ for some models are listed in the text.

	Limits						Executor phase		
Task	**Wall**	**GPU**	**Tokens**	**Dollars**	**Val.**	**Rounds**	**Calls**	**Submit**	σ
RoboDojo	24	256	150	300	20	10	80	2%	0
RoboCasa	24	256	150	300	20	10	80	2%	0
Active Search	10	12	80	120	8	15	40	2%	0
RoboLab	36	120	150	800	20	10	80	2%	7.0
LIBERO-MAX	72	256	150	300	16	15	80	2%	0
LIBERO-PRO	72	256	150	300	16	15	80	2%	0
Embodied-SWE	12	12	120	150	10	10	60	5%	2.0
Ego-Exo Memory	10	12	80	80	5	10	40	5%	2.0
Kinematics	68	272	400	600	20	60	80	2%	3.7
Mass	68	272	400	600	20	60	80	2%	1.6
ARC-AGI-1, -2, -3	10	0	30	100	40	10	80	–	0
Spatial-VSI	10	0	80	100	10	10	40	5%	1.5
Image Editing	24	256	150	300	20	15	80	100%	0.48
Multi-turn Editing	12	64	80	150	30	15	80	100%	0.55
Video Generation	12	64	80	10	30	200	80	2%	0.013
Text-to-Image	24	256	150	300	20	25	80	2%	0.20, 0.02
3D Tracking	10	12	80	100	15	15	40	2%	0
Scene from Photo	12	36	80	100	10	15	40	2%	0
Scene from Video	12	36	80	100	10	15	40	2%	0
Visual Search	10	12	80	100	10	10	40	5%	2.0
Video Quality	10	12	80	80	5	10	40	5%	2.0
VAE Decode	12	16	200	400	60	200	80	5%	0.019

C Full results

Table 13 to Table 32 give the results of every run, one table per task in the order of Table 1: the test metrics of the final solution, the test score of the starter solution where it was measured, the validation score at the start and the best validation score reached, and the number of rounds with the condition that ended the run. Table 12 reports, in the same layout, the four tasks on which only two models have been evaluated so far. Table 11 gives the average rank of every model over the tasks of the per-task tables, overall and per area.

Table 11 Leaderboard. Average rank over the tasks of the per-task tables that a model was run on, overall and per area (lower is better); models run on at least five of these tasks. Tasks: the number of these tasks the model was run on. Wins: the number of them on which it has the highest test score.

			Average rank by area						
Model	**Weights**	**Overall**	**Action**	**Reasoning**	**Creation**	**Reconstruction**	**Perception**	**Tasks**	**Wins**
Claude Opus 5.5	closed	**1.6**	**1.5**	1.8	1.7	**1.0**	**2.0**	20	**11**
Gemini 4 Argon	closed	2.0	1.7	2.0	**1.0**	3.0	3.0	9	4
GPT-6.1 Sol	closed	2.8	2.8	**1.8**	4.0	2.0	3.0	16	2
GPT-6 Astra	closed	3.8	3.5	2.8	6.5	4.3	3.0	15	2
Qwen3.8-Flash-Next	open	5.0	5.2	5.2	3.7	4.0	7.0	20	1
Qwen3.8-27B	open	5.8	6.0	5.7	4.7	6.3	6.5	19	0
Claude Sonnet 5.5	closed	6.4	6.0	2.5	7.0	11.0	9.0	8	1
Gemini 3.8 Flash	closed	6.5	6.0	6.0	9.0	8.0	5.5	8	0
DeepSeek-V4.1 Flash	open	6.7	8.0	6.0	8.0	5.0	6.5	11	0
GLM-5.3 Flash	open	7.4	6.0	7.7	9.5	4.0	8.0	10	0
Gemma-4-31B-it	open	7.5	6.4	7.8	6.7	8.0	10.5	17	0
MiniMax-M3	open	8.7	9.0	9.0	7.5	7.0	11.5	13	0

Table 12 Tasks evaluated with two models. Embodied-SWE, Spatial-VSI, Text-to-Image, and VAE Decode, in the layout of the per-task tables of this appendix. Score is the primary test metric of each task and Starter the primary test score of the starter solution; Validation gives the validation score of the starter and the best validation score reached during the run. Embodied-SWE: a robot arm in simulation must flip a slice of bread with a spatula and serve it onto a plate; the agent writes the robot program, and Score is the EmbodiedSWE-Bench rubric score, which credits each trajectory by the furthest stage of the task it reaches. Spatial-VSI: spatial questions about indoor videos, answered by a solver that calls the model; Score is the official VSI-Bench overall score. Text-to-Image: text-to-image generation on original Chinese and English prompts in two settings, one with the generator fixed to Qwen-Image-2.1 and one with any open-weight model or workflow; each image is scored by constraint completion C times image quality Q as rated by a frozen vision-language judge, $C(0.5+0.5Q)$, and Score, Starter, and Validation are the means over the two settings (test score of GPT-6 Astra 94.5 with the fixed generator and 93.4 with any model; of GPT-5.6 Sol 92.0 and 95.8), with Rounds counted per setting. VAE Decode: making the decoder of a video model faster without changing its output; Score is the speedup of the decoder over the starter on unseen inputs, and the starter has a speedup of 1.00×. On Embodied-SWE and Spatial-VSI, Qwen3.8-Flash-Next is served on our own GPUs and has no price, so its run is limited to the number of tokens used by the run of Claude Opus 5.5 (22M on Embodied-SWE and 21M on Spatial-VSI). On Text-to-Image, a pre-release version of the task, both models run at medium reasoning effort with at most 25 rounds in each setting. Scores other than the speedup are multiplied by 100.

	Model			**Hidden test**		**Validation**	**Run**	
Task	**#**	**Name**	**Weights**	**Score**	**Starter**	**Start → best**	**Rounds**	**Ended by**
Embodied-SWE	1	Claude Opus 5.5	closed	**44.4**	0.0	0.0 → 41.2	7	time budget
	2	Qwen3.8-Flash-Next	open	23.8	0.0	0.0 → 23.8	5	token budget
Spatial-VSI	1	Claude Opus 5.5	closed	**65.9**	61.0	63.6 → 70.2	10	round limit
	2	Qwen3.8-Flash-Next	open	56.1	53.4	58.0 → 60.2	6	token budget
Text-to-Image	1	GPT-6 Astra	closed	**94.0**	91.9	94.1 → 95.5	25	round limit
	2	GPT-5.6 Sol	closed	93.9	91.9	94.1 → 95.0	25	round limit
VAE Decode	1	Claude Opus 5.5	closed	**2.13×**	1.00×	1.00× → 2.14×	41	no further progress
	2	GPT-6 Astra	closed	1.96×	1.00×	1.00× → 1.95×	16	time budget

Table 13 RoboDojo. A dual-arm ARX X5 robot in simulation with the model as its policy, on 10 tasks (task score 0–100). Models are ranked by Suite on the hidden test split, the mean task score over all 10 tasks. The task columns give the test score of the final solution of each task, and Tasks > 0 the number of tasks with a test score above 0. Validation gives the suite mean of the starter's validation score and of the best validation score reached; Rounds is summed over the tasks. Every run counts up to round 10, and each task uses the model's latest run with a final test score. The starter's test score was not measured in every run and is not shown. Only models with a run on every task are listed.

Model			Hidden test metrics												Validation	Run
#	Name	Weights	Suite	Tasks > 0	Bottles	Number	Clothes	Tubes	Eggs	Conveyor	Tic-tac-toe	Xylophone	Push-T	Dolls	Start → best	Rounds
1	Claude Opus 5.5	closed	**55.1**	9	**70**	27	**68**	72	19	0	95	**100**	**80**	20	0.5 → 68.9	100
2	Gemini 4 Argon	closed	46.8	8	**70**	**71**	44	**100**	**23**	0	**100**	20	40	0	0.9 → 47.7	90
3	GPT-6.1 Sol	closed	41.5	9	14	11	36	80	4	**60**	90	80	0	**40**	0.2 → 50.5	100
4	GPT-6 Astra	closed	28.6	7	37	18	4	0	7	40	**100**	80	0	0	0.7 → 30.0	100
5	DeepSeek-V4.1 Flash	open	17.8	5	16	11	0	16	0	0	75	60	0	0	0.0 → 18.7	100
6	Muse Wonder 1.0	closed	10.6	5	6	2	4	0	4	0	90	0	0	0	0.0 → 15.4	100
7	Qwen3.8-Flash-Next	open	8.4	2	9	0	0	0	0	0	75	0	0	0	0.0 → 9.9	100
8	MiniMax-M3	open	0.0	0	0	0	0	0	0	0	0	0	0	0	0.0 → 0.0	100

Table 14 RoboCasa. A mobile manipulator in simulated kitchens (RoboCasa365) with the model as its policy, on 8 tasks (success rate over 5 test episodes per task). Models are ranked by Suite on the hidden test split, the mean success rate over the 8 tasks. The task columns give the test success rate of the final solution of each task, and Tasks > 0 the number of tasks with a test success rate above 0. Validation gives the suite mean of the starter's validation score and of the best validation score reached; Rounds is summed over the tasks. Every run counts up to round 10, and each task uses the model's latest run with a final test score. The starter's test score was not measured in every run and is not shown. A ninth task, Heat kebab sandwich, is left out: no model reached a success on it in validation or test. †Muse Spark 1.3 is accessed through an internal endpoint; it is shown for comparison and not ranked. All scores are multiplied by 100.

Model			Hidden test metrics										Validation	Run
#	Name	Weights	Suite	Tasks > 0	Coffee	Kettle	Cabinet	Drawer	Hot dogs	Stove	Microwave	Faucet	Start → best	Rounds
1	Claude Opus 5.5	closed	**57.5**	7	0	**40**	**60**	**100**	**60**	20	**100**	**80**	5.0 → 60.0	80
2	Gemini 4 Argon	closed	40.0	7	0	20	**60**	80	20	20	80	40	2.5 → 57.5	75
3	GPT-6 Astra	closed	37.5	6	**60**	20	0	80	20	**40**	80	0	0.0 → 47.5	74
4	GPT-6.1 Sol	closed	32.5	6	40	0	20	80	0	**40**	60	20	0.0 → 45.0	71
5	GLM-5.3 Flash	open	7.5	1	0	0	0	0	0	0	0	60	0.0 → 7.5	80
6	Qwen3.8-Flash-Next	open	5.0	2	0	0	0	0	0	0	20	20	2.5 → 12.5	80
7	Gemma-4-31B-it	open	2.5	1	0	0	0	0	0	0	20	0	0.0 → 1.3	78
8	MiniMax-M3	open	0.0	0	0	0	0	0	0	0	0	0	0.0 → 2.5	80
8	Qwen3.8-27B	open	0.0	0	0	0	0	0	0	0	0	0	0.0 → 3.8	80
8	DeepSeek-V4.1 Flash	open	0.0	0	0	0	0	0	0	0	0	0	0.0 → 12.5	80
–	Muse Spark 1.3†	closed	0.0	0	0	0	0	0	0	0	0	0	2.5 → 5.0	80

Table 15 **Active Search.** Active visual search in 360° scenes (H*Bench). Models are ranked by SPL on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. SPL is success weighted by path length; Success is the mean of Object SR and Path SR; Object SR and Path SR are the success rates of object-search and path-search episodes; Avg. steps is the average number of actions per episode. All scores except Avg. steps are multiplied by 100.

Model			Hidden test metrics						Validation	Run	
#	Name	Weights	SPL	Starter	Success	Object SR	Path SR	Avg. steps	Start → best	Rounds	Ended by
1	Gemini 4 Argon	closed	79.7	45.9	**100.0**	**100.0**	**100.0**	2.53	40.8 → 67.2	12	no further progress
2	GPT-6 Astra	closed	78.2	45.2	93.4	93.9	92.9	2.23	35.0 → 64.8	15	round limit
3	GPT-6.1 Sol	closed	74.4	39.3	85.3	84.9	85.7	1.98	26.6 → 59.7	15	round limit
4	Claude Opus 5.5	closed	74.2	31.1	83.2	87.9	78.6	2.21	31.6 → 66.3	15	round limit
5	Qwen3.8-27B	open	69.0	18.1	78.1	84.9	71.4	2.23	20.0 → 51.7	15	round limit
6	Gemini 3.8 Flash	closed	65.5	40.2	80.2	81.8	78.6	2.51	32.7 → 59.7	11	token budget
7	GLM-5.3 Flash	open	60.5	21.5	74.0	90.9	57.1	2.83	17.1 → 52.5	15	token budget
8	Qwen3.8-Flash-Next	open	59.3	14.1	65.9	81.8	50.0	1.98	16.7 → 50.3	12	time budget
9	Claude Sonnet 5.5	closed	55.8	26.0	62.3	81.8	42.9	2.23	26.1 → 53.0	15	round limit
10	Gemma-4-31B-it	open	52.2	24.4	59.3	75.8	42.9	2.15	17.6 → 44.4	12	validation budget
11	DeepSeek-V4.1 Flash	open	47.1	26.6	59.3	75.8	42.9	2.38	19.9 → 35.8	6	time budget
12	MiniMax-M3	open	32.5	16.0	40.5	66.7	14.3	2.72	10.0 → 32.8	15	round limit

Table 16 **RoboLab.** One generalist tabletop manipulation policy for many tasks (pick and place, stacking, reorienting, choosing objects by color, size, kind or position), scored on tasks it never saw during development. Models are ranked by Score on the hidden test split. Test metrics are those of the final solution. Validation gives the validation score of the starter and the best validation score reached during the run. Score is the RoboLab score, which credits a successful episode fully and an unsuccessful one by the fraction of its subtasks completed; Success rate is the fraction of successful episodes. Score is on its own 0–100 scale, and Success rate is multiplied by 100.

Model			Hidden test metrics		Validation	Run	
#	Name	Weights	Score	Success rate	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	71.9	**66.7**	16.7 → 70.8	10	round limit
2	GPT-6.1 Sol	closed	62.5	56.2	18.8 → 68.8	10	round limit
3	Claude Sonnet 5.5	closed	41.0	31.2	8.3 → 47.9	10	round limit
4	GPT-5.6 Sol	closed	37.8	29.2	2.1 → 29.2	10	round limit
5	GPT-6 Astra	closed	33.0	20.8	31.2 → 47.9	10	round limit
6	Qwen3.8-27B	open	31.6	22.9	4.2 → 37.5	10	round limit
7	Qwen3.8-Flash-Next	open	17.4	6.2	4.2 → 45.8	10	round limit
8	MiniMax-M3	open	8.3	4.2	8.3 → 8.3	10	round limit
9	Gemma-4-31B-it	open	5.2	0.0	4.2 → 8.3	10	GPU budget

Table 17 LIBERO-MAX. A Franka arm in LIBERO kitchen scenes where the world changes once the episode is under way; the model's controller drives it, with a frozen X-VLA policy, cameras and depth to call on (LIBERO-MAX, dynamic success rate over paired episodes). Models are ranked by Dynamic success on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Dynamic success is the success rate on episodes in which the scene is perturbed once while the episode is under way; Base success is the success rate on the paired episodes without the perturbation. All scores are multiplied by 100.

Model			Hidden test metrics			Validation	Run	
#	Name	Weights	Dynamic success	Starter	Base success	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	**45.8**	29.2	58.3	31.2 → 50.0	15	round limit
1	Qwen3.8-Flash-Next	open	**45.8**	33.3	**62.5**	31.2 → 37.5	15	round limit
3	GPT-6.1 Sol	closed	41.7	29.2	45.8	31.2 → 37.5	15	round limit
4	Gemma-4-31B-it	open	29.2	25.0	54.2	31.2 → 43.8	15	round limit
5	Qwen3.6-35B-A3B	open	25.0	25.0	54.2	31.2 → 31.2	15	round limit
5	Qwen3.8-27B	open	25.0	25.0	54.2	31.2 → 31.2	15	round limit

Table 18 LIBERO-PRO. A Franka arm in LIBERO kitchen scenes with the task perturbed before it starts (objects, positions, wording, goal or camera); the model's controller drives it, with a frozen X-VLA policy, cameras and depth to call on (LIBERO-PRO, success rate). Models are ranked by Success rate on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Success rate is the fraction of test episodes in which the task is completed. All scores are multiplied by 100.

Model			Hidden test metrics		Validation	Run	
#	Name	Weights	Success rate	Starter	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	**63.3**	53.3	50.0 → 75.0	15	round limit
2	Gemma-4-31B-it	open	53.3	53.3	50.0 → 50.0	15	round limit
2	Qwen3.6-35B-A3B	open	53.3	53.3	50.0 → 50.0	15	round limit
2	GPT-6.1 Sol	closed	53.3	53.3	50.0 → 60.0	15	round limit
2	Qwen3.8-Flash-Next	open	53.3	53.3	50.0 → 50.0	15	round limit
6	Qwen3.8-27B	open	50.0	53.3	50.0 → 55.0	15	round limit

Table 19 Ego-Exo Memory. Memory reasoning over synchronized egocentric and exocentric videos (EgoExoMem, Ego-Exo4D part). Models are ranked by Accuracy on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Accuracy is measured on all test questions; Cooking and Bike repair give the accuracy on the questions about cooking and bike-repair videos. All scores are multiplied by 100.

Model			Hidden test metrics				Validation	Run	
#	Name	Weights	Accuracy	Starter	Cooking	Bike repair	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	**72.5**	61.7	**73.6**	**69.0**	63.7 → 73.8	10	round limit
2	Gemini 4 Argon	closed	70.0	65.0	**73.6**	58.6	75.0 → 76.2	6	validation budget
3	GPT-6.1 Sol	closed	67.5	57.5	69.2	62.1	67.5 → 75.0	10	round limit
4	GPT-6 Astra	closed	65.8	60.0	69.2	55.2	75.0 → 81.2	10	round limit
5	Gemini 3.8 Flash	closed	63.3	59.2	64.8	58.6	63.7 → 73.8	10	round limit
6	Qwen3.8-27B	open	60.8	55.0	62.6	55.2	58.8 → 72.5	10	round limit
7	Qwen3.8-Flash-Next	open	59.2	54.2	59.3	58.6	61.3 → 67.5	10	round limit
7	Gemma-4-31B-it	open	59.2	52.5	63.7	44.8	61.3 → 66.2	5	time budget
9	GLM-5.3 Flash	open	56.7	52.5	60.4	44.8	60.0 → 65.0	8	validation budget
10	MiniMax-M3	open	51.7	52.5	52.8	48.3	47.5 → 56.2	10	round limit
11	DeepSeek-V4.1 Flash	open	50.8	55.8	53.8	41.4	66.2 → 66.2	7	validation budget

Table 20 Kinematics. Sizes, distances, velocities and accelerations in physical units from a short video and one known quantity, with a fixed open VLM (QuantiPhy). Models are ranked by MRA on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. MRA is the mean relative accuracy, which averages over a range of thresholds whether the relative error of a predicted quantity falls below the threshold. All scores are multiplied by 100. †Muse Spark 1.3 is accessed through an internal endpoint; it is shown for comparison and not ranked.

Model			Hidden test metrics		Validation	Run	
#	Name	Weights	MRA	Starter	Start → best	Rounds	Ended by
1	Claude Sonnet 5.5	closed	57.2	33.1	26.4 → 56.5	30	time budget
2	Claude Opus 5.5	closed	54.6	43.6	29.0 → 57.7	29	no further progress
3	GPT-5.6 Sol	closed	50.0	48.8	34.2 → 41.7	18	time budget
4	Qwen3.8-Flash-Next	open	49.2	37.3	22.1 → 45.6	14	time budget
5	Gemini 3.8 Flash	closed	48.2	42.5	34.0 → 41.5	31	time budget
6	Gemini 3.7 Flash	closed	46.3	39.0	29.8 → 39.8	24	no further progress
7	Qwen3.8-27B	open	43.4	42.1	20.0 → 53.5	36	time budget
8	Grok 4.6	closed	37.2	47.2	24.6 → 24.6	25	time budget
–	Muse Spark 1.3†	closed	45.5	36.4	31.7 → 48.5	45	no further progress

Table 21 Mass. A product's weight in kilograms from its catalog photos, type and dimensions, with a fixed open VLM (QuantiPhy-style, Amazon Berkeley Objects). Models are ranked by MRA on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. MRA is the mean relative accuracy, which averages over a range of thresholds whether the relative error of a predicted quantity falls below the threshold. All scores are multiplied by 100. †Muse Spark 1.3 is accessed through an internal endpoint; it is shown for comparison and not ranked.

Model			Hidden test metrics		Validation	Run	
#	Name	Weights	MRA	Starter	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	63.5	50.4	46.4 → 63.0	29	no further progress
2	GPT-5.6 Sol	closed	63.0	49.8	47.7 → 61.1	49	no further progress
3	Qwen3.8-27B	open	62.7	49.1	50.3 → 60.9	13	no further progress
4	Claude Sonnet 5.5	closed	62.5	51.0	50.2 → 61.7	60	round limit
5	Gemini 3.7 Flash	closed	60.6	51.8	46.7 → 61.1	25	no further progress
6	Qwen3.8-Flash-Next	open	60.4	51.4	49.4 → 61.1	24	time budget
7	Grok 4.6	closed	59.3	51.3	46.2 → 57.1	21	no further progress
8	Gemini 3.8 Flash	closed	57.5	50.0	44.3 → 59.1	32	time budget
–	Muse Spark 1.3†	closed	56.9	51.2	45.3 → 56.3	38	no further progress

Table 22 ARC-AGI-1. Abstract reasoning on grid puzzles from a few examples (ARC-AGI-1 public evaluation set, pass@2). Models are ranked by Accuracy on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Accuracy is the fraction of test puzzles solved within two attempts (pass@2). All scores are multiplied by 100.

Model			Hidden test metrics		Validation	Run	
#	Name	Weights	Accuracy	Starter	Start → best	Rounds	Ended by
1	GPT-6 Astra	closed	99.8	99.2	98.8 → 100.0	10	round limit
2	GPT-6.1 Sol	closed	98.8	98.8	98.8 → 98.8	10	round limit
3	Claude Opus 5.5	closed	98.0	96.0	98.8 → 100.0	10	round limit
4	DeepSeek-V4.1 Flash	open	91.0	91.0	93.8 → 93.8	2	time budget
5	Qwen3.8-Flash-Next	open	80.8	80.8	90.0 → 90.0	1	time budget
6	Qwen3.8-27B	open	77.2	75.5	83.8 → 95.0	5	time budget
7	GLM-5.3 Flash	open	63.5	63.5	76.2 → 76.2	6	time budget
8	Gemma-4-31B-it	open	54.0	54.0	67.5 → 67.5	2	time budget

Table 23 **ARC-AGI-2.** Abstract reasoning on grid puzzles from a few examples (ARC-AGI-2 public evaluation set, pass@2). Models are ranked by Accuracy on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Accuracy is the fraction of test puzzles solved within two attempts (pass@2). All scores are multiplied by 100.

Model			Hidden test metrics		Validation	Run	
#	Name	Weights	Accuracy	Starter	Start → best	Rounds	Ended by
1	GPT-6.1 Sol	closed	**95.8**	92.5	98.8 → 100.0	10	round limit
2	Claude Opus 5.5	closed	95.0	78.3	96.2 → 100.0	10	round limit
3	GPT-6 Astra	closed	93.3	93.3	98.8 → 98.8	10	round limit
4	DeepSeek-V4.1 Flash	open	46.7	46.7	92.5 → 92.5	4	time budget
5	Qwen3.8-Flash-Next	open	22.5	22.5	82.5 → 82.5	4	time budget
6	Qwen3.8-27B	open	19.2	15.0	77.5 → 80.0	3	time budget
7	GLM-5.3 Flash	open	14.2	14.2	68.8 → 68.8	2	time budget
8	MiniMax-M3	open	4.2	4.2	72.5 → 72.5	1	time budget
9	Gemma-4-31B-it	open	1.7	1.7	60.0 → 60.0	2	time budget

Table 24 **ARC-AGI-3.** Interactive puzzle games learned by playing, scored by action efficiency against humans (ARC-AGI-3 public games). Models are ranked by RHAE on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. RHAE is the official ARC-AGI-3 score, which rates each completed level by the agent's action efficiency relative to human players, weights the levels of a game by their number, and averages over games; it lies on a 0–100 scale and is not rescaled.

Model			Hidden test metrics		Validation	Run	
#	Name	Weights	RHAE	Starter	Start → best	Rounds	Ended by
1	GPT-6.1 Sol	closed	**44.2**	25.4	29.8 → 48.2	10	round limit
2	Claude Opus 5.5	closed	40.4	21.6	33.0 → 49.3	10	round limit
3	GPT-6 Astra	closed	36.5	36.5	39.0 → 39.0	9	no further progress
4	Qwen3.8-Flash-Next	open	0.9	0.0	0.0 → 4.1	7	time budget
5	DeepSeek-V4.1 Flash	open	0.6	0.6	3.8 → 3.8	3	time budget
6	Qwen3.8-27B	open	0.3	0.0	0.0 → 0.3	7	time budget
7	Gemma-4-31B-it	open	0.0	0.0	0.0 → 0.0	2	time budget

Table 25 Image Editing. Instruction-based image editing with any open model or tool, covering local, global, text, retouching and restoration, reasoning, and compound edits. Models are ranked by Overall on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Overall is the mean score over six families of edits, each edit scored by a vision-language judge and, where applicable, by programmatic checks: Local (a single object), Global (style and global changes), Text, Retouch (retouching and restoration), Reasoning (edits that require inference), and Compound (several edits in one instruction). All scores are multiplied by 100.

Model			Hidden test metrics								Validation	Run	
#	Name	Weights	Overall	Starter	Local	Global	Text	Retouch	Reasoning	Compound	Start → best	Rounds	Ended by
1	Gemini 4 Argon	closed	**74.2**	55.5	80.2	72.2	87.8	62.4	**73.7**	68.7	53.4 → 73.9	15	round limit
2	Claude Opus 5.5	closed	73.8	56.2	78.2	73.4	88.0	**66.2**	66.9	70.3	54.8 → 77.1	15	round limit
3	Qwen3.8-Flash-Next	open	70.0	56.2	**85.0**	75.2	87.4	60.8	42.1	69.3	57.6 → 67.9	15	round limit
4	GPT-6.1 Sol	closed	69.9	56.2	82.8	76.8	**90.6**	62.7	42.8	63.7	55.6 → 71.1	15	round limit
5	Qwen3.8-27B	open	69.7	55.7	70.8	70.2	86.4	53.7	66.6	70.5	56.4 → 64.8	15	round limit
6	GPT-6 Astra	closed	69.7	56.5	83.1	73.5	87.1	58.7	34.7	**81.1**	58.0 → 68.9	15	round limit
7	Gemma-4-31B-it	open	68.4	56.1	76.0	73.3	85.7	62.9	44.3	68.4	57.4 → 67.8	15	round limit
8	DeepSeek-V4.1 Flash	open	67.6	56.2	78.2	**79.3**	77.4	53.9	40.2	76.6	57.6 → 69.9	15	round limit
9	Gemini 3.8 Flash	closed	67.2	56.7	78.1	70.5	85.6	60.7	31.5	76.8	53.4 → 68.1	15	round limit
10	Claude Sonnet 5.5	closed	66.0	57.0	77.3	71.0	86.0	61.4	41.1	59.0	53.4 → 65.1	15	round limit
11	MiniMax-M3	open	64.3	55.6	70.3	69.0	83.3	57.4	33.3	72.4	57.6 → 63.2	15	round limit
12	GLM-5.3 Flash	open	60.6	56.6	69.1	65.0	80.4	53.7	29.1	66.1	57.6 → 63.9	15	round limit

Table 26 Multi-turn Editing. Edit a photo over five turns, each instruction revealed only at its turn and applied to the previous result (MSE-Bench). Models are ranked by Score on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Score is the mean per-turn score of a vision-language judge over the five turns of each session; Turn 1, Turn 3, and Turn 5 are the fractions of sessions in which every turn up to the first, third, and fifth is judged successful. All scores are multiplied by 100.

Model			Hidden test metrics					Validation	Run	
#	Name	Weights	Score	Starter	Turn 1	Turn 3	Turn 5	Start → best	Rounds	Ended by
1	Gemini 4 Argon	closed	**71.1**	64.8	**95.0**	62.5	**50.0**	68.0 → 72.2	6	time budget
2	Claude Opus 5.5	closed	70.3	64.6	85.0	60.0	37.5	68.1 → 72.3	15	round limit
3	Qwen3.8-27B	open	70.2	–	87.5	60.0	47.5	68.2 → 70.8	8	time budget
4	Claude Sonnet 5.5	closed	69.8	65.1	87.5	57.5	42.5	68.3 → 72.1	7	time budget
5	Qwen3.8-Flash-Next	open	69.2	64.8	90.0	60.0	42.5	68.3 → 70.4	8	time budget
6	GPT-6.1 Sol	closed	68.3	65.0	**95.0**	57.5	42.5	68.4 → 72.1	7	time budget
7	GPT-6 Astra	closed	67.8	65.1	**95.0**	**67.5**	**50.0**	68.0 → 71.2	15	round limit
8	Gemma-4-31B-it	open	66.3	65.0	**95.0**	**67.5**	35.0	68.3 → 69.9	6	time budget

Table 27 Video Generation. Text-to-video workflows around a frozen video model (MiniMax H3, 480p), scored on compositional prompts: attribute and action binding, numeracy, spatial relations (T2V-CompBench). Models are ranked by Score on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Score is the mean of the T2V-CompBench scores for attribute binding, action binding, numeracy, and spatial relations. All scores are multiplied by 100.

Model			Hidden test metrics		Validation	Run	
#	Name	Weights	Score	Starter	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	**93.9**	81.7	68.1 → 92.0	37	no further progress
2	GPT-6.1 Sol	closed	91.5	81.7	68.1 → 87.4	13	no further progress
3	Qwen3.8-Flash-Next	open	91.4	80.6	69.1 → 90.7	6	time budget
4	MiniMax-M3	open	87.3	80.6	69.1 → 76.5	5	time budget
5	Gemma-4-31B-it	open	87.1	80.6	69.1 → 77.8	4	time budget
6	Qwen3.8-27B	open	81.8	80.6	69.1 → 71.9	5	time budget
7	GLM-5.3 Flash	open	79.6	80.6	69.1 → 82.0	10	time budget

Table 28 3D Tracking. 3D instance tracking (IGGT). Models are ranked by T-mIoU on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. T-mIoU is the mean intersection over union of the predicted instance tracks; T-SR is the fraction of tracks recovered successfully; $\mathcal{J}$, $\mathcal{F}$, and $\mathcal{J}\&\mathcal{F}$ are region similarity, contour accuracy, and their mean. All scores are multiplied by 100.

Model			Hidden test metrics						Validation	Run	
#	Name	Weights	T-mIoU	Starter	T-SR	$\mathcal{J}$	$\mathcal{F}$	$\mathcal{J}\&\mathcal{F}$	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	**79.6**	21.2	65.0	**76.1**	**73.2**	**74.7**	19.4 → 84.7	15	round limit
2	GPT-6.1 Sol	closed	77.3	21.2	60.4	74.5	71.1	72.8	19.4 → 80.2	9	time budget
3	Gemini 4 Argon	closed	73.8	21.2	**69.6**	70.6	65.9	68.3	19.4 → 78.6	8	time budget
4	GLM-5.3 Flash	open	61.2	21.2	27.5	54.6	52.7	53.6	19.4 → 64.5	12	time budget
5	DeepSeek-V4.1 Flash	open	60.2	21.2	30.8	54.5	53.8	54.2	19.4 → 70.1	9	time budget
6	Qwen3.8-Flash-Next	open	58.9	21.2	30.8	54.6	51.8	53.2	19.4 → 65.7	12	time budget
7	Muse Wonder 1.0	closed	57.9	21.2	25.8	53.1	50.8	51.9	19.4 → 62.2	6	time budget
8	Gemini 3.8 Flash	closed	53.7	21.2	22.1	46.4	46.2	46.3	19.4 → 54.0	12	time budget
9	GPT-6 Astra	closed	52.8	21.2	20.4	48.5	46.5	47.5	19.4 → 57.3	15	round limit
10	Qwen3.8-27B	open	52.4	21.2	12.5	40.8	38.3	39.6	19.4 → 53.9	9	time budget
11	Claude Sonnet 5.5	closed	47.9	21.2	23.8	47.6	45.7	46.7	19.4 → 50.1	13	time budget
12	Gemma-4-31B-it	open	32.0	21.2	5.4	26.2	21.4	23.8	19.4 → 29.1	7	time budget
13	MiniMax-M3	open	27.9	21.2	1.7	23.1	21.5	22.3	19.4 → 29.4	15	time budget

Table 29 Scene from Photo. Image-to-3D: rebuild a furnished room from one photo as a Blender scene, every piece of furniture a whole 3D object (HARMONY300). Models are ranked by Score on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Score is the mean of three terms: Appearance compares renderings of the scene with the photo, Objects measures how well each annotated object is recovered as a complete and correctly placed object, and Geometry compares the scene geometry with the ground truth. †GPT-6 Astra runs at its highest reasoning effort, whereas every other run uses the standard effort; it is shown for comparison and not ranked. All scores are multiplied by 100.

Model			Hidden test metrics					Validation	Run	
#	Name	Weights	Score	Starter	Appearance	Objects	Geometry	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	**66.5**	7.1	**94.8**	62.5	**42.3**	7.7 → 67.8	15	round limit
2	GPT-6 Astra	closed	53.2	7.1	63.5	**62.5**	33.8	7.7 → 54.8	15	round limit
3	Qwen3.8-Flash-Next	open	40.0	7.1	78.7	16.5	24.7	7.7 → 43.2	10	time budget
4	MiniMax-M3	open	33.3	7.1	81.8	0.2	17.8	7.7 → 32.9	15	round limit
5	Qwen3.8-27B	open	29.0	7.1	63.0	4.4	19.7	7.7 → 28.6	11	time budget
6	Gemma-4-31B-it	open	15.4	7.1	21.4	0.4	24.4	7.7 → 17.2	11	time budget
–	GPT-6 Astra†	closed	61.6	7.1	81.1	60.0	43.6	7.7 → 63.1	12	time budget

Table 30 Scene from Video. Video-to-3D: reconstruct an indoor scene from its RGB frames alone as a simulation-ready scene, one textured mesh per object (HoloScene scenes, CoDimRecon setting). Models are ranked by Score on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Score is the mean of three terms: Geometry combines the F1 score and the normal consistency of the reconstructed surfaces, Appearance combines the SSIM and LPIPS of renderings, and Physics is the fraction of objects that remain stable in simulation. †GPT-6 Astra runs at its highest reasoning effort, whereas every other run uses the standard effort; it is shown for comparison and not ranked. All scores are multiplied by 100.

Model			Hidden test metrics					Validation	Run	
#	Name	Weights	Score	Starter	Geometry	Appearance	Physics	Start → best	Rounds	Ended by
1	Claude Opus 5.5	closed	**70.9**	26.7	**68.8**	**73.2**	**70.7**	17.3 → 27.9	15	round limit
2	GPT-6 Astra	closed	53.4	26.7	66.3	56.3	37.7	17.3 → 56.1	15	round limit
3	Qwen3.8-Flash-Next	open	34.5	26.7	54.7	48.7	0.0	17.3 → 17.3	15	round limit
4	MiniMax-M3	open	26.7	26.7	32.4	47.7	0.0	17.3 → 17.3	15	round limit
4	Qwen3.8-27B	open	26.7	26.7	32.4	47.7	0.0	17.3 → 17.3	11	no further progress
6	Gemma-4-31B-it	open	17.5	26.7	28.6	24.1	0.0	17.3 → 19.7	12	validation budget
–	GPT-6 Astra†	closed	54.8	26.7	70.5	58.6	35.4	17.3 → 47.8	12	time budget

Table 31 Visual Search. High-resolution visual search (V*Bench, HR-Bench 8K). Models are ranked by Accuracy on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Accuracy is measured on the whole test split; V*Bench and HR-Bench 8K give the accuracy on the questions from each source benchmark. All scores are multiplied by 100. †Muse Wonder 1.0 was added after the other runs; it is shown for comparison and not ranked, and its starter solution has no test score.

Model			Hidden test metrics				Validation	Run	
#	Name	Weights	Accuracy	Starter	V*Bench	HR-Bench 8K	Start → best	Rounds	Ended by
1	Gemini 4 Argon	closed	**98.6**	87.9	**100.0**	**97.1**	87.5 → 98.8	10	round limit
2	Claude Opus 5.5	closed	96.4	85.0	98.6	94.3	85.0 → 95.0	10	round limit
2	GPT-6.1 Sol	closed	96.4	89.3	95.7	**97.1**	87.5 → 93.8	10	round limit
4	Gemini 3.8 Flash	closed	95.0	90.0	97.1	92.9	88.8 → 95.0	10	round limit
5	GPT-6 Astra	closed	93.6	84.3	94.3	92.9	86.2 → 91.2	10	round limit
6	Qwen3.8-Flash-Next	open	88.6	78.6	97.1	80.0	78.8 → 87.5	10	round limit
7	Qwen3.8-27B	open	85.7	78.6	90.0	81.4	67.5 → 80.0	10	round limit
7	GLM-5.3 Flash	open	85.7	76.4	90.0	81.4	78.8 → 83.8	10	round limit
9	Claude Sonnet 5.5	closed	80.7	74.3	85.7	75.7	81.2 → 86.2	10	round limit
10	DeepSeek-V4.1 Flash	open	77.1	75.7	80.0	74.3	70.0 → 81.2	10	round limit
11	Gemma-4-31B-it	open	72.9	57.1	81.4	64.3	57.5 → 77.5	10	round limit
12	MiniMax-M3	open	70.7	62.1	77.1	64.3	61.3 → 76.2	10	round limit
–	Muse Wonder 1.0†	closed	56.4	–	55.7	57.1	53.8 → 85.0	10	round limit

Table 32 Video Quality. Perceptual video quality: score real user videos (YouTube-UGC MOS) and pick the better of AI-generated pairs (Veo 3.1 human votes). Models are ranked by Score on the hidden test split. Test metrics are those of the final solution, except Starter, which is the primary test score of the starter solution. Validation gives the validation score of the starter and the best validation score reached during the run. Score is the mean of SRCC and pairwise accuracy. SRCC and PLCC are the rank and linear correlations between predicted and human quality ratings of real videos; Pair acc. is the accuracy in choosing the better of two generated videos. All scores are multiplied by 100.

Model			Hidden test metrics					Validation	Run	
#	Name	Weights	Score	Starter	SRCC	PLCC	Pair acc.	Start → best	Rounds	Ended by
1	GPT-6 Astra	closed	**82.1**	82.6	86.7	90.4	**77.5**	84.4 → 84.4	10	round limit
2	Claude Opus 5.5	closed	82.0	80.9	86.5	89.7	**77.5**	79.9 → 85.1	10	round limit
3	DeepSeek-V4.1 Flash	open	82.0	79.1	90.5	89.7	73.5	67.3 → 82.1	7	validation budget
4	GPT-6.1 Sol	closed	80.8	81.8	92.3	**94.2**	69.4	80.4 → 89.8	10	round limit
5	Gemini 4 Argon	closed	80.6	77.0	85.6	89.7	75.5	79.8 → 86.4	10	round limit
6	Qwen3.8-27B	open	78.0	74.5	86.6	87.9	69.4	65.4 → 78.4	10	round limit
7	Gemini 3.8 Flash	closed	77.5	77.5	87.7	88.0	67.3	76.9 → 90.0	9	validation budget
8	Qwen3.8-Flash-Next	open	76.6	78.5	**94.0**	93.5	59.2	78.2 → 88.6	10	round limit
9	GLM-5.3 Flash	open	76.2	71.5	89.2	89.9	63.3	74.5 → 88.8	10	round limit
10	Gemma-4-31B-it	open	74.3	77.5	81.3	82.5	67.3	79.0 → 83.1	5	validation budget
11	MiniMax-M3	open	69.8	72.5	74.2	10.0	65.3	60.9 → 73.4	10	round limit

D Additional qualitative examples

This section extends Section 4.6 with the remaining examples of the project page, in the order of its leaderboard.

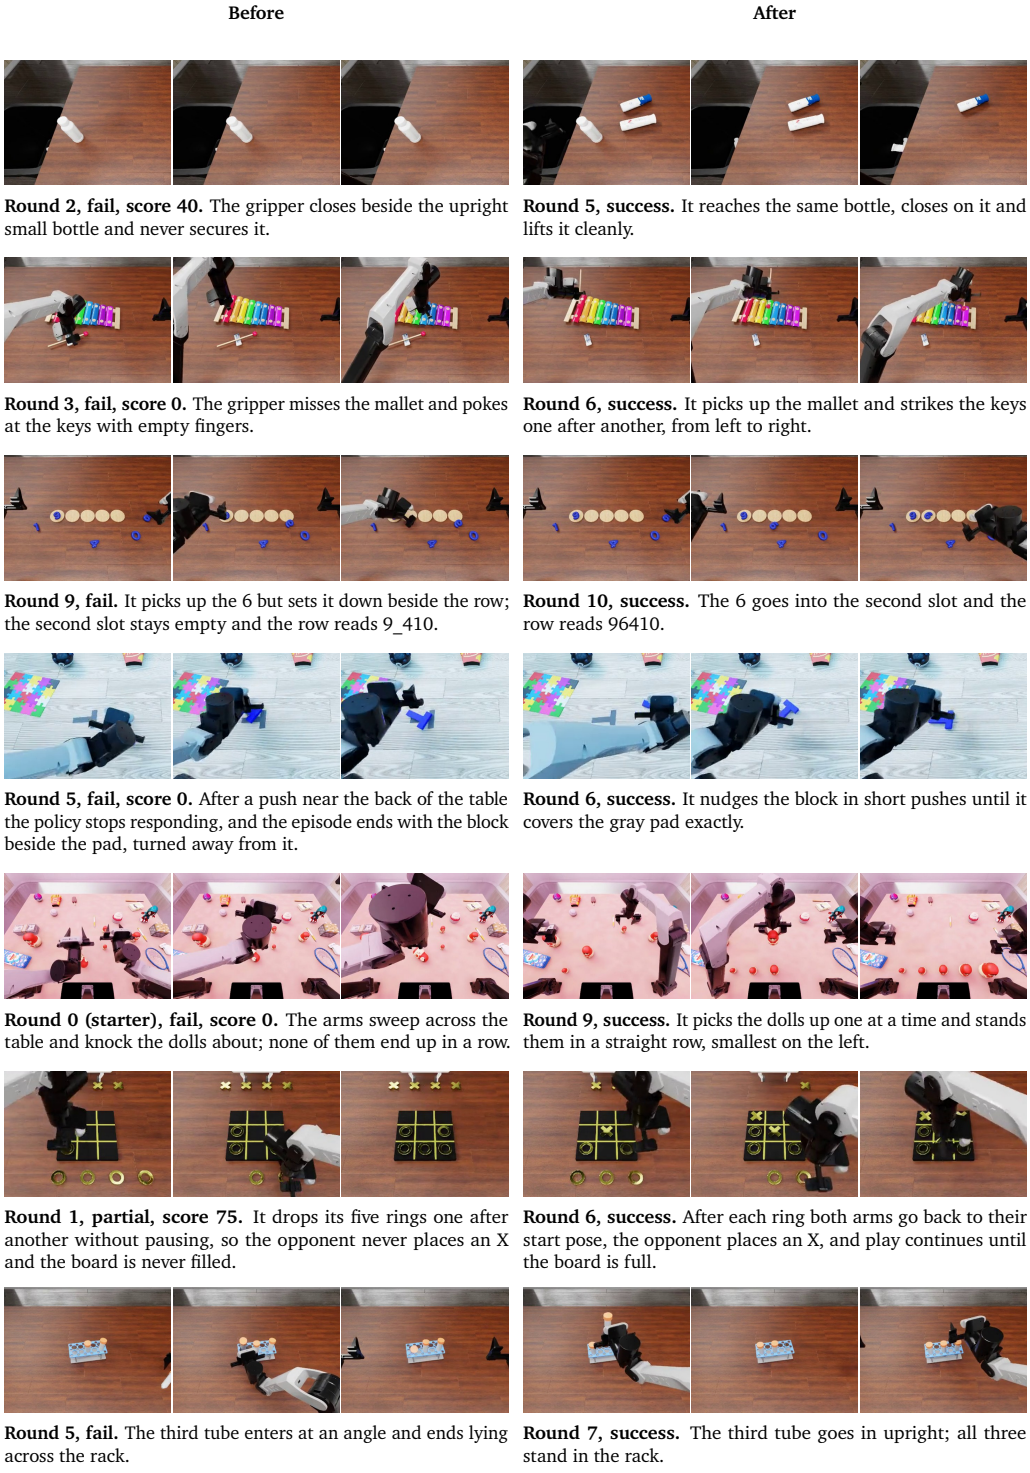

Figure 10 RoboDojo, further examples. Each row shows one episode of a robot task played by the solution of an earlier round (Before) and by that of a later round (After) of the same run, three frames each, cropped to the part of the scene where the two differ. The rows give the benchmark, the model of the run, the round of each solution and the outcome of the episode.

Figure 11 RoboCasa, further examples. Each row shows one episode of a robot task played by the solution of an earlier round (Before) and by that of a later round (After) of the same run, three frames each, cropped to the part of the scene where the two differ. The rows give the benchmark, the model of the run, the round of each solution and the outcome of the episode.

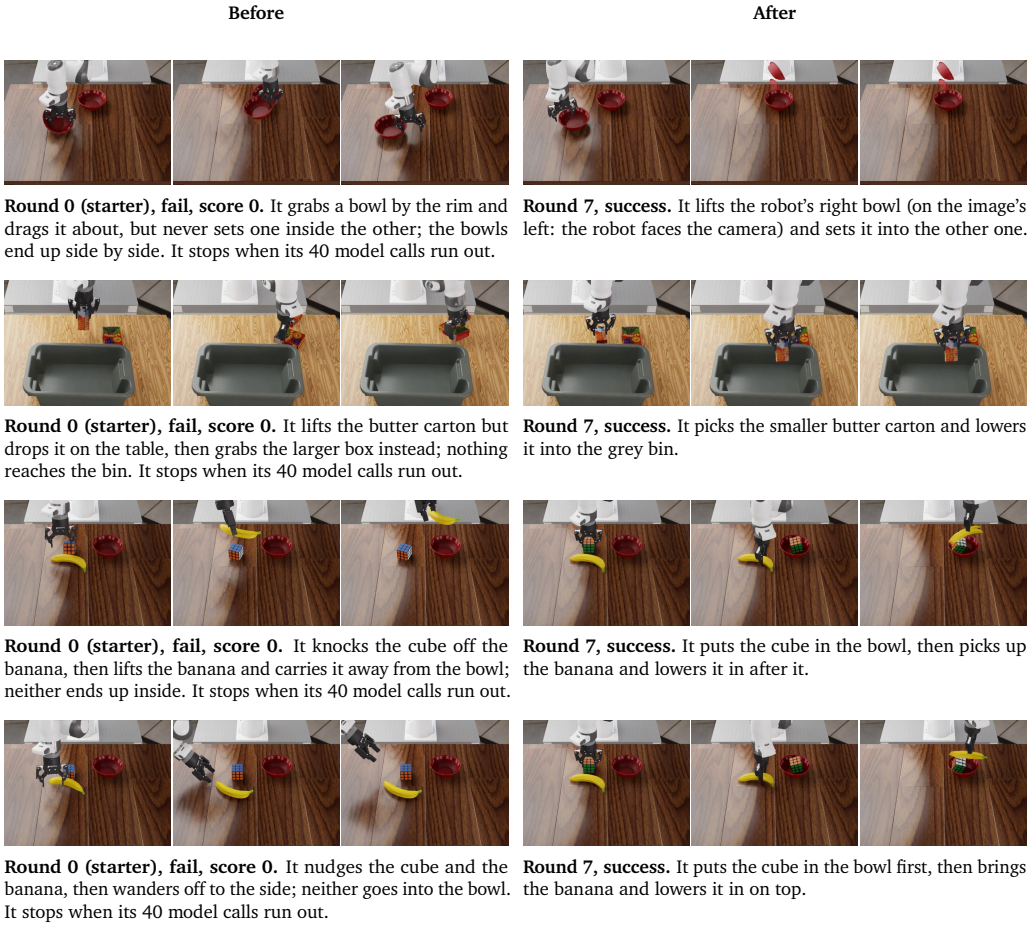

Figure 12 RoboLab, further examples. Each row shows one episode of a robot task played by the solution of an earlier round (Before) and by that of a later round (After) of the same run, three frames each, cropped to the part of the scene where the two differ. The rows give the benchmark, the model of the run, the round of each solution and the outcome of the episode. The RoboLab episodes are development tasks.

Path search. You want to read a book, where to move?

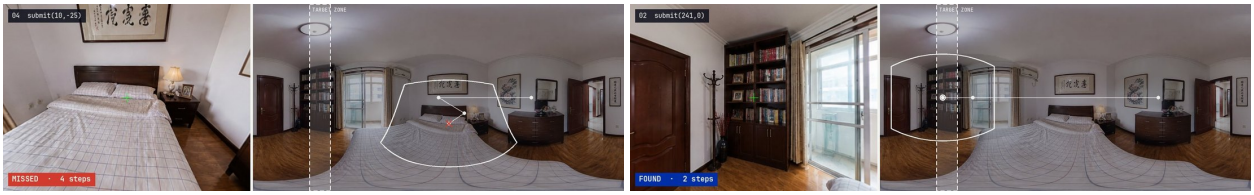

Object search. Find mannequin.

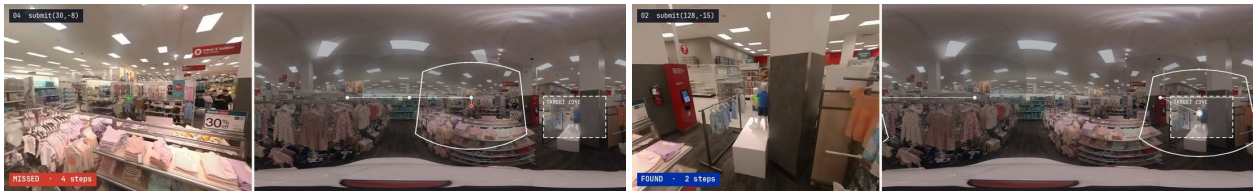

Object search. Look for the Dog Diapers.

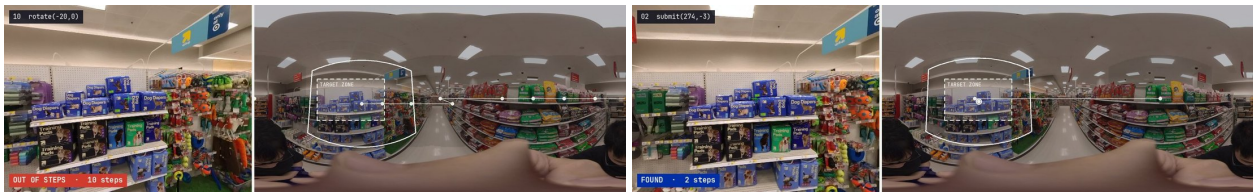

Figure 13 Active Search, further examples. Episodes of the starter solution (left of each row) and of the final solution of the run of Claude Opus 5.5 (right). Each image shows the last view of the episode (left part) and the panorama with the path of the view and the target zone (right part); the label gives the outcome and the number of actions, and the top line the last action. Object search asks for a named object, path search for the direction in which to move.

Seen by the starter: eight evenly spaced times, ego view above and exo view below **Chosen by the final solution**

Cooking Scrambled Eggs, object state. While the man in the dark short-sleeve shirt was wiping his hands at the kitchen sink, what was the state of the skillet on the stovetop island? (A) Full of clear water; (B) Covered with chopped vegetables; (C) Filled with mixed egg mixture; (D) Empty and dry Answer C; starter D, final C.

Remove a Wheel, egocentric direction. While the wearer is holding the detached black bicycle wheel over the patterned rug, where is the wheel relative to the wearer? (A) Directly in front of the wearer; (B) To the front-left of the wearer; (C) To the front-right of the wearer; (D) To the left of the wearer Answer B; starter A, final B.

Remove a Wheel, habitual location. Where is the silver combination wrench usually kept in this workshop? (A) Hanging in the lower row of the wooden tool rack filled with other metal wrenches and blue-handled tools; (B) Hanging in the rightmost section of the wooden tool rack filled with other metal wrenches and blue-handled tools; (C) Hanging in the upper row of the wooden tool rack filled with other metal wrenches and blue-handled tools; (D) Hanging in the middle row of the wooden tool rack filled with other metal wrenches and blue-handled tools Answer D; starter A, final D.

Making Sesame-Ginger Asian Salad, object state. While the man in the black T-shirt and dark glasses was getting the white bowl-like plate from the metal counter, what was the state of the pale dressing in the squeeze bottle? (A) It was still completely full and unopened; (B) It had already been poured out and was no longer full; (C) It had been spilled over the counter; (D) It was sitting empty with the cap off Answer B; starter A, final B.

Making Cucumber & Tomato Salad, instantaneous position. Where was the woman with a ponytail wearing a white floral T-shirt and dark glasses standing the FIRST time getting the cucumber? (A) At the black countertop beside the single-burner silver stove, facing the cabinets; (B) At the black countertop beside the two-burner silver stove, facing the cabinets; (C) At the black countertop beside the two-burner silver stove, facing the sink; (D) At the black countertop beside the two-burner silver stove, facing the refrigerator Answer B; starter C, final B.

Cooking an Omelet, allocentric relation. While the dark-skinned person in a black shirt pours from the small seasoning container into the pink mixing bowl, where is the white salt container relative to the seasoning container? (A) To the left of the seasoning container; (B) To the right of the seasoning container; (C) Underneath the seasoning container; (D) Directly behind the seasoning container Answer A; starter B, final A.

Figure 14 Ego-Exo Memory. Test questions, the frames that the starter solution shows the model (evenly spaced over the video), and the frames that the final solution of the run of Claude Opus 5.5 selects before answering, with the answers of both.

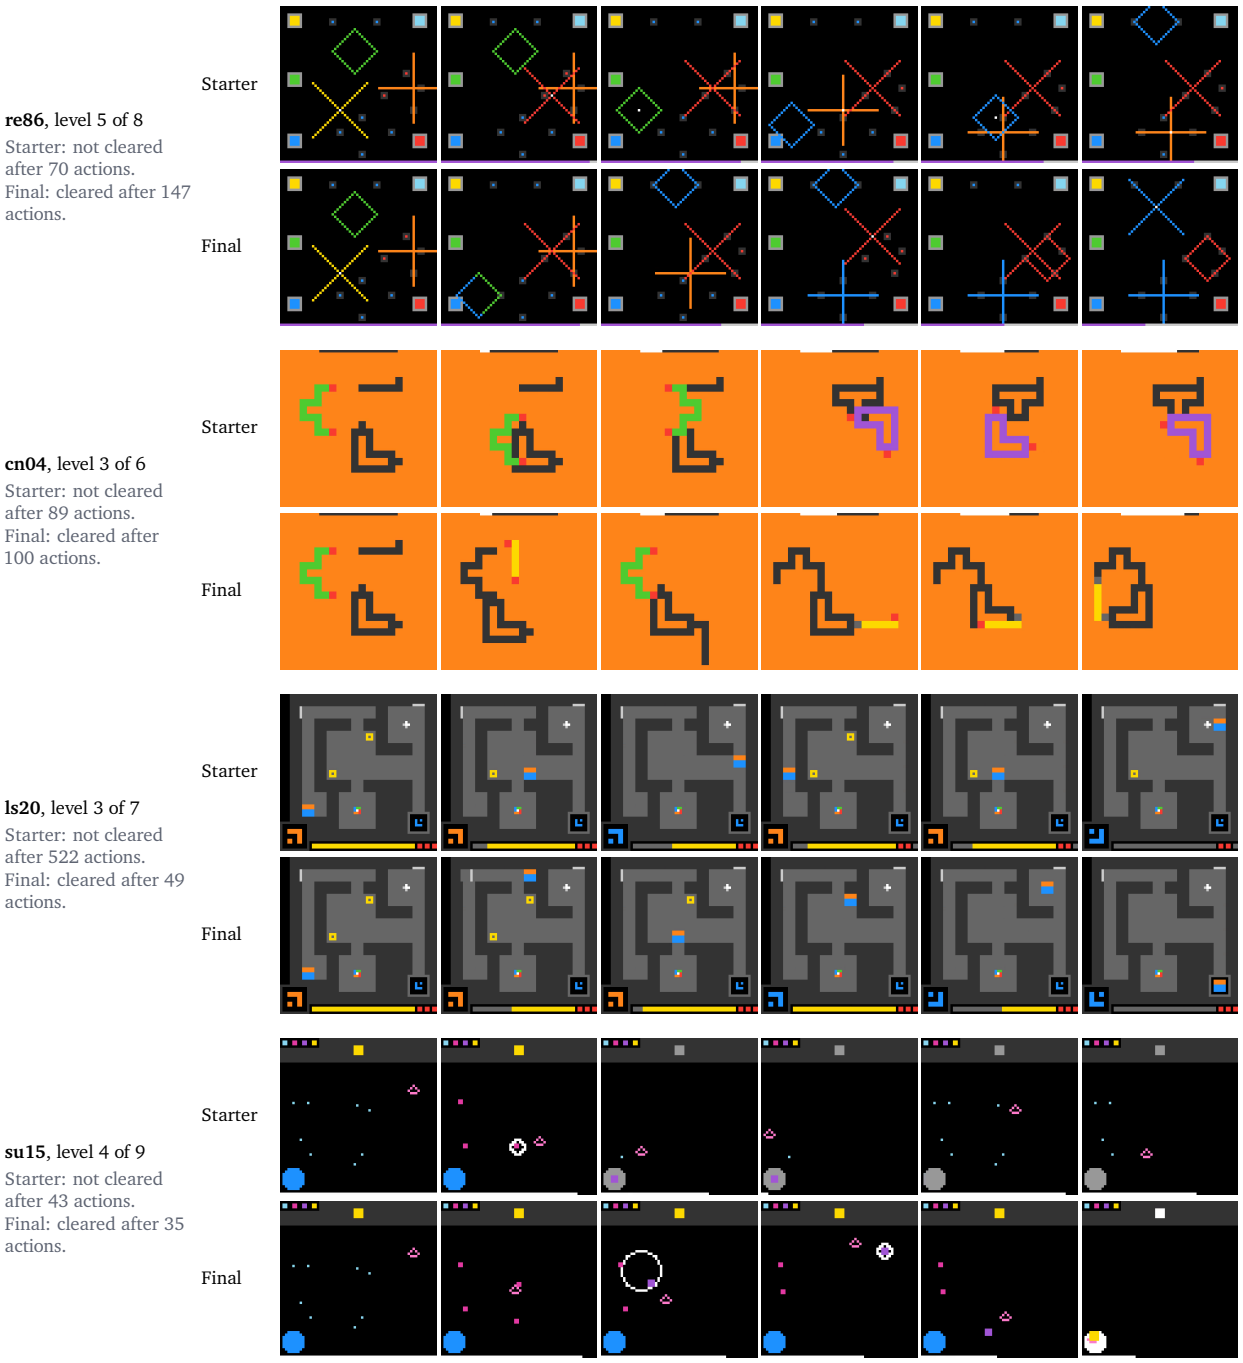

Figure 15 **ARC-AGI-3.** Four validation games, each played on one level by the starter solution and by the final solution of the run of Claude Opus 5.5. The games come without instructions: the agent finds out by acting what the controls do and what clears a level. Each row shows six frames evenly spaced from the first to the last frame of the play. The final solution tells the model which pixels each action changed, removes unavailable actions, and restarts after a game over.

Change appearance. *Change the red umbrella canopy to a yellow-and-black plaid pattern*

Input Starter: 72 Final: 89

Starter: A yellow-and-black pattern, but the tile texture around it is smoothed away. Final: Yellow-and-black plaid; the tiles and the light are identical.

Restoration, deblur. *This photo is out of focus; deblur it so the oat stalks and the clouds look sharp*

Input Starter: 0 Final: 75

Starter: Still blurry. Final: Oat stalks and clouds sharp; nothing is redrawn.

Soft-effect removal, glare. *Remove the star-burst flares and glare around the street lamps so they look like normal light sources; keep everything else unchanged*

Input Starter: 53 Final: 77

Starter: The flares soften, but the buildings and the sky change. Final: Flares gone; the lamps, the path and the buildings are unchanged.

Soft-effect removal, shadow. *Remove the large shadow cast on the brick wall; keep the brick pattern and the plants unchanged*

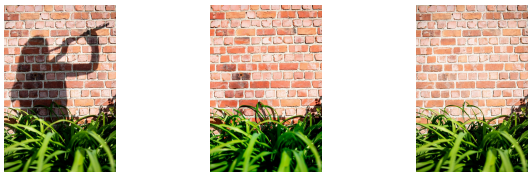

Input Starter: 66 Final: 77

Starter: The shadow goes, but the wall is redrawn and the bricks shift. Final: Shadow gone; bricks and plants line up with the original.

Remove several objects. *Tidy the desk: remove the mug, the second laptop, the green box and the papers, keeping only the laptop in the middle*

Input Starter: 48 Final: 84

Starter: The mug is gone, but the second laptop and the papers stay. Final: Only the middle laptop is left, though the desk around it changes a little.

Find the object, then edit. *Make the leftmost bottle cobalt blue, keeping the painterly brushwork*

Input Starter: 38 Final: 88

Starter: The bottle turns blue, but loses its brushwork and the background changes. Final: The leftmost bottle is cobalt blue, still painted; the rest is untouched.

Figure 16 Image Editing, further examples. For each edit instruction, the input image and the outputs of the starter solution and of the final solution of the run of Claude Opus 5.5, with the score that the task evaluator assigns to each output (multiplied by 100). Below each example, the differences between the two outputs.

Figure 17 3D Tracking. One view of each clip with the annotated instances and the instances tracked across views by the starter solution and by the final solution of the run of Claude Opus 5.5, with the tracking mIoU of each (multiplied by 100).

Figure 18 Scene from Photo, a further example. Development photos and the furniture reconstructed by the starter solution, by the solution after round 1, and by the final solution of the run of Claude Opus 5.5, next to the annotated furniture, seen from above the photo's camera; each color is one piece, and the room box is not drawn. The starter returns an empty room box; the solution of round 1 places a box template for each piece; the final solution generates the shape of each piece.

V* Bench. What is the color of the plastic stool? Answer C; starter D, final C.

V* Bench. What is the color of the watch? Answer C; starter D, final C.

V* Bench. Is the red suitcase on the left or right side of the red stool? Answer B; starter A, final B.

V* Bench. What is the color of the handbag? Answer D; starter B, final D.

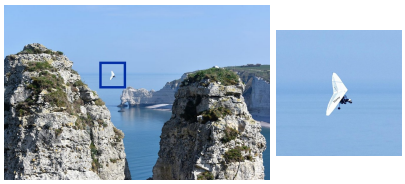

HR-Bench 8K. What word is written on the hang glider? Answer D; starter A, final D.

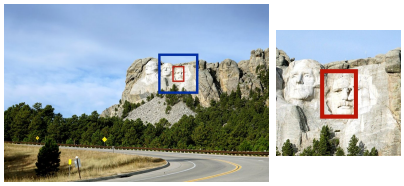

HR-Bench 8K. Which president's face is highlighted in the red rectangle on Mount Rushmore? Answer D; starter C, final D.

Figure 19 Visual Search. Test questions about small objects in high-resolution images: the image, the region that holds the answer (enlarged), and the answers of the starter solution and of the final solution of the run of Claude Opus 5.5.

Rating, music video. Human mean opinion score 3.54; starter 3.00, final 3.50 (scale 1–5).

Rating, TV clip. Human mean opinion score 1.99; starter 1.30, final 1.71 (scale 1–5).

Rating, animation. Human mean opinion score 4.52; starter 4.00, final 4.50 (scale 1–5).

Pair. Which video matches the prompt better? Prompt: A sleek, futuristic hovercar races through a neon-lit metropolis at midnight. The vehicle glides smoothly over elevated tracks, illuminating the skyline with vibrant colors as it weaves between towering skyscrapers, creating a sense of speed and advanced technology. Human choice B; starter A, final B.

Pair. Which video has fewer visual artifacts? Prompt: A daring escape through a neon-lit cyberpunk city at night. The protagonist weaves between towering skyscrapers as drones hunt, then leaps onto a hovering motorcycle and races into the vibrant skyline. Human choice B; starter A, final B.

Pair. Which video looks better overall? Prompt: A 3D animated dragon soaring above a medieval castle at sunset, its scales shimmering in the golden light. Flames flicker from its mouth as it navigates through fluffy clouds, casting dynamic shadows below. Human choice B; starter A, final B.

Pair. Which video looks better overall? Prompt: A serene mermaid glides through vibrant coral reefs, sunlight filtering through the water. Fish of every color swim around her gracefully as she dances with the ocean currents. Human choice B; starter A, final B.

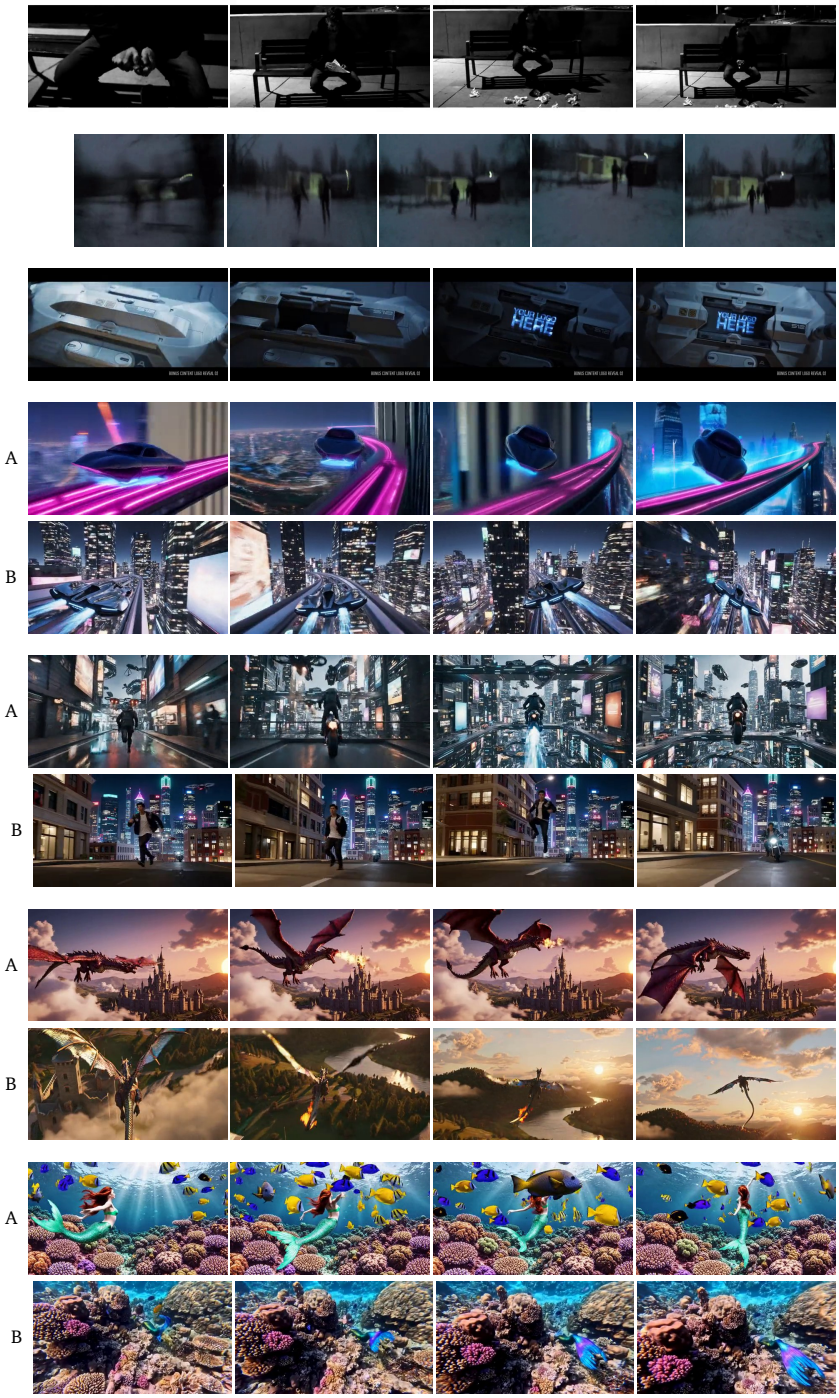

Figure 20 Video Quality. Test videos rated on a scale of 1 to 5 and pairs of generated videos compared, with the human judgement and the predictions of the starter solution and of the final solution of the run of Claude Opus 5.5. Each row shows frames evenly spaced over the video.