arXiv:2611.09454v1 [cs.LG] 6 Nov 2026

From LoRA to Language: Recovering Skill Documents from Adapters

Saranya Dintakurthy Swadha, Pavan C Shekar, Arpit Jain, and Aswanth Krishnan

QpiAI, Bengaluru, India

{saranya.d, pavan.s, arpit.j, ashwanth.krishnan}@qpiai.tech

Abstract

Agent skills are short documents that teach a language model how to handle a class of tasks. To save context, these documents are increasingly compiled into low-rank adapters (LoRA), which keep the behaviour but lose the text: an adapter is hard to read, edit, audit or move to another model. We ask whether the skill document can be written back from the adapter and its frozen base model, and we score each recovered document by the share of the original's key instructions it states and by the tasks it solves. Reading the weights alone recovered little with the methods we tried; a learned map from weights to text passed on about one bit of the document. Running the adapter worked much better. The *Executing Writer*, a small writer trained on top of the adapted model, recovers 39–44% of the key instructions of held-out skills, against 1.3–1.5% with another skill's adapter, and on adapters fitted to 41 external skills its documents restore 85–91% of the original's task gain. These strongest results are on adapters whose training stored the document; adapters trained on behaviour alone are much harder for trained writers (1–2% of instructions). For these adapters we test an agent that runs the base and adapted models and writes a skill from where they differ. On ALFWorld its skills recover 58–71% of a LoRA's gain and carry over to models that cannot load the LoRA, lifting Llama-3.1-8B from 5% to 55%. Whether such a skill comes from the adapter depends on what the adapter learned: when its behaviour cannot be guessed, the adapter is the source; when it can, a reward loop without the adapter does as well.

Keywords: low-rank adaptation; adapter interpretation; agent skills; weight-space reading; instruction recovery.

1 Introduction

An *agent skill* is a document, usually a few hundred to a few thousand words, that tells a language model how to handle a class of requests: how to review code, which conventions to follow, how to analyse a packet capture [1, 2]. Skills are written by practitioners, shared in marketplaces and placed in the model's prompt when needed. That costs context on every call, so a growing line of work compiles a skill into a *low-rank adapter* (LoRA) [3], a small set of extra weights that changes the model's behaviour with no text in the prompt [4, 5, 6, 9, 10]. The adapter is cheap to apply and to serve at scale [11], but it is a block of numbers. A text skill can be read before it is trusted; a weight skill is much harder to check.

The question. We study the inverse of compilation:

> *Given an adapter and its frozen base model, can we write the skill document that the adapter encodes?*

A recovered document would make an adapter inspectable and editable, could be re-compiled for another model, and would let a user verify an adapter before running it. Existing weight-space methods predict attributes, labels or short descriptions of adapters [12, 13, 14, 15, 16]; none of them produces a document or measures how much of one is recovered.

Why a document and not only an adapter. A skill document and a LoRA can carry the same behaviour, but they are not equally useful. A document can be read and reviewed like code, edited in one line, combined with other skills, and given unchanged to a different model or agent harness. A LoRA is tied to the weights and tokenizer of the base model it was trained on: it does not load on Llama if it was trained on Qwen, nor on a model served only through an API, and changing one rule means retraining it. Recovering the document turns a model-locked artefact back into a portable one. We therefore measure recovery in two ways: the share of the original's instructions that the recovered document states, and the share of the original's task gain that it restores.

Why it is hard. Three properties of adapters shape the methods. (i) Adapters for different skills are mostly alike: the cosine between the weight updates of two skills is 0.88, and only 11.6% of their energy is specific to the skill. (ii) Knowing which skill an adapter holds is not knowing what it says. An encoder identifies the skill from its weights for all 37 test skills, yet a learned channel from the weights to a writer passes on about one bit of the document's content. (iii) Text similarity is a weak measure, because word overlap rewards generic text. We therefore score recovered documents against frozen lists of the original's instructions, run every method with a wrong adapter as a control, and measure tasks solved wherever we can.

Our approach. The content of an adapter shows in what the network computes with it, so we let the network do the reading. We compare three readers (Figure 1, Section 4). The first treats the weights as data and learns a map from them to a prompt. The second reads the weights without running them, through the model's vocabulary and by retrieving similar instruction lines. The third, the *Executing Writer*, attaches the adapter to its frozen base model and trains one small writer adapter on top, so that the adapted model itself writes the document. Section 7 adds a fourth reader, which needs no training and suits adapters that never saw their document: an agent that runs the base and adapted models on real tasks and writes a skill from where they differ.

One distinction matters throughout. Adapters differ in whether the document text was ever a training target. A *recitation* adapter was trained to print its document, so the text is stored in it, and recovering it is extraction. A *behavioural* adapter was trained only on answers that the document shaped, so the text never appears in any training target, and recovering it is inference. The second case is the common one in practice. Our trained writers read the first kind well and the second poorly; for the second kind we turn to the executing agent.

Contributions.

1. We define adapter-to-document recovery, with coverage of frozen instruction lists, a measure of invented content, and two controls the evidence needs: a wrong adapter and retrieval of a similar document.
2. We compare three readers (Sections 4 and 6). The weights reveal a skill's identity and content words, but a learned soft-prompt channel passes on about one bit; an information measure, the *adapter gap*, shows where the difference lies.
3. The Executing Writer recovers 39–44% of a held-out skill's key instructions, about twice the best weight-only writer and 1.4–1.6 times retrieval, and only 1.3–1.5% with a wrong adapter. On adapters fitted to

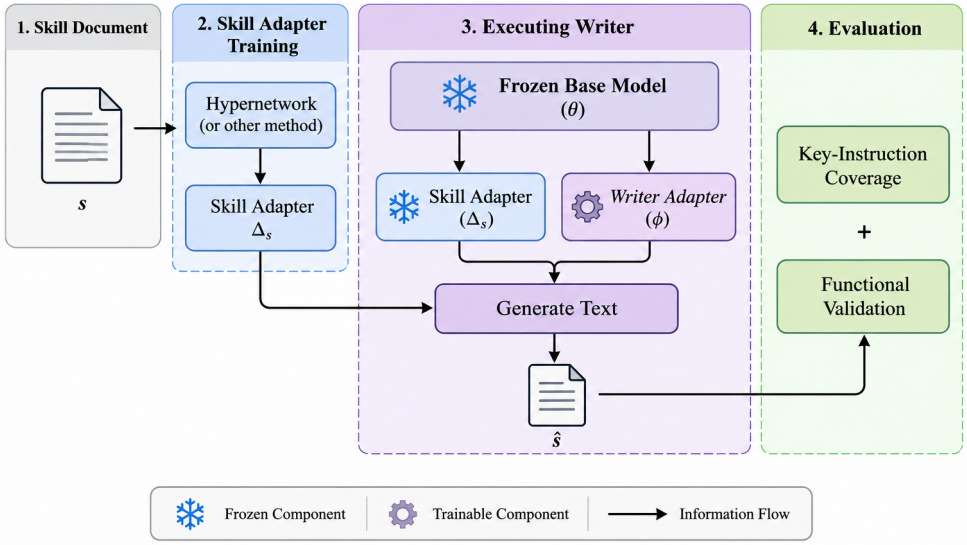

Figure 1: The pipeline. A skill document s is compiled into an adapter Δ_s. The Executing Writer attaches Δ_s (frozen) and a trainable writer ϕ to the frozen base model θ and generates a document $\hat{s}$. Recovery is scored by the share of the original's key instructions that $\hat{s}$ states and by whether $\hat{s}$ still works when given to a model.

41 external skills it recovers 90–93% of the prose at rank 64; these documents restore 85–91% of the original's task gain and help an agent on benchmark tasks almost as much as the original.

4. For behavioural adapters (Section 7) we test an executing agent on ALFWorld, a docker-compose adapter from an earlier project, synthetic tasks and adapters with hidden rules. Its skills recover 58–71% of an ALFWorld LoRA's gain and transfer to other models. Controls show that the adapter is the source of the skill only when the adapter's behaviour cannot be guessed.

2 Related Work

Agent skills. Skills package procedures and conventions as documents that an agent loads when a task calls for them [1]. SkillsBench measures how much skills help on diverse tasks with independent verifiers [2], and earlier agents kept libraries of learned skills [17] or interleaved reasoning with tool use [18, 19]. We use SkillsBench skills and tasks, which we did not write, to test whether recovered documents still work.

Compiling text into weights. LoRA [3] represents a weight update by low-rank factors, following adapters [20] and continuous prompts [21], with evidence that fine-tuning moves a model in a low-dimensional subspace [22] and that LoRA learns less than full fine-tuning [23]. Hypernetworks [24] generate weights from other inputs: HYPTER [25], Text-to-LoRA [4], Doc-to-LoRA [5] and SHINE [6] generate adapters from task descriptions or documents, and LatentSkill [9] and Parametric Skills [10] do this for agent skills. Context distillation [26, 27], gist tokens [28] and cartridges [29] also move context out of the prompt, and adapters can be composed [30] or added arithmetically [31]. All of these go from text to weights; we go the other way.

Prompt and skill optimisation. Searching over prompts with a model in the loop is an established way to improve instructions: APE [50], OPRO [51], DSPy [52], TextGrad [53] and GEPA [54] propose and score candidate prompts, and Reflexion [55], ExpeL [56] and Agent Workflow Memory [57] turn an agent's experience into reusable text. Our reward loop (Section 7) belongs to this family; what is new is the adapter

as an extra source of evidence, and the question of whether the document it yields comes from that adapter. Work on auditing models for hidden objectives [58] asks a related question about behaviour that a model was trained to hide.

Interpreting adapter weights. W2T [12] and Learning on LoRAs [32] process LoRA weights with symmetry-aware networks to predict attributes, identity and performance, and hyper-representations embed [33] or generate [34] weights. LoRAcles [13] trains a model to answer questions about an adapter from its weights, Diff Interpretation Tuning [14] trains a fine-tuned model to describe its own change, WeightWatch [15] labels the top singular vectors of a weight difference, and spectral summaries identify the training objective [16]. These produce labels, scores or short answers. We produce a document and measure how much of the original it contains.

Reading hidden states as text. The logit lens [35] and tuned lens [36] map a hidden direction to word scores with the model's own output layer, and Patchscopes [37] and SelfIE [38] ask the model itself to describe a hidden state. Models can describe behaviours they were trained on [39], and agents that experiment on a model can explain parts of it [8]. We apply the lens to an adapter's output directions after removing what adapters share, and we use the adapted model's own generation, or an agent that runs it, as the reader.

Extraction and inversion. Training data and prompts can be recovered from models and embeddings [40, 41, 42, 43]. Recovering a skill from an adapter inverts the compile step rather than attacking private data, but it raises the same question of what a released artefact reveals (Section 8).

3 Problem and Measures

Adapter families. What an adapter can reveal depends on how it was made, and Table 1 lists the three families we use. A recitation adapter exposes its text by construction, so results on recitation adapters are an upper bound on what a reader can do. They do not reflect test leakage: every writer is scored on held-out skills it never saw.

Table 1: The three adapter families used in this paper.

Family	How it is made	Document in a training target?
Document-bearing	A hypernetwork [9] compiles the document into an adapter. One checkpoint is trained to store its document (*document-fitted*); a further-trained one is tuned for agent behaviour (*task-tuned*).	yes for document-fitted, no for task-tuned
Recitation	We fit a LoRA by gradient descent to print the document after a trigger (ranks 16, 32, 64).	yes, as the output
Behavioural	The model answers requests with the document in its prompt; a LoRA is fitted to give the same answers without it (context distillation).	no

The recovery problem. Let θ be a frozen base model and s a skill document. An adapter $\Delta(s) = \{(A_b, B_b)\}_b$ adds a low-rank update $\Delta W_b = A_b B_b$ to each weight block b, giving the adapted model $q_{\theta \oplus \Delta}(\cdot \mid x)$. *Given Δ and θ but not s, produce a document $\hat{s}$ that states what s states.*

Coverage. For each original s we extract once a frozen checklist of key instructions $C(s) = \{c_1, \ldots, c_m\}$, so every method is scored against the same list. A judge model [47] assigns each c_i one of {full, partial, missing, contradicted} with respect to $\hat{s}$, and

$$\mathrm{Cov}(\hat{s}, s) = \frac{1}{m} \sum_{i=1}^{m} \Big(\mathbf{1}[\mathrm{full}_i] + \tfrac{1}{2} \mathbf{1}[\mathrm{partial}_i] \Big).$$

The primary judge returns a probability for each verdict, so coverage is an expectation. As a check, a judge from another model family scores every document again; Appendix A compares the two judges. For code-heavy skills we score a *prose* checklist (every instruction sentence) and a *structure* checklist (every code block, table and heading) separately, plus a judge-free count of identifiers such as paths, commands and flags.

Invented content. Coverage ignores what a document adds. We split $\hat{s}$ into instruction units by a fixed rule and report the share that the original does not support, $\mathrm{Unsup}(\hat{s}, s)$, as the *unsupported* share.

Controls. Every method is also run with another skill's adapter (*wrong-adapter control*): a method that depends on the adapter must fall to near zero. We also compare with *retrieval* of the training document closest to the adapter, the natural way to use the weights without writing anything.

4 Method: Three Readers

The three readers differ in what they ask a learned model to work out (Figure 2). The first (R1) treats the adapter's numbers as data: it encodes the complete adapter and trains a writer to produce the document from that encoding. The second (R2) looks up which words the numbers point to and retrieves similar instruction lines, and a writer composes the document from this text. The third (R3, the Executing Writer) does not read the numbers at all: it attaches the adapter to its frozen base model and lets the adapted model say what it knows. Only R3 runs the adapter.

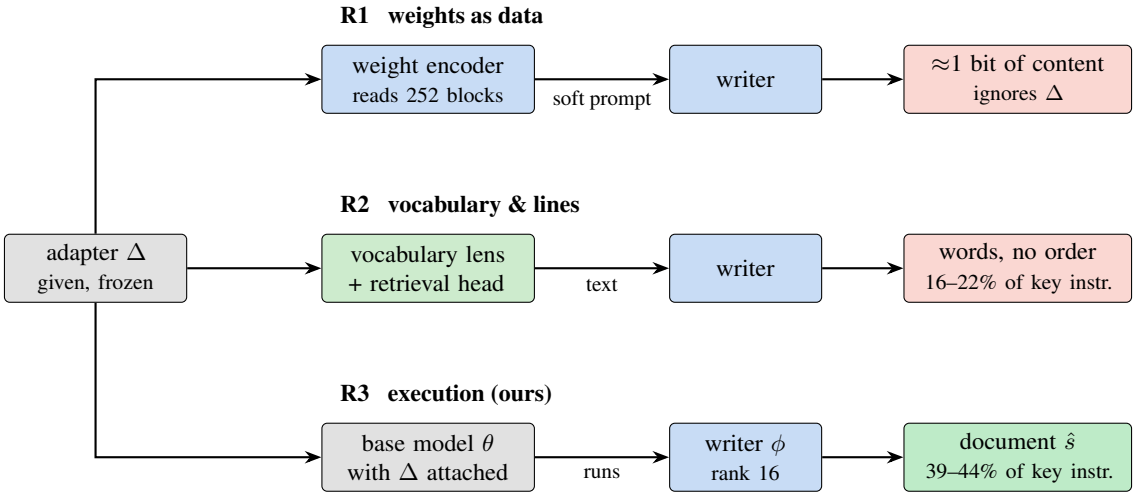

Figure 2: Three ways to read an adapter. R1 and R2 never run the adapter; R3 runs it inside its base model and trains only a small writer. Grey: given and frozen; blue: trained; green: fixed computation. Recovery figures are from Table 3.

R1: the weights as data. We encode the complete adapter (all 252 blocks, no centring or projection) with a transformer into a short learned soft prompt placed in front of a writer (Qwen3-4B or Qwen3-8B with a rank-16 adapter), which also receives the text that R2 reads. Encoder and writer are trained jointly to produce the document. We tried five variants: a pooled vector, all block states with a margin loss on the adapter gap (Section 4.1), a frozen decoder, vocabulary-aligned inputs, and the complete adapter with the 8B writer.

R2: the vocabulary and instruction retrieval. The Gram matrix of n adapters can be computed from the factors without forming ΔW: $G_{ij} = \sum_b \mathrm{tr}\left((A_b^{(i)\top} A_b^{(j)})(B_b^{(i)} B_b^{(j)\top})\right)$. Over 20 test adapters the mean cosine between different skills is 0.878. Centring, $\tilde{G} = HGH$ with $H = I - \frac{1}{n}\mathbf{1}\mathbf{1}^\top$, leaves 11.6% of the energy, and the centred residuals have mean cosine -0.051, the value $-1/(n-1)$ expected for unrelated vectors. Each adapter is a large shared part plus a small, distinct remainder. For each of the 72 blocks that write into the residual stream we take the output directions $v_{b,k}$ of the centred update and read them through the final norm γ and output embedding, $z_{b,k} = W_U(\gamma \odot v_{b,k})$; the top 100 words, weighted by singular value, are the adapter's keywords. Of an adapter's top 40 words, 73% occur in its own document against 11% in other test documents (40% against 8% without centring). Words carry no order, so a *retrieval encoder* (4 transformer layers, 7.5M parameters) reads compact linear summaries of the 72 blocks and is trained with a contrastive loss [44, 45] against document embeddings [46] on $10{,}500$ training skills. It names the correct document for all 37 test skills among 237 candidates. An instruction head built on the same encoder retrieves the 30 closest lines from a bank of $208{,}995$ lines taken from training documents only; these lines match 24.6% of the original's instructions, against 0.2% for a wrong adapter.

R3: the Executing Writer. Given Δ and θ, the writer is a rank-16 adapter ϕ (43.6M parameters) on the same base model, trained on document–adapter pairs with the base model and $\Delta(s)$ frozen and the loss on the document tokens only (Figure 3):

$$\min_{\phi} \; \mathbb{E}_s\Big[-\log q_{\theta \oplus \Delta(s) \oplus \phi}\big(s \mid x(s)\big)\Big].$$

The prompt $x(s)$ is the request *"Write the SKILL.md document"*, either alone (EW-A, adapter only) or preceded by the keywords and retrieved lines of R2 with the lines of s itself removed (EW-T, adapter and text). EW-T is trained on document-fitted adapters and EW-A on task-tuned ones. Documents are decoded greedily with a repetition penalty of 1.1.

4.1 Why running the adapter helps: the adapter gap

An adapted layer updates the hidden state as $h^{(l+1)} = h^{(l)} + F_l(h^{(l)}; W_l + \Delta W_l)$, so Δ affects any output only through the products $\Delta W_b\, h$. A reader of the numbers must learn a map from ΔW to text that implicitly simulates what the network does with ΔW, from a few thousand examples. Running the adapter hands that simulation to the frozen network, which already contains it, and asks the writer only for a query that makes the adapted model say what it holds. We measure how much each reader takes from the adapter with the *adapter gap*:

$$\mathcal{G}(\mathcal{R}) = \frac{1}{|s|}\,\mathbb{E}_{s,s'}\Big[\log q^{\mathcal{R}}\big(s \mid x; \Delta(s)\big) - \log q^{\mathcal{R}}\big(s \mid x; \Delta(s')\big)\Big],$$

in nats per token, with s' another skill. It is 0 for a reader that ignores the adapter, and $|s|\,\mathcal{G}/\ln 2$ is the number of bits the adapter adds to the document's probability. It is a diagnostic, not a bound on mutual information. For scale, naming one of our $11{,}000$ training skills takes $\log_2 11{,}000 \approx 13.4$ bits.

(a) Test time: write the document from an adapter

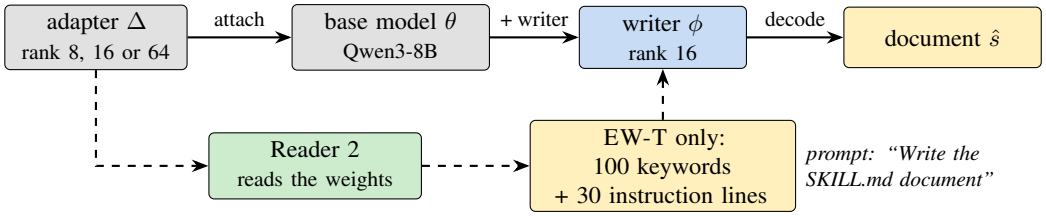

(b) Training: only ϕ is updated

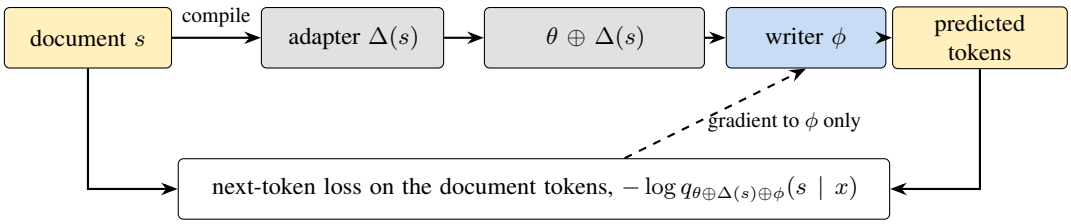

Figure 3: The Executing Writer. (a) At test time the adapter is attached to its frozen base model and a small writer ϕ makes the adapted model write the document; EW-T also receives the text read by R2 (dashed). (b) Training updates only ϕ.

Table 2: Experiments, models and metrics. Numbers from different rows are comparable only where the setting matches.

Experiment	Base model	Adapters	Skills / tasks	Metric
Three readers (§6.1)	Qwen3-8B	document-bearing, rank 8	37 skills	coverage
Transfer and rank (§6.3)	Qwen3-8B	recitation, rank 16–64	37 + 41 skills	coverage
Behavioural adapters (§6.4)	Qwen3-8B	behavioural, rank 16, 64	41 skills, 323 tasks	coverage, task score
Agent with documents (§6.5)	Claude Haiku 5.5 agent	recitation, rank 64	4 benchmark tasks	verifier tests
Executing agent (§7)	Qwen2.5-7B, Qwen3.5-4B	behavioural, rank 16–32	ALFWorld, 6 other tasks	task success

5 Experimental Setup

Models. All base and actor models are open-weight models: Qwen3-8B and Qwen3-4B [49], Qwen2.5-7B-Instruct [59] and Qwen3.5-4B (Alibaba), and Llama-3.1-8B-Instruct [60] (Meta). Agents and judges are proprietary models used through their providers: Claude Haiku 5.5 and Claude Sonnet 5.5 (a small and a mid-sized model from Anthropic), run inside Claude Code, Anthropic's command-line agent harness. The primary coverage judge is JEV-1.13, a judge model that returns a probability for each verdict; Gemini 2.5 Pro and Gemini 3 Flash (Google) serve as second judges. Open models are served with vLLM [48]. We refer to models by these names throughout; later versions may behave differently, so we give all prompts and tool definitions in the released code.

Adapters and data. Document-bearing adapters come from a public skill-to-adapter hypernetwork [9]: rank 8 on all seven projections of all 36 layers of Qwen3-8B (252 blocks, 21.8M numbers), generated on demand. From its corpus of 171,258 skill documents we draw 10,500 training documents with strict near-duplicate filtering against the evaluation set and add 500 documents from an earlier training set (11,000 in all), plus 300 validation and 200 further held-out documents. The test set has 37 skills (9 from the corpus, 5 hand-written and 23 from a split the hypernetwork never saw); seven have near-copies in the training corpus, so we also report the other 30.

Baselines. *Retrieval*: the training document closest to the adapter. *Keyword writer* and *instruction writer*: Qwen3-4B with a rank-16 adapter trained to write the document from the keywords, or from keywords and 30 retrieved lines. *Soft-prompt writers*: R1 with a Qwen3-4B or Qwen3-8B writer.

Training and evaluation. EW-T and EW-A are trained for one pass over the training set (1,312 steps at effective batch 8, learning rate 2×10^{-4}, rank 16, $\alpha = 32$). Differences below about four coverage points are within noise on 37 skills.

6 Results on Document-Bearing and Recitation Adapters

6.1 Which reader works?

Table 3: Recovery on the 37 test skills (percent). Coverage by two judges, coverage with a wrong adapter, coverage on the 30 skills without near-copies, and the unsupported share (lower is better). †: runs the adapter. EW-T reads document-fitted and EW-A task-tuned adapters; the other rows use document-fitted adapters.

Method	Judge 1	Judge 2	Wrong adapter	30 skills	Unsupported
Retrieval of a similar document	26.9	35.2	—	15.3	66
Keyword writer (weights only)	16.1	23.5	3.3	14.6	—
Instruction writer (weights only)	22.3	33.3	1.6	20.4	54
Soft-prompt writer, 4B (weights only)	22.7	33.2	22.5	18.2	55
Soft-prompt writer, 8B (weights only)	20.9	31.3	23.1	18.2	55
EW-T† (adapter and text)	**43.7**	**61.5**	1.3	**39.9**	**39**
EW-A† (adapter only)	38.7	59.7	1.5	37.3	**39**

Table 3 shows that running the adapter nearly doubles recovery. The Executing Writer recovers 38.7% and 43.7% of the key instructions, against 22.3% for the best weight-only writer and 26.9% for retrieval. This gain comes from the adapter. With another skill's adapter attached, the executing writers fall to 1.3–1.5%, although EW-T still receives the correct text, so it follows the adapter over the text, as expected for adapters that store their document. EW-A receives no text at all, so everything it writes comes from the adapter. The soft-prompt writers, in contrast, score the same with a wrong adapter, so they do not use it. The executing writers also invent less: 39% of their units are unsupported by the original, against 54–66% for the other methods. On the 30 skills without near-copies, retrieval falls to 15.3%, while the executing writers keep 37–40%. Appendix A compares the two judges: the second is more generous to every method, by 7 to 21 points, but under both the executing writers lead.

6.2 What the weights contain, and why a learned channel stays weak

The weights do contain the skill's identity and its content words. The retrieval encoder names the correct document for all 37 test skills and reaches 100% top-1 on 300 validation skills, and a single threshold on the

keyword score separates an adapter's own document from the most similar other document with AUC 0.905 on 537 unseen adapters. This gives a training-free check that an adapter matches a claimed skill, though it checks the skill and not its individual instructions.

Table 4: Adapter gap on 100 validation skills (nats per token): how much the correct adapter lowers the document's loss against a wrong one.

Reader	Loss, correct	Loss, wrong	Gap $\mathcal{G}$	Bits per 1k tokens
R1 soft prompt, complete adapter	1.220	1.221	0.001	≈ 1
R2 keywords	1.292	1.375	0.083	≈ 120
R2 keywords and instruction lines	1.271	1.398	0.127	≈ 180
R3 Executing Writer, EW-T	0.289	3.934	3.645	$\approx 5{,}260$
R3 Executing Writer, EW-A	0.323	3.936	3.613	$\approx 5{,}210$

Table 4 shows where the difficulty lies: in turning the numbers into text. Across five soft-prompt variants the gap never exceeds 0.006 nats per token, about one bit for a 1,000-token document, less than the 13.4 bits needed just to name the skill. The channel can carry information: with eight training skills the same writer learns to use the adapter completely (loss 0.000 with the correct adapter, 0.731 with a wrong one), and on training skills at scale the gap is still only 0.002–0.003, so the weakness is not overfitting. The keyword and line readers pass on a few hundred bits and the executing writers several thousand. Measured against no adapter rather than a wrong one, the executing writers' gap is 1.45 and 1.33 nats per token, about 2,000 bits per 1,000 tokens; a wrong adapter lowers the likelihood further because it points the writer at another document. Against either baseline, running the adapter passes on over ten times more than any weight-only reader. Capacity is not the cause either: a rank-16 adapter on 252 blocks has 43.6M parameters, four orders of magnitude more than the roughly 2.5 bits per token (the writer's loss with no adapter) that a document needs beyond the base model. Better weight encoders may close part of this gap; with the encoders we tried, running the adapter was far more effective.

6.3 Transfer to other adapters, and the effect of rank

To test transfer, we fitted new Qwen3-8B adapters to recite each of the 37 test skills at ranks 16, 32 and 64 and applied both writers with no retraining. We did the same for 41 skills from SkillsBench [2], a public benchmark whose skills and tasks we did not write and which no reader saw during training. All fitted adapters reproduce their document exactly, so rank does not limit storage.

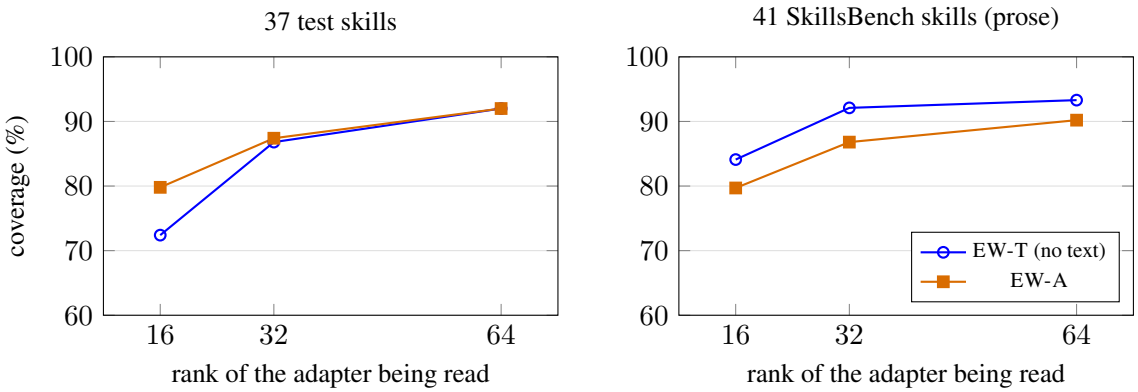

Figure 4: Recitation adapters are read back more accurately at higher rank. Both writers have rank 16; only the adapter being read changes. With a wrong adapter coverage is about 2% (left) and 9–13% (right). Full tables in Appendix B.

Figure 4 shows that both writers transfer. On the 37 skills they recover 72–80% of the instructions at rank 16 and 92% at rank 64; on the 41 external skills, 80–84% of the prose at rank 16 and 90–93% at rank 64, with the same numbers for code and structure and for a judge-free count of identifiers. Most of the gain is in place by rank 32, close to the effective rank of these updates (about 40 directions by participation ratio). These adapters were fitted to recite their document, so the experiment shows transfer, not that behaviour alone can be read.

6.4 Behavioural adapters

To test the case that matters in practice, we fitted adapters on the 41 external skills *without* the document in any training text. For each skill the base model writes about 240 requests the skill would help with and answers them with the skill in its prompt; a LoRA is trained to give the same answers from the request alone. The fitted adapters reach an answer loss close to the prompted model's. We applied both writers to them unchanged.

Table 5: Document fidelity on the 41 external skills (percent). "Same" is a second judge comparing whole documents.

		Coverage of original			Precision		Whole doc.
Method	Rank	Prose	Struct.	Ident.	Prose	Struct.	Same
Recitation adapters							
EW-T	16	84.1	85.1	85.1	85.8	65.5	83
EW-T	64	**93.3**	**94.0**	**94.1**	**94.9**	**72.0**	**90**
EW-A	16	79.7	80.1	78.0	84.5	63.2	73
EW-A	64	90.2	90.2	91.2	94.2	70.0	86
Behavioural adapters							
EW-T	16	0.9	2.4	–	19.3	15.3	0
EW-T	64	1.1	2.7	–	19.0	14.1	1
EW-A	16	1.3	4.3	–	20.9	11.0	1
EW-A	64	1.9	4.8	–	20.1	13.5	2
Wrong adapter (both writers)	64	8.3	12.2	–	13.7	11.0	5

Table 5 shows that the trained writers do not read behavioural adapters. From them both writers recover 1–2% of the prose and 2–5% of the structure, below the wrong-adapter control, and the second judge finds 0–2% of the instructions present. These writers were trained only on adapters that recite their document. Appendix D reports a writer trained on behavioural adapters instead: it depends on the adapter (a gap of 0.7 nats per token) but still recovers only about 7% with 45 training adapters, which it memorises.

To see what the documents and adapters *do*, we wrote 323 tasks (8 per skill) that each require 3 to 6 specific behaviours taken from the original, and scored the base model's answers under each condition with the primary judge. Table 6 reports the share of the original's gain over no skill that each condition restores.

Documents recovered from recitation adapters restore 74% to 91% of the original's gain on all tasks, and 85–91% at rank 64; the standard error of each paired difference to the original is 0.7–1.3 points. The hardest 91 tasks, those with a no-skill score of at most 0.4, show the same pattern. Behavioural adapters behave differently. Attached with no document, they restore 47–58% of the gain, so they carry the skill; but the documents the writers produce from them restore none of it (-4 to $+1$%). The rank-16 adapters restore somewhat more than rank 64 (58% against 47%); we did not study why.

Table 6: What documents and adapters do on 323 skill-dependent tasks (base-model score, percent). Gain recovered is relative to the original.

Condition	Rank	All tasks	Hardest 91 tasks	Gain recovered (%)
No skill	–	55.5	23.2	0
Original skill in prompt	–	77.0	65.9	100
Recovered document in prompt, from recitation adapters				
EW-T	16	73.6	59.2	84
EW-T	64	75.1	60.8	**91**
EW-A	16	71.5	56.0	74
EW-A	64	73.8	59.1	85
Recovered document in prompt, from behavioural adapters				
EW-T	64	55.7	30.1	1
EW-A	64	54.6	28.7	-4
Behavioural adapter attached, no document				
Adapter itself	16	68.0	50.8	58
Adapter itself	64	65.6	50.0	47

6.5 An agent using recovered documents

Tasks we wrote ourselves might favour our documents, so we also tested them on SkillsBench tasks, whose verifiers the benchmark authors wrote. The agent is Claude Haiku 5.5 in the benchmark's own Docker sandbox. Of 17 pilot tasks we kept the four that Haiku fails without a skill and that the reference solution passes, and ran each 8 times per arm with the full skill text in the prompt (Table 7). On the 3D-model parser task, every skill takes the agent from $0/8$ to $8/8$. On intrusion detection, the original and the EW-T reconstruction both reach $5/8$, while EW-A gets $0/8$. Over all runs, the share of verifier tests passed rises from 59.7% with no skill to 85.0% with the original and to 82.5–83.1% with the reconstructions; the differences from the original, about -2 points, have intervals of about -12 to $+7$. Because the agent is a different model from the adapter's base, these runs also show a recovered document working across models. Four tasks, two of which respond to a skill, make a small test. Appendix C adds 41 tasks we wrote: where no skill was needed the reconstructions never did worse than the original, and on one of the three skill-dependent tasks EW-T fully restored the original's effect.

Table 7: Claude Haiku 5.5 as an agent on four SkillsBench tasks: runs passing all verifier tests (of 8) and the mean share of tests passed.

Skill in prompt	Intrusion det.	3D parser	Finance QA	PDDL plan	All runs	Tests (%)
None	0/8	0/8	0/8	3/8	3/32	59.7
Original	5/8	8/8	2/8	2/8	17/32	**85.0**
EW-T reconstruction	5/8	8/8	1/8	1/8	15/32	82.5
EW-A reconstruction	0/8	8/8	1/8	4/8	13/32	83.1

Findings

- Running the adapter is what makes the difference: the Executing Writer recovers 39–44% of a held-out skill's key instructions, about twice the best weight-only writer, and falls to 1.3–1.5% with a wrong adapter.
- The weights reveal which skill an adapter holds and its content words; a learned map from weights to text passed on about one bit.
- Recovered documents work: 85–91% of the original's task gain at rank 64, and 82–83% of verifier tests for an agent against 85% with the original.
- Behavioural adapters remain much harder for trained writers (1–2% of instructions, about 7% for a writer trained on them), although the adapters themselves carry the skill.

7 Reading Behavioural Adapters with an Executing Agent

A behavioural adapter never saw its document, so there is no stored text to extract; what it holds is a change in behaviour. This section tests a reader that needs no training: an agent that runs the base and adapted models on real task instances, looks at where they act differently, and writes a skill from those differences. We ask three questions, all judged by tasks solved. Does the skill recover the adapter's benefit? Does it come from the adapter, or would the agent have written it anyway? Does it work on models the adapter cannot be attached to? Every experiment in this section, including its pass criteria, was fixed in a written plan before it ran.

7.1 Setup

The agent. The writer is a Claude model running headless in Claude Code, in a sandbox where its only command is a tool, `fx`, that wraps the models. With it the agent can play a development game with the base or the adapted model, play both on the same game and see the first step where they differ, ask both models for their next action at a chosen state, try a draft skill on the base model, and score a draft. All of these tools observe the models' outputs; two further tools that expose internal signals (Appendix G) were available but never used, so this reader works from behaviour. Test instances never pass through the tool, and a canary file placed beside the test results never appeared in any transcript. Skills are limited to about 800 words. We call one independent run of the agent, from an empty sandbox to a finished skill, a *session*. Sessions of the same arm can differ a lot, so we run several per arm and report intervals that resample both test items and sessions.

Scoring signals. In the main one-session arm the agent never sees a reward: its score is how often the base model with its draft picks the same action as the adapted model at about 500 development states, so anything it learns about the task it learns from the adapter. In the *reward loop*, the writer instead sees task success on 60 development games, revises its drafts over up to five rounds, and the final round is chosen on 80 further development games that the writer never sees; the test games are played once per final skill. A linter rejects drafts that name specific objects ("drawer 2"), copy long passages or exceed 800 words. An audit of the test wins, described in Appendix G.1, found no exploits.

Controls. The agent is also run with a *wrong adapter* (from another task), with the *task hidden*, with only the *task text*, and with ten expert *demonstrations* instead of the adapter. If a control scores like the real arm, the skill did not come from the adapter.

Tasks. The main benchmark is ALFWorld [7], a text-based household game in which an agent must, for example, find a mug, heat it and put it on a shelf: 134 unseen test games and separate development games.

The actor is Qwen2.5-7B-Instruct, playing greedily. Its LoRA (rank 16, one epoch on 1,477 expert games) wins 81.3% of test games against 25–26% for the base model. We also use a docker-compose adapter that we had trained earlier for a separate project, with no thought of reading it back (Appendix F), five synthetic and tool-calling tasks on Qwen3.5-4B (Appendix E), and ALFWorld LoRAs trained to follow hidden rules (Section 7.3). We report the share of the LoRA's gain a skill recovers, $R = (\text{skill} - \text{base})/(\text{LoRA} - \text{base})$.

7.2 How much of the LoRA's gain does a skill recover?

Table 8: ALFWorld, 134 unseen games, Qwen2.5-7B-Instruct as actor. Each skill is given to the base model. R: share of the LoRA's gain recovered.

Arm	Sessions	Success	Range	R
References				
LoRA	–	**0.813**		1.00
Base model (one session / loop protocol)[a]	–	0.254 / 0.261		0.00
Three in-context expert games[a]	–	0.709 / 0.679		0.81 / 0.76
One session, Haiku writer, no reward				
Agent reads the adapter	6	0.581	0.42–0.72	0.58
with dev reward instead of agreement	3	0.592	0.55–0.62	0.60
Agent reads 10 demonstrations	6	0.588	0.46–0.74	0.60
Agent with a wrong adapter	6	0.521	0.30–0.64	0.48
Agent with the task hidden	6	0.496	0.02–0.66	0.43
Agent with the task text only	6	0.113	0.07–0.15	−0.25
Reward loop, Sonnet writer, five rounds				
Loop with the adapter	12	0.653	0.54–0.69	0.71
Loop without any adapter	12	**0.673**	0.59–0.81	**0.75**
Loop with 10 demonstrations	3	0.649	0.58–0.72	0.70
Loop with a wrong adapter	2	0.660	0.60–0.72	0.72
Loop with the adapter, Haiku writer	4	0.618	0.56–0.69	0.65
One round with the adapter, agreement score	4	0.483	0.27–0.58	0.40
Loop skill + three in-context games (adapter / none)	12 / 12	0.795 / 0.778		0.97 / 0.94

[a] The two groups were run separately with step caps of 50 and 52; base model and in-context references are given for each. The LoRA scores 0.813 under both. "Range" is the lowest and highest session.

Table 8 shows that running the adapter gives a good skill and that a reward loop gives a better one. In one session with no reward, the agent lifts the base model from 0.254 to 0.581, 58% of the LoRA's gain. Its skills move the base model's actions measurably closer to the LoRA's (agreement at matched states 0.646, against 0.581 for the base model), and watching the models matters: an agent that sees only the task text and the reward writes skills that make the actor worse (0.113). With a reward loop, a stronger writer and five rounds, the skill alone reaches 0.653–0.673 ($R = 0.71$–0.75). Selection on unseen development games carried over to the test games, which score only about 3 points lower. The pre-set goal of $R \geq 0.80$ for a skill alone was not reached, and no skill reached the LoRA. Combining a loop skill with three in-context games reaches 0.78–0.80 ($R \approx 0.95$), against 0.68 for the games alone; this is a different protocol, since it adds demonstrations to the prompt.

The controls come close, however. In one session, a wrong adapter (0.521), a hidden task (0.496; 0.59 without one session whose skill made the actor repeat `look`) and ten demonstrations (0.588) all score near the real arm. In the reward loop, the same loop without any adapter does as well as with it: the difference is −2.0 points, with interval $[-7.6, +3.8]$ (the pre-set goal was +5). This LoRA imitates an expert, so it holds what any competent household agent does, and a capable writer can infer most of that from the task and the

reward. Session spread is large: one loop session without the adapter reached 0.806, close to the LoRA, but sessions of the same arm span 22 points.

7.3 Does the skill come from the adapter?

Table 9: Skills written from the real adapter versus a control, across tasks. Intervals resample items and sessions; ✓: excludes 0.

Task (control)	What the adapter learned	Real	Control	Difference [95% CI]
Docker run→compose (wrong)	arbitrary output conventions	0.212	0.000	$+0.212$ [$+0.127$, $+0.315$] ✓
S4 ticket triage (wrong)	arbitrary rule	0.873	0.645	$+0.228$ [$+0.148$, $+0.323$] ✓
S2 fee calculation (wrong)	arithmetic rule	0.467	0.394	$+0.072$ [-0.034, $+0.176$]
S1 house format (wrong)	learnable from inputs	0.908	0.893	$+0.015$ [-0.028, $+0.063$]
ALFWorld, one session (wrong)	expert household play	0.581	0.521	$+0.060$ [-0.061, $+0.184$]
ALFWorld, loop (no adapter)	expert household play	0.653	0.673	-0.020 [-0.076, $+0.038$]
ALFWorld hidden rules[b] (decoy)	one-step habits	0.029	0.000	$+0.029$ [$+0.006$, $+0.061$] ✓

[b]Rule success, pooled over two rules: the share of games won while following the rule.

Table 9 shows that the adapter is the source of the skill when its behaviour cannot be guessed. On docker, skills from the real adapter reach 0.212 exact match against 0.000 in every wrong-adapter session: the reference output follows conventions, such as a fixed project name and a null default network, that the base model and the agent have no way to know. On the synthetic triage task S4 the real adapter gives 0.873 against 0.645. Where the behaviour is guessable or learnable from the inputs (ALFWorld, S1, S2) the intervals include zero.

Hidden rules. To test the unguessable case on ALFWorld itself, we trained two more LoRAs on the same expert games with one arbitrary step inserted: LoRA-A says `inventory` just before the winning action, and LoRA-B says `look` right after the first successful `take`. They follow their rules in 87% and 85% of test games, while the base model never does. Each skill is scored under both rules, so a skill written from LoRA-A is the real arm for rule A and a competent decoy for rule B. The agents found the rules: the skill text stated the right rule in 10 of 12 sessions and the wrong rule in none, while agents reading demonstrations or the base model never stated either. But the actor rarely carried them out: rule success was only 0.5% for rule A and 5.3% for rule B. Even a skill that states the rule outright makes the 7B actor follow it in only 3% (A) and 19% (B) of games, so this test is limited by the actor more than by the reader.

Rules the actor can follow. We then searched for rules the actor does follow from text and found two (65% and 53% compliance when stated): `inventory`, or examining the object, right after the first take. With LoRAs for these rules and the reward loop, skills reached 9–41% rule success for the first rule and 0–22% for the second, against 78% for the LoRAs and 0% for the decoy rule in every session (details in Appendix H). The agent can read an arbitrary rule from a LoRA, but turning it into a skill that a small actor follows reliably remains hard.

7.4 Portability and what the skills add

Table 10: The same skill documents, unchanged, on other actors (ALFWorld, 134 unseen games). The LoRA cannot be attached to any of them.

Actor	Base	One-session skill	Loop, adapter	Loop, no adapter	Three games
Qwen2.5-7B (the LoRA's base)	0.254	0.612	0.676	0.658	0.709
Llama-3.1-8B-Instruct	0.052	**0.545**	0.454	0.405	0.269
Qwen3.5-4B	0.022	**0.664**	0.385	0.564	0.396
Claude Haiku 5.5 (18 games)	0.944	1.000	n.r.	n.r.	n.r.

One-session skill: the session chosen on development games. Loop columns: mean of the first five sessions of each arm, with no selection. n.r.: not run.

Table 10 shows the main advantage of a document: it moves to any model. The one-session skill lifts Llama-3.1-8B from 0.052 to 0.545 and Qwen3.5-4B from 0.022 to 0.664, gains of 49 and 64 points, larger than its 36 on the home model; three in-context games transfer much less ($+22$ and $+37$). Part of this is formatting: both models fail mostly through invalid commands, which any skill repairs. Skills from the reward loop also transfer, though less than the one-session skill (Llama-3.1-8B 0.41–0.45, Qwen3.5-4B 0.39–0.56, where two loop sessions with the adapter failed), perhaps because the loop fits them to the home actor. Loop skills written without any adapter transfer as well as those written with one, so what moves between models is the document itself. Claude Haiku is already at ceiling, so its row says little.

Table 11: Success by ALFWorld task type on the home actor (reward-loop skills, mean of 12 sessions). Games per type in parentheses.

	Pick (24)	Examine (18)	Clean (31)	Heat (23)	Cool (21)	Pick two (17)
Base model	0.50	0.06	0.16	0.39	0.19	0.24
LoRA	**1.00**	0.56	**1.00**	**1.00**	**1.00**	0.00
Loop skill, adapter	0.84	**0.64**	0.69	0.75	0.54	0.35
Loop skill, no adapter	0.83	0.61	0.70	0.73	0.65	**0.43**

Table 11 shows that the skills and the LoRA succeed on different games. The LoRA solves every pick, clean, heat and cool game and no game that asks for two objects; the skills solve 35–43% of those and more of the examine-in-light games than the LoRA, while losing a quarter to a half of the clean, heat and cool games. A written rule such as "after placing the first object, go back for the second" is easy to state and easy for the actor to follow; a long fixed sequence of steps is where the LoRA's training pays off.

Text on top of the LoRA. Adding a skill to the LoRA helps only if it is short. A 148-word skill raises the LoRA from 0.813 to 0.858 ($+4.5$ points, interval over games $[+1.5, +8.2]$), and three in-context games give 0.836; the full 847-word skill lowers it to 0.716, and the reward-loop skills to 0.74–0.77. The LoRA was trained with no text in its prompt, and a long text moves it away from what it saw in training; Appendix G gives every placement we tried.

7.5 What the weights show

The agent above runs the models. We also asked whether the hidden rules can be seen inside the model. A logit lens on the singular directions of the weight update is not specific: the rule words rank first for the plain ALFWorld LoRA too, as do most command words. A lens on the *activations* is specific. We take the difference between the adapted and base model's hidden states at 200 development states, subtract the mean over the three ALFWorld LoRAs, and read it through the output embedding: `inventory` ranks first at

layer 26 for LoRA-A (447 and 5,113 for the other two), and `look` first at layer 24 for LoRA-B (222 and 4,400); the layer was chosen after the fact. Ablations place the rule in the MLP weights of the middle and upper layers: keeping only the MLP updates of LoRA-A preserves 92% rule success on development games, attention alone 7%, and the first nine layers alone none. These lenses show that a hidden rule leaves a specific trace inside the adapted model, and where that trace sits; turning such traces into a complete, usable skill without running the model is left for future work (Appendix I gives the details).

Findings

- An agent that runs the base and adapted models writes skills that recover 58% of an ALFWorld LoRA's gain in one session and 71% with a reward loop, and its skills transfer to models the LoRA cannot be attached to (Llama-3.1-8B $0.05 \rightarrow 0.55$, Qwen3.5-4B $0.02 \rightarrow 0.66$).
- The adapter is the source of the skill when its behaviour cannot be guessed (docker 0.21 against 0.00; an arbitrary rule 0.87 against 0.65). On ALFWorld a reward loop with no adapter does as well (-2.0 points, $[-7.6, +3.8]$).
- Hidden one-step rules are read by the agent (10 of 12 sessions) but rarely followed by a 7B actor from text; rules the actor can follow are partly transferred (9–41% against 0% for the decoy).
- Skills and the LoRA fail on different task types; short text helps the LoRA and long text hurts it. A centred activation lens finds the hidden rule's word; a lens on the weight update alone does not.

8 Discussion

What the results support. An adapter's document can be recovered to a useful degree from the adapter and its base model, and the most effective route we found is to let the adapted model do the reading. The wrong-adapter control, the retrieval baseline and the unsupported-content measure show that the recovered text depends on the adapter and is more faithful than the alternatives. Recovery is far from exact on the hypernetwork's own adapters (44% of key instructions, 39% of units unsupported) and high on adapters that recite their document (90–93% of prose at rank 64), where higher rank helped most.

Why recitation and behaviour differ. A recitation adapter holds the document's tokens, so the adapted model can be asked to produce them. A behavioural adapter holds the effect of the document, and nothing in its training asked it to write the document down. Our trained writers saw only the first kind, so their weak result on the second is a statement about their training data as much as about the method; in Appendix D, a writer trained on a small bank of behavioural adapters improved slightly and then overfitted. The executing agent sidesteps training and reads behaviour directly, but on a task like ALFWorld a capable writer can produce the same skill without the adapter. The adapter's information is most valuable exactly where behaviour is arbitrary.

Why recovery is worth having. Even where a LoRA scores higher than any document, the document keeps properties the adapter lacks: it can be read before use, edited without retraining, merged with other skills and given to any model. A skill an agent wrote from a Qwen2.5 LoRA lifted Llama-3.1-8B from 5% to 55%, and a skill written for a 4B model's fee-calculation task raised Claude Haiku from 0% to 96% (Appendix E), a model on which the LoRA cannot be loaded at all. Recovery is also a way to move a skill off the model it was compiled for.

Limitations. Most positive results are on adapters that store their document; the hypernetwork's adapters are one family and Qwen3 the only base family for the trained writers. Coverage relies on LLM judges with no human labels: the two judges agree on 272 of 485 verdicts and the second is 7 to 21 points more

generous, so absolute numbers are less certain than the order of methods. The main skill set has 37 skills and the benchmark agent test four tasks. The executing-agent results rest on one main benchmark and one home actor, with 2 to 12 sessions per arm, and the spread between sessions is often larger than the gap between arms. The 7B actor limits what any skill can show for hidden rules. Finally, the original document is an imperfect target: an adapter may not have learned every instruction, and different wording can have the same effect.

Future work. The main open problem is behaviour that is both arbitrary and followable: an adapter whose behaviour cannot be guessed is where reading it matters most, yet a small actor often fails to carry out the rule once it is written down. Three directions follow. First, breadth: more benchmarks (SkillsBench, WebShop, ScienceWorld, τ-bench), more adapters trained by others, and a second base-model family, with human labels for a sample of documents. Second, the unguessable case: adapters whose behaviour is arbitrary but easy to follow (a format, a required field, a naming convention), the reward loop run on such adapters to see whether the adapter then adds what it did not add on ALFWorld, and stronger actors. Third, better readers: writers trained on several hundred behavioural adapters with early stopping, agents that combine execution with activation lenses on the middle layers, and searches that use the adapter's own answers to score candidate documents.

Broader impact. Recovering a document from an adapter helps audit and edit adapters before they are run, and lets a user check that an adapter matches its stated skill. The same capability lets someone who holds a released adapter reconstruct the skill it compiles, which may be proprietary, so releasing only weights should not be treated as keeping a skill private. Recovered documents are partial and contain unsupported content, and should be checked before they are trusted.

9 Conclusion

Skills compiled into adapters are cheap to run but lose what makes a skill useful: a document that can be read, edited, audited and given to another model. We asked whether that document can be written back, and judged the answer by the instructions it recovers and the tasks it solves. Reading the weights alone recovered a skill's identity and its words, but little of its instructions. Running the adapter worked. A small writer trained on top of the adapted model recovers 39–44% of a held-out skill's key instructions, about twice the best weight-only writer, and almost nothing with a wrong adapter. On adapters that store their document it recovers 90–93% of the prose, and the recovered documents restore 85–91% of the original's task gain and help an agent on benchmark tasks nearly as much as the original. For adapters trained on behaviour alone, an agent that runs the base and adapted models writes skills that recover 58–71% of an ALFWorld LoRA's gain, read rules that no one could guess, and carry the skill to models the LoRA cannot be loaded on, lifting Llama-3.1-8B from 5% to 55%. Together these results show that an adapter need not be a sealed artefact: run inside its base model, it can be turned back into a document that people can read and other models can use.

References

[1] Anthropic. Equipping agents with Agent Skills. Anthropic Engineering blog, 2025. https://www.anthropic.com/engineering/equipping-agents-for-the-real-world-with-agent-skills

[2] X. Li et al. SkillsBench: Benchmarking How Well Agent Skills Work Across Diverse Tasks. 2026. arXiv:2602.12670.

[3] E. J. Hu et al. LoRA: Low-Rank Adaptation of Large Language Models. *ICLR*, 2022. arXiv:2106.09685.

[4] R. Charakorn et al. Text-to-LoRA: Instant Transformer Adaption. *ICML*, 2025. arXiv:2506.06105.

[5] R. Charakorn, E. Cetin, S. Uesaka and R. T. Lange. Doc-to-LoRA: Learning to Instantly Internalize Contexts. *ICML*, 2026. arXiv:2602.15902.

[6] Y. Liu, X. Wang, Y. Mao, Y. Gelberg, H. Maron and M. Zhang. SHINE: A Scalable In-Context Hypernetwork for Mapping Context to LoRA in a Single Pass. *ICML*, 2026. arXiv:2602.06358.

[7] M. Shridhar, X. Yuan, M.-A. Côté, Y. Bisk, A. Trischler and M. Hausknecht. ALFWorld: Aligning Text and Embodied Environments for Interactive Learning. In *ICLR*, 2021.

[8] T. R. Shaham, S. Schwettmann, F. Wang, A. Rajaram, E. Hernandez, J. Andreas and A. Torralba. A Multimodal Automated Interpretability Agent. In *ICML*, 2024.

[9] A. Yu et al. LatentSkill: From In-Context Textual Skills to In-Weight Latent Skills for LLM Agents. 2026. arXiv:2606.06087.

[10] X. Zhao et al. Parametric Skills. 2026. arXiv:2606.30015.

[11] Y. Sheng et al. S-LoRA: Serving Thousands of Concurrent LoRA Adapters. *MLSys*, 2024. arXiv:2311.03285.

[12] X. Han, F. Neri, Z. Jiang, F. Wu, Y. Ye, L. Yin and Z. Wang. W2T: LoRA Weights Already Know What They Can Do. 2026. arXiv:2603.15990.

[13] C. De Schamphelaere, J. Bauer, N. Nanda, E. Ong. LoRAcles: Self-Supervised Weight-Space Interpretability at Scale. *NeurIPS*, 2026.

[14] A. Goel, Y. Kim, N. Shavit and T. T. Wang. Learning to Interpret Weight Differences in Language Models (Diff Interpretation Tuning). *ICLR*, 2026. arXiv:2510.05092.

[15] Z. Zhong and A. Raghunathan. Watch the Weights: Unsupervised Monitoring and Control of Fine-tuned LLMs. *ICLR*, 2026. arXiv:2508.00161.

[16] R. Paul. Spectral Geometry of LoRA Adapters Encodes Training Objective and Predicts Harmful Compliance. 2026. arXiv:2604.08844.

[17] G. Wang et al. Voyager: An Open-Ended Embodied Agent with Large Language Models. *TMLR*, 2024. arXiv:2305.16291.

[18] S. Yao et al. ReAct: Synergizing Reasoning and Acting in Language Models. *ICLR*, 2023. arXiv:2210.03629.

[19] J. Yang et al. SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. *NeurIPS*, 2024. arXiv:2405.15793.

[20] N. Houlsby et al. Parameter-Efficient Transfer Learning for NLP. *ICML*, 2019. arXiv:1902.00751.

[21] X. L. Li, P. Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. *ACL*, 2021. arXiv:2101.00190.

[22] A. Aghajanyan, L. Zettlemoyer, S. Gupta. Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning. *ACL*, 2021. arXiv:2012.13255.

[23] D. Biderman et al. LoRA Learns Less and Forgets Less. 2024. arXiv:2405.09673.

[24] D. Ha, A. Dai, Q. V. Le. HyperNetworks. *ICLR*, 2017. arXiv:1609.09106.

[25] Q. Ye, X. Ren. Learning to Generate Task-Specific Adapters from Task Description (HYPTER). *ACL*, 2021. arXiv:2101.00420.

[26] A. Askell et al. A General Language Assistant as a Laboratory for Alignment. 2021. arXiv:2112.00861.

[27] C. Snell, D. Klein, R. Zhong. Learning by Distilling Context. 2022. arXiv:2209.15189.

[28] J. Mu, X. L. Li, N. Goodman. Learning to Compress Prompts with Gist Tokens. *NeurIPS*, 2023. arXiv:2304.08467.

[29] S. Eyuboglu et al. Cartridges: Lightweight and General-Purpose Long Context Representations via Self-Study. 2025. arXiv:2506.06266.

[30] C. Huang et al. LoraHub: Efficient Cross-Task Generalization via Dynamic LoRA Composition. 2023. arXiv:2307.13269.

[31] G. Ilharco et al. Editing Models with Task Arithmetic. *ICLR*, 2023. arXiv:2212.04089.

[32] T. Putterman et al. Learning on LoRAs: GL-Equivariant Processing of Low-Rank Weight Spaces for Large Finetuned Models. 2024. arXiv:2410.04207.

[33] K. Schürholt et al. Hyper-Representations as Generative Models: Sampling Unseen Neural Network Weights. 2022. arXiv:2209.14733.

[34] K. Wang et al. Neural Network Diffusion. 2024. arXiv:2402.13144.

[35] nostalgebraist. Interpreting GPT: the logit lens. LessWrong, 2020. https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru

[36] N. Belrose et al. Eliciting Latent Predictions from Transformers with the Tuned Lens. 2023. arXiv:2303.08112.

[37] A. Ghandeharioun et al. Patchscopes: A Unifying Framework for Inspecting Hidden Representations of Language Models. *ICML*, 2024. arXiv:2401.06102.

[38] H. Chen et al. SelfIE: Self-Interpretation of Large Language Model Embeddings. *ICML*, 2024. arXiv:2403.10949.

[39] J. Betley et al. Tell me about yourself: LLMs are aware of their learned behaviors. 2025. arXiv:2501.11120.

[40] N. Carlini et al. Extracting Training Data from Large Language Models. *USENIX Security*, 2021. arXiv:2012.07805.

[41] M. Nasr et al. Scalable Extraction of Training Data from (Production) Language Models. 2023. arXiv:2311.17035.

[42] J. X. Morris et al. Text Embeddings Reveal (Almost) As Much As Text. *EMNLP*, 2023. arXiv:2310.06816.

[43] J. X. Morris et al. Language Model Inversion. *ICLR*, 2024. arXiv:2311.13647.

[44] A. van den Oord, Y. Li, O. Vinyals. Representation Learning with Contrastive Predictive Coding. 2018. arXiv:1807.03748.

[45] A. Radford et al. Learning Transferable Visual Models From Natural Language Supervision. *ICML*, 2021. arXiv:2103.00020.

[46] Y. Zhang et al. Qwen3 Embedding: Advancing Text Embedding and Reranking Through Foundation Models. 2025. arXiv:2506.05176.

[47] L. Zheng et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. *NeurIPS Datasets and Benchmarks*, 2023. arXiv:2306.05685.

[48] W. Kwon et al. Efficient Memory Management for Large Language Model Serving with PagedAttention. *SOSP*, 2023. arXiv:2309.06180.

[49] A. Yang et al. Qwen3 Technical Report. 2025. arXiv:2505.09388.

[50] Y. Zhou et al. Large Language Models Are Human-Level Prompt Engineers. *ICLR*, 2023. arXiv:2211.01910.

[51] C. Yang et al. Large Language Models as Optimizers. *ICLR*, 2024. arXiv:2309.03409.

[52] O. Khattab et al. DSPy: Compiling Declarative Language Model Calls into Self-Improving Pipelines. *ICLR*, 2024. arXiv:2310.03714.

[53] M. Yuksekgonul et al. Optimizing generative AI by backpropagating language model feedback (TextGrad). *Nature*, 2025. arXiv:2406.07496.

[54] L. A. Agrawal et al. GEPA: Reflective Prompt Evolution Can Outperform Reinforcement Learning. 2025. arXiv:2507.19457.

[55] N. Shinn et al. Reflexion: Language Agents with Verbal Reinforcement Learning. *NeurIPS*, 2023. arXiv:2303.11366.

[56] A. Zhao et al. ExpeL: LLM Agents Are Experiential Learners. *AAAI*, 2024. arXiv:2308.10144.

[57] Z. Z. Wang et al. Agent Workflow Memory. 2024. arXiv:2409.07429.

[58] S. Marks et al. Auditing Language Models for Hidden Objectives. 2025. arXiv:2503.10965.

[59] Qwen Team. Qwen2.5 Technical Report. 2024. arXiv:2412.15115.

[60] A. Grattafiori et al. The Llama 3 Herd of Models. 2024. arXiv:2407.21783.

A Evaluation details

Judges. The primary judge is a judge model (JEV-1.13) that returns a probability for each verdict on each instruction, so coverage is an expectation. The second judge (Gemini 2.5 Pro) returns one verdict per instruction. The two agree on 272 of 485 instruction verdicts. Most disagreements are a "full" from the second judge against a "partial" from the primary; in a random audit of 12 disagreements the second judge was too generous in about 10 (for example, accepting "cherry-pick supporting data" as stating "avoid cherry-picking"), so its coverage runs about 20 points higher. Both judges were reached through a commercial API router.

Frozen requirements. The key instructions of each original are extracted once, before any method is scored, and reused for every method and judge. For the external skills no model extracts requirements: the prose list is every instruction sentence of the original (6–60 words, with code, tables and headings removed) and the structure list is every code block, table and heading. Requirements are judged in chunks of 15 with the whole original and the whole candidate in each call (817 prose and 931 structure units over 41 skills). Units are not weighted by importance and may repeat each other, so these numbers compare methods on identical questions rather than giving an absolute share of the skill.

Whole-document judge. Table 5 also uses a judge that needs no checklist: Gemini 3 Flash (temperature 0) reads the complete original and reconstruction, lists every instruction, command, parameter or warning of the original, marks each as same, partial, missing or contradicted, and states whether an agent given the reconstruction would behave like one given the original (76–90% of skills for recitation adapters at rank 64).

Unsupported content. Units of the candidate are split by rule, not extracted by a model, and the judge is asked whether the original supports each sampled unit.

Intervals. For the executing agent, intervals are a paired bootstrap that resamples test items and, within each arm, sessions. Intervals over items alone are narrower and are not used for claims. All plotted values are drawn from the tables; Figures 1 and 2–3 are schematic.

B Transfer and Rank: Full Results

Table 12: Executing Writers on adapters fitted to recite each of the 37 test skills (coverage, percent; whole-document match probability).

	Coverage			Wrong adapter			Whole-doc. match		
Writer	$r = 16$	$r = 32$	$r = 64$	$r = 16$	$r = 32$	$r = 64$	$r = 16$	$r = 32$	$r = 64$
EW-T (no text)	72.4	86.8	**92.0**	2.3	1.9	1.7	60.7	79.3	82.7
EW-A	79.8	87.4	**92.0**	1.9	1.9	2.0	52.2	69.3	76.1

Table 13: Executing Writers on 41 SkillsBench skills with recitation adapters (percent). Identifier recall is counted by rule, with no judge.

Writer	Rank	Prose cov.	Structure cov.	Identifier recall	Match prob.
EW-T (no text)	16	84.1	85.1	85.1	74.3
EW-T (no text)	32	92.1	93.4	**94.4**	79.1
EW-T (no text)	64	**93.3**	**94.0**	94.1	**81.2**
EW-A	16	79.7	80.1	78.0	52.1
EW-A	32	86.8	86.9	87.7	64.8
EW-A	64	90.2	90.2	91.2	73.1

The fitted adapters reach 100% token accuracy at every rank ($r = 16$ already stores a 1,451-token document exactly after 25 steps). Coverage rises with rank and levels off near the effective rank of the updates (participation ratio 11.6 at $r = 16$ and 40.7 at $r = 64$). The wrong-adapter control is higher on the external skills (9–13%) because these code-heavy documents share boilerplate. The three checklists agree within about one point, so leaving code out of the prose list did not inflate the scores.

C Examples, Self-Authored Agent Tasks and Cost

Original (Zustand store skill)

````
## Basic Store
```typescript
import { create } from "zustand";
interface CounterStore {
  count: number;
  increment: () => void;
  decrement: () => void;
  reset: () => void;
}
const useCounterStore = create<CounterStore>((set...
````

EW-A, coverage 68%

````
## Basic Store
```typescript
import { createStore } from "zustand";
interface Counter {
  count: number;
  increment: () => void;
}
const useCounter = createStore<Counter>()((set) =>...
```
## Persist Middleware
````

Original (GSAP routing skill)

```
# GSAP Router
Routes to 4 specialized skills based on animation...
## Routing Protocol
1. **Classify** -- Identify animation type from u...
2. **Match** -- Apply signal matching rules below
3. **Combine** -- Production animations often nee...
4. **Load** -- Read matched SKILL.md files before...
```

EW-A, coverage 27%

```
# GS Router
Routes animations to specialized skills based on ...
## Routing Protocol
### Tier 1: Core Requirements (from signal)
1. **Class** -- Identify animation type from `ani...
2. **Combine** -- Apply combined SKILL routing be...
3. **Match** -- Production signals often need prod...
```

Figure 5: Original and recovered openings for two test skills with no near-copy in the training corpus (long lines truncated). Top: structure and code are kept, one API name differs and a method is dropped. Bottom: the steps are reworded and a second tier is invented; the document is almost twice as long as the original. At rank 64 on recitation adapters, recovery is near-verbatim for hundreds of lines.

Self-authored agent tasks. We also wrote one task per pilot skill (41 tasks with our own prompts, verifiers and reference solutions) and ran the same Haiku agent: four tasks 8 times per arm and 37 tasks 3 times. Because we wrote the verifiers, this is a weaker check than the benchmark tasks. On 36 tasks the agent passes with or without any skill, and the original and both reconstructions match it; two fail in every arm. On the three tasks where the skill matters, the original passes 18 of 19 runs, EW-T 12 and EW-A 11, against 1 with no skill; on threat triage EW-T reproduces the original exactly ($8/8$ against $0/8$ with no skill) while EW-A fails. Over all 142 runs per arm: no skill 117, original 134, EW-T 130, EW-A 125 strict passes.

Benchmark task screening. Of the 17 pilot SkillsBench tasks we screened 13 without a skill and kept six that the agent failed; two were dropped after checking each with its reference solution. One needs a live download blocked in the sandbox; in the other, the expected value contradicts the workbook's own rule (the filled cells of the same row are six-year averages, matching 6721.8, 859.3 and 1100.2 exactly, while the verifier expects a five-year average). With the skill given as a file in the sandbox instead of in the prompt (3 runs per task), the original and both reconstructions passed 4, 4 and 5 of 12 runs, but the agent opened the skill in only about half of them.

Cost. One document at a time the Executing Writer decodes about 8 tokens per second (69 minutes for 37 documents). The adapter layer accepts a batch of adapters, so documents can be written together with each row using its own adapter; rows do not influence each other. With 8 documents at a time the 37 rank-64 adapters take 18.5–19.9 minutes (3.4–3.7$\times$ faster) with unchanged coverage.

D Training Writers on Behavioural Adapters

Section 6.4 showed that writers trained on recitation adapters do not read behavioural ones. Here the writer is trained on behavioural adapters instead (EW-B). This is a smaller setting (Qwen3-4B base, judge-free metrics), so its numbers are not comparable with the main text.
```

Setup. For each skill, the base model writes about 60 requests and answers them with the document in its prompt; a rank-16 LoRA is fitted to give those answers without the document. All 78 fitted adapters lowered held-out answer loss by at least 10% (median 69%). 54 were used for training and the rest for testing, together with 4 skills made with the main paper's recipe. EW-B is the Executing Writer of Section 4, trained on 11, 22 and 45 adapters. We compared it with training-free methods: generic and domain-targeted probes, vocabulary-lens words, *amplified answers* (sampling from $\log p_{\text{adapter}} + \alpha(\log p_{\text{adapter}} - \log p_{\text{base}})$, $\alpha = 1$), best-of-n sampling, self-report and retrieval. Metrics are judge-free: *behavioural fidelity* M1 is the share of the original document's help to the likelihood of the prompted model's answers that a candidate provides (1 for the original, 0 for none), and coverage counts the original's units that have an embedding neighbour in the candidate (cosine ≥ 0.757, calibrated on documents that restate each other).

Table 14: Judge-free recovery on held-out behavioural adapters (Qwen3-4B). Paper recipe: adapters made as in Section 6.4.

Set	Method	n	M1, own	M1, wrong	Cov. own	Cov. wrong
Held-out skills	Generic probes	8	−0.14	−0.05	0.0	0.0
	Vocabulary lens	8	−0.04	−0.09	6.6	0.0
	Amplified answers	7	**0.19**	−0.07	4.5	0.0
	Plain request	8	0.13	0.02	0.7	0.0
	Best of n	7	0.11	0.00	5.6	0.1
	Retrieval	8	0.03	−0.07	1.9	0.7
	EW-B, 45 adapters	8	0.11	0.00	**7.3**	1.2
Paper recipe	Amplified answers	3	−0.01	−0.04	0.0	0.0
	Plain request	6	−0.00	−0.09	0.0	0.0
	EW-B, 45 adapters	4	−0.02	−0.05	1.1	0.0

Results. The writer depends on the adapter: the correct adapter lowers validation loss by about 0.7 nats per token against a wrong one, against 0.001 for the soft-prompt writer of Table 4. But it overfits within 30 of 120 steps and memorises its training documents. It has the best coverage on held-out skills (7.3%, interval $[2.4, 13.5]$, against 1.2% for a wrong adapter; sign test $p = 0.36$, $n = 8$), far below the 20% we had set as success, and 1.1% on the paper-recipe adapters (Table 14). Its documents name the right area and miss the specifics (for an architecture-diagram skill: "AI Assistant for Code Review and Documentation"). Coverage rises from 0.0% with 11 training adapters to 5.3% with 45, but only 2 to 14 documents were scored per size, so this is a direction, not a measurement. No training-free method beats the wrong-adapter control on coverage by more than 4 points; amplified answers are the exception on behavioural fidelity (0.19 against −0.07, $p = 0.011$, $n = 10$), so they carry some adapter-specific behaviour without stating its rules. As a check of the stacking mechanism, the public Diff Interpretation Tuning adapter [14] loaded in the writer slot named a hidden topic from a weight difference in 11 of 12 cases at a reduced scale. A bank of several hundred behavioural adapters with early stopping is the natural next step.

E Skills versus LoRA on Synthetic and Tool-Calling Tasks

We compare a skill written by a Claude Haiku agent with a LoRA on six tasks with Qwen3.5-4B as the base model: four synthetic rule-following tasks with programmatic scorers and known true skills (S1 house format, S2 fee calculation, S3 code table, S4 ticket triage) and two real tasks, R1 text-to-SQL and R3 tool calling (the data of the public `doosra` adapters). Each task has 100 development items, visible to the agent only through a scoring command, and 200 test items. LoRAs are rank 16 on all projections, trained on 2,000 examples;

R1's LoRA gained only $+0.035$ on development and was not read. Each skill is placed in the system prompt of the base model. Arms: *P*, task text only; *D*, task text and 32, 128 or 400 training examples; *A*, the adapter only, scored by how often base plus skill reproduces the adapter's output on development prompts; *A-wrong*, a different task's adapter; *A+D*, adapter and data. One session per arm, so differences below about 0.05 are within noise.

Table 15: Test score (200 items) per task. Skills are applied to the frozen base model unless stated. "–": not defined or not run.

Arm	S1	S2	S3	S4	R1	R3
Base model	0.000	0.194	0.000	0.225	0.735	0.140
LoRA (2,000 examples)	0.891	**0.698**	**1.000**	**1.000**	**0.770**	**0.895**
True skill	0.898	0.447	0.998	0.902	–	–
Skill P (task text)	0.909	0.404	0.000	0.860	0.730	0.190
Skill D (400 examples)	0.955	0.472	**1.000**	0.993	0.665	0.245
Skill A (adapter only)	0.927	0.454	0.970	0.935	–	0.200
Skill A-wrong	0.938	0.409	0.922	0.730	–	0.145
Skill A+D	**0.964**	0.515	**1.000**	0.955	–	0.215
Claude Haiku 5.5 with a skill	1.000	0.958	1.000	0.983	0.730	0.630

Results. On the rule-like tasks S1, S3 and S4 a skill written from 400 examples matches the LoRA trained on 2,000, and with 32 examples the skill beats a LoRA trained on the same 32 on five of the six tasks (most strongly on S3, 0.999 against 0.302; R3 is the exception). The LoRA wins where the small model must execute something a prompt does not teach: on S2 (fee arithmetic) the LoRA scores 0.698 and the best skill 0.531, although Claude Haiku with the same kind of skill scores 0.958, so the document was right and the 4B model could not carry it out; on R3 (tool calling) the LoRA scores 0.895 and every skill 0.40 or less. Skills written from the adapter alone recover 92–104% of the LoRA's gain on S1, S3 and S4, but a skill from another task's adapter recovers nearly as much on S1 and S3, because the agent infers those tasks from the description and the development feedback. With repeated sessions (Table 9), the adapter clearly matters on S4. Giving the agent both adapter and data did not beat data alone.

F An Existing Adapter: Docker Run to Compose

To test reading on an adapter that was not built for this study, we used the `docker_compose_gen` LoRA that we had trained earlier for a separate project and released in the `adaptive-minds-loras` collection (rank 32, Qwen2.5-7B-Instruct). The task is to convert a `docker run` command into the exact compose file that `docker compose` would print. Data: 200 test items over ten flag families (20 each), 100 development and 400 training items, all disjoint. A response is exact if its parsed YAML equals the reference; the soft score is a field-level F1. Writers here are Claude Sonnet agents with the arms of Appendix E, plus *A-deep* (at least 25 probes before the first score) and *A+D-probe* (at least 20 adapter probes before reading any example).

Results. With no training data, after two to four probes of the adapter, the agent wrote skills that take the base model from 0.000 to 0.263 exact on average (0.235–0.300) and recover about 84% of the LoRA's soft gain. A skill from the wrong adapter and a skill from the task text alone both stay at 0.000, because the reference follows conventions the base model cannot guess (a fixed project name, a null default network, a `command` list for tokens after the image). Probing harder did not help, but forcing probes before reading data gave the best skill (0.360). Twenty in-context examples (0.435) beat every skill, and the LoRA (0.595)

Table 16: Docker-to-compose, 200 test items: exact match and soft score. Rows with several sessions give the session mean.

Arm	Exact	Soft
Base model (Qwen2.5-7B-Instruct)	0.000	0.186
LoRA `docker_compose_gen`	**0.595**	**0.885**
Wrong adapter (`dockerfile_gen`)	0.000	0.079
Skill P (task text)	0.000	0.340
Skill D (training data), 2 sessions	0.313	0.819
Skill A (adapter only), 3 sessions	0.263	0.771
Skill A-deep	0.250	0.764
Skill A-wrong, 2 sessions	0.000	0.278
Skill A+D-probe	0.360	0.859
5 / 20 in-context examples (no agent)	0.300 / 0.435	0.820 / 0.872
LoRA + skill P / A / A+D	0.450 / 0.605 / **0.690**	0.866 / 0.890 / **0.937**

beats everything; a good skill on top of the LoRA adds up to 0.095, and a vague one costs up to 0.145. Per flag family, adapter-only skills come close to the LoRA on logging (0.38 against 0.40) but recover little on numeric and structured fields (health checks 0.08 against 0.70). Renaming the service in 199 test items left the LoRA unchanged (0.593), so it had not memorised the test items. The one-session Haiku agent of Section 7 gave 0.212 with the real adapter and 0.000 with a wrong one.

G ALFWorld: Further Details

Environment and adapter. Text version of ALFWorld (`alfworld` 0.4.2). Test: the 134 `valid_unseen` games, never shown to any agent. Development: games from `valid_seen`. The actor is Qwen2.5-7B-Instruct with plain chat prompts, greedy decoding and a step cap of 50 (one-session study) or 52 (reward loop and hidden rules). The LoRA (rank 16, $\alpha = 32$, all linear layers) was trained for one epoch on $1{,}477$ expert games and passed a pre-set gate of $+15$ points on development (0.70 against 0.10).

Agent tools. `task`; `rollout` (play a development game with the base or adapted model); `compare` (both models on the same game, first divergence marked); `probe` (next action of both models at a written state); `counterfactual` and `logitdiff` (edit a state, or show the tokens the adapter boosts); `try`; `score`; `examples` (demonstrations, for arms that have them). Across all one-session ALFWorld runs the agents played 416 rollouts, 93 comparisons, 104 probes and 138 scores, and never used `counterfactual` or `logitdiff`. Agents given the wrong adapter probed far more (75 probes against 13), because comparing the two models showed them nothing.

Further one-session arms. One retrieved expert game in the prompt gives 0.522. Reading the adapter and then data gives 0.602; skills written from 50 and from all demonstrations give 0.552 and 0.527, so more data did not help. The LoRA asked to write its own skill gives 0.149, and an automatic miner of base–adapter divergences followed by one writing call gives 0.256 (sessions 0.57, 0.06, 0.13). Giving the agent a mechanical list of first divergences made the agent name rule A in every hidden-rule session and gave the steadiest win rate of any arm (session sd 0.026).

Text on top of the LoRA.

Table 17: The ALFWorld LoRA with text added to its prompt (134 unseen games, step cap 52).

Added to the LoRA	Success	Change
Nothing	0.813	–
Short skill (148 words), system prompt	**0.858**	+0.045
Short skill, after the task text	0.843	+0.030
Three in-context games	0.836	+0.023
Three games + full skill, system prompt	0.821	+0.008
Reward-loop skills, system prompt (mean of 12)	0.744–0.767	−0.05 to −0.07
Full skill (847 words), system prompt	0.716	−0.097
Full skill, after the task text	0.597	−0.216
Three games + full skill, after the task text	0.425	−0.388

Adapter dependence against a task-hidden control. Real adapter minus the task-hidden agent: ALFWorld $+0.085$ $[-0.081, +0.296]$; docker -0.028 $[-0.137, +0.082]$ (the "hidden" docker arm still sees `docker run` in its prompts, so the task is not really hidden); S1 -0.002 $[-0.021, +0.017]$. This control was not run on S2, S4 and R3. On S3 one task-hidden session scored 0.63, so S3 is dominated by session luck and we draw no conclusion from it. R3 sessions hit the time limit on a shared server and have no valid result.

Portability details. The one-session skill used on other actors is the session chosen on development success. A skill written from ten demonstrations and chosen the same way gives 0.739 on the home actor, 0.485 on Llama-3.1-8B, 0.731 on Qwen3.5-4B and 0.889 on Claude Haiku (18 games, three per task type). On docker, transfer is uneven: the data-based skill gives exact match 0.96 on Claude Haiku, 0.56 on Qwen3.5-4B and 0.14 on Llama, and the adapter-derived skill 0.44, 0.11 and 0.02.

G.1 Reward-loop protocol and audit

The 60 games the writer sees and the 80 selection games come from `valid_seen` (shuffled, seed 0); on them the base model scores 0.133 and 0.163 and the LoRA 0.767 and 0.813. Each round is a fresh call of the writer in a persistent sandbox with a notes file as memory; up to three drafts per round are played on the 60 games and the best is played on the 80. The linter rejected 43 of 101 draft attempts with the adapter and 60 of 115 without, mostly for naming specific objects; every rejected draft could be fixed and resubmitted. One session lost two rounds to an unclear linter message; dropping it moves the adapter-minus-no-adapter gap from -2.0 to -1.9 points. One further session produced no skill and is excluded. From selection to test the mean score moved from 0.683 to 0.653 with the adapter and 0.698 to 0.673 without. No test win took fewer than three steps, and a judge model read 20 sampled wins per arm and found 39 of 40 legitimate; the exception ends as won after a heating step without a placing step, using only valid commands, which we read as a scoring quirk of the environment.

H Hidden-Rule Experiments

Rules and scoring. Rule A holds if the game is won and the action just before the winning action is exactly `inventory`. Rule B holds if the game is won and the action right after the first successful `take` is exactly `look`. The original LoRA, the base model and the expert games score 0% on both rules; the expert games with the rule step inserted score 100%. LoRA-A and LoRA-B use the recipe of the ALFWorld LoRA with the loss on the inserted step weighted three times, and passed a pre-set gate after one epoch (development rule success 87.5% and 70.0%, base 0%). For the adapter arms the agent's score is agreement with that adapter; no prompt, tool output or file the agent can read contains either rule.

Table 18: Hidden-rule LoRAs on ALFWorld (134 unseen games). Rule columns give rule success: won and the rule followed.

Arm	Sessions	Win	Rule A	Rule B	Skill states its rule
Base model	–	0.261	0.000	0.000	–
LoRA-A	–	0.896	**0.866**	0.000	–
LoRA-B	–	0.851	0.000	**0.851**	–
Agent reading LoRA-A	6	0.604	0.005	0.000	5/6 (never rule B)
Agent reading LoRA-B	6	0.565	0.000	0.053	5/6 (never rule A)
Agent reading LoRA-A with a divergence list	6	0.636	0.001	0.000	6/6
Agent reading 10 rule-A demonstrations	2	0.563	0.007	0.000	0/2
Agent with the base model only	2	0.090	0.000	0.000	0/2
Skill stating rule A outright	–	0.239	0.030	0.000	yes
Skill stating rule B outright	–	0.209	0.007	0.194	yes

Contrasts. Real minus decoy is $+0.005$ $[0.000, +0.016]$ for rule A and $+0.054$ $[+0.010, +0.119]$ for rule B, pooled $+0.029$ $[+0.006, +0.061]$, far below the pre-set $+15$ points. Counting every game, won or not, the LoRA-B actor says `look` right after its first take in 99.3% of games, the actor with the outright skill in 67.9%, and actors with agent-written skills in 10.3%. The actor with the outright rule-A skill says `inventory` somewhere in 84.3% of games but just before the last step in only 7.5%. The actor can be made to say the word, but not at the right moment.

Rules the actor can follow. We tested six candidate rules by appending each to a competent skill and measuring compliance on development games: `look` first (25%), `inventory` first (45%), examine the taken object right after the first take (53%), `inventory` right after the first take (65%), `look` after the first move (20%), and `inventory` as the second action (0%). The two above 50% were used. LoRAs trained for them win 78% of development games and follow their rule in 78%, and are scored under each other's rule as decoys. With three reward-loop sessions per rule (three rounds each), test rule success was 0.09, 0.41, 0.25 for `inventory` after the first take and 0.00, 0.22, 0.00 for examining the object, against 0.00 for the decoy rule in all six sessions; win rates were 0.59–0.74.

I Reading Weights

Lenses. The weight lens takes the top singular directions of each LoRA update in the attention-output and MLP-down projections and reads them through the final norm and output embedding, with and without centring over the three ALFWorld LoRAs. For every LoRA, most ALFWorld command words reach rank 1 in some layer, direction and sign, so the best rank is not specific. The activation lens averages, over 200 development states, the difference between the adapted and base model's last-token hidden state at each layer, subtracts the mean of this difference over the three LoRAs, and reads it through the output embedding.

Ablations. Keeping the update of LoRA-A only in some layers or modules, on 60 development games (win / rule A): full 0.90 / 0.90; layers 0–8 0.48 / 0.00; layers 9–18 0.83 / 0.52; layers 19–27 0.50 / 0.48; MLP only 0.92 / 0.92; attention only 0.77 / 0.07. For the ALFWorld LoRA, layers 9–18 alone give 0.93 and MLP only 0.97 on development, above the full adapter's 0.77; differences of a few games are within noise.